



A.A. 2010/2011

Calcolo Scientifico per l'Informatica - Laboratory Class - 1

Davide Fugazza
fugazza@elet.polimi.it

Outline

- 1 Introduction to MATLAB
- 2 Floating-Point: IEEE-754
- 3 Accuracy of the Arithmetic
- 4 Loss-of-Significance errors
- 5 Homework

Outline

- 1 Introduction to MATLAB
- 2 Floating-Point: IEEE-754
- 3 Accuracy of the Arithmetic
- 4 Loss-of-Significance errors
- 5 Homework

MATLAB[®]: a quick review

A list of some basics:

- MATLAB[®] user interface;
- variable definition: scalars, vectors and matrices, chars;
- vector and data matrix manipulation;
- code modularity: functions and scripts;
- data plots and simple graphics commands;
- I/O: loading and saving files, variables, and data.

These Labs are not intended to provide an introduction to MATLAB[®]. The student is supposed to have already developed a basic knowledge of the environment throughout his/her academic carrier.

Outline

- 1 Introduction to MATLAB
- 2 Floating-Point: IEEE-754**
- 3 Accuracy of the Arithmetic
- 4 Loss-of-Significance errors
- 5 Homework

Introductory examples

Cleve Moler's code sequence (From [4])

```
>>format long
>> a = 4/3
a =
    1.333333333333333

>> b = a - 1
b =
    0.333333333333333

>> c = 3*b
c =
    1.000000000000000

>> e = 1 - c
e =
    2.220446049250313e-016
```

Introductory examples, cont'd

Compare the results of the following code segments:

Example 1

```
>> x = 0;  
>> while x ~= 1  
>> x = x + 1/16;  
>> [x, sqrt(x)]  
>> end
```

Example 2

```
>> x = 0;  
>> while x ~= 1  
>> x = x + 0.1;  
>> [x, sqrt(x)]  
>> end
```

Is everything as expected?

We are dealing with floating point arithmetic!

Floating-Point Arithmetic

Technical computing environments use floating-point arithmetic, which involves a *finite set of numbers* with *finite precision*. The typical implementation represents the number in the *normalized form*:

$$\bar{x} = (-1)^{sign} \times significand \times base^{exponent}$$


that can be equivalently written as:

$$\bar{x} = (-1)^s \times (1 + f) \times b^e$$

where

- s is the sign: 0 (+) or 1 (-);
- b is the *base* of the number system (*e.g.* for *hex* format $b = 16$);
- e is the *exponent*, $e_{min} \leq e \leq e_{max}$;
- $(1 + f)$ is the *mantissa* or *fractional part*.


IEEE Std 754: Binary encodings



The overall representation is encoded in k bits

- 1 bit for the sign S
- w -bits for the biased exponent $E = e + bias$
- t -bits trailing significand field digit string $T = d_1 d_2 d_3 \dots d_{p-1}$, where the leading bit (d_0) is implicitly encoded in the biased exponent E .


IEEE Std 754: Binary encodings



Exponent E : stored as unsigned number

- $L \leq e \leq U$, where $U = 2^{w-1} - 1$ and $L = 2 - 2^{w-1}$.
- $E = e + bias$, where the *bias* is equal to $2^{w-1} - 1$.
- $E = 0$ is reserved to encode ∓ 0 ($T = 0$) and subnormal numbers ($T \neq 0$).
- $E = 2^w - 1$ is reserved to encode $\mp \infty$ ($T = 0$) and *NaNs* ($T \neq 0$).
- w is a limitation on the achievable *range* of the arithmetic.


IEEE Std 754: Binary encodings



Exponent E : normal and subnormal numbers

- If $1 \leq E \leq 2^{w-1} - 2$, then $x = (-1)^S \times 2^{E-bias} \times (1.F)$, where $1.F$ is intended to represent the binary number created by prefixing F with an implicit leading 1 and binary point.
- If $E = 0$, then $x = (-1)^S \times 2^L \times (0.F)$. These numbers are called non-normalized or subnormal numbers.

IEEE Std 754: Binary encodings



Trailing significand field T (*Mantissa*)


T is a string of bits in the form:

$$\frac{d_1}{2^1} + \frac{d_2}{2^2} + \frac{d_3}{2^3} + \cdots + \frac{d_{p-1}}{2^t}.$$

Note that:

- note that d_0 remains implicit and that d_1 is always $\neq 0$.
- p represents a limitation on the *precision* of the arithmetic.

IEEE Std 754: Binary encodings



Parameter	half	single	double	quad	binary(k)
k, storage width in bits	16	32	64	128	$1 + w + t$
p, precision in bits	11	24	53	113	
exponent bias in bits	15	127	1023	16383	$2^{k-p-1} - 1$
Sign bit	1	1	1	1	1
w, exp. field width in bits	5	8	11	15	
t, trail. sign. field width in bits	10	23	52	112	$k - w - 1$

From IEEE Std 754TM - 2008 documentation.

IEEE Std 754: examples

for single precision ($k = 32\text{bits}$)

$$0\ 10000000\ 000000000000000000000000 = +1 \cdot 2^{128-127} \cdot 1.0_2 = 2$$

$$0\ 10000010\ 100010000000000000000000 = +1 \cdot 2^{130-127} \cdot 1.10001_2 = 12.25$$

$$0\ 00000000\ 000000000000000000000000 = 0$$

$$0\ 11111111\ 000000000000000000000000 = +\textit{Infinity}$$

$$1\ 11111111\ 000000000000000000000000 = -\textit{Infinity}$$

$$0\ 11111111\ 000000000000100000010000 = \textit{NaN}$$

$$0\ 00000001\ 000000000000000000000000 = +1 \cdot 2^{1-127} \cdot 1.0_2 = 2^{-126}$$

$$0\ 00000000\ 100000000000000000000000 = +1 \cdot 2^{-126} \cdot 0.1_2 = 2^{-127}$$

$$0\ 00000000\ 000000000000000000000001 = +1 \cdot 2^{-126} \cdot 0.00 \dots 01_2 = 2^{-149}$$

IEEE Std 754: hex format

In MATLAB every computation is performed in *double precision*.

- Use `format` and the appropriate options to set the output format (Note: this does not affect the way the computations are performed).

```
>> format long
```

```
>> x = 1
```

```
x =
```

```
1
```

```
>> format hex
```

```
>> x = 1
```

```
x =
```

```
3ff0000000000000
```

- Use `single` (`double`) commands to convert a number to single (double) precision.
- Use `num2hex` (`hex2num`) ... for conversion in the representations.

Refer to the related helps (`help format,...`) for further details!

IEEE Std 754: examples

The minimum and the maximum normalized machine's number can be obtained as follows:

```
>> minimum_real = realmin
```

```
>> maximum_real = realmax
```

- Note that `realmin` is different from `eps`.
- Any computation trying to produce a value larger than `realmax`, it is said to overflow. The result is called `Inf`.
- Any computation trying to produce a value smaller than `realmin`, it is said to underflow.
- For any computation trying to produce a value that is undefined the result is called `NaN` (e.g. `0/0`, `Inf/Inf`, `Inf-Inf`).

For double precision: $\text{realmin} = 2^{-1022}$, $\text{realmax} = (2 - \text{eps}) * 2^{1023}$, $\text{eps} = 2^{-52}$.

Outline

- 1 Introduction to MATLAB
- 2 Floating-Point: IEEE-754
- 3 Accuracy of the Arithmetic**
- 4 Loss-of-Significance errors
- 5 Homework

Machine epsilon: Exercise

Exercise 1 (From [2])

Evaluate the *epsilon machine* in the following three ways:

- 1 Implementing an *ad hoc* routine;
- 2 Executing the following code sequence (by Cleve Moler):

```
>> a = 4/3  
>> b = a - 1  
>> c = b + b + b  
>> e = 1 - c
```

- 3 Using `eps` command.

Comment and justify the obtained results.

Floating-Points: Comments

Note that floating-point numbers are rational. The base determines the fractions that can be represented. For instance $1/10$ or $1/5$ can be exactly represented only in the decimal format.

```
>> x = 0.1
```

The mathematical value stored in x is not exactly 0.1 , because the fraction $1/10$ in binary would require an infinite series!!

$$\frac{1}{10} = \frac{1}{2^4} + \frac{1}{2^5} + \frac{0}{2^6} + \frac{0}{2^7} + \frac{1}{2^8} + \frac{1}{2^9} + \frac{0}{2^{10}} + \frac{0}{2^{11}} + \dots$$

Let's verify:

```
>> b = num2bin(quantizer('double'),0.1)
b =
10111111111111001100110011001100110011001100110011001100110011010
```

Floating-Points: Exercise

Exercise 2

Consider a binary floating-point arithmetic with $k = 6\text{bits}$ ($w = 3\text{bits}$ and $t = 2\text{bits}$), with rounding.

- 1 How many numbers can be represented in this arithmetic? What about the error due to binary representation?
- 2 Represent the following numbers: 2, 1.6, -1, -14.01. Check the values with MATLAB.
- 3 Compare the results with the *double precision* format (64bits).
- 4 Which are the correspondent representations in hexadecimal format?
- 5 Convert the obtained results to *single precision*.

See also `help quantizer` and `help num2bin`. Comment the results.

Outline

- 1 Introduction to MATLAB
- 2 Floating-Point: IEEE-754
- 3 Accuracy of the Arithmetic
- 4 Loss-of-Significance errors
- 5 Homework

Numerical Cancellation

Subtracting two nearly equal numbers, we lose a great deal of accuracy!

Example

```
>> a = 9.901020304e-9;  
>> b = 9.0990055e-9;  
>> c = 1.8500185e8;  
>> s1 = (a+b) + c;  
>> s2 = a + (b+c);  
>> e = abs(s1-s2)  
e =  
    2.980232238769531e-008
```

Note: as a consequence of the loss of significant digits, the associative property is no more longer valid within floating-point arithmetic: *The order in which the operations are performed sometimes matters!*

Loss-of-Significance: Exercise

Exercise 3 (From [5])

Consider the following function:

$$f(x) = \frac{e^x - 1}{x}$$

- 1 Use De L'Hôpital's and Taylor's theorems to get an estimation of $f(x)$ around $x = 0$.
- 2 Experimentally evaluate $f(x)$ for values of x near zero (try with 2^{-k} , $k \in [1, 30]$): what do you obtain for single and double precision? Explain the results.
- 3 Propose an approach to fix the problem (Hint: Use Taylor's expansion).
- 4 How many terms in a Taylor's expansion are needed to get single precision accuracy (7 decimal digits) for all $x \in [0, 1/2]$? How many terms are needed for double precision accuracy (14 decimal digits) over the same range?

Outline

- 1 Introduction to MATLAB
- 2 Floating-Point: IEEE-754
- 3 Accuracy of the Arithmetic
- 4 Loss-of-Significance errors
- 5 Homework

Homework

Exercise H1.1

- 1 Use Taylor polynomial approximation to avoid the loss-of-significance errors in the following function formula when x approaches 0:

$$f(x) = \frac{1 - \cos(x)}{x^2}$$

- 2 Reformulate the following function $g(x)$ to avoid the loss-of-significance error in its evaluation for increasing values of x towards $+\infty$:

$$g(x) = x(\sqrt{x+1} - \sqrt{x})$$

Exercise H1.2 (From [5])

What is the machine's epsilon for a computer that uses binary arithmetic, 24bits for the fraction, and rounds? What if it chops?

Homework

Exercise H1.3 (From [5])

We can compute e^{-x} using Taylor polynomials in two ways, either using:

$$e^{-x} \approx 1 - x + \frac{1}{2}x^2 - \frac{1}{6}x^3 + \dots$$

or using

$$e^{-x} = \frac{1}{e^x} \approx \frac{1}{1 + x + \frac{1}{2}x^2 + \frac{1}{6}x^3 + \dots}$$

Discuss which approach is the more accurate. Which one is less susceptible to rounding error?

Homework

Exercise H1.4 (From [2])

Consider the following command sequence:

```
x = 1e-15
```

```
y = ((1+x)-1))/x
```

- 1 What can you say about y ? Give reason for the result through the direct analysis of errors propagation in the sequence.
- 2 Explain the obtained results by calculating the conditioning of the problem.