



Authorization and access control



Access control: deceptively simple

- ❑ A binary decision:
 - ❑ Access either allowed or denied
- ❑ However:
 - ❑ How do we *design* the access rules?
 - ❑ How do we *express* the access rules?
 - ❑ How do we *store* them
 - ❑ How do we appropriately *apply* them?
- ❑ Access control policies
- ❑ Access control models
 - ❑ Design and expression of policies; storage (in part)
 - ❑ A model should be *complete* and *consistent*
- ❑ Access control mechanisms and implementation
 - ❑ Storage (in part), application at runtime



Access control policies

- ❑ They can be roughly divided in
 - ❑ Discretionary Access Control (DAC)
 - ❑ Mandatory Access Control (MAC)
 - ❑ Role-Based Access Control (RBAC)
- ❑ The key difference between *discretionary* and *mandatory* access control is that in discretionary system the *owner* of a resource can assign privileges on it, including ownership, to others. In mandatory systems, only a security administrator can set privileges
- ❑ All of the common COTS operating systems are DAC based
 - ❑ MAC extensions are available for Linux, BSD, and other systems



Examples of well-known DAC systems

- ❑ Unix
 - ❑ Subjects: users, groups
 - ❑ Objects: files
 - ❑ Actions: read, write, execute
- ❑ Windows (in the NT family, inc. business versions of XP and Vista)
 - ❑ Subjects and objects as above but with “roles” instead of groups, multiple ownership of users and roles over files
 - ❑ Added actions: delete, change permissions, change ownership



Modeling access controls in DAC

- ❑ We need to model the following entities:
 - ❑ **Subjects** who can exercise privileges
 - ❑ **Objects** on which privileges are exercised and
 - ❑ **Actions** which can be exercised
- ❑ A **protection state** is commonly defined (Lampson 71, Graham and Denning 72, Harrison, Ruzzo and Ullmann 76) by a triple (S, O, A) where A is represented by a **matrix** with S on rows and O on columns. $A[s,o]$ represents the privileges of subject s over object o

	<i>Alpha</i>	<i>Beta</i>	<i>Gamma</i>
<i>Alice</i>	Read	Write	
<i>Bob</i>		Read/Write	
<i>Charlie</i>	Write		Read/Write



States and transitions in the HRU model

❑ Basic operations

- ❑ create subject $\langle s \rangle$; create object $\langle o \rangle$; enter $\langle \text{permission} \rangle$ into $[s, o]$; delete r from $[s, o]$; destroy subject s ; destroy object o

❑ Transitions are combinations of commands

- ❑ E.g. create file (subject u ; file f) means: create object f ; enter "own" into $[u, f]$; enter "read" into $[u, f]$; enter "write" into $[u, f]$

- ❑ Is this right? No, we need to check if f existed before, otherwise u would be stealing it away!

- ❑ We need an "if" construct into transitions

❑ Enter the safety problem



Safety problem

- ❑ Given a set of transitions and an initial configuration, can a certain subject attain a certain right on a certain object?
 - ❑ Obviously, yes: if the owner allows it!
 - ❑ But, if the owner does not allow it or does not exist?
 - ❑ If it happens the set of commands is *unsafe by design* !
- ❑ More formally, given a configuration Q , is there any sequence of transitions which leaks a certain right r (for which the owner is removed) into the access matrix somewhere?
 - ❑ If not, then the system is safe wrt right r
 - ❑ In a general HRU model **this is undecidable**
 - ❑ Decidable only in mono-operational systems, which are substantially useless (e.g. you cannot create a file and own it), or if subjects/objects are finite



Take-grant model

- ❑ Restricts the model to be able to prove safety
- ❑ We model privileges as a directed **graph** where S and O are nodes, and privileges are edges:
 - ❑ Read, write
 - ❑ Take: a node with take privilege wrt another can take (not in the sense of remove) any privilege from the latter
 - ❑ Grant: a node with grant privilege wrt another can give it any privilege it owns
- ❑ Operations are take, grant, create (a new object with certain privileges) and remove (certain privileges from an object)
- ❑ Safety here is decidable and $O(n)$ wrt a specific privilege, and $O(n^3)$ globally
- ❑ A pity this is substantially useless in real world...



Common DAC implementations

- ❑ DAC implementations are substantially reproduction of HRU models
- ❑ Access matrix is a sparse matrix = lots of wasted space
- ❑ Alternative implementations:
 - ❑ **Authorizations table**: records non-null triples S-O-A, typically used in DBMS
 - ❑ **Access Control Lists**: records by column (i.e. for each object, the list of subjects and authorizations)
 - ❑ **Capabilities Lists**: records by row (i.e. for each subject, the list of objects and authorizations)
- ❑ ACLs and capabilities obviously have different pros and cons



ACL vs capability lists

- ❑ ACL needs *subjects* to be securely authenticated
- ❑ ACL is more efficient if privilege assignments and revocations happen *per object*
- ❑ This is what usually happens, and this is why most systems (POSIX, Win) use ACLs
- ❑ Some systems (e.g. POSIX) use abbreviated ACLs
- ❑ Capability needs complex and verified methods for propagation and appropriate identification of objects
- ❑ Capabilities more efficient if privilege assignments and revocations happen *per subject*
- ❑ Capabilities are optional in POSIX (Linux and BSD)



Generic DAC shortcomings

- ❑ Cannot prove safety
- ❑ Control access to objects but not to information inside objects
 - ❑ E.g. susceptible to *trojan horse* problems even if users are trustworthy: if a program executes malicious actions, it will do so with the privileges of the user, and will be able to tamper with any file owned by the user
- ❑ Problems of scalability and management, since each user can potentially compromise security of the system with his/her own decisions



Mandatory access control policies

- ❑ Generic approach: define a *classification* of subjects (clearance) and objects (sensitivity)
- ❑ A classification is generally a partial order relationship, with a dominance concept which we will denote as \geq
- ❑ Often the classification is composed of
 - ❑ A strictly ordered **secretcy level** (e.g. the US classification: Unclassified < FOUO < Secret < Top Secret; the NATO levels: Unclassified < NATO Confidential < COSMIC Secret < COSMIC Top Secret)
 - ❑ A set of labels (e.g. "Policy", "Energy", "Finance"; or ATOMAL)
- ❑ Dominance is usually defined as
 - ❑ $\{C1, L1\} \geq \{C2, L2\} \Leftrightarrow C1 \geq C2 \text{ and } C2 \subseteq C1$



Just to get all formal ...

- ❑ This relationship is a lattice:
 - ❑ Reflexive
 - ❑ Transitive
 - ❑ Anti-symmetric
- ❑ For each couple of objects I can identify a least upper bound (an object which dominates both, and which in turn is dominated by any other object dominating both) and a greatest lower bound (an object which is dominated by both, and which in turn dominates any other object which is dominated by both)



Secrecy constraints; Bell-LaPadula model

- ❑ Subject s can read object o iff s dominates o (no read up property, aka simple security property)
- ❑ Subject s can write object o iff o dominates s (no write down property, aka *-property)
 - ❑ For this reason, a user may be allowed to connect with any security level which he currently dominates
- ❑ Tranquillity property: security classifications cannot be changed (clearance can)
- ❑ This would create a monotonic flow of information towards higher classification



Limitations of a BLP MAC model

- ❑ Sanitization: the output of a process can be of lower sensitivity than the inputs
- ❑ Declassification: at times, information is not secret anymore (think of WW2 battle plans today)
- ❑ The *association* of two values may have a higher sensitivity than each
- ❑ The *aggregation* of some data may be more sensitive than each data point singularly
- ❑ *Covert channels* may still be an issue
 - ❑ e.g. a TS process which cannot be directly observed by a S user can be identified by the response time of the CPU
 - ❑ System design must take care of noninterference between processes at different security levels



Integrity constraints: BIBA

- ❑ Dual of BLP for integrity
 - ❑ Instead of secrecy level, integrity level
 - ❑ User integrity: how much trust we have in user and correctness of his information
 - ❑ Object integrity: how much value the correctness of the information has, and how much it is trusted
- ❑ Dual rules
 - ❑ Subject s can write object o iff s dominates o (no write up property aka *-property)
 - ❑ Subject s can read object o iff o dominates s (no read down aka simple security property)
- ❑ Biba and BLP can coexist but with **different** sets of classifications !



Relational multilevel systems

- ❑ Multilevel security can be applied also to RDBMS
- ❑ Each tuple, or even each attribute of a tuple, can have an associated security classification
- ❑ Each user will see a view comprising only those tuples and attribute (s)he dominates
- ❑ Each user will write at his/her own security level
- ❑ This leads to the problem of poli-instantiation, because a user at lower level can instantiate a tuple with the same primary key of an existing tuple which is classified at a level he cannot access



Access control mechanisms

- ❑ Access control usually based on a *reference monitor*
 - ❑ Tamper proof
 - ❑ No bypass
 - ❑ Small enough to be verifiable
 - ❑ In other words, a *security kernel*
- ❑ A number of potential vulnerabilities
 - ❑ Storage channels: re-assigned resources such as RAM should be cleared before reassignment
 - ❑ Side channels: sometimes timing, or other issues, can reveal hidden information