



# **Intrusion Detection Systems, Honeynets**

---

A basic introduction to the  
concepts of detection and  
reaction



# A troublesome world

---

- ❑ Recap of unique challenges in security:
  - ❑ Vulnerabilities are discovered on a daily basis
  - ❑ A firewall cannot control everything
  - ❑ Access control and privilege management does not scale to large networks
  - ❑ Unauthenticated users access unsafe services
  - ❑ Trusted but untrustworthy machines are exposed to the network perils
  - ❑ TCP/IP packets are unauthenticated and unencrypted, allowing all sorts of manipulation
  - ❑ ...



# Crouching tiger, hidden dragon

---

- ❑ Crackers **do not** use brutal attacks against our defenses, they rely on skill, subtlety, and treachery
- ❑ The classic paradigm of security of authentication and authorization, “Who are you? What can you do?”, makes it necessary to build layers upon layers of complex defense systems
- ❑ Remember an historical lesson: King Dario at Gaugamela plains against Alexander Magnus
- ❑ For computer scientists it’s the KISS principle: Keep It Simple, Stupid
- ❑ We need a complementary paradigm



## A paradigm of resistance, detection and reaction

---

- ❑ Murphy's law is unescapable
  - ❑ Trusted components and people are untrustworthy
  - ❑ Secure programs are insecure
  - ❑ Reliable systems do fail
- ❑ The only difference between systems that can fail and systems that cannot is that, when a system which cannot fail, fails, it fails in a totally devastating way
- ❑ We must prepare for the fact that security incidents will – not can – happen
  - ❑ Build systems that fail gracefully
  - ❑ Build systems that rely on others as little as possible
  - ❑ Build systems that are tamper-evident
  - ❑ Design system for recovery



## Detecting intrusions: why are you doing this ?

---

- ❑ We want to be able to *detect* security incidents, in order to *react* as soon as possible. How can we do this?
- ❑ Back to basics: the goal of information security is to ensure that a system is used for a specific set of *goals* following a proper *procedure*, and that it cannot be used otherwise
- ❑ A violation of the security paradigm is evident, because the system does something else than what it's intended to do
- ❑ From the "Who are you ? What can you do ?" mindset to the "Why are you doing this?" mindset
- ❑ INTRUSION DETECTION SYSTEMS



# The IDS mantras

---

- ❑ An IDS is a system, not a software
  - ❑ A skilled human looking at logs is an IDS
  - ❑ A skilled network admin looking at TCPdump is an IDS
  - ❑ A company maintaining and monitoring your firewall is an IDS
  - ❑ A box bought by a vendor and plugged into the network is **not** an IDS by itself
- ❑ An IDS is not a panacea, it's a component
  - ❑ Does not substitute a firewall (despite what Gartner thinks)
  - ❑ It's the last component to add to a security architecture, not the first
- ❑ Detection without reaction is a no-no
  - ❑ Like burglar alarms with no guards!
- ❑ Reaction without human supervision is a dream
  - ❑ "Network, defend thyself !"



# Terminology and taxonomies

---

- ❑ Different types of software involved in IDS
  - ❑ Logging and auditing systems
  - ❑ Correlation systems
  - ❑ So-called “IDS” software
  - ❑ Honeypots / honeytokens
- ❑ The logic behind an IDS is always the same: those who access a system for illegal purposes act differently than normal users
- ❑ Two main detection methods:
  - ❑ Anomaly Detection: we try to describe what is normal, and flag as anomalous anything else
  - ❑ Misuse Detection: we try to describe the attacks, and flag them directly

# Anomaly vs. misuse

---



## Anomaly Detection Model

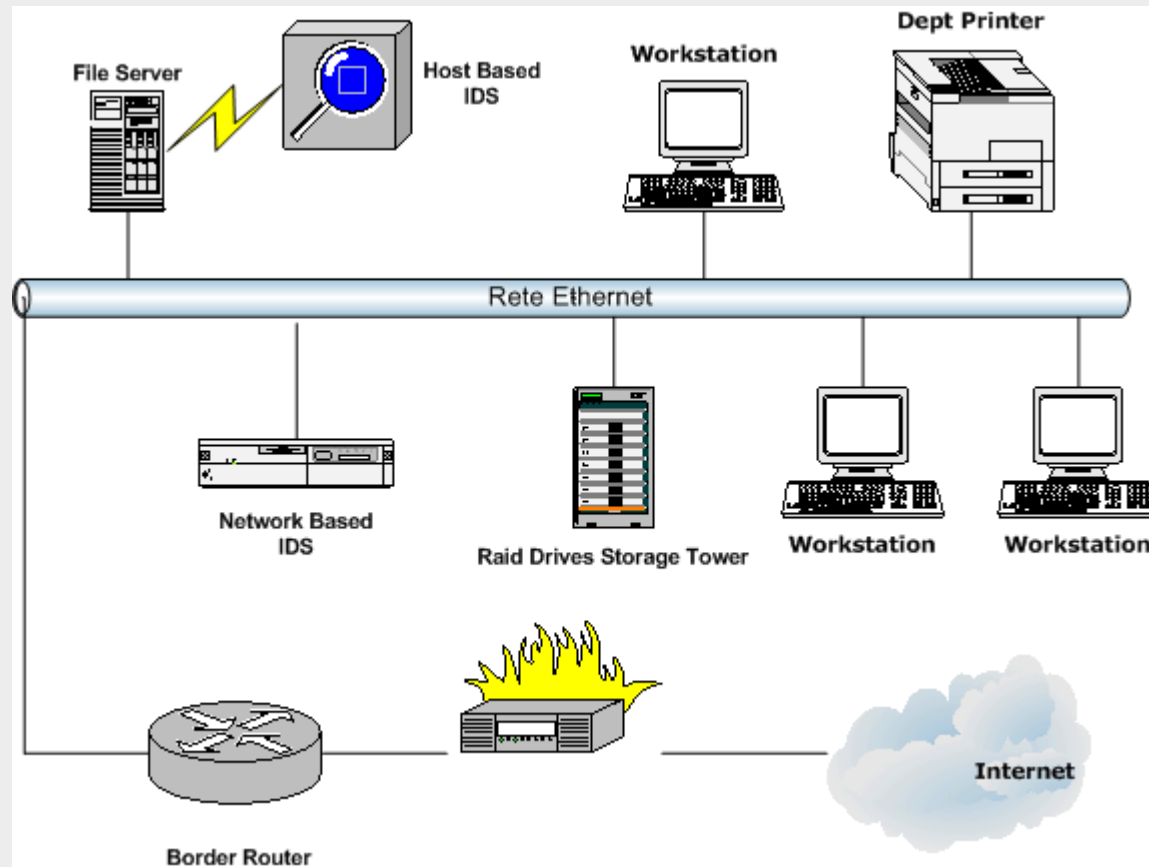
- ❑ Describes normal behaviour, and flags deviations
- ❑ Uses statistical or machine learning models of behaviour
- ❑ Theoretically able to recognize any attack, also 0-days
- ❑ Strongly dependent on the model, the metrics and the thresholds
- ❑ Generates statistical alerts: "Something's wrong"

## Misuse Detection Model

- ❑ Uses a knowledge base to recognize the attacks
- ❑ Can recognize only attacks for which a "signature" exists in the KB
- ❑ When new types of attacks are created, the language used to express the rules may not be expressive enough
- ❑ Problems for polymorphism
- ❑ The alerts are precise: they recognize a specific attack, giving out many useful informations



# Host based vs. Network based





# Misuse detection alone is an awful idea

---

- ❑ Misuse detection systems rely on a knowledge base (think of the anti-virus example, if it's easier to grasp)
- ❑ Updates continuously needed, and not all the attacks become known (as opposed to viruses)
- ❑ Signature engineering problems (a NAV update, a couple of years ago, recognized *itself* as a virus...)
- ❑ Attacks are polymorphs, more than computer viruses: ADMutate, UTF encoding...
- ❑ Attacks are not atomic, and interleaving helps in avoiding detection!



# Anomaly Detection, perhaps not better

---

- ❑ We must describe the behaviour of a system
  - ❑ Which features/variables/metrics do we use?
  - ❑ Which model do we choose to fit them?
- ❑ Thresholds must be chosen to minimize false positive vs. detection rate: a difficult process
- ❑ The base model is fundamental: if the attack shows up only in variables we discarded, the system is blind on it!
- ❑ Any type of alert is more or less equivalent to "hey, something's wrong here". What? Your guess!

# SNORT: a Lightweight NIDS

---





# What's Snort ?

---

- ❑ Snort is a “multi-mode packet analysis tool”, fundamentally a network based IDS sensor, which operates fundamentally on misuse detection (signatures)
- ❑ Developed by Martin Roesch
- ❑ Snort is available at [www.snort.org](http://www.snort.org)
- ❑ Snort-based commercial appliances and support: [www.sourcefire.com](http://www.sourcefire.com)



# Characteristics of snort

---

- ❑ It's a lightweight, high-performance packet sniffer
- ❑ Based on Libpcap libraries
- ❑ Can work in 4 basic modes
  - ❑ Sniffer Mode (similar to tcpdump)
  - ❑ Packet Logger Mode (as above, but more output options)
  - ❑ NIDS Mode (activates attack rulesets)
  - ❑ Forensic Data Analysis Mode (as above, but takes a network dump as input)
- ❑ A plug-in system adds flexibility



## Snort plug-ins

---

- ❑ Preprocessor: examines and mangles packets before detection stage
- ❑ Detection plugin: executes tests on a field or aspect of packet
- ❑ Output plugin: transform analysis data for representation, visualization and logging



## A complete NIDS sensor

---

- ❑ Snort rules are widely available, a true standard: they are even inserted in advisories sometimes. Even ISS added snort rules format support to RealSecure
- ❑ Detection Plug-ins extend rules expressivity, allowing to write rules not only matching packet content and header, but also adding statistical informations and protocol anomaly detection
- ❑ Snort preprocessors: portscan detection, IP defragmentation, TCP stream reassembly, application layer analysis, etc.





# Snort outputs

---

- ☐ Database
  - ☐ MySQL
  - ☐ PostgreSQL
  - ☐ Oracle
- ☐ XML (SNML DTD by CMU/CERT)
- ☐ tcpdump/libpcap format
- ☐ Unified format (Snort's own format)
- ☐ ASCII
- ☐ Syslog
- ☐ WinPopup (SMB)



# Snort 1.x Detection Engine

---

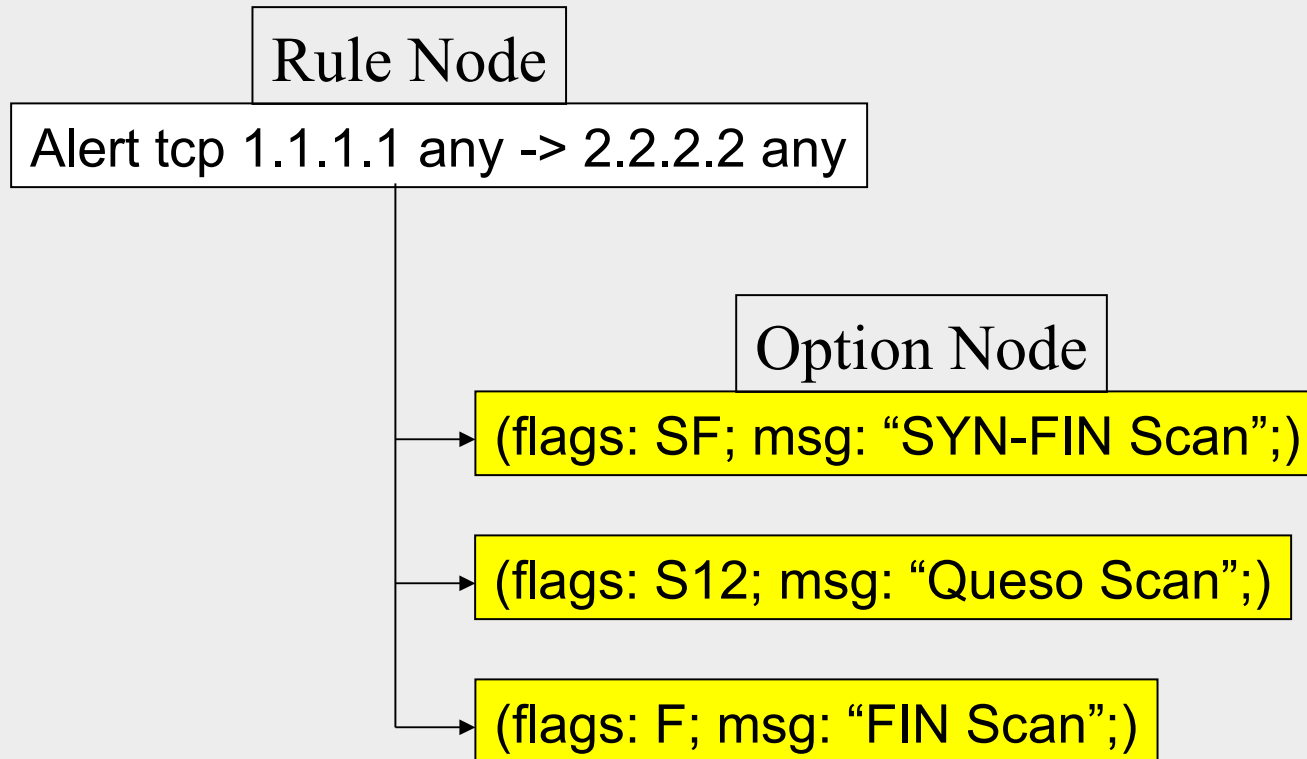
- ❑ Tridimensional linked list of rules
  - ❑ The first and the second dimension contain the parameters to test on the packets
  - ❑ Third dimension contains pointers to functions that must be used on packets
  - ❑ The engine is recursively applied on packets
  - ❑ First match detection: as soon as one rule matches, packet is discarded and next packet is processed
- ❑ Snort reaches with few problems 100Mb/sec of throughput on really common hardware



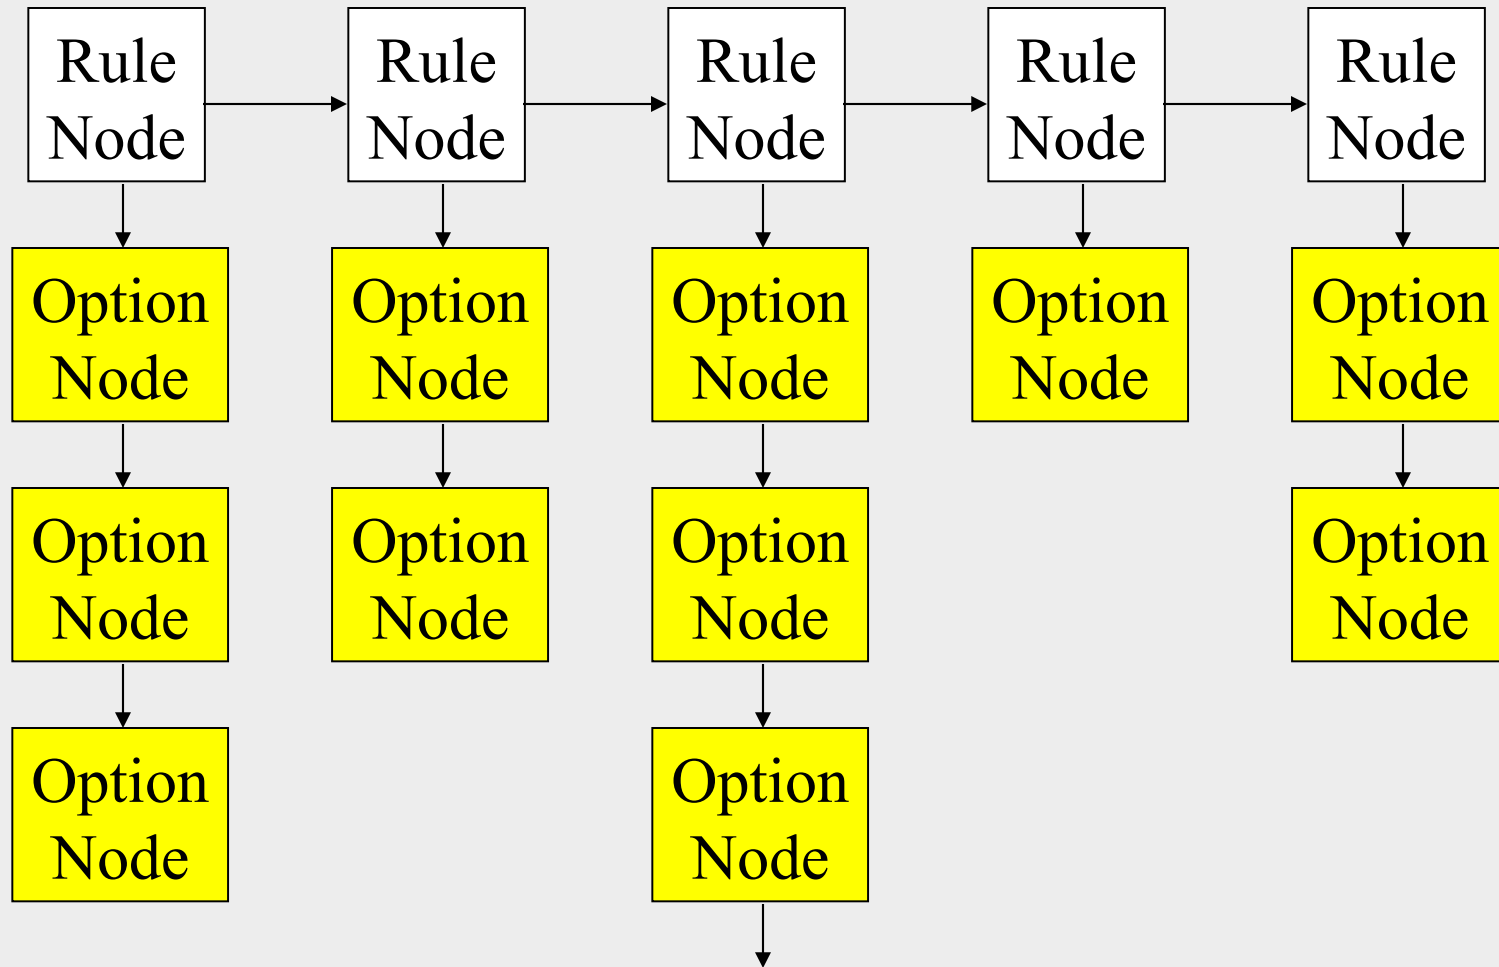
# Three sample rules

Rule Header	Rule Options
Alert tcp 1.1.1.1 any -> 2.2.2.2 any	(flags: SF; msg: "SYN-FIN Scan");
Alert tcp 1.1.1.1 any -> 2.2.2.2 any	(flags: S12; msg: "Queso Scan");
Alert tcp 1.1.1.1 any -> 2.2.2.2 any	(flags: F; msg: "FIN Scan");

# Their representation...



# You got the idea ?





# Snort 1.x architectural limits

---

- ❑ Focused on the single IP packet
- ❑ IP defragmentation and TCP stream reassembly as preprocessors... often becoming bottleneck
- ❑ Output plugins, too, were wasting resources
- ❑ New protocol support was not modular and not easy to add
- ❑ Application layer not managed by the rule engine, relayed on the shoulders of the rule writer. So, simple to write rules for IP/TCP/UDP/ICMP/IGMP... but complex to describe HTTP, RPC, SMTP...



## Snort 2.0: a new architecture

---

- ❑ Lots of new plug-in types
  - ❑ Acquires various data formats
  - ❑ An extensible traffic protocol decoder (protocol verification, stream analysis on multi-path routing...)
  - ❑ Multiple formats for rules (from a DBMS, XML format...)
  - ❑ Detection engine is a plug-in: Standard NIDS, Target-based IDS, Statistical IDS, Host-based IDS...
- ❑ Improved pattern matching:
  - ❑ Aho-Corasick; Boyer-Moore; set-wise Boyer-Moore-Horspool algorithms
  - ❑ in human terms, a speedup of five times!
- ❑ Output plugins are now attached to a secondary process ("barnyard") which acts as a buffer, avoiding Snort the output bottleneck



## Regarding snort configuration

---

- ❑ Configuration files with explanations and comments, for 1.9 and 2.0:  
<http://www.0xdeadbeef.info/>
  - ❑ Lots of other goodies, there, too :)
- ❑ The user manual explains everything about rule writing:  
[www.snort.org/docs/SnortUsersManual.pdf](http://www.snort.org/docs/SnortUsersManual.pdf)
- ❑ “Intrusion Signatures and Analysis” (Northcutt, Cooper, Fearnow, Frederick): a great book, generic, but using snort as a specific example!



- ❑ An host based IDS sensor uses system auditing auditing resources to gather information
  - ❑ System logs
  - ❑ System events (e.g. login attempts)
  - ❑ System calls (file access, devices...)
- ❑ File integrity checking is a form of HIDS: see Tripwire, [www.tripwire.com](http://www.tripwire.com)
- ❑ Anomaly detection is often prominent here!
- ❑ A compromised system is NOT a great source of information, as you may guess...



# What is an “honeypot” ?

---

- ❑ *An honeypot is a security resource whose value lies in being probed, attacked, or compromised*
- ❑ Simple concept: a resource that expects no data, so any traffic to or from it is most likely unauthorized activity
- ❑ A form of “anomaly detection” !
- ❑ Honeypots *do not solve a specific problem*, but can be a tool that contributes to your overall security architecture
- ❑ Their value, and the problems they help solve, depends on how you build, deploy, and use them



# How do they fit into the IDS scenario ?

---

## ☐ Prevention

- ☐ Honeypots, in general are not effective prevention mechanisms.
- ☐ “Deception, Deterrence, Decoy”: psychological weapons. They do NOT work against worms or auto-rooters

## ☐ Detection

- ☐ They are a great anomaly IDS
- ☐ Unique advantages

## ☐ Response

- ☐ Honeypots can be used to help respond to an incident
- ☐ They can be pulled offline and analyzed (unlike production systems)
- ☐ Little to no data pollution (all the data on a honeypot is attack data)



## Also Research Honeypots

---

- ❑ Used by universities and research groups
- ❑ Early Warning and Prediction
- ❑ Discovery of new Tools and Tactics
- ❑ They help to understand motives, behavior, and organization of the criminals
- ❑ Used for training and development of analysis and forensic skills
- ❑ E.g. the HoneyNet project (born in 1999)
- ❑ E.g. WOMBAT project ([www.wombat-project.eu](http://www.wombat-project.eu))

## Example: Motives and Behavior

---



- ❑ Captured IRC session between hacker and pal on a private channel

**J4ck:** why don't you start charging for packet

attacks?

**J4ck:** "give me x amount and I'll take bla bla offline

for this amount of time"

**J1LL:** it was illegal last I checked.

**J4ck:** heh, then everything you do is illegal. Why not make money off of it?

**J4ck:** I know plenty of people that'd pay exorbitant amounts for packeting.



# Level of Interaction

---

- ❑ Level of Interaction determines amount of functionality a honeypot provides: “how real and complex it looks”
- ❑ The greater the interaction, the more you can learn
- ❑ The greater the interaction, the more complexity and risk
- ❑ Chances are that an attacker can use your honeypot to harm, attack, or infiltrate other systems or organizations
- ❑ Which is best? It depends on your objectives



# Low Interaction vs. High Interaction

---

- ❑ Lowly interactive honeypots:
  - ❑ Provide Emulated Services
  - ❑ No operating system for attacker to access
  - ❑ Information limited to transactional information and attackers activities with emulated services
  - ❑ Many attackers will realize that services are fake after a while
- ❑ Highly interactive honeypots:
  - ❑ Provide Actual Operating Systems
  - ❑ Learn extensive amounts of information
  - ❑ High risk: a subverted honeypot can be turned against you

# Honeyd

---



```
create default
```

```
set default personality "FreeBSD 2.2.1-STABLE"
```

```
set default default action open
```

```
add default tcp port 80 "sh /usr/local/honeyd/scripts/web.sh"
```

```
add default tcp port 22 "sh /usr/local/honeyd/scripts/test.sh"
```

```
add default tcp port 113 reset
```

```
add default tcp port 1 reset
```

```
create windows
```

```
set windows personality "Windows NT 4.0 Server SP5-SP6"
```

```
set windows default action reset
```

```
add windows tcp port 80 "sh /usr/local/honeyd/scripts/web.sh"
```



# Honeytokens

---



- ❑ Honeytoken: a honeypot that is not a computer
  - ❑ A honeytoken can be a credit card number, Excel spreadsheet, PowerPoint presentation, a database entry, or even a bogus login.
  - ❑ Just as a honeypot computer has no authorized value, no honeytoken has any authorized use.
- ❑ Map-making companies inserting bogus cities or roads into maps to spot spies
- ❑ Example:
  - ❑ Add a bogus credit card number into an e-commerce DB
  - ❑ E.g., the credit card number 4356974837584710
  - ❑ Add a Snort rule: alert ip any any -> any any (msg:"Honeytoken Access"; content:"4356974837584710";)
- ❑ Minimal cost and complication!

# Honeynets

---



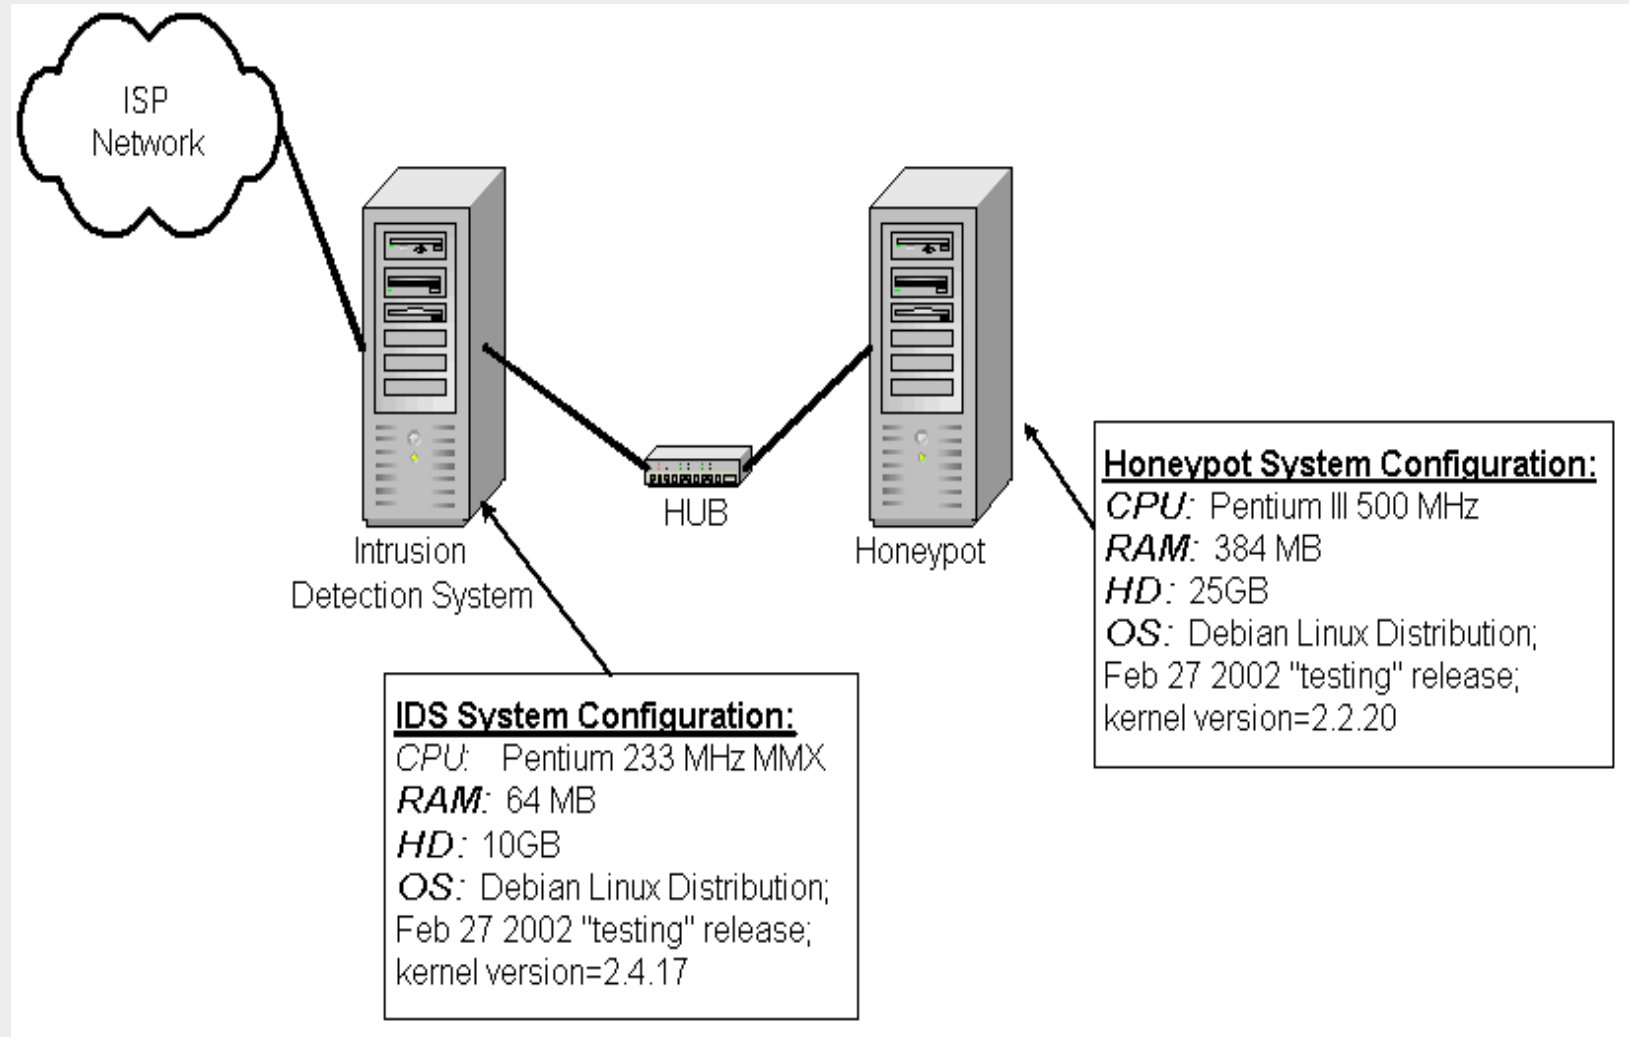
- ❑ <http://project.honeynet.org>
- ❑ An extension of a Honeypot
- ❑ Network topology provides many advantages over standard honeypot
  - ❑ Covert logging
  - ❑ More points of attack for a blackhatter
  - ❑ Looks realistic from the outside
- ❑ In a typical honeynet, there may be one or more honeypot machines spanning several operating systems

# Basics of an honeynet

---

- ❑ Containment: once a honeypot within the Honeynet is compromised, we have to contain the activity and ensure the honeypots are not used to harm non-Honeynet systems. We must control how traffic can flow in and out of the Honeynet
- ❑ Data Capture: capture all activity within the Honeynet and the information that enters and leaves the Honeynet
- ❑ Data Collection: Once data is captured, it is securely forwarded to a centralized data collection point
- ❑ ... All of these avoiding that the “bad guys” detect us!

# A sample Honeynet architecture





# Limitations of HIH and LIH

---

- ❑ LIH are limited to known services and specific types of attacks
- ❑ HIH are labour-intensive and prone to disasters, e.g. being taken over by an attacker and abused
- ❑ Many projects (e.g. sgnet) aim to blend the two types of systems, e.g. by automated learning