



Association Rules: Advanced Topics

Data Mining and Text Mining (UIC 583 @ Politecnico di Milano)

- ❑ Frequent patterns without candidate generation
- ❑ Multilevel association rules
- ❑ Correlation rules
- ❑ Sequential rules

Is Apriori Fast Enough?

Performance Bottlenecks

- ❑ The core of the Apriori algorithm
 - ▶ Use frequent $(k-1)$ -itemsets to generate candidate frequent k -itemsets
 - ▶ Use database scan and pattern matching to collect counts for the candidate itemsets
- ❑ The bottleneck of Apriori: candidate generation
- ❑ Huge candidate sets:
 - ▶ 10^4 frequent 1-itemset will generate 10^7 candidate 2-itemsets
 - ▶ To discover a frequent pattern of size 100, e.g., $\{a_1, a_2, \dots, a_{100}\}$, needs to generate $2^{100} \approx 10^{30}$ candidates.
- ❑ Needs multiple $(n+1)$ scans of database with n is the length of the longest pattern

- ❑ How to generate candidates?
 - ▶ Step 1: self-joining L_k
 - ▶ Step 2: pruning

- ❑ How to count supports of candidates?

- ❑ Example of Candidate-generation
- ❑ Example:
 - ▶ $L_3 = \{abc, abd, acd, ace, bcd\}$
 - ▶ Self-joining: $L_3 * L_3$ generates $abcd$ from abc and abd
 - ▶ Generates also $acde$ from acd and ace
 - ▶ Pruning: $acde$ is removed because ade is not in L_3
 - ▶ $C_4 = \{abcd\}$

How to Generate Candidates?

- ❑ Suppose the items in L_{k-1} are listed in an order
- ❑ Step 1: self-join L_{k-1}
 - ▶ INSERT INTO C_k
 - ▶ SELECT $p.item_1, p.item_2, \dots, p.item_{k-1}, q.item_{k-1}$
 - ▶ FROM $L_{k-1} \ p, L_{k-1} \ q$
 - ▶ WHERE $p.item_1 = q.item_1, \dots, p.item_{k-2} = q.item_{k-2},$
 $p.item_{k-1} < q.item_{k-1}$
- ❑ Step 2: pruning
 - ▶ For each itemset c in C_k do
 - For each $(k-1)$ -subsets s of c do if (s is not in L_{k-1}) then delete c from C_k

Mining Frequent Patterns Without Candidate Generation

- ❑ Compress a large database into a compact, Frequent-Pattern tree (FP-tree) structure
 - ▶ Highly condensed, but complete for frequent pattern mining
 - ▶ Avoid costly database scans
- ❑ Develop an efficient, FP-tree-based frequent pattern mining method
- ❑ A divide-and-conquer methodology: decompose mining tasks into smaller ones
- ❑ Avoid candidate generation: sub-database test only

FP-growth...

- ❑ Leave the generate-and-test paradigm of Apriori
- ❑ Data sets are encoded using a compact structure, the **FP-tree**
- ❑ Frequent itemsets are extracted directly from the FP-tree

- ❑ Major Steps to mine FP-tree
 - ▶ Construct conditional pattern base for each node in the FP-tree
 - ▶ Construct conditional FP-tree from each conditional pattern-base
 - ▶ Recursively mine conditional FP-trees and grow frequent patterns obtained so far
 - ▶ If the conditional FP-tree contains a single path, simply enumerate all the patterns

Building the FP-tree: First Scan to Compute Frequent Items

10

TID	Items
1	{A,B}
2	{B,C,D}
3	{A,C,D,E}
4	{A,D,E}
5	{A,B,C}
6	{A,B,C,D}
7	{A}
8	{A,B,C}
9	{A,B,D}
10	{B,C,E}



Item	Count
A	8
B	7
C	6
D	5
E	3

minsup=2


Building the FP-tree: Second Scan to Build the Tree

11


TID	Items
1	{A,B}
2	{B,C,D}
3	{A,C,D,E}
4	{A,D,E}
5	{A,B,C}
6	{A,B,C,D}
7	{A}
8	{A,B,C}
9	{A,B,D}
10	{B,C,E}

- Items are sorted in decreasing support counts

After reading TID=1:




After reading TID=2:



TID	Items
1	{A,B}
2	{B,C,D}
3	{A,C,D,E}
4	{A,D,E}
5	{A,B,C}
6	{A,B,C,D}
7	{A}
8	{A,B,C}
9	{A,B,D}
10	{B,C,E}


After reading TID=3:



TID	Items
1	{A,B}
2	{B,C,D}
3	{A,C,D,E}
4	{A,D,E}
5	{A,B,C}
6	{A,B,C,D}
7	{A}
8	{A,B,C}
9	{A,B,D}
10	{B,C,E}

Item	Pointer
A	
B	
C	
D	
E	

- Pointers are used to assist frequent itemset generation
- The FP-tree is typically smaller than the data



Decompose the frequent item generation
into multiple subproblems


Start backward


Find frequent itemsets ending in E

Find frequent itemsets ending in D


...

- Start from the bottom, from E
- Compute the support count, by adding the counts associated to E
- Since, E is frequent we compute the frequent itemsets from the FP-tree







(a) Paths containing node e



(b) Paths containing node d



(c) Paths containing node c



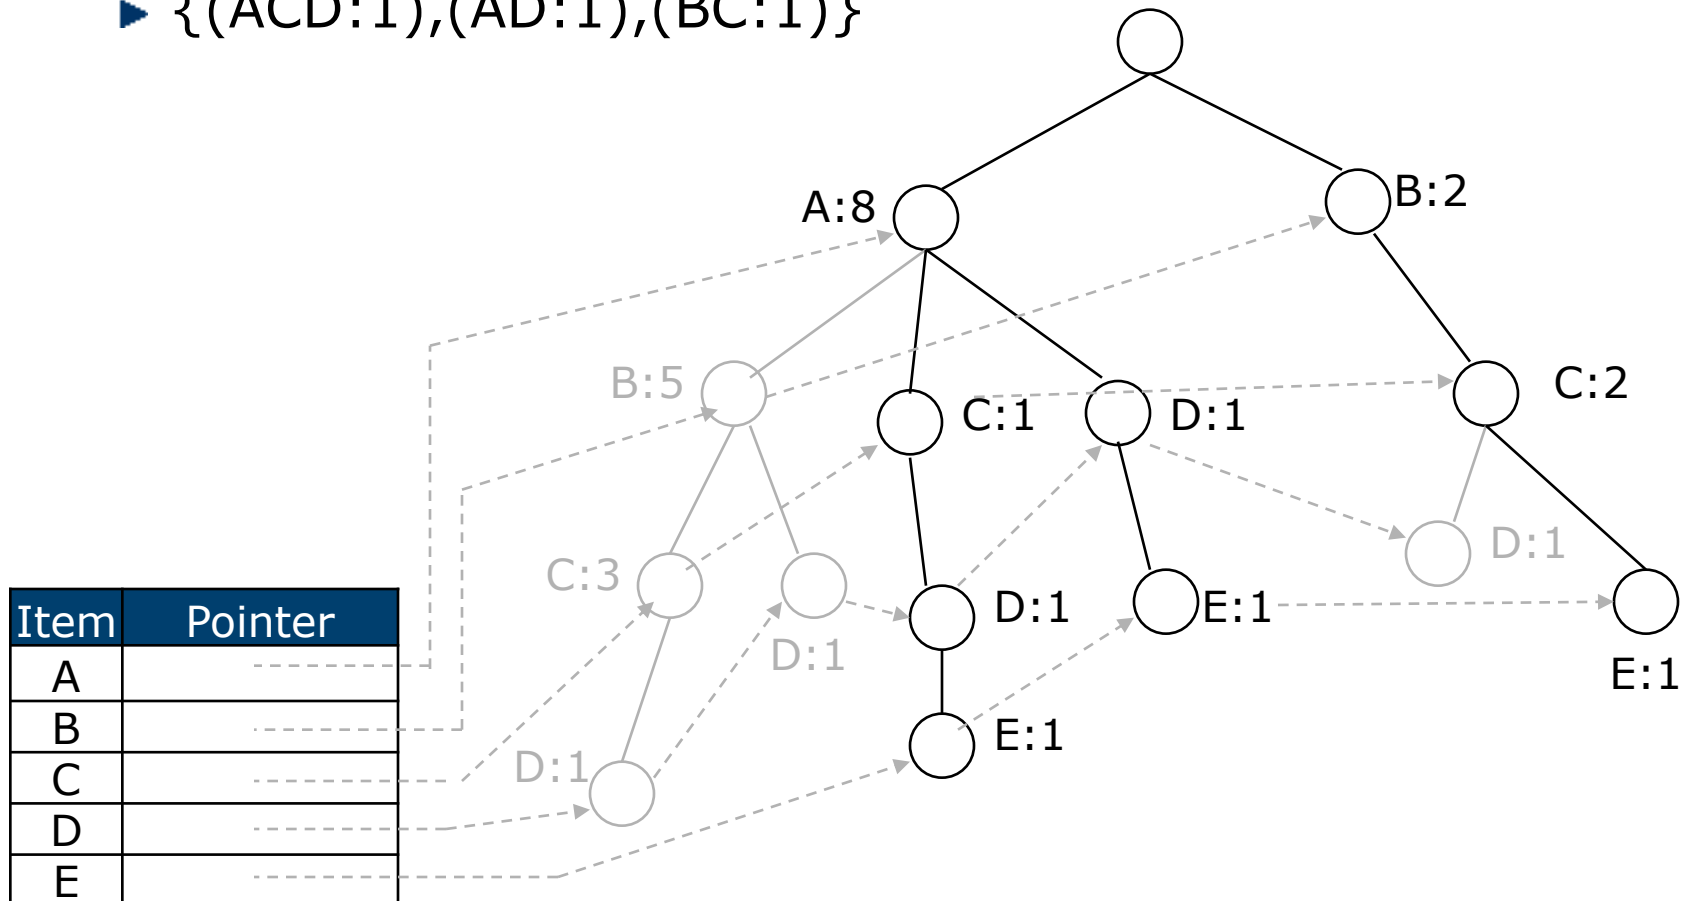
(d) Paths containing node b



(e) Paths containing node a



- ❑ Considers the path to E, it is frequent
- ❑ Extracts the conditional pattern base
 - ▶ $\{(ACD:1), (AD:1), (BC:1)\}$

minsup=2




- From the conditional pattern base
 $\{(ACD:1), (AD:1), (BC:1)\}$
build the conditional FP-tree
- Delete unfrequent items

minsup=2




- Use the conditional FP-tree to extract the frequent itemsets ending with **DE**, **CE**, and **AE**

minsup=2




Conditional FP-tree for E




Prefix paths ending in DE

- ❑ Consider the suffix **DE**, it is frequent
- ❑ Build the conditional FP-tree for **DE**
- ❑ The last tree contains only **A** which is frequent
- ❑ So far we have obtained three frequent itemsets, {E, DE, ADE}

minsup=2




Prefix paths ending in de




Conditional FP-tree for de

- After **DE**, the suffix **CE** is considered
- CE** is frequent and thus added to the frequent itemsets
- Then, we search for the itemsets ending with **AE**
- At the end the frequent itemsets ending with E are {E, DE, ADE, CE, AE}

minsup=2



Prefix paths ending in CE



Prefix paths ending in AE

- ❑ General idea (divide-and-conquer)
 - ▶ Recursively grow frequent pattern path using the FP-tree
- ❑ Method
 - ▶ For each item, construct its conditional pattern-base, and then its conditional FP-tree
 - ▶ Repeat the process on each newly created conditional FP-tree
 - ▶ Until the resulting FP-tree is empty, or it contains only one path (single path will generate all the combinations of its sub-paths, each of which is a frequent pattern)


- ❑ FP-growth is an order of magnitude faster than Apriori, and is also faster than tree-projection

- ❑ Why?
 - ▶ No candidate generation, no candidate test
 - ▶ Use compact data structure
 - ▶ Eliminate repeated database scan
 - ▶ Basic operation is counting and FP-tree building

- ❑ Start at the last item in the table
- ❑ Find all paths containing item
 - ▶ Follow the node-links
- ❑ Identify conditional patterns
 - ▶ Patterns in paths with required frequency
- ❑ Build conditional FP-tree C
- ❑ Append item to all paths in C, generating frequent patterns
- ❑ Mine C recursively (appending item)
- ❑ Remove item from table and tree

FP-growth vs. Apriori: Scalability With the Support Threshold

25



❑ Completeness

- ▶ Preserve complete information for frequent pattern mining
- ▶ Never break a long pattern of any transaction

❑ Compactness

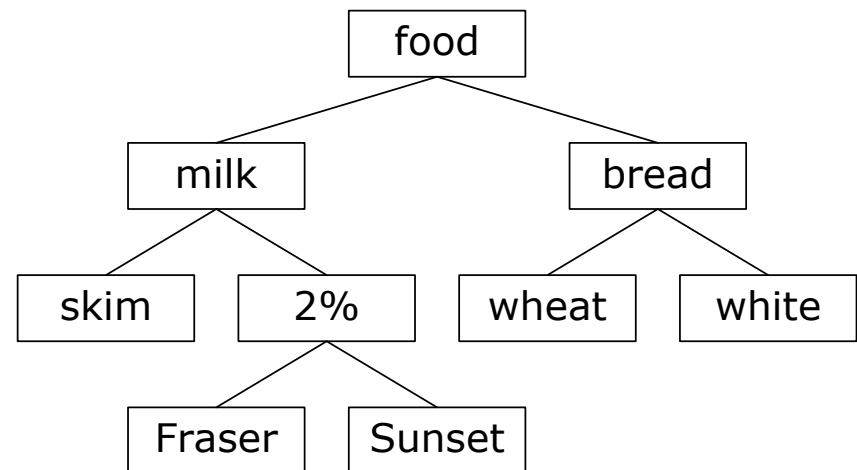
- ▶ Reduce irrelevant info—infrequent items are gone
- ▶ Items in frequency descending order: the more frequently occurring, the more likely to be shared
- ▶ Never be larger than the original database (not count node-links and the count field)
- ▶ For Connect-4 DB, compression ratio could be over 100

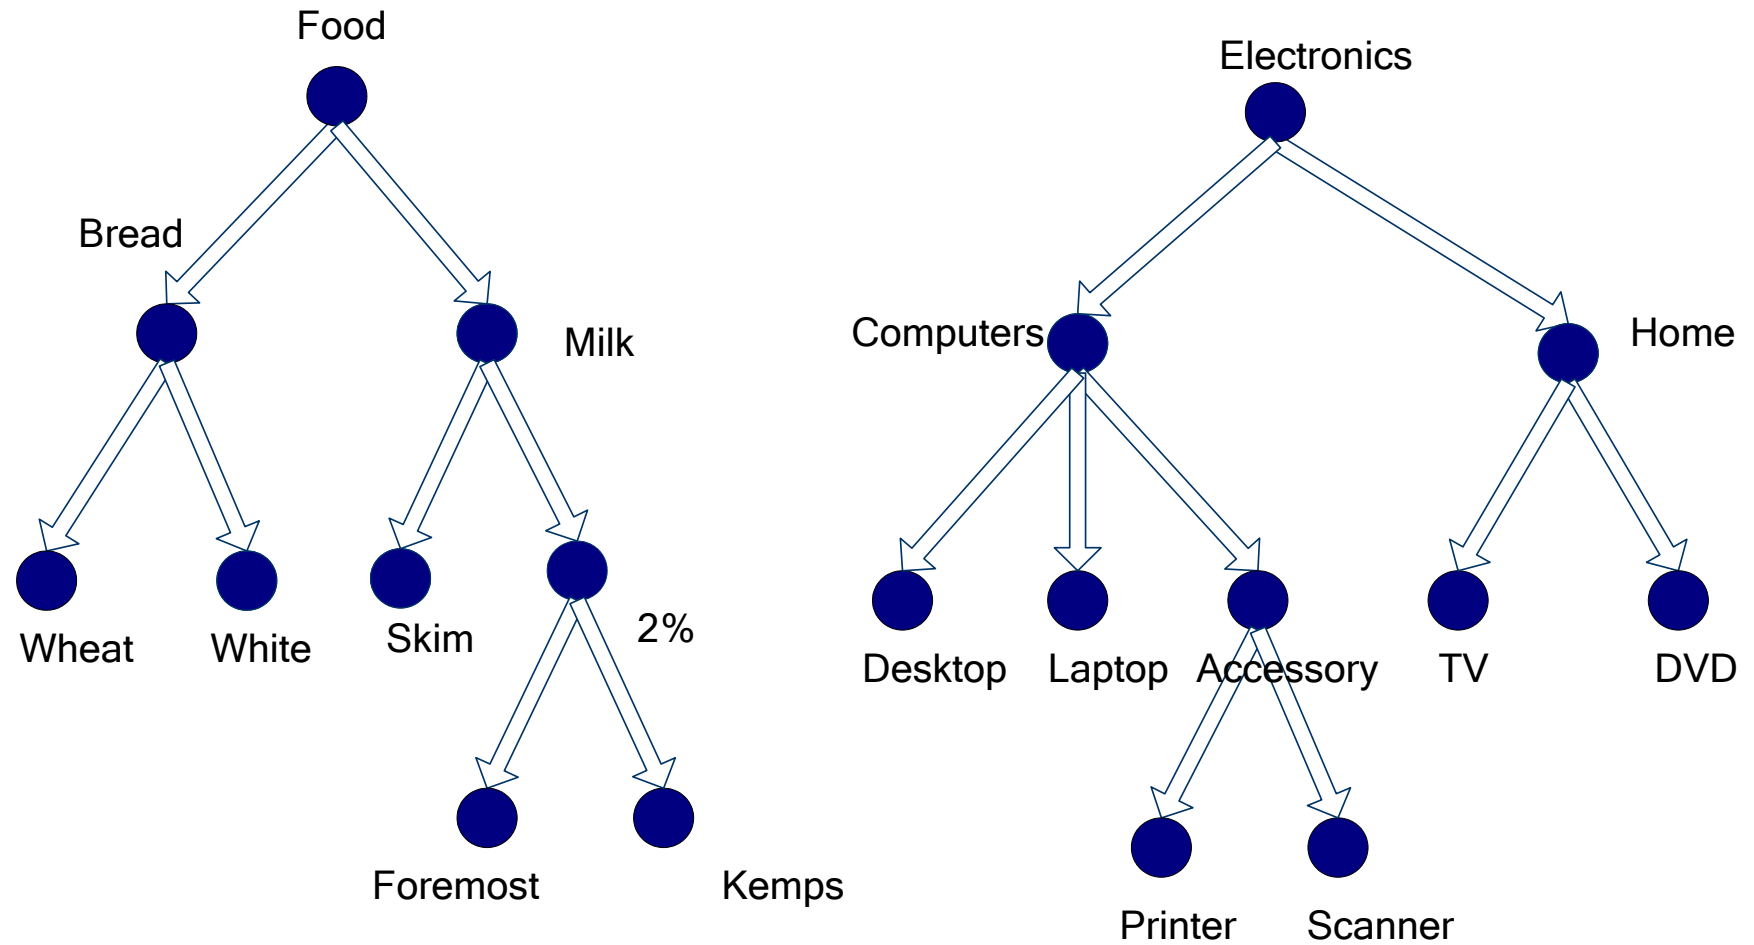
- ❑ Divide-and-conquer:
 - ▶ Decompose both the mining task and DB according to the frequent patterns obtained so far
 - ▶ Leads to focused search of smaller databases

- ❑ Other factors
 - ▶ No candidate generation, no candidate test
 - ▶ Compressed database: FP-tree structure
 - ▶ No repeated scan of entire database
 - ▶ Basic ops—counting local freq items and building sub FP-tree, no pattern search and matching

Multi-level Association Rules

- ❑ Items often form hierarchy
- ❑ Items at the lower level are expected to have lower support
- ❑ Rules regarding itemsets at appropriate levels could be quite useful
- ❑ Transaction database can be encoded based on dimensions and levels
- ❑ We can explore shared multi-level mining





- ❑ Rules at lower levels may not have enough support to appear in any frequent itemsets
- ❑ Rules at lower levels of the hierarchy are overly specific,

$\{2\% \text{ milk}\} \Rightarrow \{\text{wheat bread}\}$

is indicative of association between milk and bread

- How do support and confidence vary as we traverse the concept hierarchy?
 - ▶ If X is the parent item for both $X1$ and $X2$, then
 $\sigma(X) \leq \sigma(X1) + \sigma(X2)$
 - ▶ If $\sigma(X1 \cup Y1) \geq \text{minsup}$,
and X is parent of $X1$, Y is parent of $Y1$
then $\sigma(X \cup Y1) \geq \text{minsup}$, $\sigma(X1 \cup Y) \geq \text{minsup}$
 $\sigma(X \cup Y) \geq \text{minsup}$
 - ▶ If $\text{conf}(X1 \Rightarrow Y1) \geq \text{minconf}$,
then $\text{conf}(X1 \Rightarrow Y) \geq \text{minconf}$

❑ Approach 1:

- ▶ Extend current association rule formulation by augmenting each transaction with higher level items
- ▶ Original Transaction: {skim milk, wheat bread}
- ▶ Augmented Transaction:
{skim milk, wheat bread, milk, bread, food}

❑ Issues:

- ▶ Items that reside at higher levels have much higher support counts
- ▶ If support threshold is low, too many frequent patterns involving items from the higher levels
- ▶ Increased dimensionality of the data

❑ Approach 2:

- ▶ Generate frequent patterns at highest level first
- ▶ Then, generate frequent patterns at the next highest level, and so on

❑ Issues:

- ▶ I/O requirements will increase dramatically because we need to perform more passes over the data
- ▶ May miss some potentially interesting cross-level association patterns

- ❑ A top down, progressive deepening approach
- ❑ First find high-level strong rules:

$$\{\text{milk}\} \Rightarrow \{\text{bread}\} \quad [20\%, 60\%]$$


- ❑ Then find their lower-level “weaker” rules:

$$\{2\% \text{ milk}\} \Rightarrow \{\text{wheat bread}\} \quad [6\%, 50\%]$$

- ❑ Some rules may be redundant due to “ancestor” relationships between items
- ❑ A rule is redundant if its support is close to the “expected” value, based on the rule’s ancestor.
- ❑ Example
 - ▶ milk \Rightarrow wheat bread
[support = 8%, confidence = 71.5%]
 - ▶ 2% milk \Rightarrow wheat bread
[support = 7.5%, confidence = 72%]
- ❑ The first rule is an ancestor of the second rule.
- ❑ One of the two rules is redundant: both rules provide the same information, in fact, they have similar support and confidence


Level 1
min_sup = 5%

Level 2
min_sup = 5%



Level 1
min_sup = 5%

Level 2
min_sup = 3%



Multi-level Association: Uniform Support vs. Reduced Support

39

- ❑ Uniform Support: the same minimum support for all levels
 - ▶ + One minimum support threshold. No need to examine itemsets containing any item whose ancestors do not have minimum support.
 - ▶ – Lower level items do not occur as frequently. If support threshold
 - too high \Rightarrow miss low level associations
 - too low \Rightarrow generate too many high level associations
- ❑ Reduced Support: reduced minimum support at lower levels
 - ▶ There are 4 search strategies:
 - Level-by-level independent
 - Level-cross filtering by k-itemset
 - Level-cross filtering by single item
 - Controlled level-cross filtering by single item

Multi-Level Mining: Progressive Deepening

40

- ❑ A top-down, progressive deepening approach
- ❑ First mine high-level frequent items:
milk (15%), bread (10%)
- ❑ Then mine their lower-level “weaker” frequent itemsets:
2% milk (5%), wheat bread (4%)
- ❑ Different min support threshold across
multi-levels lead to different algorithms:
 - ▶ If adopting the same min_support across multi-levels
then toss t if any of t 's ancestors is infrequent.
 - ▶ If adopting reduced min_support at lower levels
then examine only those descendents whose ancestor's
support is frequent/non-negligible.

Correlation Rules

- Suppose perform association rule mining

	c	Not c	row
t	20	5	25
Not t	70	5	75
col	90	10	100

$\{\text{tea}\} \Rightarrow \{\text{coffee}\} [20\% \ 80\%]$

- But the 90% of the customers buy coffee anyway!
- A customer who buys tea is less likely (10% less) to buy coffee than a customer about whom we have no information

$A \Rightarrow B$ [support, confidence, correlation]

- ❑ One measure is to calculate correlation
- ❑ If two random variables A and B are independent,

$$\text{lift}(A, B) = \frac{P(A \wedge B)}{P(A)P(B)} = 1$$

- ❑ For tea and coffee,

$$\frac{P(t \wedge c)}{P(t)P(c)} = 0.89$$

this means that the presence of tea is negatively correlated with the presence of coffee

- The χ^2 statistic is defined as

$$\chi^2 = \sum_{r \in R} \frac{(O(r) - E[r])^2}{E[r]}$$

- Look up for significance value in a statistical textbook
 - ▶ There are k-1 degrees of freedom
 - ▶ If test fails cannot reject independence, otherwise contingency table represents dependence.

- ❑ Use the χ^2 statistical test for testing independence
- ❑ Let n be the number of baskets
- ❑ Let k be the number of items considered
- ❑ Let r be a rule consisting of item values (i_1, i_2, \dots, i_k)
- ❑ Let $O(r)$ represent the number of baskets containing rule r (i.e. frequency)
- ❑ Let $E[i_j] = O(i_j) / n$ for a single item (negation is $n - E[i_j]$)
- ❑ $E[r] = n * E[r_1] / n * \dots * E[r_k] / n$

- ❑ Back to tea and coffee

	c	Not c	row
t	20	5	25
Not t	70	5	75
col	90	10	100

- ❑ $E[t] = 25$, $E[\text{Not } t]=75$, $E[c]=90$, $E[\text{Not } c]=10$
- ❑ $E[tc]=100 * 25/100 * 90 /100=22.5$
- ❑ $O(tc) = 20$
- ❑ Contrib. to $\chi^2 = (20 - 22.5)^2 / 22.5 = 0.278$
- ❑ Calculate for the rest to get $\chi^2=2.204$
- ❑ Not significant at 95% level (3.84 for $k=2$)
- ❑ Cannot reject independence assumption

- ❑ If χ^2 test shows significance, then want to find most interesting cell(s) in table
- ❑ $I(r) = O(r)/E[r]$
 - ▶ Look for values far away from 1
 - ▶ $I(t\ c) = 20/22.5 = 0.89$
 - ▶ $I(t\ \text{not } c) = 5/2.5 = 2$
 - ▶ $I(\text{not } t\ c) = 70/67.5 = 1.04$
 - ▶ $I(\text{not } t\ \text{not } c) = 5/7.5 = 0.66$

- ❑ Upward closed
 - ▶ If a k-itemset is correlated, so are all its supersets
 - ▶ Look for minimal correlated itemsets, no subset is correlated
- ❑ Can combine a-priori and χ^2 statistic efficiently
 - ▶ No generate and prune as in support-confidence

Sequential Rules

- ❑ Association rule mining does not consider the order of transactions
- ❑ In many applications such orderings are significant
- ❑ In market basket analysis, it is interesting to know whether people buy some items in sequence
- ❑ Example: buying bed first and then bed sheets later
- ❑ In Web usage mining,
 - ▶ it is useful to find navigational patterns of users in a Web site from sequences of page visits of users

❑ Web sequence

- ▶ < {Homepage} {Electronics} {Digital Cameras}
{Canon Digital Camera} {Shopping Cart} {Order
Confirmation} {Return to Shopping} >

❑ Sequence of initiating events causing the nuclear accident at 3-mile Island:

<{clogged resin} {outlet valve closure} {loss of feedwater}
{condenser polisher outlet valve shut} {booster pumps trip}
{main waterpump trips} {main turbine trips} {reactor
pressure increases}>

❑ Sequence of books checked out at a library:

- ▶ <{Fellowship of the Ring} {The Two Towers}
{Return of the King}>

- Let $I = \{i_1, i_2, \dots, i_m\}$ be a set of items.
- A sequence is an **ordered** list of itemsets.
- We denote a sequence s by $\langle a_1, a_2, \dots, a_r \rangle$, where a_i is an itemset, also called an element of s .
- An element (or an itemset) of a sequence is denoted by $\{x_1, x_2, \dots, x_k\}$, where $x_j \in I$ is an item.
- We assume without loss of generality that items in an element of a sequence are in lexicographic order

□ Size

- ▶ The size of a sequence is the number of elements (or itemsets) in the sequence

□ Length

- ▶ The length of a sequence is the number of items in the sequence
- ▶ A sequence of length k is called k -sequence

□ Given two sequences, $s_1 = \langle a_1 a_2 \dots a_r \rangle$ and $s_2 = \langle b_1 b_2 \dots b_v \rangle$

□ s_1 is a subsequence of s_2 if there exist integers $1 \leq j_1 < j_2 < \dots < j_{r-1} < j_r \leq v$ such that $a_1 \subseteq b_{j_1}, a_2 \subseteq b_{j_2}, \dots, a_r \subseteq b_{j_r}$

□ A sequence $s_1 = \langle a_1 a_2 \dots a_r \rangle$ is a subsequence of another sequence $s_2 = \langle b_1 b_2 \dots b_v \rangle$, or s_2 is a supersequence of s_1 , if there exist integers $1 \leq j_1 < j_2 < \dots < j_{r-1} < j_r \leq v$ such that $a_1 \subseteq b_{j_1}, a_2 \subseteq b_{j_2}, \dots, a_r \subseteq b_{j_r}$. We also say that s_2 contains s_1 .

- ❑ Let $I = \{1, 2, 3, 4, 5, 6, 7, 8, 9\}$.
- ❑ The sequence $\langle \{3\}\{4, 5\}\{8\} \rangle$ is contained in (or is a subsequence of) $\langle \{6\} \{3, 7\}\{9\}\{4, 5, 8\}\{3, 8\} \rangle$
- ❑ In fact, $\{3\} \subseteq \{3, 7\}$, $\{4, 5\} \subseteq \{4, 5, 8\}$, and $\{8\} \subseteq \{3, 8\}$
- ❑ However, $\langle \{3\}\{8\} \rangle$ is not contained in $\langle \{3, 8\} \rangle$ or vice versa
- ❑ The size of the sequence $\langle \{3\}\{4, 5\}\{8\} \rangle$ is 3, and the length of the sequence is 4.

- ❑ The input is a set S of input data sequences (or sequence database)
- ❑ The problem of mining sequential patterns is to find all the sequences that have a user-specified minimum support
- ❑ Each such sequence is called a frequent sequence, or a sequential pattern
- ❑ The support for a sequence is the fraction of total data sequences in S that contains this sequence

- ❑ Itemset
 - ▶ Non-empty set of items
 - ▶ Each itemset is mapped to an integer
- ❑ Sequence
 - ▶ Ordered list of itemsets
- ❑ Support for a Sequence
 - ▶ Fraction of total customers that support a sequence.
- ❑ Maximal Sequence
 - ▶ A sequence that is not contained in any other sequence
- ❑ Large Sequence
 - ▶ Sequence that meets minisup

Customer ID	Transaction Time	Items Bought
1	June 25 '93	30
1	June 30 '93	90
2	June 10 '93	10,20
2	June 15 '93	30
2	June 20 '93	40,60,70
3	June 25 '93	30,50,70
4	June 25 '93	30
4	June 30 '93	40,70
4	July 25 '93	90
5	June 12 '93	90

Customer ID	Customer Sequence
1	< (30) (90) >
2	< (10 20) (30) (40 60 70) >
3	< (30) (50) (70) >
4	< (30) (40 70) (90) >
5	< (90) >

Maximal seq with support > 40%
< (30) (90) >
< (30) (40 70) >

Note: Use Minisup of 40%, no less than two customers must support the sequence
 < (10 20) (30) > Does not have enough support (Only by Customer #2)
 < (30) >, < (70) >, < (30) (40) > ... are not maximal.

- ❑ Apriori-based Approaches
 - ▶ GSP
 - ▶ SPADE
- ❑ Pattern-Growth-based Approaches
 - ▶ FreeSpan
 - ▶ PrefixSpan

- ❑ The Apriori property for sequential patterns
 - ▶ If a sequence S is not frequent, then none of the super-sequences of S is frequent
 - ▶ For instance, if $\langle hb \rangle$ is infrequent so do $\langle hab \rangle$ and $\langle (ah)b \rangle$
- ❑ GSP (Generalized Sequential Pattern) mining algorithm
 - ▶ Initially, every item in DB is a candidate of length-1
 - ▶ for each level (i.e., sequences of length- k) do
 - scan database to collect support count for each candidate sequence
 - generate candidate length- $(k+1)$ sequences from length- k frequent sequences using Apriori
 - ▶ repeat until no frequent sequence or no candidate can be found
- ❑ Major strength: Candidate pruning by Apriori

The Generalized Sequential Pattern (GSP) Mining Algorithm

60

□ Step 1:

- ▶ Make the first pass over the sequence database D to yield all the 1-element frequent sequences

□ Step 2:

Repeat until no new frequent sequences are found

▶ **Candidate Generation:**

- Merge pairs of frequent subsequences found in the $(k-1)th$ pass to generate candidate sequences that contain k items

▶ **Candidate Pruning:**

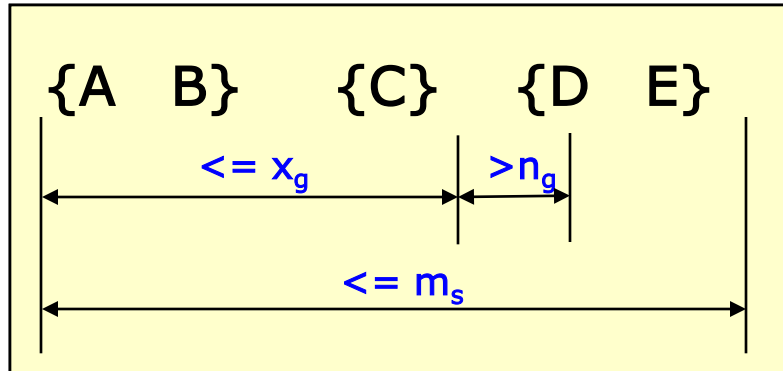
- Prune candidate k -sequences that contain infrequent $(k-1)$ -subsequences

▶ **Support Counting:**

- Make a new pass over the sequence database D to find the support for these candidate sequences

▶ **Candidate Elimination:**

- Eliminate candidate k -sequences whose actual support is less than *minsup*



x_g : max-gap

n_g : min-gap

m_s : maximum span

$x_g = 2, n_g = 0, m_s = 4$

Data sequence	Subsequence	Contain?
$\langle \{2,4\} \{3,5,6\} \{4,7\} \{4,5\} \{8\} \rangle$	$\langle \{6\} \{5\} \rangle$	Yes
$\langle \{1\} \{2\} \{3\} \{4\} \{5\} \rangle$	$\langle \{1\} \{4\} \rangle$	No
$\langle \{1\} \{2,3\} \{3,4\} \{4,5\} \rangle$	$\langle \{2\} \{3\} \{5\} \rangle$	Yes
$\langle \{1,2\} \{3\} \{2,3\} \{3,4\} \{2,4\} \{4,5\} \rangle$	$\langle \{1,2\} \{5\} \rangle$	No

- ❑ Approach 1:
 - ▶ Mine sequential patterns without timing constraints
 - ▶ Postprocess the discovered patterns

- ❑ Approach 2:
 - ▶ Modify GSP to directly prune candidates that violate timing constraints

Sequential Rules

- ❑ FP-growth
- ❑ Multi-level association rules
- ❑ Correlation rules
- ❑ Sequential pattern mining