



Mining Data Streams

Data Mining and Text Mining (UIC 583 @ Politecnico di Milano)

- ❑ What is stream data?
 - ❑ Why Stream Data Systems?
 - ❑ Stream data management systems: Issues and solutions
 - ❑ Stream data cube and multidimensional OLAP analysis
 - ❑ Stream frequent pattern analysis
 - ❑ Stream classification
 - ❑ Stream cluster analysis
 - ❑ Research issues
-
- ❑ Reference: Jiawei Han and Micheline Kamber, "Data Mining: Concepts and Techniques", The Morgan Kaufmann Series in Data Management Systems (Second Edition)
Chapter 8, part 1

Data Streams

❑ Data Streams vs DBMS

- ▶ **Data streams**
continuous, ordered, changing, fast, huge amount
- ▶ **Traditional DBMS**
data stored in finite, persistent data sets

❑ Characteristics

- ▶ Huge volumes of continuous data, possibly infinite
- ▶ Fast changing and requires fast, real-time response
- ▶ Data stream captures nicely our data processing needs of today
- ▶ Random access is expensive
single scan algorithm (can only have one look)
- ▶ Store only the summary of the data seen thus far
- ▶ Most stream data are at pretty low-level or multi-dimensional in nature, needs multi-level and multi-dimensional processing

- ❑ Telecommunication calling records
- ❑ Business: credit card transaction flows
- ❑ Network monitoring and traffic engineering
- ❑ Financial market: stock exchange
- ❑ Engineering & industrial processes: power supply & manufacturing
- ❑ Sensor, monitoring & surveillance: video streams, RFIDs
- ❑ Security monitoring
- ❑ Web logs and Web page click streams
- ❑ Massive data sets (even saved but random access is too expensive)

- ❑ Persistent relations
- ❑ One-time queries
- ❑ Random access
- ❑ “Unbounded” disk store
- ❑ Only current state matters
- ❑ No real-time services
- ❑ Relatively low update rate
- ❑ Data at any granularity
- ❑ Assume precise data
- ❑ Access plan determined by query processor, physical DB design
- ❑ Transient streams
- ❑ Continuous queries
- ❑ Sequential access
- ❑ Bounded main memory
- ❑ Historical data is important
- ❑ Real-time requirements
- ❑ Possibly multi-GB arrival rate
- ❑ Data at fine granularity
- ❑ Data stale/imprecise
- ❑ Unpredictable/variable data arrival and characteristics

Ack. From Motwani's PODS tutorial slides


The Architecture of Stream Query Processing

7

SDMS
(Stream Data Management System)

Continuous
Query

Multiple streams




User/Applications

Stream Query
Processor

Results

Scratch Space
(Main memory and/or Disk)



What are Challenges of Stream Data Processing?

- ❑ Multiple, continuous, rapid, time-varying, ordered streams
- ❑ Main memory computations
- ❑ Queries are often continuous
 - ▶ Evaluated continuously as stream data arrives
 - ▶ Answer updated over time
- ❑ Queries are often complex
 - ▶ Beyond element-at-a-time processing
 - ▶ Beyond stream-at-a-time processing
 - ▶ Beyond relational queries (scientific, data mining, OLAP)
- ❑ Multi-level/multi-dimensional processing and data mining
 - ▶ Most stream data are at low-level or multi-dimensional in nature

❑ Query types

- ▶ One-time query vs. **continuous query** (being evaluated continuously as stream continues to arrive)
- ▶ **Predefined query** vs. ad-hoc query (issued on-line)

❑ Unbounded memory requirements

- ▶ For real-time response, **main memory algorithm** should be used
- ▶ Memory requirement is unbounded if one will join future tuples

❑ Approximate query answering

- ▶ With bounded memory, it is not always possible to produce exact answers
- ▶ **High-quality approximate answers** are desired
- ▶ Data reduction and synopsis construction methods: Sketches, random sampling, histograms, wavelets, etc.

What the Methodologies for Stream Data Processing?

10

- ❑ Major challenges
 - ▶ Keep track of a large universe, e.g., pairs of IP address, not ages
- ❑ Methodology
 - ▶ **Synopses** (trade-off between accuracy and storage)
 - ▶ Use synopsis data structure, much smaller ($O(\log^k N)$ space) than their base data set ($O(N)$ space)
 - ▶ Compute an approximate answer within a small error range (factor ϵ of the actual answer)
- ❑ Major methods
 - ▶ Random sampling
 - ▶ Histograms
 - ▶ Sliding windows
 - ▶ Multi-resolution model
 - ▶ Sketches
 - ▶ Radomized algorithms

- ❑ **Random sampling** (but without knowing the total length in advance)
 - ▶ Reservoir sampling: maintain a set of s candidates in the reservoir, which form a true random sample of the element seen so far in the stream. As the data stream flow, every new element has a certain probability (s/N) of replacing an old element in the reservoir.
- ❑ **Sliding windows**
 - ▶ Make decisions based only on recent data of sliding window size w
 - ▶ An element arriving at time t expires at time $t + w$
- ❑ **Histograms**
 - ▶ Approximate the frequency distribution of element values in a stream
 - ▶ Partition data into a set of contiguous buckets
 - ▶ Equal-width (equal value range for buckets) vs. V-optimal (minimizing frequency variance within each bucket)
- ❑ **Multi-resolution models**
 - ▶ Popular models: balanced binary trees, micro-clusters, and wavelets

- ❑ Sliding windows
 - ▶ Only over sliding windows of **recent** stream data
 - ▶ Approximation but often more desirable in applications
- ❑ Batched processing, sampling and synopses
 - ▶ **Batched** if update is fast but computing is slow
 - Compute periodically, not very timely
 - ▶ **Sampling** if update is slow but computing is fast
 - Compute using sample data, but not good for joins, etc.
 - ▶ **Synopsis** data structures
 - Maintain a small synopsis or sketch of data
 - Good for querying historical data
- ❑ Blocking operators, e.g., sorting, avg, min, etc.
 - ▶ **Blocking** if unable to produce the first output until seeing the entire input

- ❑ Research projects and system prototypes
 - ▶ STREAM (Stanford): A general-purpose DSMS
 - ▶ Cougar (Cornell): sensors
 - ▶ Aurora (Brown/MIT): sensor monitoring, dataflow
 - ▶ Hancock (AT&T): telecom streams
 - ▶ Niagara (OGI/Wisconsin): Internet XML databases
 - ▶ OpenCQ (Georgia Tech): triggers, incr. view maintenance
 - ▶ Tapestry (Xerox): pub/sub content-based filtering
 - ▶ Telegraph (Berkeley): adaptive engine for sensors
 - ▶ Tradebot (www.tradebot.com): stock tickers & streams
 - ▶ Tribeca (Bellcore): network monitoring
 - ▶ MAIDS (UIUC/NCSA): Mining Alarming Incidents in Data Streams

- ❑ Stream mining—A more challenging task in many cases
 - ▶ It shares most of the difficulties with stream querying
 - ▶ But often requires less “precision”, e.g., no join, grouping, sorting
 - ▶ Patterns are hidden and more general than querying
 - ▶ It may require exploratory analysis, not necessarily continuous queries

- ❑ Stream data mining tasks
 - ▶ Multi-dimensional on-line analysis of streams
 - ▶ Mining outliers and unusual patterns in stream data
 - ▶ Clustering data streams
 - ▶ Classification of stream data

- ❑ Most stream data are at pretty low-level or multi-dimensional in nature: needs ML/MD processing
- ❑ Analysis requirements
 - ▶ Multi-dimensional trends and unusual patterns
 - ▶ Capturing important changes at multi-dimensions/levels
 - ▶ Fast, real-time detection and response
 - ▶ Comparing with data cube: Similarity and differences
- ❑ Stream (data) cube or stream OLAP: Is this feasible?
 - ▶ Can we implement it efficiently?

❑ Analysis of Web click streams

- ▶ Raw data at low levels: seconds, web page addresses, user IP addresses, ...
- ▶ Analysts want: changes, trends, unusual patterns, at reasonable levels of details
- ▶ E.g., Average clicking traffic in North America on sports in the last 15 minutes is 40% higher than that in the last 24 hours."

❑ Analysis of power consumption streams

- ▶ Raw data: power consumption flow for every household, every minute
- ▶ Patterns one may find: **average** hourly power consumption **surges up 30%** for **manufacturing companies** in Chicago in the **last 2 hours today** than that of the same **day a week ago**

Architectures

❑ A tilted time frame

- ▶ Different time granularities:
second, minute, quarter, hour, day, week, ...

❑ Critical layers


- ▶ Minimum interest layer (m-layer)
- ▶ Observation layer (o-layer)
- ▶ User: watches at o-layer and occasionally needs to drill-down down to m-layer

❑ Partial materialization of stream cubes

- ▶ Full materialization: too space and time consuming
- ▶ No materialization: slow response at query time
- ▶ Partial materialization: what do we mean “partial”?


□ Natural tilted time frame:

- ▶ Example: Minimal: quarter, then 4 quarters \rightarrow 1 hour, 24 hours \rightarrow day, ...



□ Logarithmic tilted time frame:

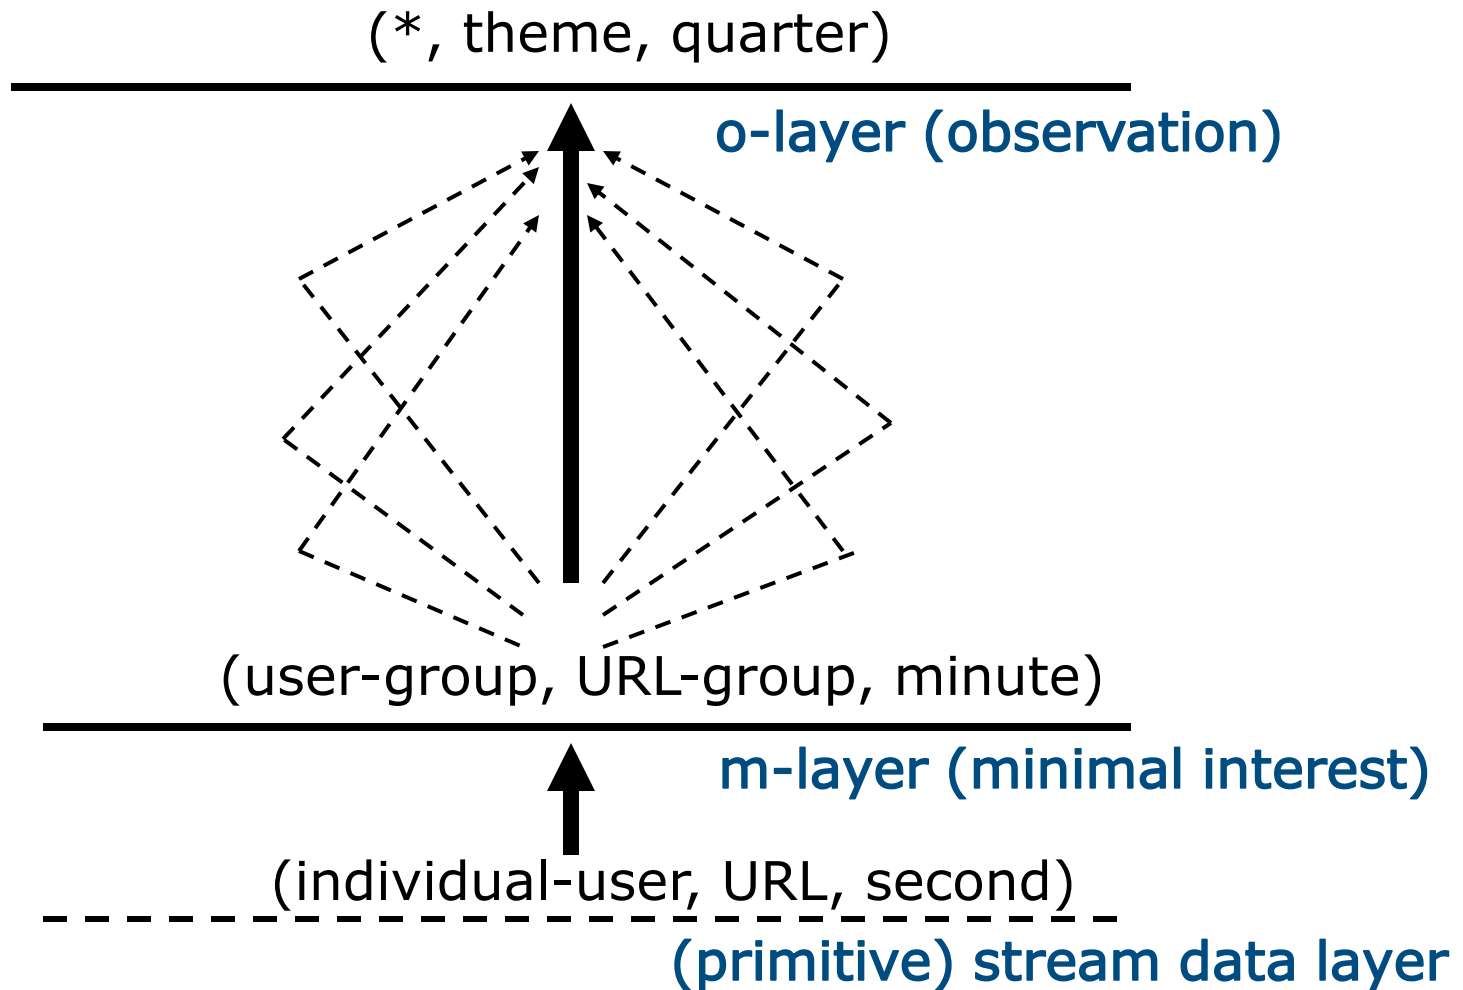
- ▶ Example: Minimal: 1 minute, then 1, 2, 4, 8, 16, 32, ...



❑ Pyramidal tilted time frame

- ▶ Example: Suppose there are 5 frames and each takes maximal 3 snapshots
- ▶ Given a snapshot number N , if $N \bmod 2d = 0$, insert into the frame number d . If there are more than 3 snapshots, “kick out” the oldest one.

Frame no.	Snapshots (by clock time)
0	69 67 65
1	70 66 62
2	68 60 52
3	56 40 24
4	48 16
5	64 32



❑ On-line materialization


- ▶ Materialization takes precious space and time
 - Only incremental materialization (with tilted time frame)
- ▶ Only materialize “cuboids” of the critical layers?
 - Online computation may take too much time
- ▶ Preferred solution:
 - popular-path approach: Materializing those along the popular drilling paths
 - H-tree structure: Such cuboids can be computed and stored efficiently using the H-tree structure


❑ Online aggregation vs. query-based computation

- ▶ Online computing while streaming:
aggregating stream cubes
- ▶ Query-based computation:
using computed cuboids

Stream Cube Structure: From m-layer to o-layer

25





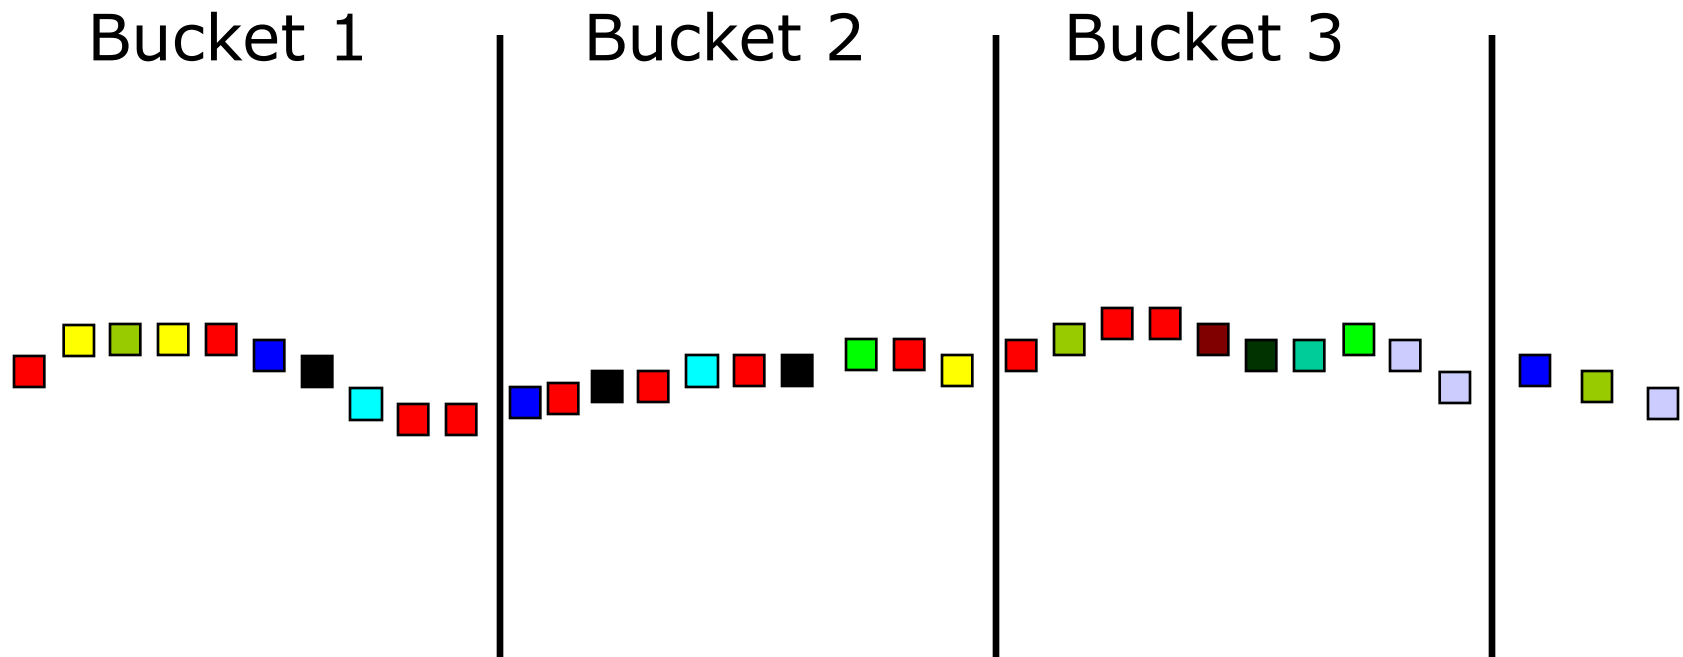
- ❑ H-tree and H-cubing
 - ▶ Developed for computing data cubes and ice-berg cubes
 - ▶ Fast cubing, space preserving in cube computation

- ❑ Using H-tree for stream cubing
 - ▶ Space preserving: intermediate aggregates can be computed incrementally and saved in tree nodes
 - ▶ Facilitate computing other cells and multi-dimensional analysis
 - ▶ H-tree with computed cells can be viewed as stream cube

Frequent patterns


- ❑ Frequent pattern mining is valuable in stream applications
 - ▶ e.g., network intrusion mining
- ❑ Mining precise freq. patterns in stream data: unrealistic
 - ▶ Even store them in a compressed form, such as FPtree
- ❑ How to mine frequent patterns with good approximation?
 - ▶ Approximate frequent patterns (Manku & Motwani VLDB'02)
 - ▶ Keep only current frequent patterns? No changes can be detected
- ❑ Mining evolution freq. patterns
(C. Giannella, J. Han, X. Yan, P.S. Yu, 2003)
 - ▶ Use tilted time window frame
 - ▶ Mining evolution and dramatic changes of frequent patterns
- ❑ Space-saving computation of frequent and top-k elements
(Metwally, Agrawal, and El Abbadi, ICDT'05)

- ❑ Mining precise freq. patterns in stream data: unrealistic
 - ▶ Even store them in a compressed form, such as FPtree
- ❑ Approximate answers are often sufficient (e.g., trend/pattern analysis)
- ❑ Example: a router is interested in all flows:
 - ▶ whose frequency is at least 1% (σ) of the entire traffic stream seen so far
 - ▶ and feels that 1/10 of σ ($\varepsilon = 0.1\%$) error is comfortable
- ❑ How to mine frequent patterns with good approximation?
 - ▶ Lossy Counting Algorithm (Manku & Motwani, VLDB'02)
 - ▶ Major ideas: not tracing items until it becomes frequent
 - ▶ Adv: guaranteed error bound
 - ▶ Disadv: keep a large set of traces




Divide Stream into 'Buckets' (bucket size is $1/\epsilon = 1000$)

Empty
(summary)




At bucket boundary, decrease all counters by 1




At bucket boundary, decrease all counters by 1


- ❑ Given
 - ▶ (1) support threshold: σ ,
 - ▶ (2) error threshold: ε , and
 - ▶ (3) stream length N
- ❑ Output: items with frequency counts exceeding $(\sigma - \varepsilon) N$
- ❑ How much do we undercount?
 - ▶ If (stream length seen so far = N)
and (bucket-size = $1/\varepsilon$)
then frequency count error $\leq \text{\#buckets} = \varepsilon N$
- ❑ Approximation guarantee
 - ▶ No false negatives
 - ▶ False positives have true frequency count at least $(\sigma - \varepsilon)N$
 - ▶ Frequency count underestimated by at most εN
 - ▶ The space requirement is limited to $1/\varepsilon \log(\varepsilon N)$

Divide Stream into 'Buckets' as for frequent items
But fill as many buckets as possible in main memory one time





If we put 3 buckets of data into main memory one time,
Then decrease each frequency count by 3



Itemset () is deleted.

That's why we choose a large number of buckets
– delete more



If we find itemset () is not frequent itemset,
Then we needn't consider its superset

- ❑ Strength
 - ▶ A simple idea
 - ▶ Can be extended to frequent itemsets

- ❑ Weakness:
 - ▶ Space Bound is not good
 - ▶ For frequent itemsets, they do scan each record many times
 - ▶ The output is based on all previous data. But sometimes, we are only interested in recent data

- ❑ A space-saving method for stream frequent item mining
 - ▶ Metwally, Agrawal and El Abbadi, ICDT'05

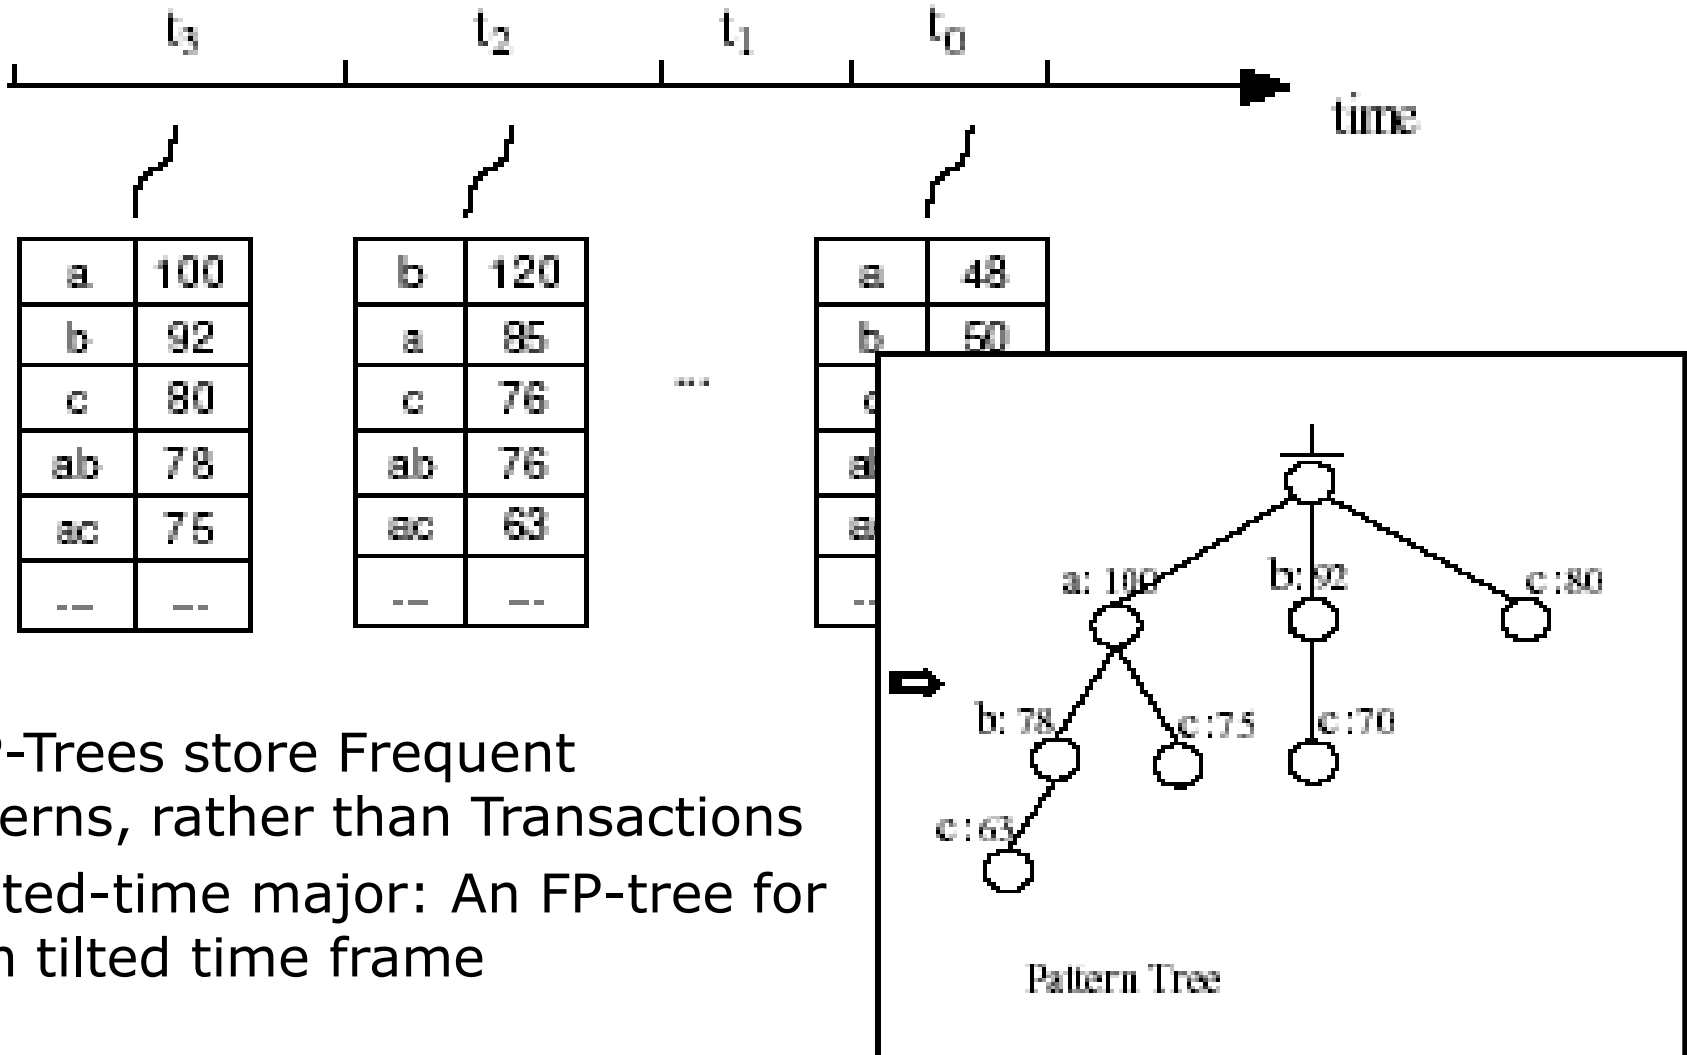
- ❑ Approximate frequent patterns (Manku & Motwani VLDB'02)
 - ▶ Keep only current frequent patterns:
No changes can be detected

- ❑ Mining evolution and dramatic changes of frequent patterns (Giannella, Han, Yan, Yu, 2003)
 - ▶ Use tilted time window frame
 - ▶ Use compressed form to store significant (approximate) frequent patterns and their time-dependent traces

- ❑ Note: To mine precise counts, one has to trace/keep a fixed (and small) set of items

Two Structures for Mining Frequent Patterns with Tilted-Time Window (1)

40




- ❑ FP-Trees store Frequent Patterns, rather than Transactions
- ❑ Tilted-time major: An FP-tree for each tilted time frame

Two Structures for Mining Frequent Patterns with Tilted-Time Window (2)

41

□ The second data structure:

- ▶ Observation: FP-Trees of different time units are similar
- ▶ Pattern-tree major: each node is associated with a tilted-time window



Classification

Classification in Data Streams

What are the issues?

43

- ❑ It is impossible to store the whole data set, as traditional classification algorithms require
- ❑ It is usually not possible to perform multiple scans of the input data
- ❑ Data streams are time-varying!
There is concept drift.

- ❑ Decision tree induction for stream data classification
 - ▶ VFDT (Very Fast Decision Tree)/CVFDT

- ❑ Is decision-tree good for modeling fast changing data, e.g., stock market analysis?

- ❑ Other stream classification methods
 - ▶ Instead of decision-trees, consider other models:
 - Naïve Bayesian,
 - Ensemble (Wang, Fan, Yu, Han. KDD'03)
 - K-nearest neighbors (Aggarwal, Han, Wang, Yu. KDD'04)
 - ▶ Tilted time framework, incremental updating, dynamic maintenance, and model construction
 - ▶ Comparing of models to find changes

- ❑ Initially introduced to analyze click-streams
- ❑ With high probability, classifies tuples the same
- ❑ Only uses small sample
 - ▶ Based on Hoeffding Bound principle

- ❑ Hoeffding Bound (Additive Chernoff Bound)
 - ▶ r : random variable representing the attribute selection method
 - ▶ R : range of r
 - ▶ n : # independent observations
 - ▶ Mean of r is at least $r_{\text{avg}} - \varepsilon$, with probability $1 - \delta$

$$\varepsilon = \sqrt{\frac{R^2 \ln(1 / \delta)}{2n}}$$

- The algorithm uses the bound to determine, with high probability the smallest number N of examples needed at a node to select the splitting attribute

□ Hoeffding Tree Input


- ▶ S : sequence of examples
- ▶ X : attributes
- ▶ $G(\cdot)$: evaluation function
- ▶ δ : desired accuracy

```
for each example in S
    retrieve  $G(X_a)$  and  $G(X_b)$     //two highest  $G(X_i)$ 
    if (  $G(X_a) - G(X_b) > \epsilon$  )
        split on  $X_a$ 
        recurse to next node
    break
```

- Complexity is $O(ldvc)$ where l is the depth, d is the number of attributes, v is the maximum number of attributes, c is the number of classes

Decision-Tree Induction with Data Streams

48



From Gehrke's SIGMOD tutorial slides

Hoeffding Tree: Strengths and Weaknesses

49

Strengths

- ❑ Scales better than traditional methods
 - ▶ Sublinear with sampling
 - ▶ Very small memory utilization
- ❑ Incremental
 - ▶ Make class predictions in parallel
 - ▶ New examples are added as they come

Weaknesses

- ❑ Could spend a lot of time with ties
- ❑ Memory used with tree expansion
- ❑ Number of candidate attributes

- ❑ Modifications to Hoeffding Tree
 - ▶ Near-ties broken more aggressively
 - ▶ G computed every n_{\min}
 - ▶ Deactivates certain leaves to save memory
 - ▶ Poor attributes dropped
 - ▶ Initialize with traditional learner (helps learning curve)
- ❑ Compare to Hoeffding Tree: Better time and memory
- ❑ Compare to traditional decision tree
 - ▶ Similar accuracy
 - ▶ Better runtime with 1.61 million examples
 - 21 minutes for VFDT
 - 24 hours for C4.5
- ❑ Still does not handle concept drift

❑ Concept Drift

- ▶ Time-changing data streams
- ▶ Incorporate new and eliminate old

❑ CVFDT

- ▶ Increments count with new example
- ▶ Decrement old example
 - Sliding window
 - Nodes assigned monotonically increasing IDs
- ▶ Grows alternate subtrees
- ▶ When alternate more accurate, then replace old
- ▶ $O(w)$ better runtime than VFDT-window

- ❑ H. Wang, W. Fan, P. S. Yu, and J. Han, "Mining Concept-Drifting Data Streams using Ensemble Classifiers", KDD'03.
- ❑ Method (derived from the ensemble idea in classification)

`train K classifiers from K chunks`

`for each subsequent chunk`

`train a new classifier`

`test other classifiers against the chunk`

`assign weight to each classifier`

`select top K classifiers`

Clustering


Clustering Evolving Data Streams

What methodologies?

54

- ❑ Compute and store summaries of past data
- ❑ Apply a divide-and-conquer strategy
- ❑ Incremental clustering of incoming data streams
- ❑ Perform microclustering as well as macroclustering analysis
- ❑ Explore multiple time granularity for the analysis of cluster evolution
- ❑ Divide stream clustering into on-line and off-line processes

- ❑ Base on the k-median method
 - ▶ Data stream points from metric space
 - ▶ Find k clusters in the stream s.t. the sum of distances from data points to their closest center is minimized
- ❑ Constant factor approximation algorithm
In small space, a simple two step algorithm:
 1. For each set of M records, S_i ,
find $O(k)$ centers in S_1, \dots, S_l
Local clustering: Assign each point in S_i to its closest center
 2. Let S' be centers for S_1, \dots, S_l with each center weighted by number of points assigned to it
Cluster S' to find k centers



❑ Method

- ▶ Maintain at most m level- i medians
- ▶ On seeing m of them, generate $O(k)$ level- $(i+1)$ medians of weight equal to the sum of the weights of the intermediate medians assigned to them

❑ Drawbacks

- ▶ Low quality for evolving data streams (register only k centers)
- ▶ Limited functionality in discovering and exploring clusters over different portions of the stream over time

- ❑ Network intrusion detection: one example
 - ▶ Detect bursts of activities or abrupt changes in real time—by on-line clustering

- ❑ The methodology by C. Agarwal, J. Han, J. Wang, P.S. Yu, VLDB'03
 - ▶ Tilted time frame work:
o.w. dynamic changes cannot be found
 - ▶ Micro-clustering: better quality than k-means/k-median
 - incremental, online processing and maintenance)
 - ▶ Two stages: micro-clustering and macro-clustering
 - ▶ With limited “overhead” to achieve high efficiency, scalability, quality of results and power of evolution/change detection

□ Design goal

- ▶ High quality for clustering evolving data streams with greater functionality
- ▶ While keep the stream mining requirement in mind
 - One-pass over the original stream data
 - Limited space usage and high efficiency

□ CluStream: A framework for clustering evolving data streams

- ▶ Divide the clustering process into online and offline components
- ▶ Online component: periodically stores summary statistics about the stream data
- ▶ Offline component: answers various user questions based on the stored summary statistics

❑ Micro-cluster

- ▶ Statistical information about data locality
- ▶ Temporal extension of the cluster-feature vector
 - Multi-dimensional points $X_1 \dots X_k \dots$
with time stamps $T_1 \dots T_k \dots$
 - Each point contains d dimensions, i.e., $X = (x^1 \dots x^d)$
 - A micro-cluster for n points is defined as a $(2.d + 3)$ tuple

$$\left(\overline{CF2^x}, \overline{CF1^x}, CF2^t, CF1^t, n \right)$$

❑ Pyramidal time frame

- ▶ Decide at what moments the snapshots of the statistical information are stored away on disk

- ❑ Snapshots of a set of micro-clusters are stored following the pyramidal pattern
- ❑ They are stored at differing levels of granularity depending on the recency
- ❑ Snapshots are classified into different orders varying from 1 to $\log(T)$
 - ▶ The i -th order snapshots occur at intervals of α^i where $\alpha \geq 1$
 - ▶ Only the last $(\alpha + 1)$ snapshots are stored

❑ Online micro-cluster maintenance

- ▶ Initial creation of q micro-clusters
 - q is usually significantly larger than the number of natural clusters
- ▶ Online incremental update of micro-clusters
 - If new point is within max-boundary, insert into the micro-cluster
 - O.w., create a new cluster
 - May delete obsolete micro-cluster or merge two closest ones

❑ Query-based macro-clustering

- ▶ Based on a user-specified time-horizon h and the number of macro-clusters K , compute macroclusters using the k-means algorithm

Stream Data Mining:

What are the Research Issues?

63

- ❑ Mining sequential patterns in data streams
- ❑ Mining partial periodicity in data streams
- ❑ Mining notable gradients in data streams
- ❑ Mining outliers and unusual patterns in data streams

- ❑ Stream clustering
 - ▶ Multi-dimensional clustering analysis?
Cluster not confined to 2-D metric space, how to incorporate other features, especially non-numerical properties
 - ▶ Stream clustering with other clustering approaches?
 - ▶ Constraint-based cluster analysis with data streams?

Summary

- ❑ Stream Data Mining is a rich and on-going research field
- ❑ Current research focus in database community:
 - ▶ DSMS system architecture
 - ▶ Continuous query processing
 - ▶ Supporting mechanisms
- ❑ Stream data mining and stream OLAP analysis
 - ▶ Powerful tools for finding general and unusual patterns
 - ▶ Effectiveness, efficiency and scalability:
lots of open problems
- ❑ Philosophy on stream data analysis and mining
 - ▶ A multi-dimensional stream analysis framework
 - ▶ Time is a special dimension: Tilted time frame
 - ▶ What to compute and what to save?—Critical layers
 - ▶ Partial materialization and precomputation
 - ▶ Mining dynamics of stream data

- ❑ C. Aggarwal, J. Han, J. Wang, P. S. Yu. A Framework for Clustering Data Streams, VLDB'03
- ❑ C. C. Aggarwal, J. Han, J. Wang and P. S. Yu. On-Demand Classification of Evolving Data Streams, KDD'04
- ❑ C. Aggarwal, J. Han, J. Wang, and P. S. Yu. A Framework for Projected Clustering of High Dimensional Data Streams, VLDB'04
- ❑ S. Babu and J. Widom. Continuous Queries over Data Streams. SIGMOD Record, Sept. 2001
- ❑ B. Babcock, S. Babu, M. Datar, R. Motwani and J. Widom. Models and Issues in Data Stream Systems", PODS'02. ([Conference tutorial](#))
- ❑ Y. Chen, G. Dong, J. Han, B. W. Wah, and J. Wang. "Multi-Dimensional Regression Analysis of Time-Series Data Streams, VLDB'02
- ❑ P. Domingos and G. Hulten, "Mining high-speed data streams", KDD'00
- ❑ A. Dobra, M. N. Garofalakis, J. Gehrke, R. Rastogi. Processing Complex Aggregate Queries over Data Streams, SIGMOD'02
- ❑ J. Gehrke, F. Korn, D. Srivastava. On computing correlated aggregates over continuous data streams. SIGMOD'01
- ❑ C. Giannella, J. Han, J. Pei, X. Yan and P.S. Yu. Mining frequent patterns in data streams at multiple time granularities, Kargupta, et al. (eds.), Next Generation Data Mining'04

- ❑ S. Guha, N. Mishra, R. Motwani, and L. O'Callaghan. Clustering Data Streams, FOCS'00
- ❑ G. Hulten, L. Spencer and P. Domingos: Mining time-changing data streams. KDD 2001
- ❑ S. Madden, M. Shah, J. Hellerstein, V. Raman, Continuously Adaptive Continuous Queries over Streams, SIGMOD02
- ❑ G. Manku, R. Motwani. Approximate Frequency Counts over Data Streams, VLDB'02
- ❑ A. Metwally, D. Agrawal, and A. El Abbadi. Efficient Computation of Frequent and Top-k Elements in Data Streams. ICDT'05
- ❑ S. Muthukrishnan, Data streams: algorithms and applications, Proceedings of the fourteenth annual ACM-SIAM symposium on Discrete algorithms, 2003
- ❑ R. Motwani and P. Raghavan, Randomized Algorithms, Cambridge Univ. Press, 1995
- ❑ S. Viglas and J. Naughton, Rate-Based Query Optimization for Streaming Information Sources, SIGMOD'02
- ❑ Y. Zhu and D. Shasha. StatStream: Statistical Monitoring of Thousands of Data Streams in Real Time, VLDB'02
- ❑ H. Wang, W. Fan, P. S. Yu, and J. Han, Mining Concept-Drifting Data Streams using Ensemble Classifiers, KDD'03