8th European Conference on Interactive TV and Video 1

Tampere, Finland – June 2010, 9th

Recommender Systems for Interactive TV

Tutorial

Roberto Turrin, Paolo Cremonesi roberto.turrin@polimi.it paolo.cremonesi@poliimi.it



- Introduction
- Recommender systems
 - Content-based and collaborative recommendations
 - Advanced algorithms
- Evaluation of recommender systems

- Case study: a recommender system for IPTV/VOD provider
- Recommender system demo



Introduction

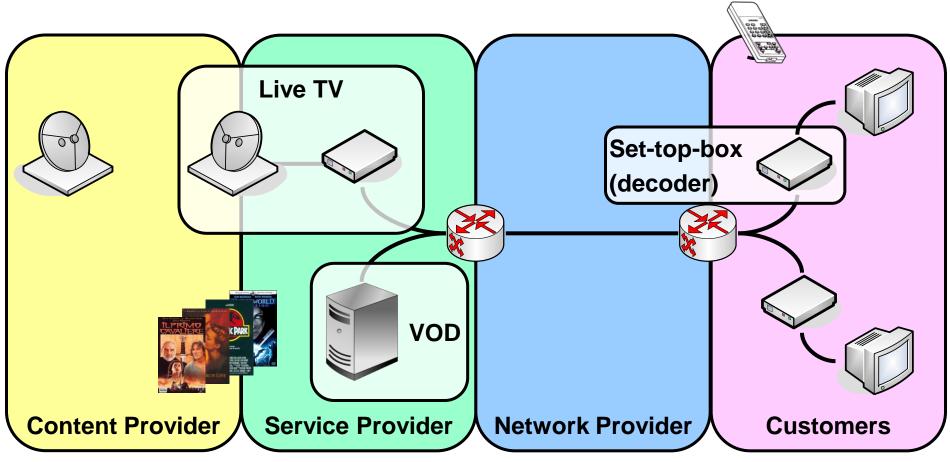
- Recommender systems
 - Content-based and collaborative recommendations
 - Advanced algorithms
- Evaluation of recommender systems

- Case study: a recommender system for IPTV/VOD provider
- Recommender system demo



A scenario: IPTV







IPTV Platform: Now









CUSTOMERS FACE DIFFICULTIES FINDING THE "RIGHT" CONTENT





IPTV Platform: with a recommender system

From this....



Today recommendations, based on your personal taste, are:

To this.



Peculiarities in the settings of VoD/IPTV

- Large number of items (movies, TV series, music, etc)
- Slow channel switching
- The available resources increase continuously so each item gets less and less visibility
- The amount of resources interesting for a single user is relatively low
- The interaction between user & system is a remote control



- Introduction
- Recommender systems
 - Content-based and collaborative recommendations
 - Advanced algorithms
- Evaluation of recommender systems

- Case study: a recommender system for IPTV/VOD provider
- Recommender system demo



Recommender systems

A recommender system is meant to alleviate the problem of searching for personalized resources by proposing to a user a very limited set of items that he is probably interested in.





- Improve user satisfaction, surprise users
 - Users pay for a better service
- Selling new content to users (increase sales)
 - VOD (Video On Demand)
 - Pay-per-view channels
- Targeted advertisement



How recommender systems works

Main **inputs**:

- User tastes?
 - User ratings (profile)



Required

- Explicit
- Implicit
- content preferences
- Item content information

Main outputs:

- Predicted ratings
- Customized list of items



Main interest



Explicit Ratings: each user can actively express his judgement about each item, for example by assigning a numerical value to it.

- © Easy computation of recommendation
- B Users might get annoyed and leave hasty ratings or don't leave a rating at all
- Ratings are subjected to personal interpretation

Implicit Ratings: the recommendation system implicitly infers the rating by observing the user behaviour (time spent watching a particular TV show, purchased items history...)

- © Transparent to the users (not annoyed)
- Often offers effective recommendation...
- ...but not so often perceived as such by the user
- Requires complex computations/inference



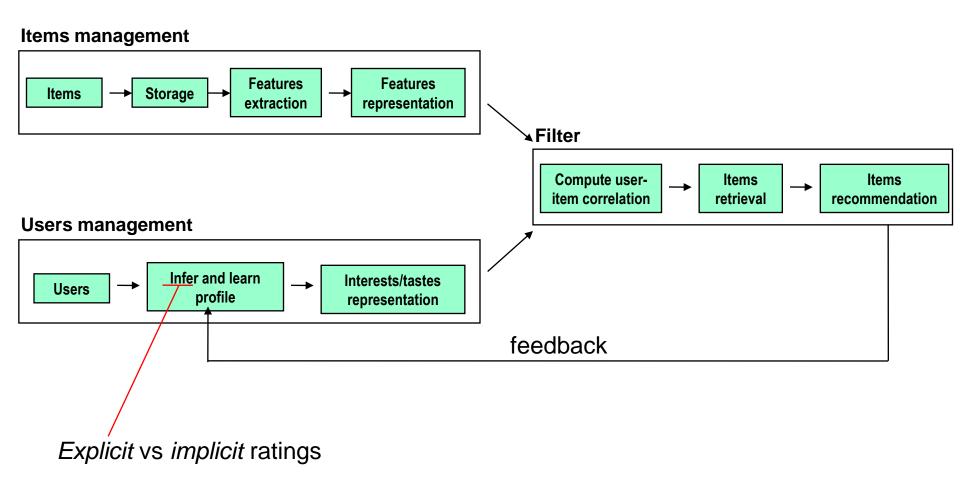
In several domains <u>we typically dispose only of an implicit dataset</u> or, sometimes, of very few explicit ratings that coexist in a mainly-implicit dataset, for instance because users get annoyed by leaving too many ratings.



We have a dataset composed by a mixture of explicit and implicit ratings



Recommender architecture





Recommender architecture: requirements

We would like the recommender system to guarantee:

- High quality of recommendations:
 - the capability to suggest personalized items which are potentially interesting to the user and to discard items which are disliked by the user
 - Other aspects of quality...
- Capability to recommend any user at any time (e.g., even if his profile has just been added)
- Scalability and real-time performance: a recommender system deals with a large amount of data, so it is crucial to apply algorithms which are capable to satisfy real-time response times

The architecture of a recommender system is generally divided into two asynchronous components:

- The *batch* or "off line" process
- The *real-time* or "on line" process

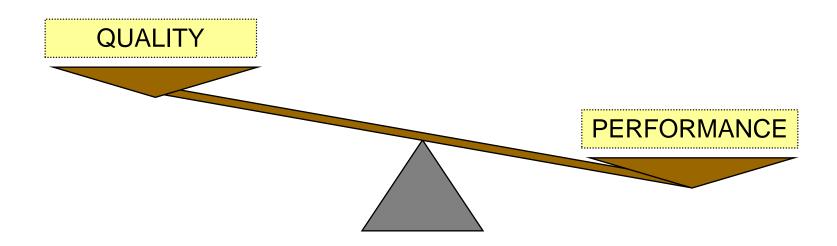
PURPOSE: move as many operations as possible in the off-line part in order to lighten the on-line part.

NEPTUNY



➡ The batch component faces with the most expensive part of the computation with the highest memory and computing requirements

➡ The real-time component faces the real-time requests for recommendation from the client interface, with rigid time constraints





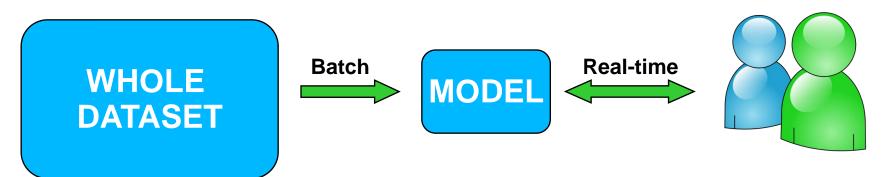
Item-, model-based algorithms

Using item-, model-based algorithms!

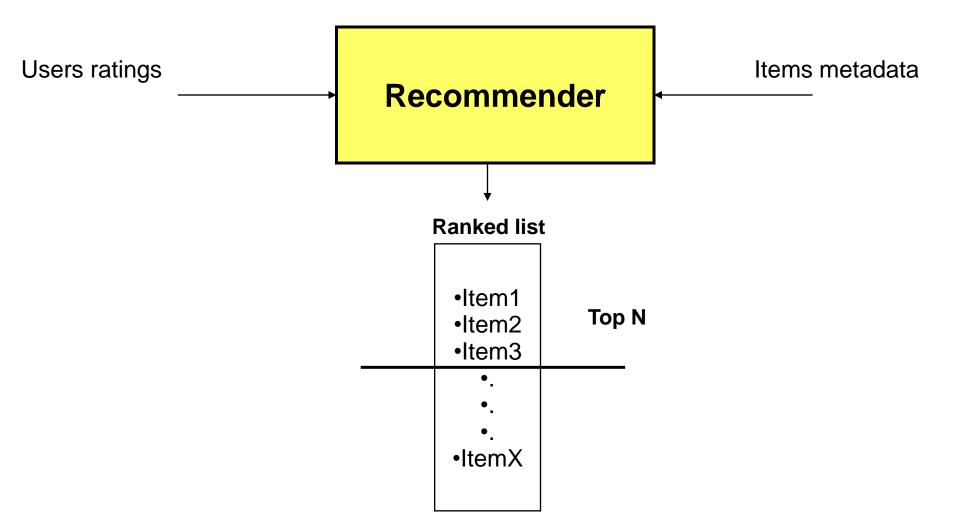
MODEL: a representation of the whole dataset which captures the principal information

item-based: only items are explicitly modeled, user model is implicitly derived

- The real-time component can work on the model instead of the whole dataset
- Generating the model is VERY expensive in terms of calculations.









Road to personalization

- One of the main goals of recommender systems is providing personalized suggestions to customers
 - Higher personalization means better technique
- We can then classify today's solutions in ascending order of personalization effectiveness as follows:
 - Non-personalized: a system which only gives generic suggestions, identical for every viewer
 - Roughly personalized: personalization is trivial (e.g., based on preferred category)
 - Personalized: advanced techniques are used in order to get a strong personalization, this can be obtained via contentbased or collaborative filtering



Non- and roughly personalized systems

- Examples of *non-personalized* systems are recommendations based on advertising, popularity (e.g., top sellers) or highest average rating
- Very easy to implement!
- Suggestions are not accurate and thus may be irrelevant

- The main example of *rough personalization* is that based on the so-called *demographic profiles*: (e.g., age, gender, location, ...)
- The suggestions are not-so-generic
- No need to provide info about preferences to get suggestions
- Bigh personalization impossible, due to exiguity data about profiles

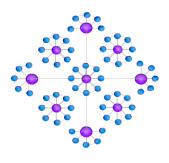


Non- and roughly personalized systems: baselines

Ratings are affected by subjectivity



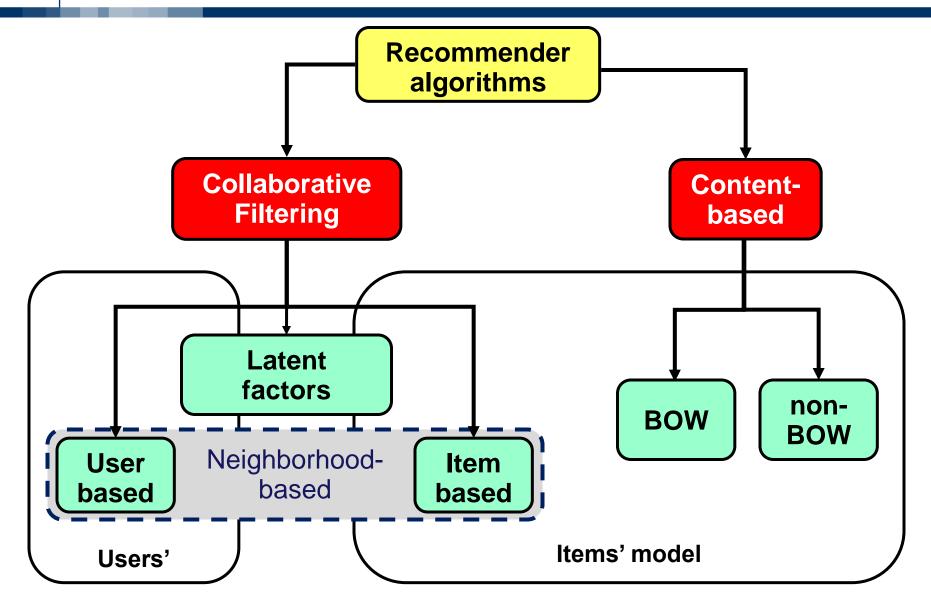




$$\widetilde{r}_{ui} = r_{ui} - (b_u + b_i + \mu)$$



Personalized systems





- Introduction
- Recommender systems
 - Content-based and collaborative recommendations
 - Advanced algorithms
- Evaluation of recommender systems

- Case study: a recommender system for IPTV/VOD provider
- Recommender system demo



Content-based and collaborative systems

Content-based

- based on the analysis of the content of the items
- Items are modelled by sets of features representing their content, and the assumption is that those features can somehow capture the meaning and the relevance of the item

Collaborative

 based on the preferences expressed by other users: it is not needed to extract any explicit features from the items, since recommendations are based only on the community behaviour

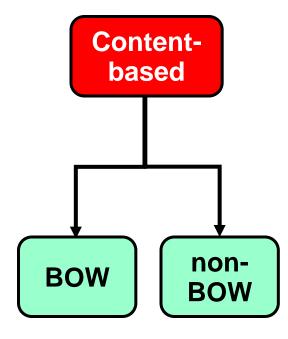
NEPTUNY



Content-based systems

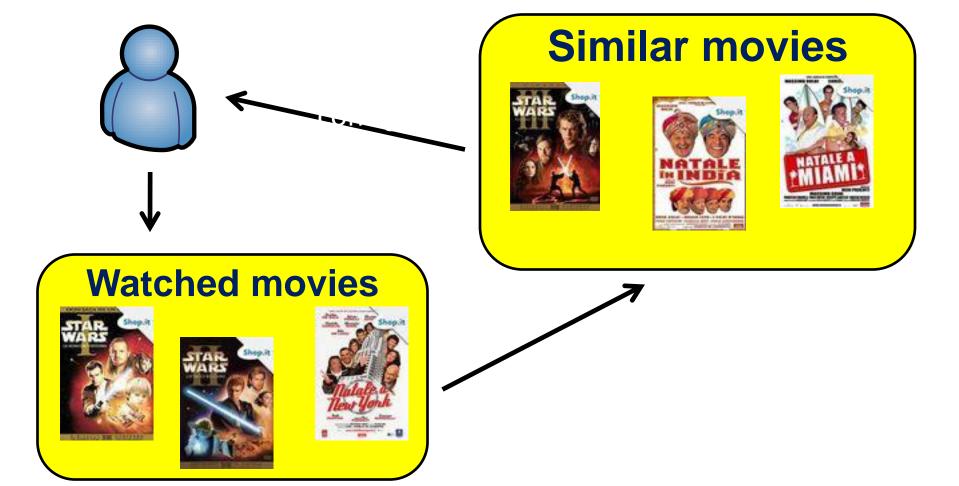
Content-based

- based on the analysis of the content of the items
- Items are modelled by sets of features representing their content, and the assumption is that those features can somehow capture the meaning and the relevance of the item

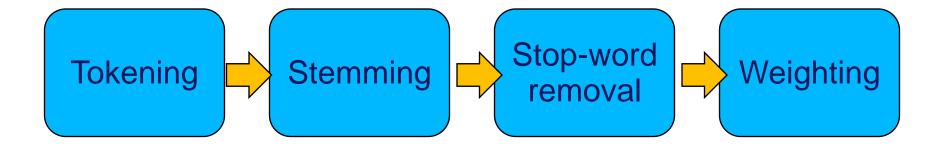




Content-based systems

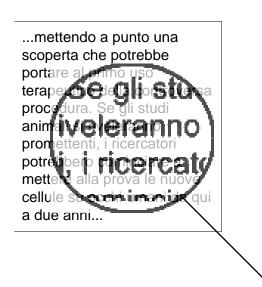


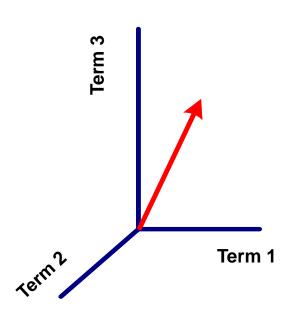
BoW: Bag of Words





Content-based filtering

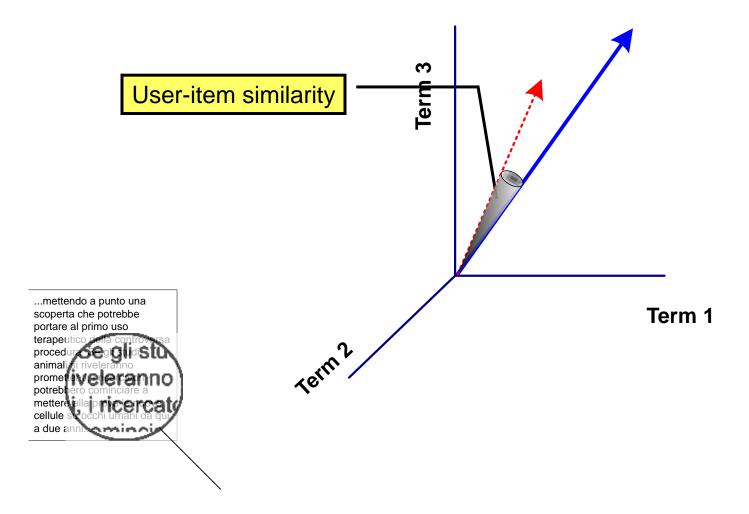




- Similar items contain the same terms (features)
- The more a term occurs in an item, the more representative it is
- The more a term occurs in the collection, the less representative it is (i.e. it is less important in order to distinguish a specific item)

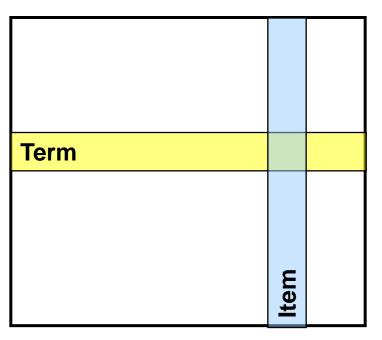


Content-based filtering: prediction





Content-based filtering: Item-Content Matrix



- Each column represents an item, each row a term considered a tag or a keyword indexing a feature of the items
- An element r_{ij} at the *i-th* row and *j-th* column is equal to 1 if the item $\frac{1}{2}$ contains the term $\frac{1}{2}$, 0 otherwise.

- The Item-Content Matrix will then be a nxm matrix where n
 is the number of terms and m the number of items in the
 dataset
- Big and sparse matrix: → matrix factorization



Content-based filtering: pro & cons

Pro:

- No need for data about other users' preferences
- No cold-start or sparsity problems, neither first-rater
 - Able to recommend users with unique tastes
 - Able to recommend new and unpopular items
- Can provide explanations about recommended items
- Well-known technology
- Can integrate explicit (content) preferences

Cons:

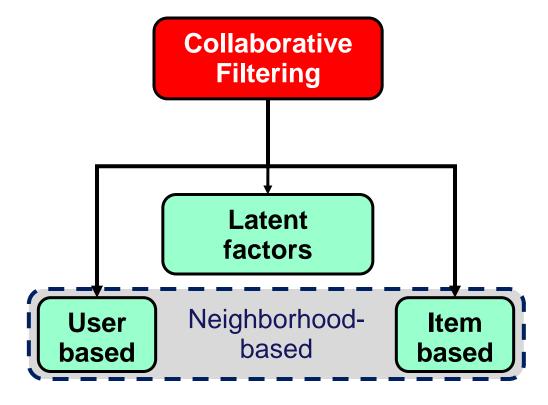
- Requires a structured content
- Lower accuracy (than collaborative)
- Users' tastes must be represented as a function of the content
- Unable to exploit quality judgments of other users
 - Cannot distinguish between a "valuable" item and a "bad" one
- Cannot generate new interests in users



Collaborative systems

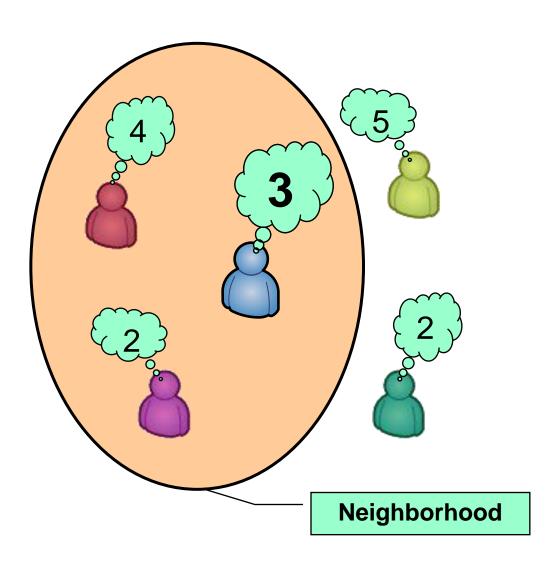
Collaborative

 based on the preferences expressed by other users: it is not needed to extract any explicit features from the items, since recommendations are based only on the community behaviour





Collaborative filtering: user-based

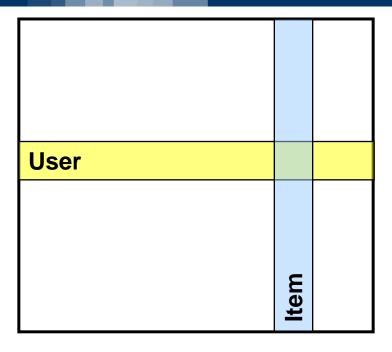


User-based

similar users rate an item similarly: the *neighborhood* is the set of users who are considered similar to a certain user according to such definition



Collaborative filtering: User-Rating Matrix

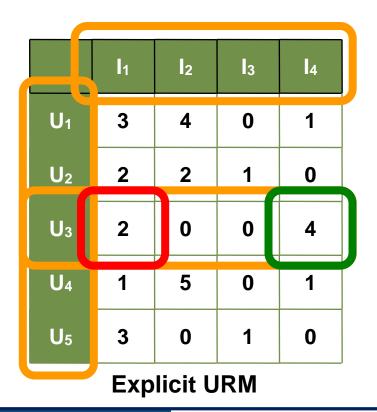


- Each column represents an item, each row a user profile
- An element r_{ij} at the *i-th* row and *j-th* column is the rating of user profile i about item
- An element r_{ij} is null (zero) if the user did not interact with the item
- The *User-Rating Matrix* contains info about the user interaction with the system and is the only info required by collaborative algorithms. It is a **n**x**m** matrix where **n** is the number of users and **m** the number of items in the dataset
- Big and sparse matrix: → factorization



Explicit and implicit URM

• The simplest form of implicit rating is a binary classification: r_{ij} is equal to one if the user interacted with an item, zero otherwise



	I ₁	l ₂	l 3	I 4
U ₁	0	1	0	1
U ₂	0	0	1	1
U ₃	1	1	0	0
U ₄	1	1	0	1

Implicit URM

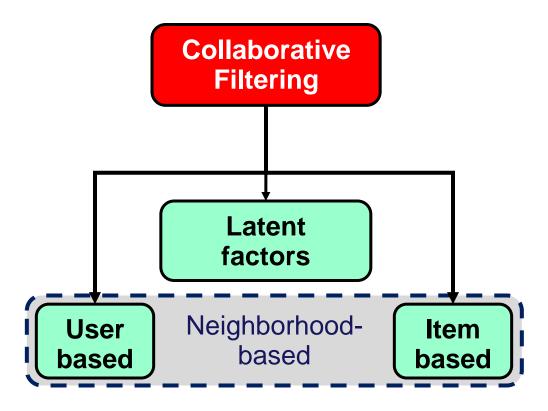


Collaborative Filtering: pro & cons

- Pro:
- There is no need for content
- Cons:
 - Cold Start: we need to have enough users in the system to find a match.
 - Sparsity: when the user/ratings matrix is sparse it is hard to find a neighborhood.
 - First Rater: cannot recommend an item that has not been previously rated by anyone else
 - Popularity Bias: cannot recommend items to someone with unique tastes. Tends to recommend popular items, since they are the most rated (dataset coverage)

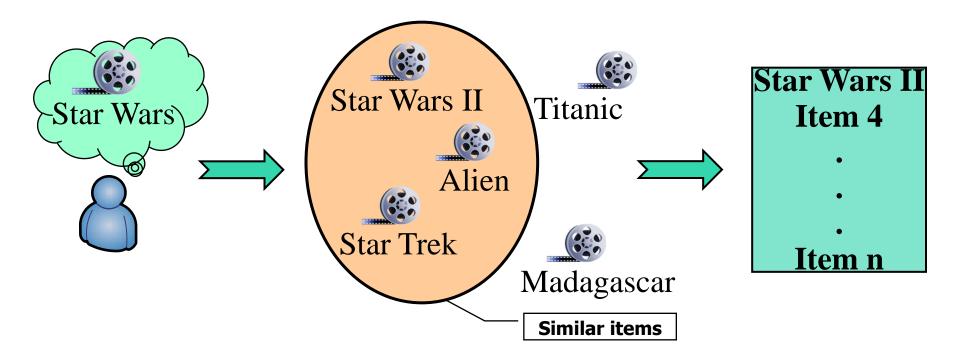


Item-based (neighborhood) collaborative filtering





Item-based collaborative filtering

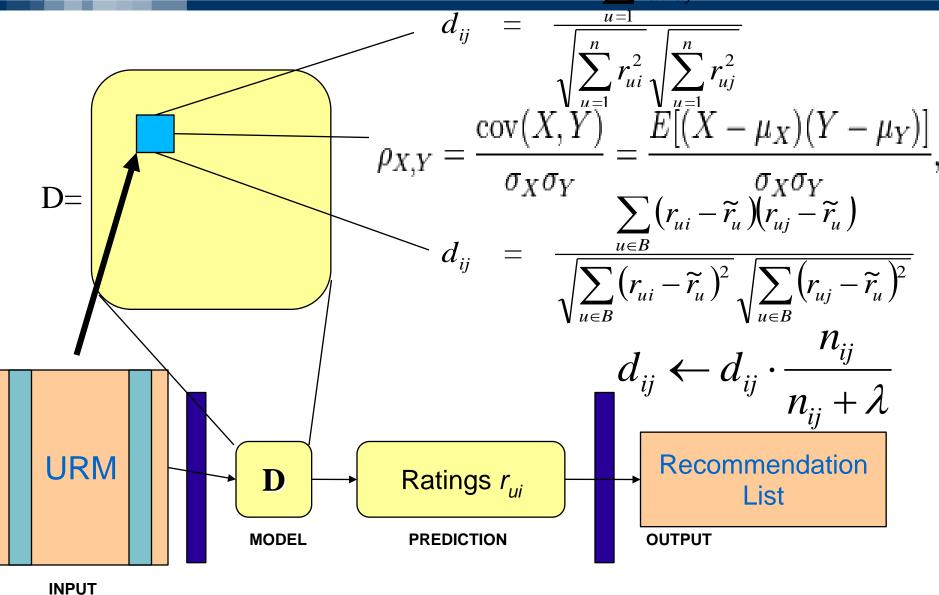


Item-based

similar items are rated by a user similarly: the *neighborhood* is the set of items which are considered similar to a certain item according to the community's ratings

Item-based algorithms







Item-based algorithms

• The estimated rating r_{ui} of an item i for a user u is computed as follows:

$$\hat{r}_{ui} = K + \frac{\sum_{j \in S^k(i;u)} d_{ij}}{Q} \qquad Q = \sum_{j \in S^k(i;u)} (d_{ij})$$

$$\hat{r}_{ui} = b_{ui} + |R^{k}(i;u)|^{-\frac{1}{2}} \sum_{j \in R^{k}(i;u)} w_{ij} (r_{uj} - b_{uj}) + |N^{k}(i;u)|^{-\frac{1}{2}} \sum_{j \in N^{k}(i;u)} c_{uj}$$



Item-based algorithms

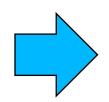
 The model – W, C, b_u, b_i – is learnt by minimizing the error (RMSE) between known ratings and predicted ratings

$$\min \sum (r_{ui} - \hat{r}_{ui})^{2}$$

$$\min \sum (r_{ui} - \hat{r}_{ui})^{2} + \lambda (b_{u}^{2} + b_{i}^{2} + \sum w_{ij}^{2} + \sum c_{ij}^{2})$$

prediction rule

$$\hat{r}_{ui} = b_{ui} + \left| R^k(i;u) \right|^{-\frac{1}{2}} \sum_{j \in R^k(i;u)} w_{ij} (r_{uj} - b_{uj}) + \left| N^k(i;u) \right|^{-\frac{1}{2}} \sum_{j \in N^k(i;u)} c_{uj}$$

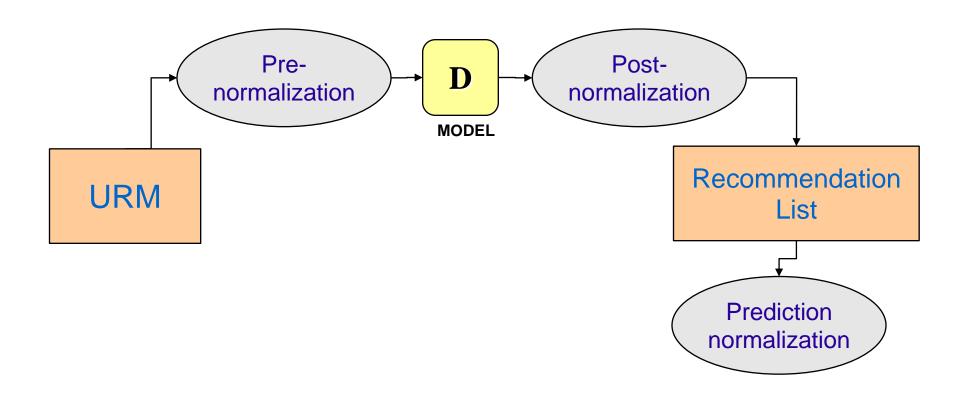


Gradient descent



Observation: Normalization

 The adjusted cosine similarity metric shows a first example of normalization applied to the dataset



NEPTUNY



Item-based algorithms: managing implicit URM

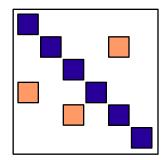
- Note: not all metrics for explicit datasets can be used
 - E.g., adjusted cosine does not work: subtracting the average of all user ratings would reset to zero all the user ratings!
- Similarity metric is usually *frequency-based*
 - For instance, a high similarity between item *i* and item *j* means that if user *u* bought item *i* he will most likely buy also item *j*
- Cosine similarity can still be used, but a post-normalization process is desirable.
 - Furthermore, with regard to implicit ratings, cosine is just a special case of a more general approach that we refer to as Direct Relations (DR)



Item-based algorithms: Direct Relations

 Given the (implicit) URM R, the item-item matrix D used with DR is computed as

$$D = R^{\mathsf{T}} \cdot R$$



- The element d_{ii} on the principal diagonal represents the total number of ratings for item i
- The element d_{ij} represents instead the number of users that have watched both movie i and movie j

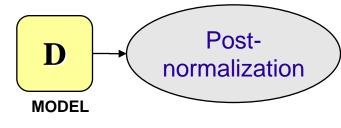
NEPTUNY



Item-based algorithms: post-normalization

 Given the (implicit) URM R, the item-item matrix D used with DR is computed as





General definition of post normalization:

$$d_{ij} \leftarrow \frac{d_{ij}}{d_{ij}^{\beta}d_{jj}^{\gamma}}$$

- β and γ are constant parameters whose optimal value depends on the dataset
- both parameters set to 0.5 corresponds to the cosine

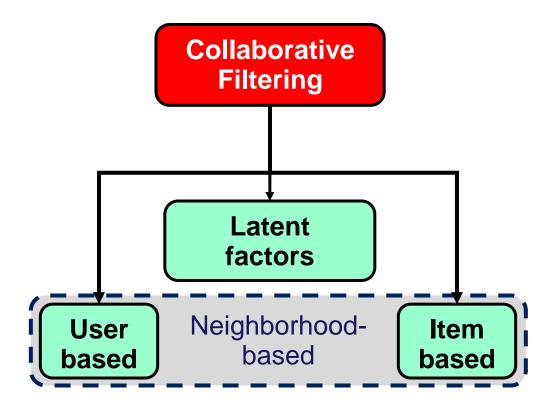


- Introduction
- Recommender systems
 - Content-based and collaborative recommendations
 - Advanced algorithms
- Evaluation of recommender systems

- Case study: a recommender system for IPTV/VOD provider
- Recommender system demo



Collaborative filtering: Latent factor models





Latent factor algorithms

 The dataset (users and items) is described by means of a limited set of f hidden features

Item feature vector

$$\hat{r}_{ui} = p_u^T q_i$$

User feature vector



- Error minimization (Gradient descent)
 - Matrix factorization



Latent factor algorithms: gradient descent

$$\hat{r}_{ui} = b_{ui} + p_u^T q_i \qquad \min \sum (r_{ui} - \hat{r}_{ui})^2 + \lambda (b_u^2 + b_i^2 + \sum |p_u|^2 + \sum |q_i|^2)$$

$$\begin{array}{c} b_{u}=b_{u}+\gamma(e_{ui}-\lambda b_{u})\\ \\ e_{ui}=\hat{r}_{ui}-r_{ui}\\ \\ b_{i}=b_{i}+\gamma(e_{ui}-\lambda b_{i})\\ \\ \vdots\\ \\ p_{u}=p_{u}+\gamma(e_{ui}q_{i}-\lambda p_{u}) \end{array}$$



- **Error minimization** Matrix factorization (gradient descent)



Latent factor algorithms: best NETFLIX algos

$$\hat{r}_{ui} = b_{ui} + p_u^T q_i$$
Roughly person.
$$\hat{r}_{ui} = b_{ui} + q_i^T \left(p_u + |N(u)|^{-\frac{1}{2}} \sum_{j \in N(u)} y_j \right) + \left| R^k(i;u) \right|^{-\frac{1}{2}} \sum_{j \in N^k(i;u)} w_{ij} (r_{uj} - b_{uj}) + \left| N^k(i;u) \right|^{-\frac{1}{2}} \sum_{j \in N^k(i;u)} c_{uj}$$
Integrated Model



- Error minimization (gradient descent)
- Matrix factorization



Latent factor algorithms: ..a step back to model-based

$$\hat{r}_{ui} = b_{ui} + q_i^T \left(p_u + |N(u)|^{-\frac{1}{2}} \sum_{j \in N(u)} y_j \right)$$

$$p_{u} = |R(u)|^{-\frac{1}{2}} \sum_{j \in R(u)} (r_{uj} - b_{uj}) x_{j}$$

AsySVD



- Error minimization (gradient descent)
 - Matrix factorization



Latent factor algorithms: ..SVD

$$\hat{r}_{ui} = p_u^T q_i$$
 R
 $= U$
 V^T

SVD

$$M = U \cdot S \cdot V^T$$


$$U^{\mathsf{T}} \cdot \mathsf{U} = \mathsf{I}$$
$$\mathsf{V}^{\mathsf{T}} \cdot \mathsf{V} = \mathsf{I}$$




- Error minimization (Gradient descent)
- Matrix factorization



Singular Value Decomposition





$$p_u = U \cdot \sqrt{S}$$

$$\hat{r}_{ui} = p_u q_i$$

Item feature vector

$$q_i = \sqrt{S} \cdot V^T$$



PureSVD: some steps over...

$$R = U \cdot S \cdot V^{T}$$
 $R \cdot (V) = U \cdot S \cdot V^{T} \cdot (V)$
 $R \cdot V = U \cdot S$

$$u_u \cdot S = r_u \cdot V$$



PureSVD: some nice properties

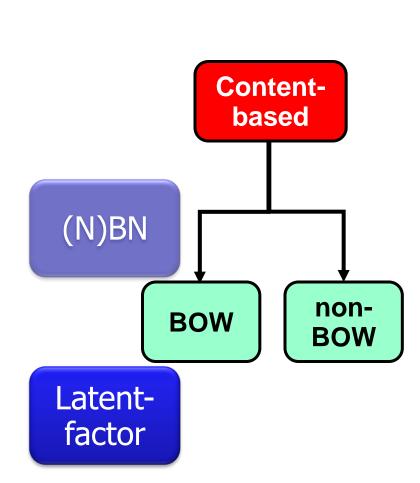
$$\hat{r}_{ui} = r_u \cdot V \cdot v_i^T$$

$$\hat{r}_u = r_u \cdot V \cdot V^T$$

$$u_u \cdot S = r_u \cdot V$$


- V-V^T is a mxm item-item matrix!
- PureSVD is a model-based algorithm!







Content-based Filtering: LSA rank-reduced SVD





Further observations about IPTV recommenders

- Multi-language content
 - (e.g., Switzerland)
- New user problem
 - (user-based algorithms)
- New item problem
 - (all collaborative algorithms)
- Semantic problem
 - (e.g., house and home)



- Introduction
- Recommender systems
 - Content-based and collaborative recommendations
 - Advanced algorithms
- Evaluation of recommender systems

- Case study: a recommender system for IPTV/VOD provider
- Recommender system demo



The multiple facets of recommender system evaluation

- Relevance
 - Error
 - Classification/accuracy
- Confidence
- "Explainability"
- Serendipity
- Diversity
- Novelty
- Coverage
- Training rate



- Relevance
 - Error
 - Classification/accuracy

Many works do not describe clearly the methods used for performance evaluation and model comparison

Different dataset partition methodologies and evaluation metrics lead to divergent results

The Netflix prize has **improperly** focused the research attention on **Hold-out RMSE**



Dataset partitioning

In order to evaluate the capability of a recommender system in suggesting items to users the URM (User Rating Matrix) must be partitioned into two different sets:

- **Training set**: used to generate the model of the system (it can be further divided into a model and a validation set for tuning some parameters of the algorithm)
- **Test set**: used to test the model generated into the training step

IMPORTANT!

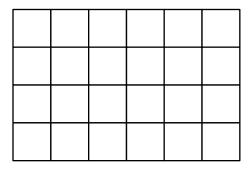
Test set must be different and independent from training set



Partitioning techniques

Different techniques can be used to partition the URM:

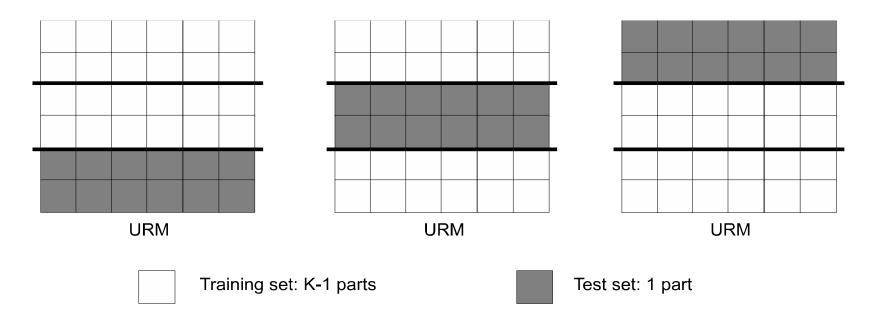
- K-fold cross validation
- Leave-one-out
- Hold-out





K-fold cross validation

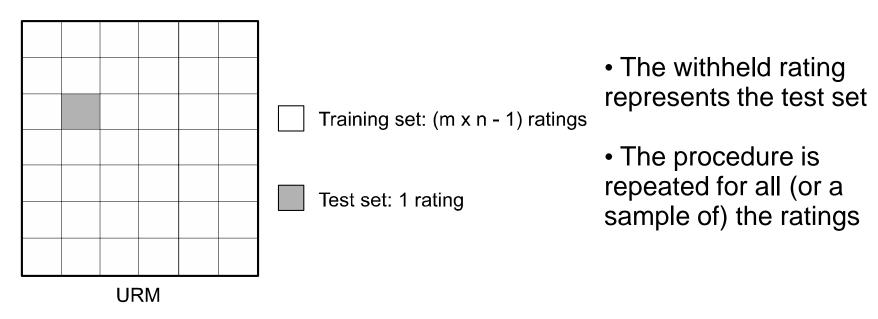
The K-fold cross validation divides the users of the URM into K distinct partitions and at each step K-1 partitions are used to generate the model and the remaining partition si used for the tests.



- The number of partitions suggested to have a robust test is 10.
- The tested users are unknown to the system because they are not used to build the model.



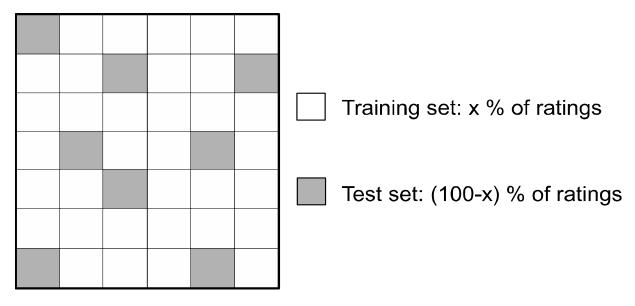
During the construction of the model we withhold one non null user rating of the URM (es: set it to 0), while the remaining ratings are used to generate the model.



- © Leave-one-out is attractive because almost all the ratings are available to build the model
- However it suffers from overfitting problems and it has a high computational complexity



A random set of ratings is withheld from the URM and it is used as test set. The remaining part of the ratings are used as training set.



URM

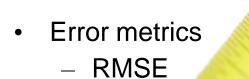
- Can suffer the same overfitting problems as leave-one-out, because the tested users are not totally unknown to the model (especially if the test set is too small).
- Moreover this technique modifies the user profiles and this can lead to erroneous results (particularly perceptible in case of datasets with short user profiles.



RELEVANCE: Evaluation metrics

The **evaluation metrics** are techniques used to evaluate a recommender system and they are so important when we decide which recommender system is better to adopt for a specific application domain.

Quality estimation

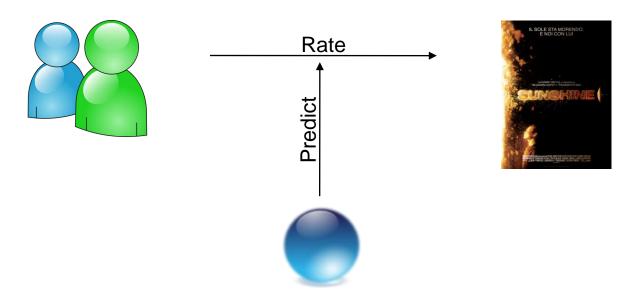


- Classification metrics
 - Recall, precision, f-measure





Error metrics evaluate the error made in the prediction of the rating given by user *u* for item *i*.

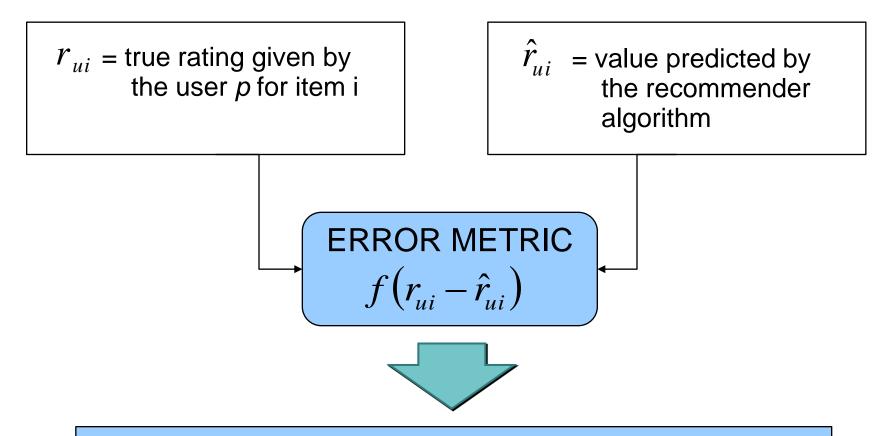


IMPORTANT!

Error metrics can be applied **only** to explicit datasets since in the implicit URM we have no information about the level of preference of users!



Error metrics details



An error metric estimates the difference between r_{ui} and \hat{r}_{ui}



Error metrics details

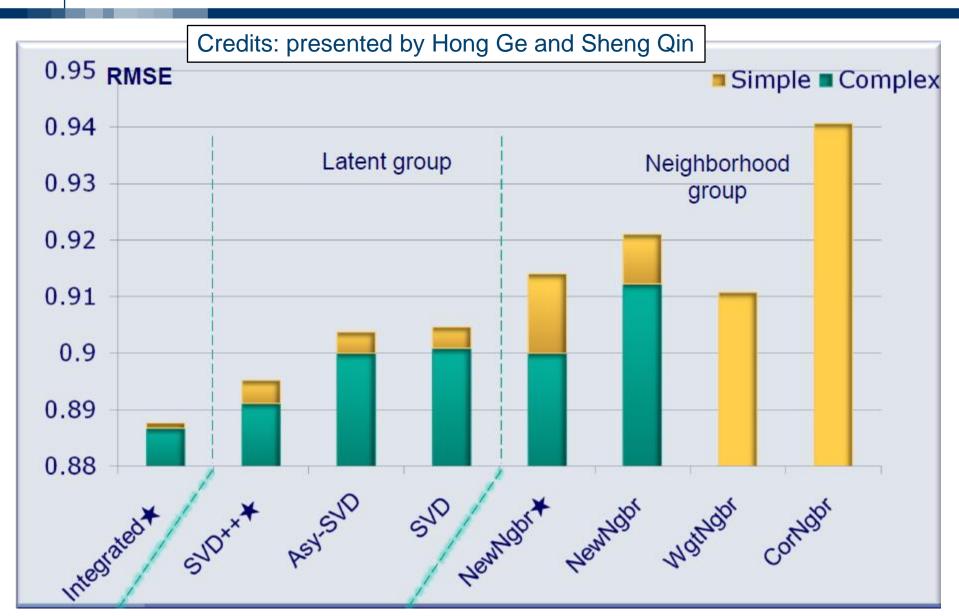
Among error metrics, the most used in the evaluation of reccommender systems are *root mean squared error* (RMSE), *mean squared error* (MSE) and *mean absolute error* (MAE).

$$MSE = \frac{1}{n} \sum_{i=1}^{n} (\hat{r}_{ui} - r_{ui})^2$$
 $MAE = \frac{1}{n} \sum_{i=1}^{n} |\hat{r}_{ui} - r_{ui}|$

$$RMSE = \sqrt{\frac{1}{n} \sum_{u,i} (\hat{r}_{ui} - r_{ui})^2}$$



From "Factorization Meets the Neighborhood: a Multifaceted Collaborative Filtering Model" - Y. Koren, 2008 –





Pro:

© RMSE, MSE and MAE are error metrics easy to compute but...

Cons:

- ...they are suitable only for explicit datasets and...
- :..these metrics can suffer the *outlier problem* so they are sensible to ratings with large prediction errors.

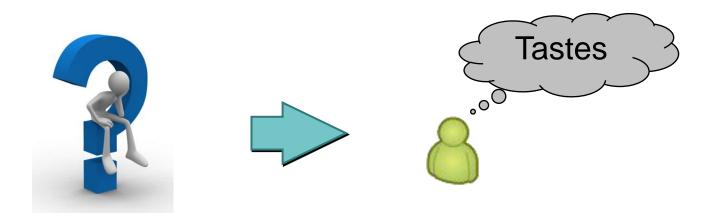
Note: with respect to MSE and RMSE, MAE has a more intuitive interpretation since it represents the average distance between true rating and predicted one.



Accuracy Metrics

Accuracy metrics (also called classification metrics) allow to measure the capacity of recommender systems in suggesting interesting or uninteresting items to the user.

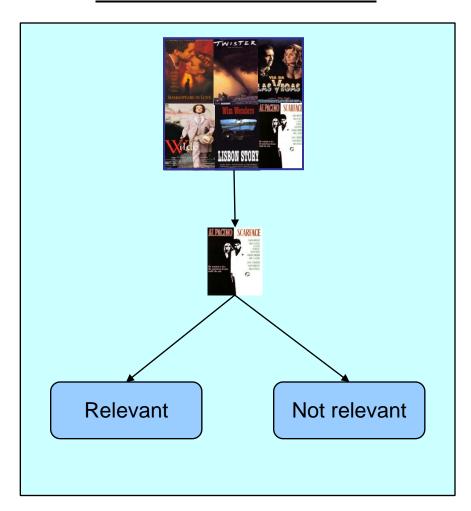
So the purpose of these metrics is not to predict the exact value of the ratings but they allow to understand whether a suggested item will be interesting or not to the user.



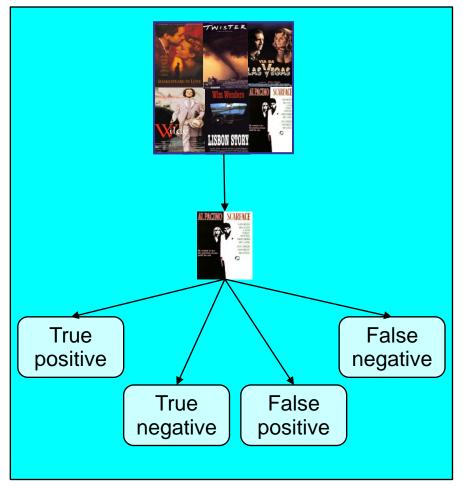


Classification of items into the URM

Classification of items for information retrieval



Classification of items according to accuracy metrics





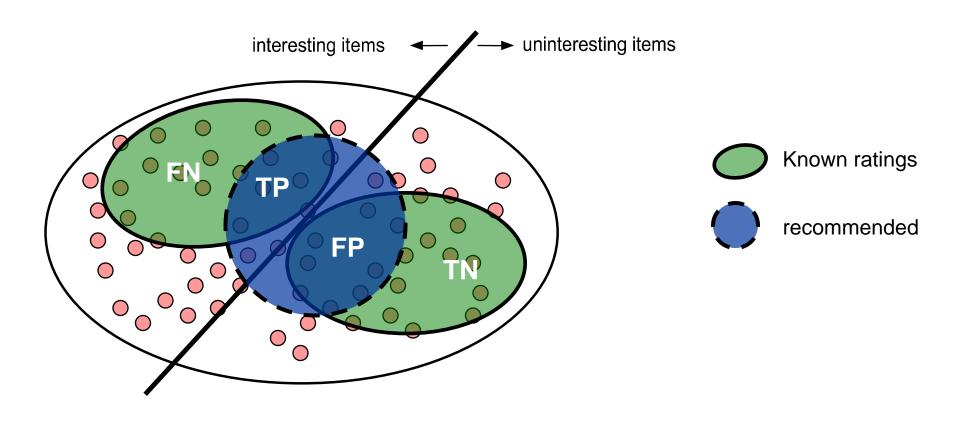
Classification details

- For a given user, items can be classified either in:
 - relevant
 - Non-relevant
- For a given user and a given recommendation rule, items can be classified either in:
 - True positive (TP): the system suggests a relevant item to the user
 - True negative (TN): the system doesn't suggest that the user dislikes
 - False positive (FP): the system suggests an item that the user dislikes
 - False negative (FN): the system doesn't suggest an item interesting for the user

NEPTUNY



Classification metrics





Recall and Fallout

The *recall* metric evaluates the capability of the system in recommending interesting items to the user (this metric is also referred to as true positive rate.

$$T = \frac{TP}{TP + FN}$$

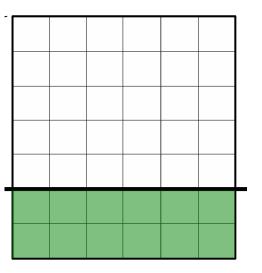
On the other side, *fallout*, also referred to as <u>false positive rate</u>, is a metric complementary to recall and measures how frequently the recommender system suggests non-interesting items to the user.

$$F = \frac{FP}{TN + FP}$$



A possible methodology

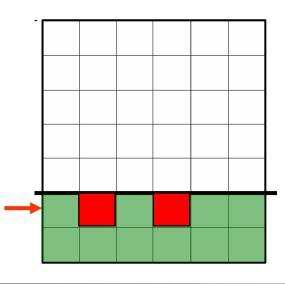
- 10-fold cross validation
 - 9 folds: training set
 - 1 fold: test set
 - Leave-one-out





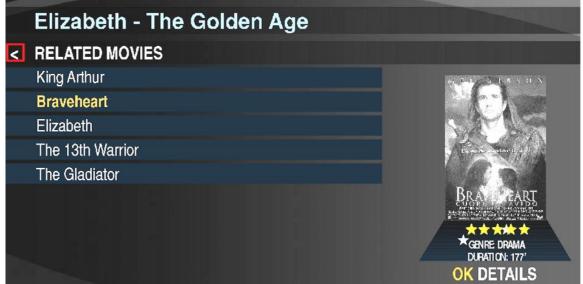


- 10-fold cross validation
 - 9 folds: training set
 - 1 fold: test set
 - Leave-one-out → test item



Test item recommended in top-5?

- Positive rating → TP
- Negative rating → FP





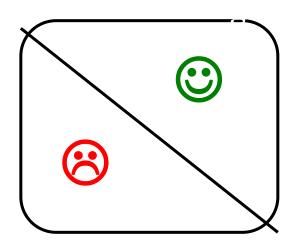


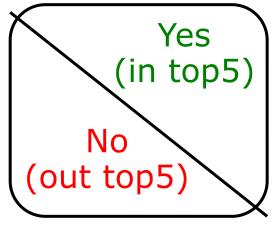
- 10-fold cross validation
 - 9 folds: training set
 - 1 fold: test set
 - Leave-one-out → test item

Test item recommended in top-5?

- Positive rating → TP
- Negative rating → FP

- Metrics
 - Recall
 - Precision
 - F-measure







Accuracy metrics: reformulation


The most popular accuracy metrics are:

- Recall (T)
- Fallout (F)
- Precision (P)

And can be redefined as:

$$T = \frac{TP}{TP + FN} = \frac{\# hits^{+}}{|Test|}$$

$$P = \frac{TP}{TP + FP} = \frac{\# hits^{+}}{|N \cdot |Test|} = \frac{T}{N}$$



$$F = \frac{FP}{TN + FP} = \frac{\# hits^{-}}{|tests|}$$

Where N is the number of recommended items.



The definition of precision is the following:

$$P = \frac{TP}{TP + FP} = \frac{T}{N}$$

...and when applied in the settings of recommender algorithms considers unrated items as non-interesting items.



So we assume that all recommended items for which we have no rating information are to be considered as **non-relevant** and **they count negatively in computing the precision**.

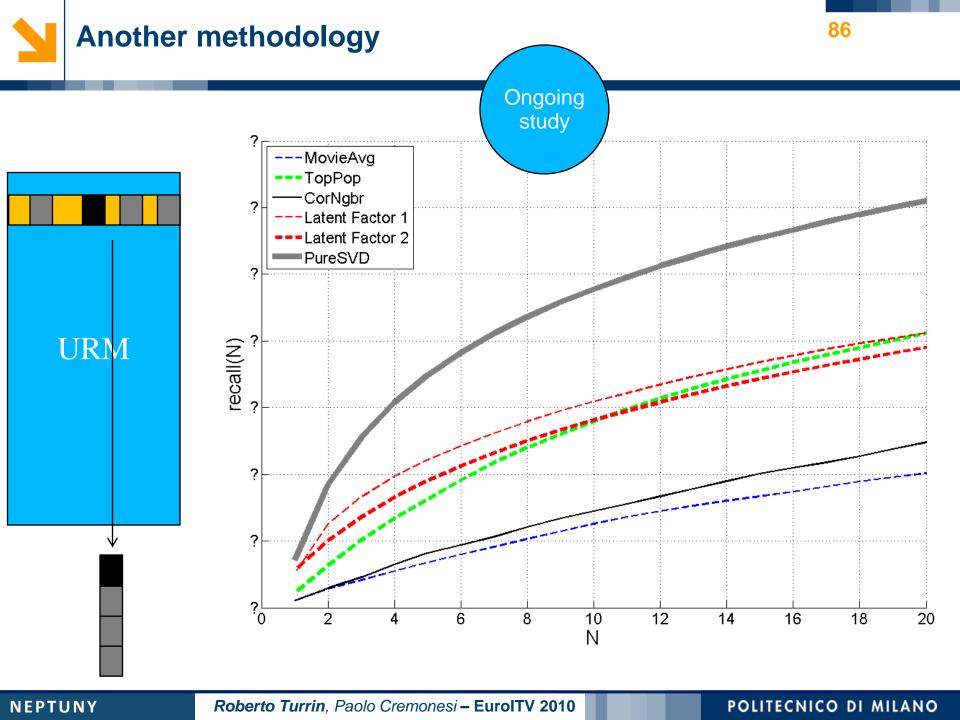


Reducing error in precision metric

- One way to reduce the error in the estimate of the precision is to increase the number of ratings withheld from the user profile.
- → Moving from a Leave-one-out to and hold-out...



This implies reducing the number of ratings used to describe a user which negatively affects the quality of the recommender algorithm.





Recall/fallout and precision: some considerations

- Recall,... are suitable in the evaluation of the global behavior of a recommender algorithm in top-N tasks!
- Both precision and recall heavily depend on the number of rated items per user and, thus, their values should not be interpreted as absolute measures but only to compare different algorithms on the same dataset.
- Measuring precision and recall is difficult because it is often unknown, for a certain user, how many relevant items there exist in the catalogue.
- Any catalogue probably contains a large number of items that meet
 a specific user taste. If that user has rated only a small percentage
 of such items a recommender algorithm might appear to have a
 low *recall*, since the system may recommend unrated, relevant
 items (which are considered as irrelevant).



Netflix dataset: test user profile

Title	Rating
King Kong	5
The Village	5
Pocahontas	5
Men in Black II	5
Little Women	5
The Lord of the Rings: the Return of the king	g 1





Title	Rating
King Kong	5
The Village	5
Pocahontas	5
Men in Black II	5
Little Women	5
The Lord of the Rings: the Return of the king	, 1

Title

The Lord of the Rings: the Fellowship ... Ring The Lord of the Rings: the Two Towers

Lost: Season 1

The Shawshank Redemption

Arrested Development

RMSE: 0.95 Recall: 1%

F-measure: 0.01



Netflix dataset: Item-based neighborhood

Title	Rating
King Kong	5
The Village	5
Pocahontas	5
Men in Black II	5
Little Women	5
The Lord of the Rings: the Return of the king	g 1

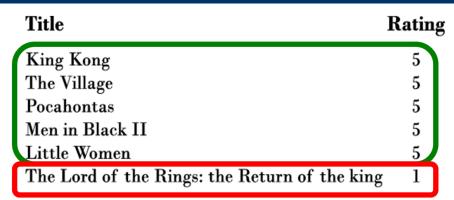
Title

Dinosaur Planet
Isle of Man TT 2004 Review
Character
Paula Abdul's Get Up & Dance
The Rise and Fall of ECW

RMSE: 1.6 Recall: 8%

F-measure: 0.16





Title

I Robot
Independence Day
Men in Black
Harry Potter and the Prisoner of Azkaban
The Day After Tomorrow

RMSE: 2.7 Recall: 17% F-measure: 0.28



Metrics comparison summary

Error metrics

- Mean Square Error (MSE)
- Root Mean Square Error (RMSE)
- Mean Absolute Error (MAE)
- Only for explicit datasets
- Top-N recommender systems

Accuracy metrics

- Recall
- Precision
- Fallout

Both implicit and explicit datasets

There is NOT a "universal" metric perfect for all application domains but we must choose the most suitable one for each situation.

User-perceived vs algorithmic quality

93

88 screened subjects

JUST
PRELIMINARY
RESULTS!
(with few
subjects)

USER-PERCEIVED QUALITY

VS.

ALGORITHMIC METRICS

		Statistical matrice		User-centered metrics		
Algorithm	Profile	Statistical metrics		Mean feedback		
Algorium	length	RMSE	Recall	All	Novel	Serendipity
				movies	movies	
Collaborative	Short	1.76	5.71	3.63	2.55	14%
Conaborative	Long	1.73	5.13	3.77	2.77	10%
Content-based	Short	2.47	3.77	3.17	2.75	46%
Comem-vasea	Long	2.42	2.33	3.29	2.65	36%



- Introduction
- Recommender systems
 - Content-based and collaborative recommendations
 - Advanced algorithms
- Evaluation of recommender systems

- Case study: a recommender system for IPTV/VOD provider
- Recommender system demo

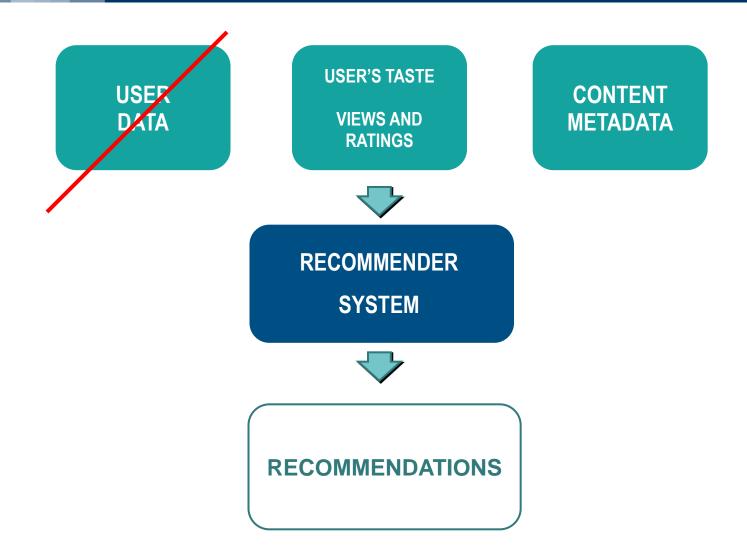


Large-scale European IPTV/VOD provider

- IP-based broadband TV (2001)
 - Hundreds of thousands of customers
 - Catalog of thousands of multimedia contents
- Recommender system (Content Wise) integrated
 - Small list of recommended items
 - Screen definition
 - Reduced navigation capabilities
 - Real-time constraints
 - Scalability (30000 recommendations per day)
 - Live boradcast channels
 - User identification
 - Quality of content information




Recommender System





Model-based architecture





Proposed approach

Batch system

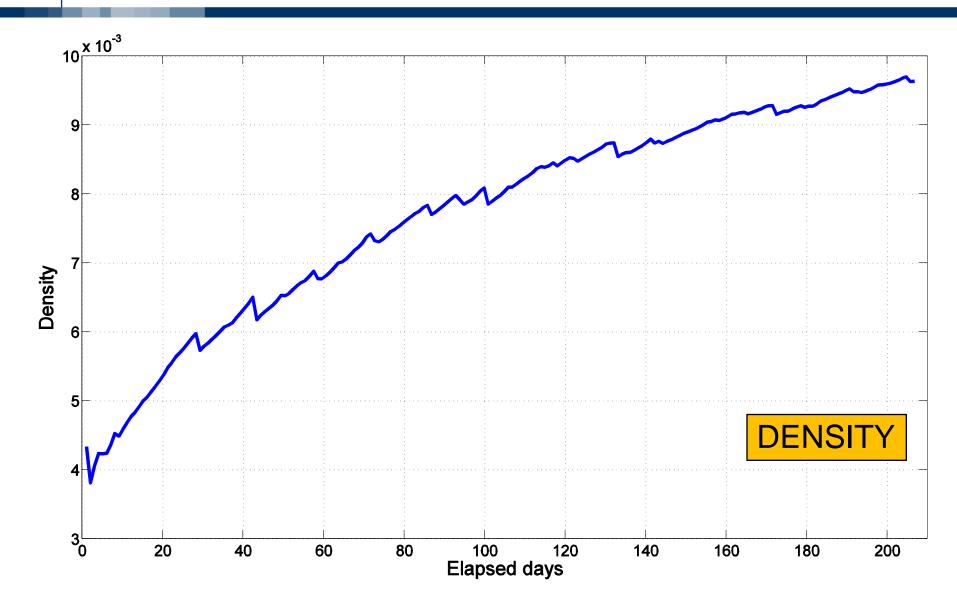
- Statistical analysis of the dataset
- Definition of a number of models
- Accuracy evaluation for different user profiles

Run-time system

- User profile analysis
- Selection of best candidate model
- Recommendation

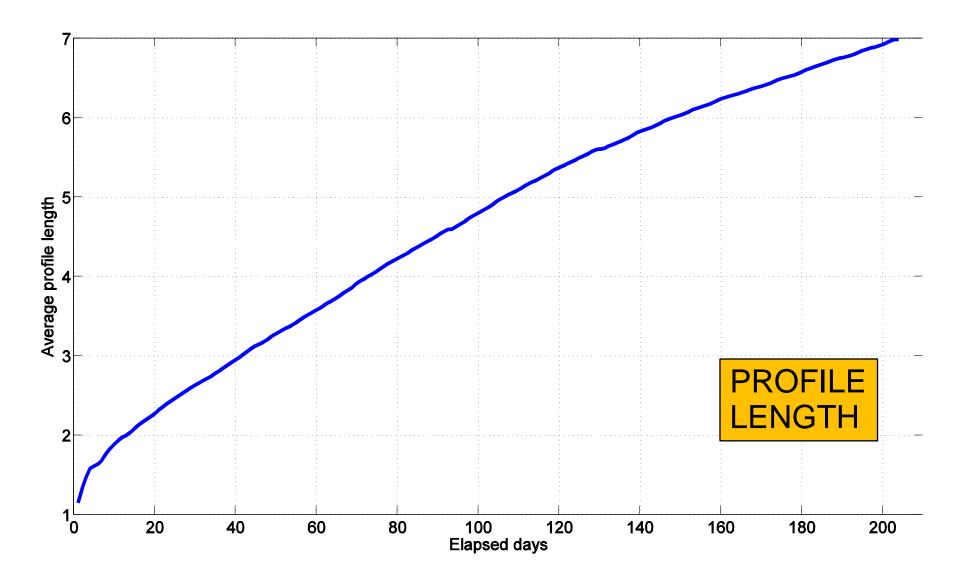


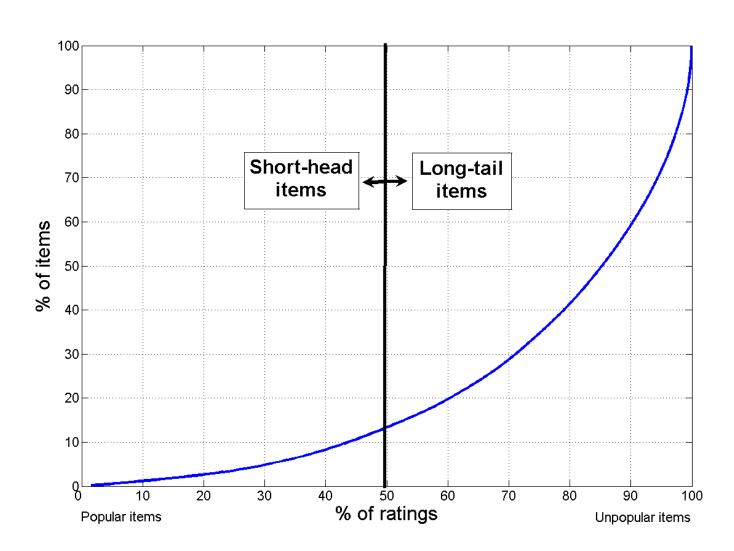
Some data...













Recommendation: relevance on VoD

Algorithm	Darameter	Recall		
Algorithm	Parameter	3 months	6 months	
	k = 10	16.8%	14.9%	
Item-based-CF	k = 50	18.7%	16.4%	
Heili-based-Cr	k = 100	19.0%	16.6%	
	k = 150	18.8%	16.5%	
	l = 5	15.1%	12.7%	
SVD-CF	I = 15	12.6%	13.3%	
SVD-CF	l = 25	10.9%	11.5%	
	I = 50	9.3%	9.9%	
	I = 100	6.3%	8.0%	
	l = 50	1.9%	1.7%	
LSA-CB	I = 100	2.3%	2.3%	
	I = 150	2.4%	2.4%	
	I = 200	2.5%	2.5%	
Top-rated		12.2%	7.7%	



Recommendation: relevance on VoD

		Recall non-top-10		Recall non-top-50%	
Algorithm	Parameter	3 months	6 months	3 months	6 months
	k = 10	14.0%	13.2%	7.7%	9.6%
Item-based-CF	k = 50	14.0%	13.8%	6.8%	9.0%
nem-based-Cr	k = 100	13.8%	13.5%	6.2%	8.3%
	k = 150	13.5%	13.2%	6.1%	7.9%
	l = 5	6.6%	6.8%	0.7%	1.4%
SVD-CF	I = 15	11.5%	10.2%	1.2%	3.5%
SVD-CF	l = 25	12.6%	12.0%	2.2%	4.9%
	I = 50	11.4%	11.2%	4.8%	7.8%
	I = 100	7.6%	9.3%	9.8%	11.8%
	l = 50	2.1%	1.8%	1.8%	1.7%
LSA-CB	I = 100	2.3%	2.3%	2.0%	2.5%
	I = 150	2.5%	2.5%	2.1%	2.5%
	I = 200	2.6%	2.6%	2.2%	2.6%
Top-rated		0.4%	1.0%	0%	0%



Recommendation: relevance on VoD

	2 hours	24 hours	7 days
All	17.0%	19.8%	24.7%
Top 10	5.1%	7.0%	10.6%
Non-top 10	24.2%	27.6%	32.1%
Top 50%	9.4%	11.5%	16.2%
Non-top 50%	28.4%	32.2%	36.1%



- Introduction
- Recommender systems
 - Content-based and collaborative recommendations
 - Advanced algorithms
- Evaluation of recommender systems

- Case study: a recommender system for IPTV/VOD provider
- Recommender system demo

- P.Cremonesi, Y. Koren, and R.Turrin, "Performance of Recommender Algorithms on Top-N Recommendation Tasks".
 To appear on Proc. of the 4th International Conference on Recommender Systems (RecSys), ACM, 2010, Barcelona, Spain.
- P.Cremonesi, and R.Turrin, "Time-evolution of IPTV Recommender Systems".
 Proc. of the 8th European Conference on Interactive TV and Video, ACM, 2010, Tampere, Finland.
- R. Bambini, P.Cremonesi, and R.Turrin, "A Recommender System for an IPTV Service Provider: a Large-Scale Production Environment".
 Book Chapter on Recommender Systems Handbook (P.B.Kantor, F.Ricci, L.Rokach, and B. Shapira), Springer, 2010.
- Herlocker, J.; Konstan, J.; Terveen, L. & Riedl, J. "Evaluating collaborative filtering recommender systems".
 ACM Transactions on Information Systems (TOIS), ACM Press New York, NY, USA, 2004, 22, 5-53
- Adomavicius, G. & Tuzhilin, A. "Toward the next generation of recommender systems: a survey
 of the state-of-the-art and possible extensions".
 Knowledge and Data Engineering, IEEE Transactions on, 2005, 17, 734-749



- Adomavicius, G. & Tuzhilin, A. "Toward the next generation of recommender systems: a survey
 of the state-of-the-art and possible extensions".
 Knowledge and Data Engineering, IEEE Transactions on, 2005, 17, 734-749
- P.Cremonesi, and R.Turrin, "Analysis of cold-start recommendations in IPTV systems".
 Proc. of the 3rd International Conference on Recommender Systems (RecSys), ACM, 2009, New York, USA
- P.Cremonesi, E. Lentini, M. Matteucci, and R.Turrin, "An evaluation methodology for recommender systems".
 Proc. of the 4th International Conference on Automated Solutions for Cross Media Content and Multi-channel Distribution (AXMEDIS), IEEE, 2008, Firenze, Italy.
- Bell, R. M. & Koren, Y."Scalable Collaborative Filtering with Jointly Derived Neighborhood Interpolation Weights".
 7th IEEE Int. Conf. on Data Mining, IEEE International Conference on Data Mining (ICDM'07), 2007, 43-52
- Koren, Y. "Factorization meets the neighborhood: a multifaceted collaborative filtering model".
 KDD '08: Proceeding of the 14th ACM SIGKDD international conference on Knowledge discovery and data mining, ACM, 2008, 426-434
- E. Campochiaro, R. Casatta, P.Cremonesi, and R.Turrin, "Do metrics make recommender algorithms?".
 Proc. of the International Symposium on Mining and Web (MAW), IEEE, 2009, Bradford, UK.