



---

# *Software Lab*

*Synchronization, IPC and ItC*

Roberto Farina  
roberto.farina@cefriel.it



- Synchronization and ItC
  - ▶ Global variables
  - ▶ Mutex
  - ▶ Semaphores
  - ▶ Condition variables
- Synchronization and IPC
  - ▶ Signals
  - ▶ Shared memory
  - ▶ Semaphores
  - ▶ Mapped memory
  - ▶ Pipes
  - ▶ Sockets

# Race condition



- Most of bugs in multithread programs happens because **different** threads access **same** data
  - ▶ The critical operation is the modification
- **Race condition**
  - ▶ “Race to be the first who modifies data”
  - ▶ Correctness depends on the thread scheduling
  - ▶ Typical errors:
    - Different threads do the same job
    - Segmentation fault
- Solution: make operations **atomic**



- **MUT**ual **EX**clusion locks
- Special “lock” that allows a single thread to access a resource in mutual exclusion
  - ▶ E.g.: bathroom door
- **BLOCKING** mechanism
  - ▶ If a thread tries to “close” an already closed lock it goes in a blocked state until the lock is reopened



- Use the data type `pthread_mutex_t`
- Pass a pointer to `pthread_mutex_init`
  - ▶ The second argument is a pointer to a mutex attribute object; NULL for default values
- It is possible to use `PTHREAD_MUTEX_INITIALIZER`
- `pthread_mutex_lock` to acquire the access right
  - ▶ BLOCKING if the mutex is already locked
- `pthread_mutex_unlock` to unlock the mutex waking up blocked threads
  - ▶ It must be invoked by the same thread which acquired the lock



- Every lock must have the corresponding unlock
  - ▶ If a thread terminates before unlocking a mutex, it will remain blocked
- **Deadlock**
  - ▶ One or more threads are waiting for an event that will never occur
- It is possible to use non blocking calls
  - ▶ `pthread_mutex_trylock` on a blocked mutex returns **EBUSY**

# Three mutex types



- **Fast** mutex (default)
  - ▶ Can generate deadlock if the thread tries to lock a mutex it has already blocked
- **Recursive** mutex
  - ▶ It keeps track of the number of locks
  - ▶ It is necessary to execute the same number of unlock to free the mutex
- **Error-checking** mutex
  - ▶ A second call to `pthread_mutex_lock` returns **EDEADLK**

# Changing mutex type



- Use the `pthread_mutexattr_t` object
- Initialize it with `pthread_mutexattr_init`
- `pthread_mutexattr_setkind_np` to change the mutex type
  - ▶ `PTHREAD_MUTEX_RECURSIVE_NP`
  - ▶ `PTHREAD_MUTEX_ERRORCHECK_NP`
  - ▶ `pthread_mutexattr_settype`
- Invoke `pthread_mutex_init` with the pointer to the mutex attribute object
- `pthread_mutexattr_destroy` to destroy the object
- NP means Not Portable: GNU/Linux specific!



# Thread semaphores



- **Counter** used for thread synchronization
- It is granted that reading and writing a semaphore value is secure
- Two basic operations
  - ▶ **Wait**: decrease the semaphore value by 1. If the value is 0, the operation is BLOCKING
  - ▶ **Post**: increase the semaphore value by 1, causing the awakening of a blocked thread if any
- Two implementations
  - ▶ POSIX for thread communication
  - ▶ Specific for process communication

# Semaphore usage



- Use the data type `sem_t`
- Initialize with `sem_init`
  - ▶ The second parameter should be 0
    - GNU/Linux doesn't support sharing of this kind of semaphore
  - ▶ The third argument is the initial value
- `sem_wait` to wait for a semaphore
  - ▶ `sem_trywait` returns `EAGAIN` if the value is 0
- `sem_post` to increase the semaphore value
- `sem_getvalue` to read the semaphore value
  - ▶ Don't use the value to take decisions



- A synchronization mechanism to implement more complex execution conditions
  - ▶ When the condition is false, the thread doesn't poll but transits in a blocked state
  - ▶ When the condition becomes true, the thread is awakened
- Similar to semaphores in the waiting mechanism
- Different from semaphores because **it doesn't have memory**
  - ▶ If the condition is signaled but there's no thread waiting the signal is lost



- The thread checks the execution flag
  - ▶ If it is not set, waits for the c.v.
- The thread which sets the c.v., modifies the flag and then it signals the new condition
- Problem: **race condition**
  - ▶ If the first thread is interrupted by the scheduler during control...
- Solution: lock on the flag and the c.v. with a single mutex
  - ▶ In GNU/Linux, every c.v. is used in conjunction with a mutex



- Mutex lock and check of the value
- If it is set, mutex unlock and execution
- If it is not set, mutex unlock and wait for the execution condition
  - ▶ Possible problem if those operations are not performed **atomically**
    - Another thread could change the flag value and signal the condition between the test and the wait
    - GNU/Linux allows to perform this step atomically

# Condition variables



- Use the data type `pthread_cond_t`
  - ▶ In conjunction with `pthread_mutex_t`
- Initialization with `pthread_cond_init`
  - ▶ The second argument is ignored in GNU/Linux
- `pthread_cond_signal` to signal the condition
  - ▶ `pthread_cond_broadcast` to unblock EVERY thread that is waiting for the specific c.v.
- `pthread_cond_wait` to wait a c.v.
  - ▶ The second argument is the mutex to unlock
  - ▶ When the condition is signaled, the lock is gained again



- Synchronization and ItC
  - ▶ Global variables
  - ▶ Mutex
  - ▶ Semaphores
  - ▶ Condition variables
- Synchronization and IPC
  - ▶ Signals
  - ▶ Shared memory
  - ▶ Semaphores
  - ▶ Mapped memory
  - ▶ Pipes
  - ▶ Sockets



- **sigaction** alternative: use the **signal** system call to install a new signal handler
  - ▶ **SIG\_IGN**, **SIG\_DFL**, specific handler
  - ▶ Returns a pointer to the old handler or **SIG\_ERR**
  - ▶ Non standard behavior
    - In some cases the handler is reset to the default value
- It is possible to mask nested signals
  - ▶ **sigemptyset**, **sigaddset**, **sigdelset**
    - Work on the **sa\_mask** field of **sigaction** (**sigset\_t**)
- **pause** to wait for a signal



# Shared memory



- It is possible to share a memory region between processes
- Upon creation, a process allocates the segment, other processes open it (`attach`)
  - ▶ Trying to allocate an existing segment results in a reference to it
  - ▶ Segments are integer multiples of the page dimension defined by the system
    - In Linux usually it is 4KB, `getpagesize()`
- Upon deletion, every process closes the segment (`detach`), a single process has to deallocate it

# Creation of a segment



- **shmget** to allocate it (**SH**ared **M**emory **GET**)
  - ▶ It is necessary to specify an **integer key** to be able to access it
    - **IPC\_PRIVATE** to grant unique keys
  - ▶ Returns the **shmid**
  - ▶ It is necessary to specify the segment dimension
    - Automatically rounded (paging)
  - ▶ The third argument is a bitwise OR of options:
    - **IPC\_CREAT** to create the segment
    - **IPC\_EXCL** to abort if the segment already exists
    - **Mode flags**: access rights similar to those for files
      - **S\_IRUSR, S\_IWUSR, ...** (**sys/stat.h**)
      - Exec rights ignored

# Attach and detach



- **shmat** (**SH**ared **M**emory **AT**tach)
  - ▶ It is necessary to specify the id key
  - ▶ Pointer to specify where to map the segment
    - NULL to let Linux choose an available one
  - ▶ Third argument
    - **SHM\_RND** to automatically align with the page
    - **SHM\_RDONLY** to indicate it will be only read
  - ▶ Returns a pointer to the segment
- **shmdt** (**SH**ared **M**emory **DeT**ach) to close
  - ▶ Automatic when calling **exit** or **exec**
  - ▶ The segment is removed if it is already deallocated and the last process closes it



- **shmctl** (SHared Memory ConTrol)
  - ▶ To obtain information use **IPC\_STAT** as the second param and a pointer to **shmid\_ds** as the third one
  - ▶ To remove the segment use **IPC\_RMID** as second param and NULL as the third one
    - Removed when the last process requests detach
- Segments must be explicitly deallocated to avoid reaching the max number of segments
  - ▶ **ipcs -m** to view shared memory segments
  - ▶ **ipcrm shm <id>** to remove a segment

# Shared memory



- Probably the quickest and simplest IPC mechanism
- Suitable for bidirectional communication
  - ▶ Risk of **race conditions**: an “access protocol” must be established
    - Exclusive access not granted even with **IPC\_PRIVATE**
- The problem of id key exchange...

# Process semaphores



- Also called System V semaphores
- They are organized in sets
- **semget** to allocate them
  - ▶ Identification key
  - ▶ Number of semaphores in the set
  - ▶ Flags
    - **IPC\_CREAT**, **IPC\_EXCL**, ...
- **semctl** to control and remove them
  - ▶ It is necessary to explicitly deallocate them to not saturate the system
    - Differently from the shared memory, semaphores are immediately deallocated



- It is necessary to define a `union semun`
- Invoke `semctl`
  - ▶ `semid` as the first parameter
  - ▶ 0 as the second parameter
  - ▶ `SETALL` as the third one
  - ▶ Fourth parameter: create a `union semun` with the `array` field referring to an array of `unsigned short`

# Operations on semaphores



- **semop** system call for both wait and post operations
- It is necessary to create a vector of **sembuf** structures
  - ▶ **sem\_num**: semaphore number
  - ▶ **sem\_op**: integer that specifies the operation on the semaphore
    - If positive, it is added to the current value
    - If negative, the absolute value is subtracted
      - If this operation would cause the semaphore value to become negative, the process is blocked until the semaphore value allows the operation to be performed



# Operations on semaphores



- `sem_flag`
  - ▶ `IPC_NOWAIT` to avoid to be blocked
    - If the call will take the process to a blocked state, it is not performed and an error is returned
  - ▶ `SEM_UNDO` to let Linux undo every operation on the semaphore when the process terminates
    - If the process terminates (cleanly or not) the semaphore value is restored undoing the effects of the process
- `ipcs -s` for info on semaphore sets
- `icpru sem <id>` to remove a set

# Mapped memory



- It allows the communication through a **shared file**
- Different from shared memory
  - ▶ Used both as IPC and to access the content of a file
- **Association** between a file and a process' memory
  - ▶ The content is divided in chunks with dimension equal to memory pages and it is loaded in virtual memory
  - ▶ Memory read and write operations
    - The OS handles read and write operation on the file transparently

# mmap



- **mmap** (Memory MAPped) to make the mapping
  - ▶ `#include <sys/mman.h>`
  - ▶ First parameter: address to map the file to
    - NULL lets Linux choose the first one available
  - ▶ Second parameter: mapping length in bytes
  - ▶ Third parameter: protection flag
    - `PROT_READ`, `PROT_WRITE`, `PROT_EXEC`, `PROT_NONE`
  - ▶ Fourth parameter: additional options
  - ▶ Fifth parameter: opened file descriptor
  - ▶ Sixth parameter: offset from the beginning of the file to which make the map
  - ▶ It returns the mapping address or `MAP_FAILED`



- It is necessary to remove the mapping: **munmap**
- Two parameters:
  - ▶ Mapping starting address
  - ▶ Section length
- Further references to addresses within the range generate “invalid reference” (**SIGSEGV**)
- Automatic unmapping when the process terminates
- **NO** automatic unmapping when closing the file descriptor

# mmap: additional options



- **MAP\_FIXED**

- ▶ Use the specified address for the mapping
- ▶ It must be page-aligned

- **MAP\_PRIVATE**

- ▶ Write operations will be performed on a private copy of the file, not on the original one (copy-on-write)
- ▶ Other processes are not aware of modifications
- ▶ To not be used with MAP\_SHARED

- **MAP\_SHARED**

- ▶ Shares the mapping with other processes that map the object
- ▶ Write operations are performed on the original file (**msync** or **munmap** to be sure the file is updated)

# mmap: additional options



- Mapped memory with **mmap** is preserved upon a **fork**, with same attributes
- Linux supports its own flags, not just POSIX ones
  - ▶ **MAP\_DENYWRITE**
    - Write operations fail with **ETXTBUSY** (DoS attacks)
  - ▶ **MAP\_ANONYMOUS** (**MAP\_ANON**, deprecated)
    - Not associated to a file, fd and offset params are ignored
  - ▶ **MAP\_32BIT**
    - Mapping in the first 2GB of the process address space
  - ▶ **MAP\_LOCKED**
    - Lock the pages of the mapped region into memory (**mlock**)



- Linux could buffer write operations
- **msync** forces the flush on the file
  - ▶ First two parameters similar to **munmap**
  - ▶ Third parameter:
    - **MS\_ASYNC**: the update is scheduled but not necessarily executed before the function terminates
    - **MS\_SYNC**: immediate update, blocking
    - **MS\_INVALIDATE**: invalidation request for other mappings to update with fresh values
  - ▶ It returns **EINVAL** or **ENOMEM** in case of errors



- It is necessary to establish an access protocol to avoid race conditions (semaphore)
- Not only for IPC
  - ▶ Data structures (**struct**) in mapped memory
  - ▶ When program terminates there's a file with all the data structures
  - ▶ When the program starts again, data structures are immediately available remapping the file
    - It is necessary to map to the same address to not invalidate pointers
  - ▶ Mapping of special files (**/dev/zero**)





- UNIDIRECTIONAL and SERIAL communication mechanism
- Usually used between threads or related processes
- | symbol in the shell
  - ▶ It connects the stdout of the first process to the stdin of the second one
- Limited memory
  - ▶ Writing on a full pipe or reading from an empty pipe results in a blocked state
    - Automatic synchronization



- Invoke `pipe`
- The argument is an integer array of dimension 2
  - ▶ 0: reading file descriptor
  - ▶ 1: writing file descriptor
  - ▶ Data written in 1 are read through 0
  - ▶ The descriptors created this way are valid in the current process and in its sons
- `fdopen` to convert file descriptors in `FILE*` in order to use the high level API (`printf`, `fgets`)



- It is possible to redirect stdin, stdout and stderr streams of a program
- Call to `dup2`
  - ▶ Identical file descriptors share the SAME position within the file and the same flags
    - They don't share the `close-on-exec` flag
  - ▶ `STDIN_FILENO`, `STDOUT_FILENO`, `STDERR_FILENO`
    - `dup2 (fd, STDIN_FILENO)` to redirect the standard input

# popen and pclose



- To avoid to call **pipe**, **fork**, **dup2**, **exec**, **fdopen**
- **popen** creates a pipe, executes the **fork** and invokes the shell
  - ▶ Command to be executed
  - ▶ Operation (read, "**r**", or write, "**w**")
  - ▶ It returns one end of the pipe or NULL
    - The other end is connected to the stdin of the child process
- **pclose** waits for child termination
  - ▶ It returns the exit status



- Pipe with a name in the filesystem (**named pipes**)
- Communication between unrelated processes
- **mkfifo** command for the creation
  - ▶ Can be cancelled as a normal file
- **mkfifo** function
  - ▶ **sys/types.h**, **sys/stat.h**
  - ▶ First parameter: creation path
  - ▶ Second parameter: rights
  - ▶ It returns -1 if the pipe cannot be created

# FIFO usage



- They are used as normal files
  - ▶ A process opens it for writing, the other for reading
  - ▶ Both low level primitives (`open`, `close`, ...) and high level ones (`fopen`, `fprintf`, `fscanf`, ...) can be used
- More than one process as reader/writer
- Data are atomically written up to `PIPE_BUF`
  - ▶ 4KB in Linux
  - ▶ Simultaneous read/write operations can result in interleaving sections

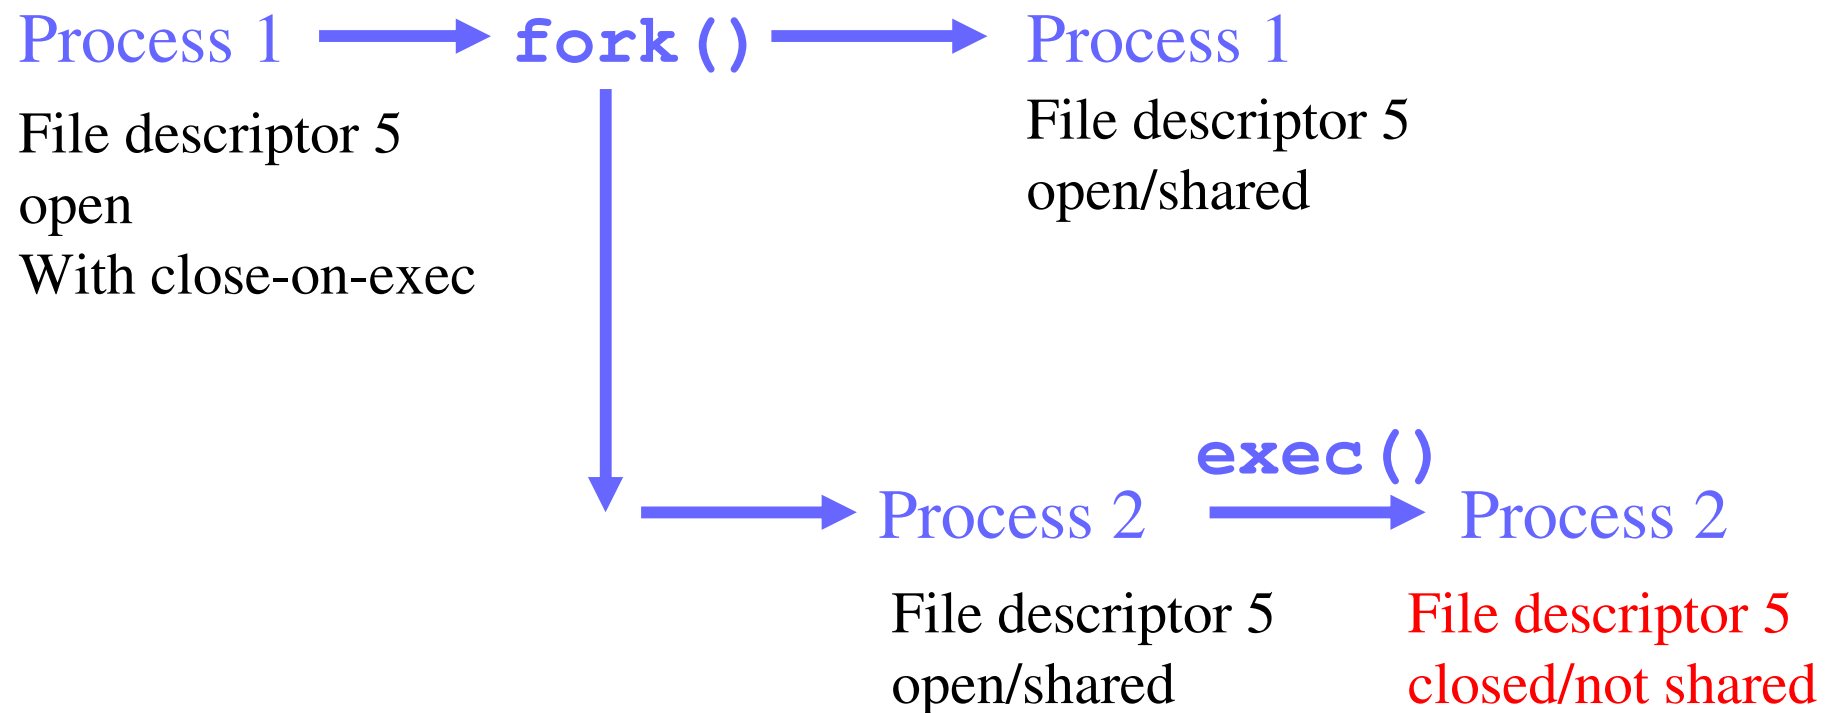
# Differences with Win32

---



- In Windows, named pipes are similar to sockets
  - ▶ It is possible to connect processes on different machines
- It is possible to avoid interleaving
- It is possible to make a bidirectional communication
- Only WinNT can create named pipes, Win9x can only create client connections

# Close-on-exec





# Bibliography



- *UNIX Network Programming, Volume 1, Second Edition: Networking APIs: Sockets and XTI*, Prentice Hall, 1998
- *UNIX Network Programming, Volume 2, Second Edition: Interprocess Communications*, Prentice Hall, 1999
- <http://mij.oltrelinux.com/devel/unixprg/>