



CEFRIEL

Consorzio per la Formazione e la Ricerca
in Ingegneria dell'Informazione

Politecnico di Milano



Scheduling for uniprocessor systems

Introduction

Lecturer:

William Fornaciari

Politecnico di Milano

fornacia@elet.polimi.it

www.elet.polimi.it/~fornacia



SUMMARY



- Basic Concepts
- Scheduling Criteria
- Scheduling Algorithms

...more to come



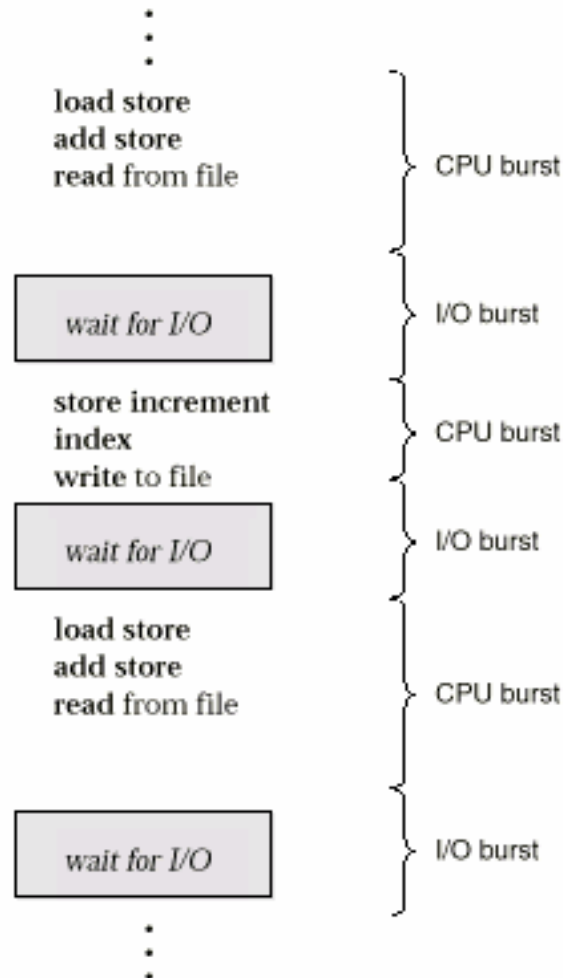
Basic Concepts



- Maximum CPU utilization obtained with multiprogramming
- CPU-I/O Burst Cycle - Process execution consists of a *cycle* of CPU execution and I/O wait.
- CPU burst distribution



Alternating Sequence of CPU And I/O Bursts





CPU Scheduler



- Selects from among the processes in memory that are ready to execute, and allocates the CPU to one of them.
- CPU scheduling decisions may take place when a process:
 1. switches from running to waiting state.
 2. switches from running to ready state.
 3. switches from waiting to ready.
 4. terminates.
- Scheduling under 1 and 4 is *nonpreemptive*.
- All other scheduling is *preemptive*.



Dispatcher



- Dispatcher module gives control of the CPU to the process selected by the short-term scheduler; this involves:
 - ▶ switching context
 - ▶ switching to user mode
 - ▶ jumping to the proper location in the user program to restart that program
- *Dispatch latency* - time it takes for the dispatcher to stop one process and start another running.



Scheduling Criteria



- CPU utilization - keep the CPU as busy as possible
- Throughput - # of processes that complete their execution per time unit
- Turnaround time - amount of time to execute a particular process
- Waiting time - amount of time a process has been waiting in the ready queue
- Response time - amount of time it takes from when a request was submitted until the first response is produced, not output (for time-sharing environment)



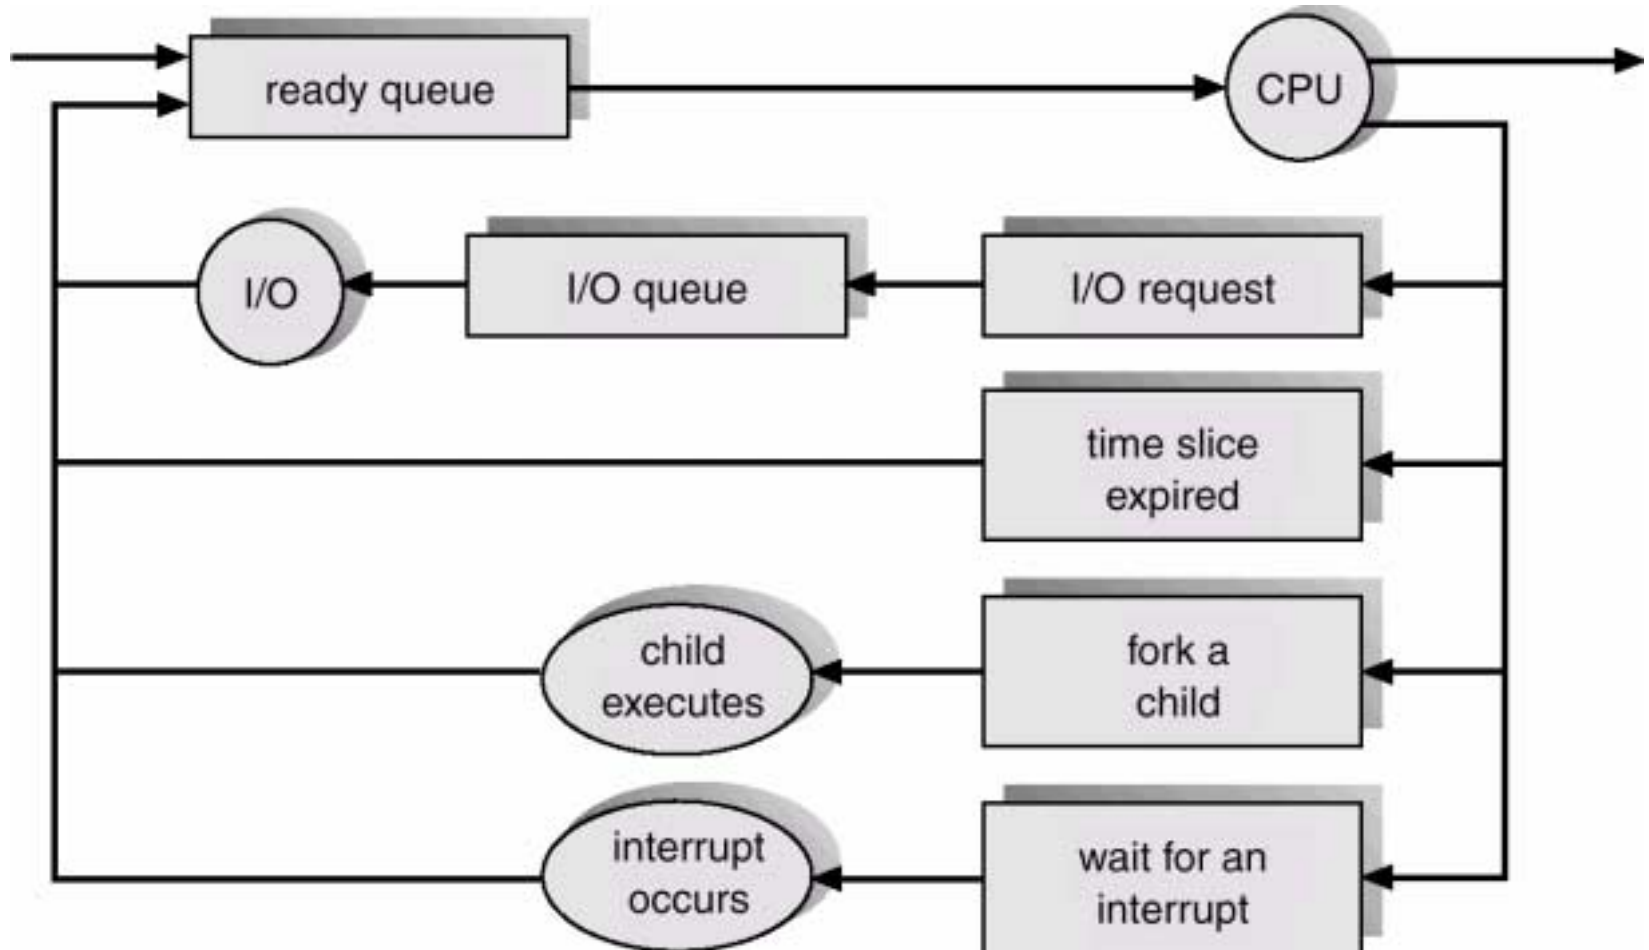
Optimization Criteria



- Max CPU utilization
- Max throughput
- Min turnaround time
- Min waiting time
- Min response time

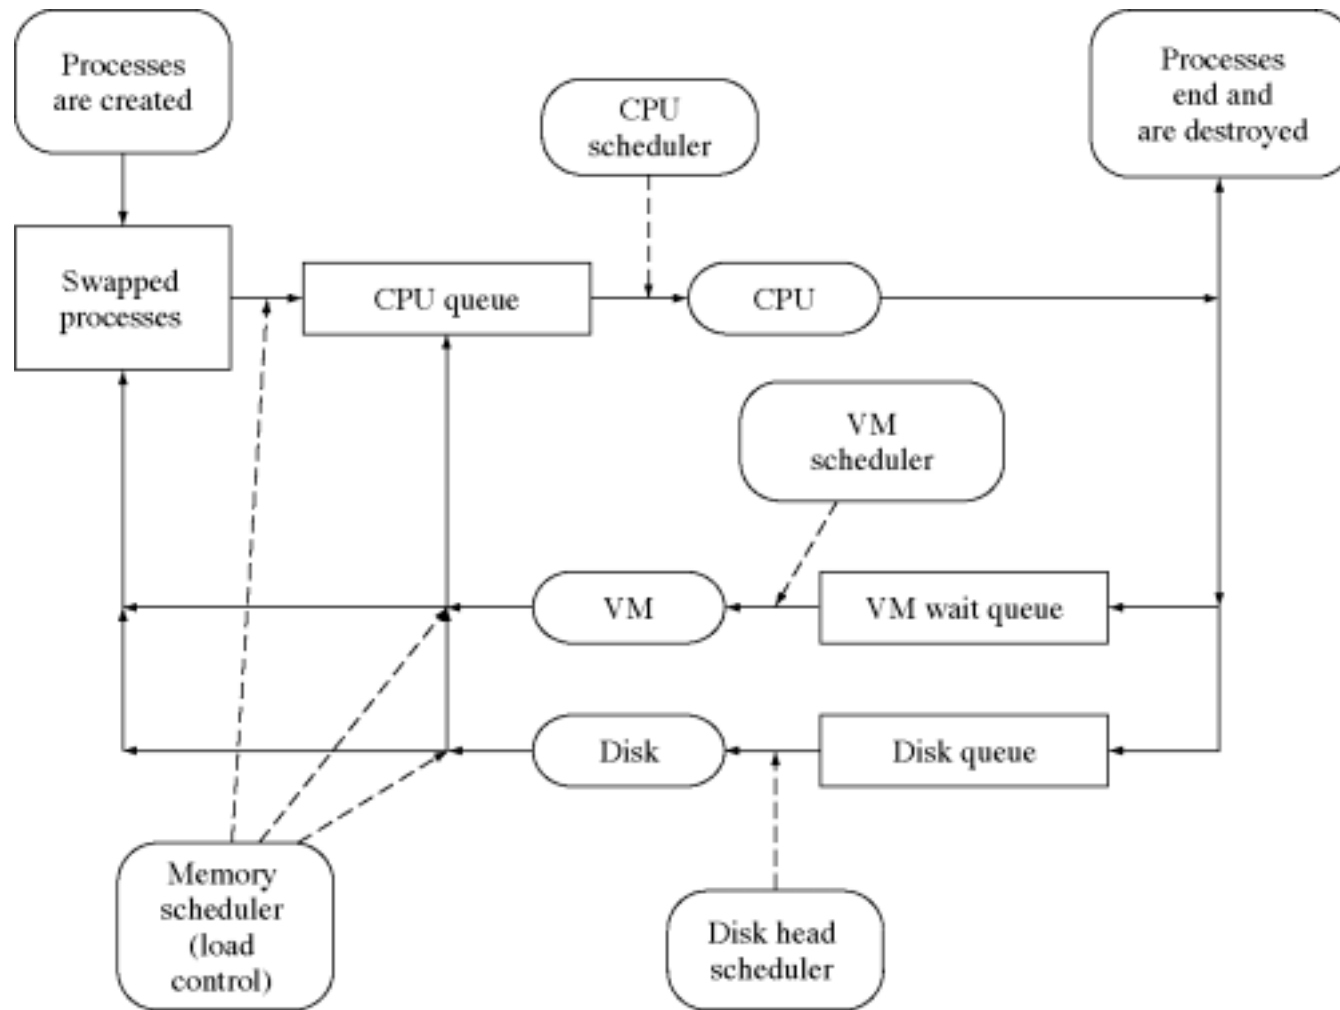


Process Scheduling: Examples



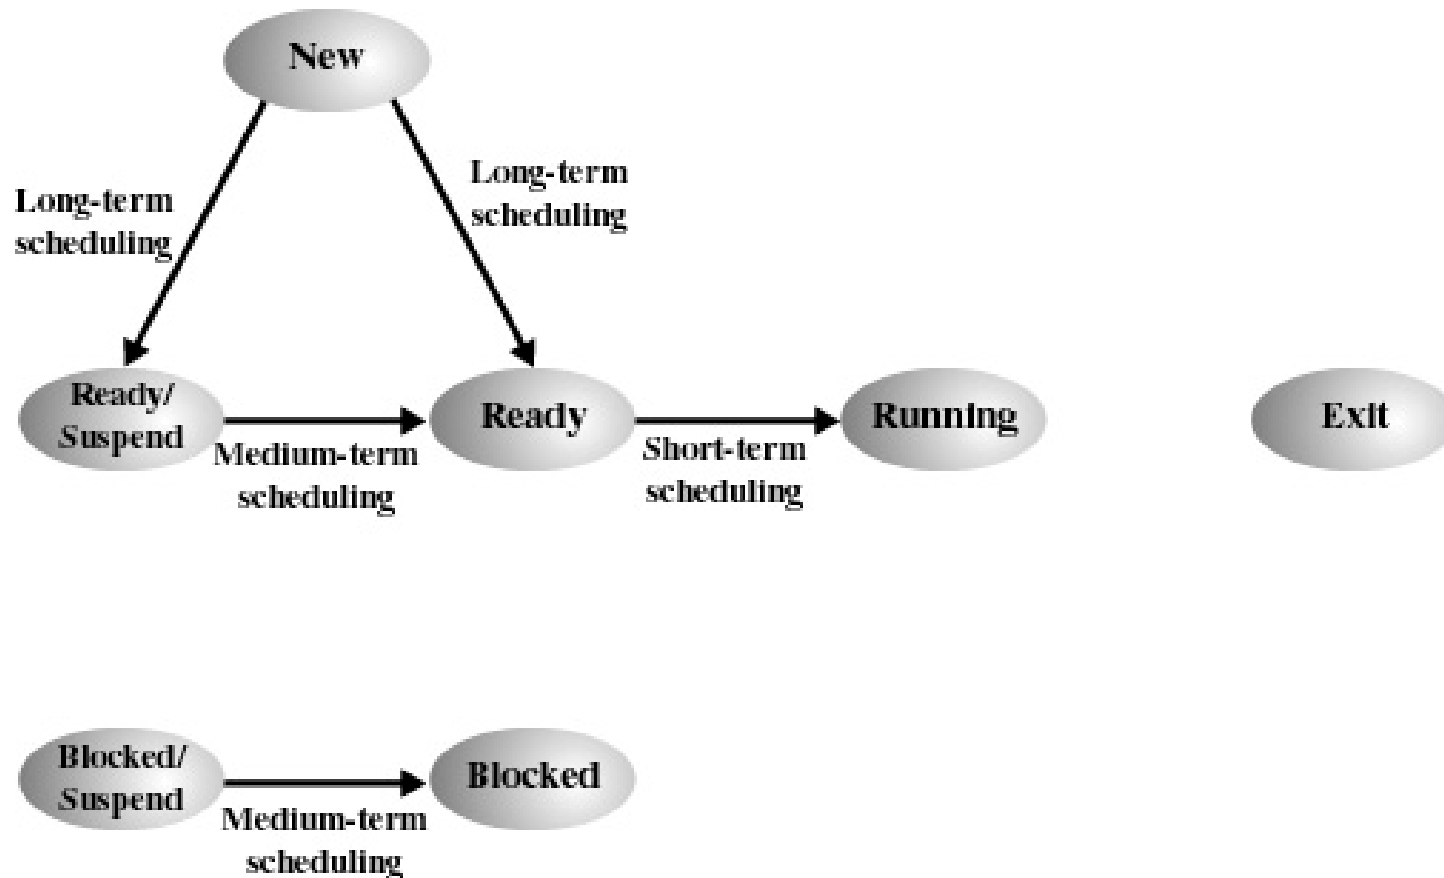


OS Schedulers



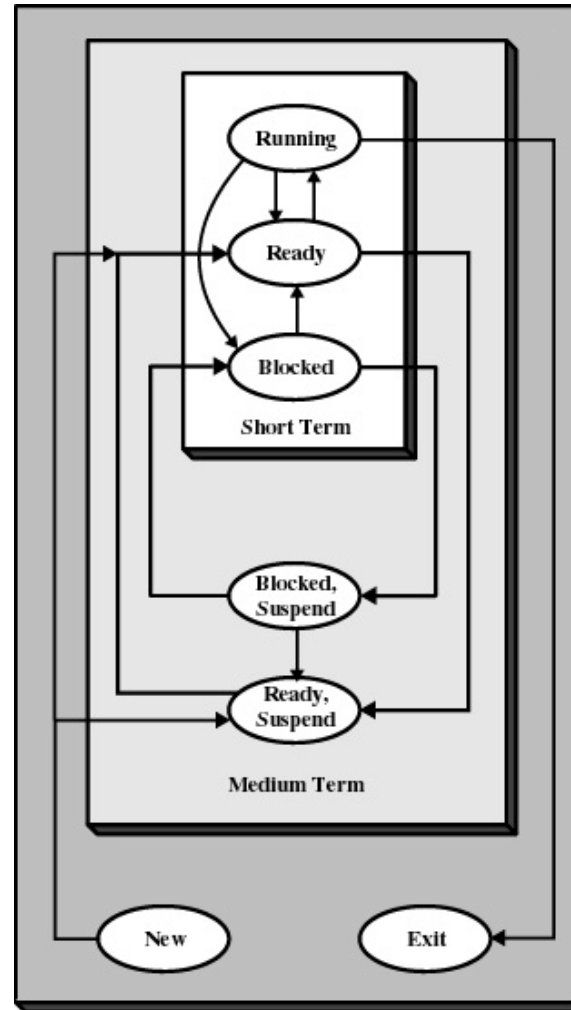


Scheduling and Process State Transition





Levels of Scheduling





Types of Scheduling



- Long-term
 - ▶ The decision to add to the pool of processes to be executed
- Medium-term
 - ▶ The decision to add to the number of processes that are partially or fully in main memory
- Short-term
 - ▶ The decision as to which available process will be executed by the processor
- I/O Scheduling
 - ▶ The decision as to which process's pending I/O request shall be handled by an available I/O device



Long-Term Scheduling



- Determines which programs are admitted to the system for processing
- Controls the degree of multiprogramming
- More processes, smaller percentage of time each process is executed



Medium-Term Scheduling



- Part of the swapping function
- Based on the need to manage the degree of multiprogramming



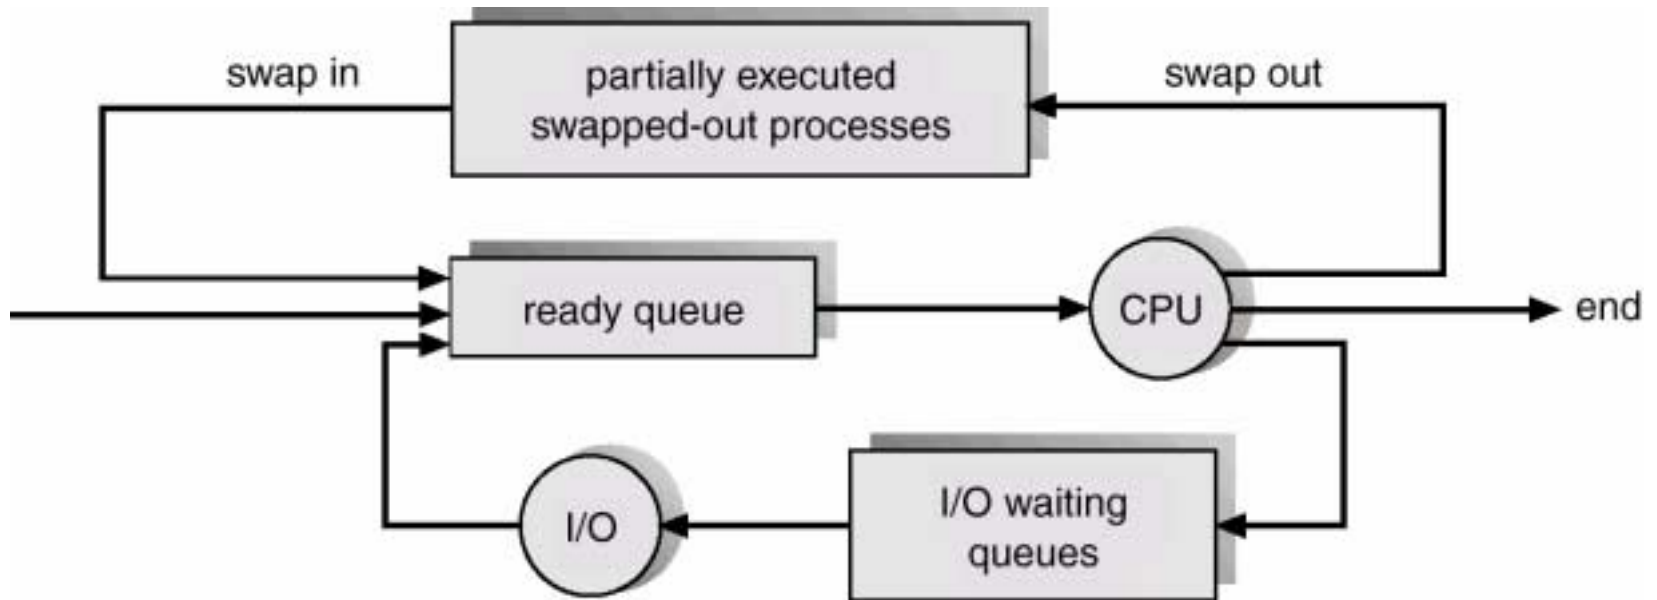
Short-Term Scheduling



- Known as the dispatcher
- Executes most frequently
- Invoked when an event occurs
 - ▶ Clock interrupts
 - ▶ I/O interrupts
 - ▶ Operating system calls
 - ▶ Signals



Addition of Medium Term Scheduling





Short-Term Scheduling Criteria



- User-oriented
 - ▶ Response Time
 - Elapsed time between the submission of a request until there is output
- System-oriented
 - ▶ Effective and efficient utilization of the processor



Short-Term Scheduling Criteria



- Performance-related
 - ▶ Quantitative
 - ▶ Measurable such as response time and throughput
- Not performance related
 - ▶ Qualitative
 - ▶ Predictability



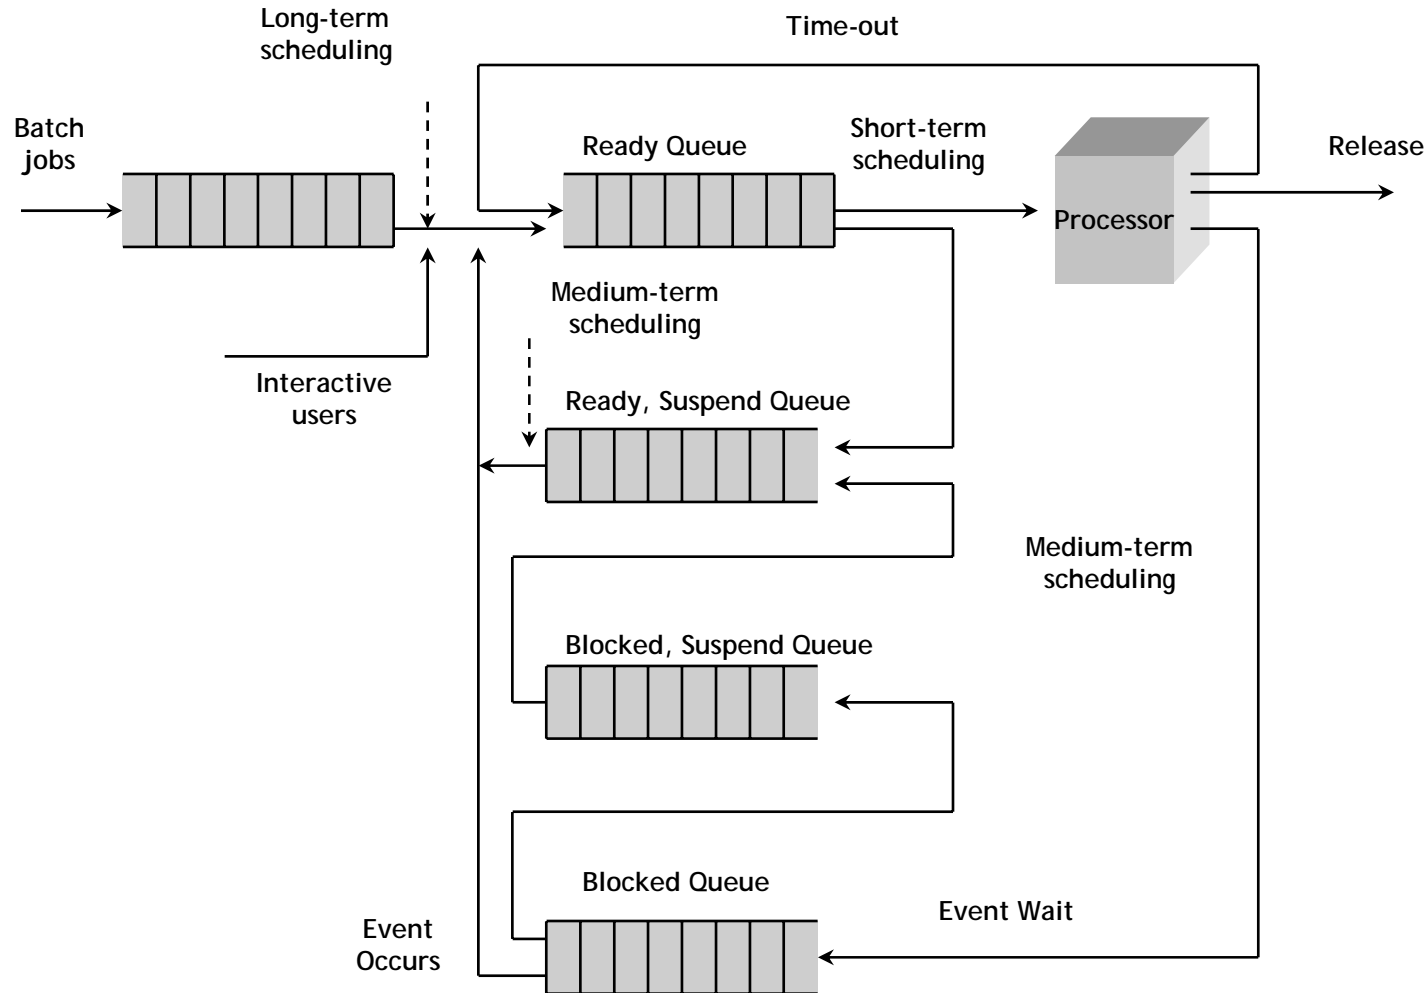
Priorities



- Scheduler will always choose a process of higher priority over one of lower priority
- Have multiple ready queues to represent each level of priority
- Lower-priority may suffer starvation
 - ▶ allow a process to change its priority based on its age or execution history

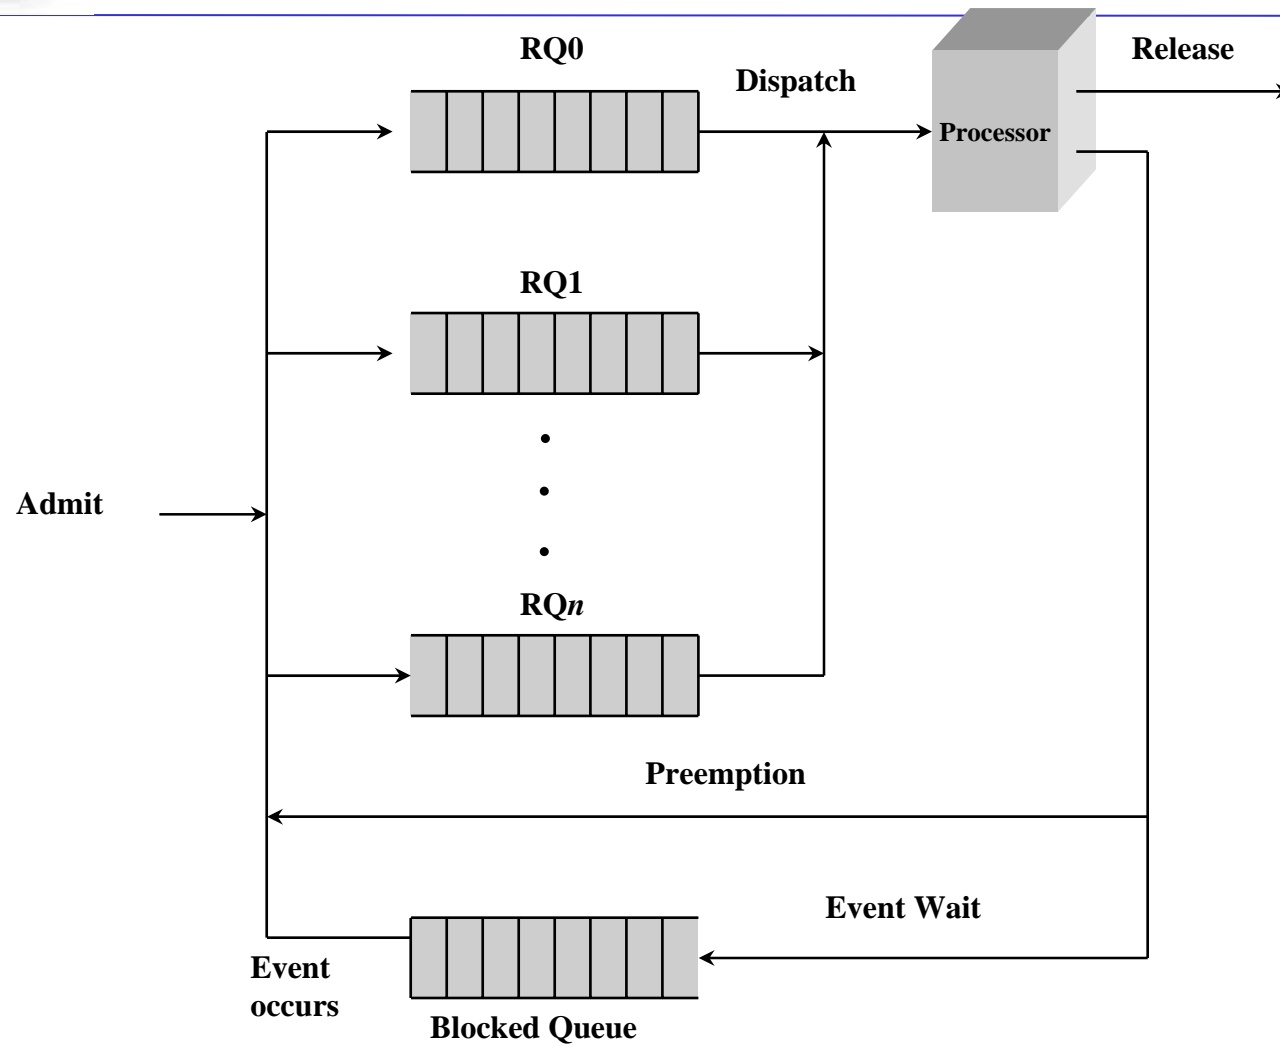


Queuing Diagram for Scheduling





Priority Queuing





Decision Mode



- Nonpreemptive
 - ▶ Once a process is in the running state, it will continue until it terminates or blocks itself for I/O
- Preemptive
 - ▶ Currently running process may be interrupted and moved to the Ready state by the operating system
 - ▶ Allows for better service since any one process cannot monopolize the processor for very long




Process Scheduling: Example



Process	Arrival Time	Service Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2



First-Come, First-Served (FCFS) Scheduling



- Each process joins the Ready queue
- When the current process ceases to execute, the oldest process in the Ready queue is selected



First-Come, First-Served (FCFS) Scheduling



- A short process may have to wait a very long time before it can execute
- Favours CPU-bound processes
 - ▶ I/O processes have to wait until CPU-bound process completes



First-Come, First-Served (FCFS) Scheduling



- Example:

<u>Process</u>	<u>Burst Time</u>
P_1	24
P_2	3
P_3	3

- Suppose that the processes arrive in the order: P_1 , P_2 , P_3
The Gantt Chart for the schedule is:



- Waiting time for $P_1 = 0$; $P_2 = 24$; $P_3 = 27$
- Average waiting time: $(0 + 24 + 27)/3 = 17$



FCFS Scheduling (Cont.)



Suppose that the processes arrive in the order

$$P_2, P_3, P_1.$$

- The Gantt chart for the schedule is:



- Waiting time for $P_1 = 6$; $P_2 = 0$; $P_3 = 3$
- Average waiting time: $(6 + 0 + 3)/3 = 3$
- Much better than previous case.
- *Convoy effect* short process behind long process



Shortest-Job-First (SJF) Scheduling



- Associate with each process the length of its next CPU burst. Use these lengths to schedule the process with the shortest time.
- Two schemes:
 - ▶ nonpreemptive - once CPU given to the process it cannot be preempted until completes its CPU burst.
 - ▶ Preemptive - if a new process arrives with CPU burst length less than remaining time of current executing process, preempt. This scheme is known as the Shortest-Remaining-Time-First (SRTF).
- SJF is optimal - gives minimum average waiting time for a given set of processes.




Example of Non-Preemptive SJF



<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0.0	7
P_2	2.0	4
P_3	4.0	1
P_4	5.0	4

- SJF (non-preemptive)



- Average waiting time = $(0 + 6 + 3 + 7)/4 = 4$




Example of Preemptive SJF



<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0.0	7
P_2	2.0	4
P_3	4.0	1
P_4	5.0	4

- SJF (preemptive)




- Average waiting time = $(9 + 1 + 0 + 2)/4 = 3$



Shortest Process Next



- Nonpreemptive policy
- Process with shortest expected processing time is selected next
- Short process jumps ahead of longer processes





Shortest Process Next




- Predictability of longer processes is reduced
- If estimated time for process not correct, the operating system may abort it
- Possibility of starvation for longer processes



Shortest Remaining Time



- Preemptive version of shortest process next policy
- Must estimate processing time






Highest Response Ratio Next (HRRN)



- Choose next process with the highest ratio

$$\frac{\text{time spent waiting} + \text{expected service time}}{\text{expected service time}}$$






Feedback




- Penalize jobs that have been running longer
- Don't know remaining time process needs to execute





Feedback Scheduling





Fair-Share Scheduling




- User's application runs as a collection of processes (threads)
- User is concerned about the performance of the application
- Need to make scheduling decisions based on process sets




Example of Fair-Share Scheduling three Processes two Groups



Time	Process A			Process B			Process C		
	Priority	Process	Group	Priority	Process	Group	Priority	Process	Group
0	60	0	0	60	0	0	60	0	0
		1	1						
		2	2						
		.	.						
		.	.						
		60	60						
1	90	30	30	60	0	0	60	0	0
					1	1			1
					2	2			2
					.	.			.
					.	.			.
					60	60			60
2	74	15	15	90	30	30	75	0	30
		16	16						
		17	17						
		.	.						
		.	.						
		75	75						
3	96	37	37	74	15	15	67	0	15
						16		1	16
						17		2	17
						.		.	.
						.		.	.
						75		60	75
4	78	18	18	81	7	37	93	30	37
		19	19						
		20	20						
		.	.						
		.	.						
		78	78						
5	98	39	39	70	3	18	76	15	18



Group 1



Group 2

Shaded rectangle represent executing process



Determining Length of Next CPU Burst



- Can only estimate the length
- Can be done by using the length of previous CPU bursts, using exponential averaging.

1. t_n = actual length of n^{th} CPU burst
2. τ_{n+1} = predicted value for the next CPU burst
3. $\alpha, 0 \leq \alpha \leq 1$
4. Define :

$$\tau_{n+1} = \alpha t_n + (1 - \alpha) \tau_n$$



Examples of Exponential Averaging



- $\alpha = 0$
 - ▶ $\tau_{n+1} = \tau_n$
 - ▶ Recent history does not count
- $\alpha = 1$
 - ▶ $\tau_{n+1} = t_n$
 - ▶ Only the actual last CPU burst counts
- If we expand the formula, we get:
$$\begin{aligned}\tau_{n+1} = & \alpha t_n + (1 - \alpha) \alpha t_{n-1} + \dots \\ & + (1 - \alpha)^j \alpha t_{n-j} + \dots \\ & + (1 - \alpha)^{n+1} t_n \tau_0\end{aligned}$$
- Since both α and $(1 - \alpha)$ are less than or equal to 1, each successive term has less weight than its predecessor



Priority Scheduling




- A priority number (integer) is associated with each process
- The CPU is allocated to the process with the highest priority (smallest integer = highest priority)
 - ▶ preemptive
 - ▶ nonpreemptive
- SJN is a priority scheduling where priority is the predicted next CPU burst time
- Problem: Starvation - low priority processes may never execute
- Solution: Aging - as time progresses increase the priority of the process



Priority Scheduling



A scheduling algorithm with four priority classes






Round-Robin



- Uses preemption based on a clock
- An amount of time is determined that allows each process to use the processor for that length of time





Round Robin (RR)



- Each process gets a small unit of CPU time (*time quantum*), usually 10-100 milliseconds. After this time has elapsed, the process is preempted and added to the end of the ready queue
- If there are n processes in the ready queue and the time quantum is q , then each process gets $1/n$ of the CPU time in chunks of at most q time units at once. No process waits more than $(n-1)q$ time units
- Performance
 - ▶ q large \Rightarrow FIFO
 - ▶ q small $\Rightarrow q$ must be large with respect to context switch, otherwise overhead is too high




Example: RR with Time Quantum = 20



<u>Process</u>	<u>Burst Time</u>
P_1	53
P_2	17
P_3	68
P_4	24


- The Gantt chart is:



- Typically, higher average turnaround than SJF, but better *response*



How a Smaller Time Quantum Increases Context Switches





Multilevel Queue



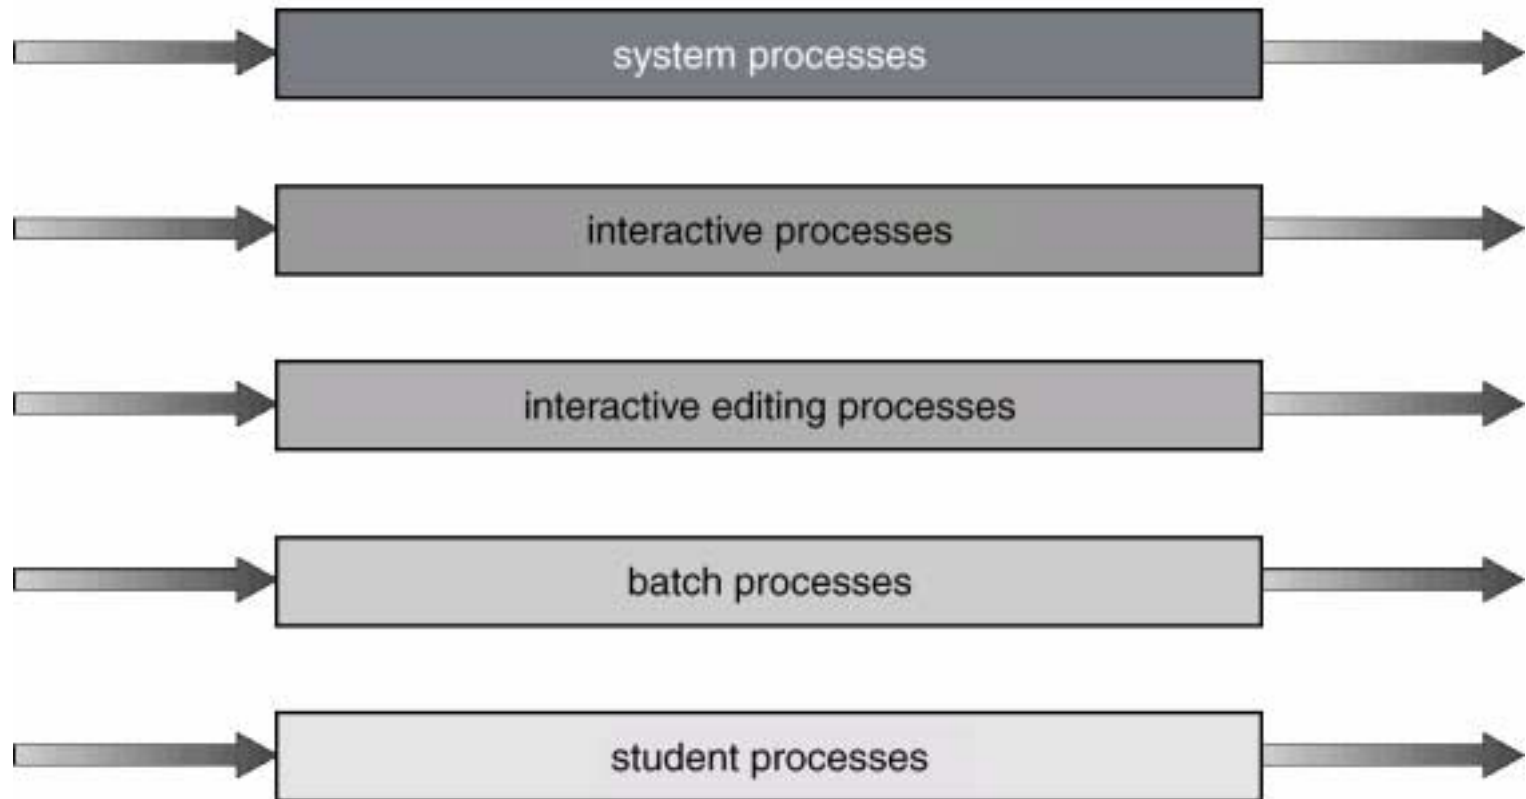
- Ready queue is partitioned into separate queues:
foreground (interactive)
background (batch)
- Each queue has its own scheduling algorithm,
foreground - RR
background - FCFS
- Scheduling must be done between the queues
 - ▶ Fixed priority scheduling; i.e., serve all from foreground then from background. Possibility of starvation.
 - ▶ Time slice - each queue gets a certain amount of CPU time which it can schedule amongst its processes; i.e.,
80% to foreground in RR
 - ▶ 20% to background in FCFS



Multilevel Queue Scheduling



highest priority



lowest priority




Multilevel Feedback Queue



- A process can move between the various queues; aging can be implemented this way
- Multilevel-feedback-queue scheduler defined by the following parameters:
 - ▶ number of queues
 - ▶ scheduling algorithms for each queue
 - ▶ method used to determine when to upgrade a process
 - ▶ method used to determine when to demote a process
 - ▶ method used to determine which queue a process will enter when that process needs service



Multilevel Feedback Queues





Example of Multilevel Feedback Queue



- Three queues:
 - ▶ Q_0 - time quantum 8 milliseconds
 - ▶ Q_1 - time quantum 16 milliseconds
 - ▶ Q_2 - FCFS
- Scheduling
 - ▶ A new job enters queue Q_0 which is served FCFS. When it gains CPU, job receives 8 milliseconds. If it does not finish in 8 milliseconds, job is moved to queue Q_1 .
 - ▶ At Q_1 job is again served FCFS and receives 16 additional milliseconds. If it still does not complete, it is preempted and moved to queue Q_2