

 POLITECNICO DI MILANO

Dipartimento di
Elettronica e Informazione

Data Stream Management Systems (DSMSs)

December 15, 2011

Davide Barbieri

Motivation

- Context:

- Sensor networks (and IT network and traffic monitoring)
- Financial applications
- Web logs and click-streams

- Issues:

- Real-time querying
- analysis of huge amounts of data
- Summarization with lossy approach instead of time-consuming processing of the whole database
- Load is unpredictable (data reduction techniques needed)

Part 1

INTRODUCTION TO DATA STREAM MANAGEMENT

Data Streams

Data Streams: Continuous, unbounded, rapid, time-varying, ordered sequence of data elements

- Continuous:
 - Data is continuously flowing
 - We *can monitor* streams, we *can not store* them
- Unbounded
 - *The information carried can be thought as infinite, we can not forecast when a stream will stop to carry data*
- Rapid, Time-varying:
 - In principle we can not assume anything on the distribution of events over time
- Ordered
 - The arrival time is crucial in most applications

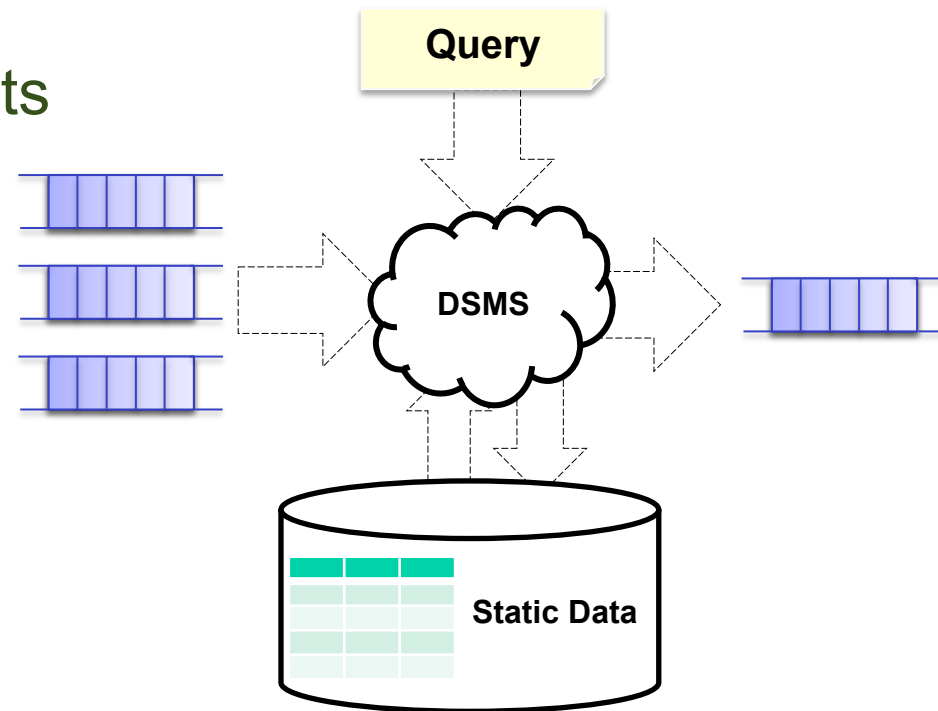
Comparison between DSMSs and DBMSs

DSMS	DBMS
Model: mainly transient data	Model: persistent data
Infinite sequence of tuples	Table: set bag of tuples
Access: sequential	Access: random
Updates: append only	Updates: all primitives
Query: persistent (continuous) queries	Query: transient queries
Query Answer: often approximate	Query Answer: typically exact
Query Eval. one-pass	Query Eval. multi-pass
Operators: unblocking only	Operators: blocking
Query Plan: adaptive	Query Plan: fixed

DSMS Generic Architecture

A four step process:

1. Stream registration
2. Static data import (optional)
3. Query registration
4. Subscription to results

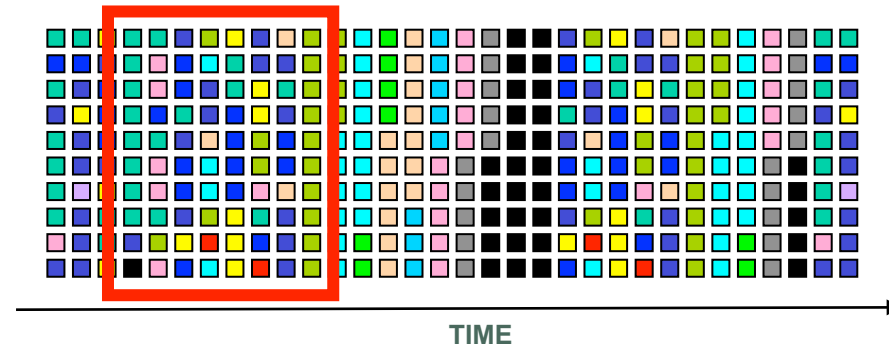


From Database Systems to DSMSs

- 1980s - Relational systems & query languages
- 1990s - Active rule systems
- 2000s - Continuous query engines

DSMS name	Year	Querying paradigm/ model	Architecture
Gigascop	2002	GSQL (SQL-like syntax)	Centralized
TelegraphCQ	2003	SQL-like language with programmatic definition of selection windows	Centralized
Aurora	2003	Algebraic, WYSIWYG Query Plan Editor	Centralized
STREAM	2004	CQL (SQL-like language)	Centralized
Borealis	2005	Same as Aurora	Distributed
Wavescope	2008	Programmatic approach, Wavescript (Ocaml-like language)	Generates distributed applications

Selection over Data Streams: Windows



- Landmark window: starts with the stream, unbounded
 - good for materialization of “slow” streams
- Window features:
 - Size
 - Physical (by number of elements)
 - Logical (by time)
 - Stride (pace)
 - Tumbling (not overlapping)
 - Sliding

Join and Aggregation

- *Approximate join* is needed
 1. Join windows from each stream
 2. merge the result
- *Aggregation*: first-class citizen in streaming applications
 - Data has to be summarized, can not be stored
 - Usually synopses are sufficient to accomplish application tasks
 - Synopses can be stored as opposed to entire streams
- Load management challenges
 - Quality of service (regular DBs weak here)
 - Multi-query optimization
 - Careful resource allocation & use
 - Graceful load shedding
 - Sampling

Part 2

C-SPARQL: BRIDGING STREAMING DATA AND REASONING

Motivations

- There are many scenarios requiring a fast solution to complex reasoning problems
 - Traffic jam detection
 - Routing traffic to avoid congestion
- Querying very large RDF knowledge bases with SPARQL has known scalability issues
- Combination of static and streaming RDF data leads to *stream reasoning* [1]
 - Enabling current reasoners to also consider rapidly changing data

[1] E. Della Valle, S. Ceri, D. Barbieri, D. Braga, A. Campi: **A First Step Towards Stream Reasoning**
In Proceedings of **FIS Future Internet Symposium**. Vienna, Austria, September 2008

Reasoning on Streaming Data

- **RDF data streams:** new data format set at the confluence of relational data streams and RDF data, defined as

“an ordered sequence of pairs of an RDF triple and its timestamp”

($\langle \text{subj}, \text{pred}, \text{obj} \rangle, \text{ts}$)

Timestamps are not required to be unique

Querying RDF streams: C-SPARQL

- Continuous SPARQL:
 - minimal extension to SPARQL [2]
 - compliant with version 1.1 [3]
 - operational formal semantics
 - C-SPARQL engine:
 - an optimized execution environment
 - extensible plug-in architecture
 - Efficient maintenance of ontological entailments
 - Validated on social data use cases
-
- [2] D. Barbieri, D. Braga, S. Ceri, E. Della Valle, M. Grossniklaus: **C-SPARQL: SPARQL for Continuous Querying**, WWW 2009
 - [3] D. Barbieri, D. Braga, S. Ceri, E. Della Valle, M. Grossniklaus: **Querying RDF Streams with C-SPARQL - ACM SIGMOD Record**, 2010

Query and Stream Registration

- All queries over RDF data streams are continuous
 - Registered through the **REGISTER QUERY** statement
 - Run at a frequency specified by the optional **COMPUTED EVERY** clause
- The output of queries is in the form of tables of variable bindings or RDF graphs
- If it is a **CONSTRUCT** or **DESCRIBE** query, can be registered as a new RDF stream (through the **REGISTER STREAM** statement)
- **Composability:**
 - Query results registered as streams can feed other registered queries like any other RDF stream

Selection Over Streams

- In a C-SPARQL query, data streams are associated with an IRI
 - As in standard **FROM** clause
- Windows over streams are specified through the **FROM STREAM** clause
 - Can be physical or logical, tumbling or sliding
 - **FROM STREAM** <http://streams.org/p.trdf> [**TRIPLES** 100]
 - **FROM STREAM** <http://streams.org/l.trdf> [**RANGE** 30m **STEP** 5m]
- C-SPARQL queries can combine triples from more than one RDF stream
 - More than one **FROM STREAM** clause in a query are supported
 - Every **FROM STREAM** can have its own window definition

Another C-SPARQL Example

```
REGISTER STREAM CarsEnteringCityCenterPerDistrict COMPUTED EVERY 5m  
AS
```

```
CONSTRUCT {?district t:has-entering-cars COUNT(?car) AS ?passages}
```

```
FROM <http://city.org/info.rdf>
```

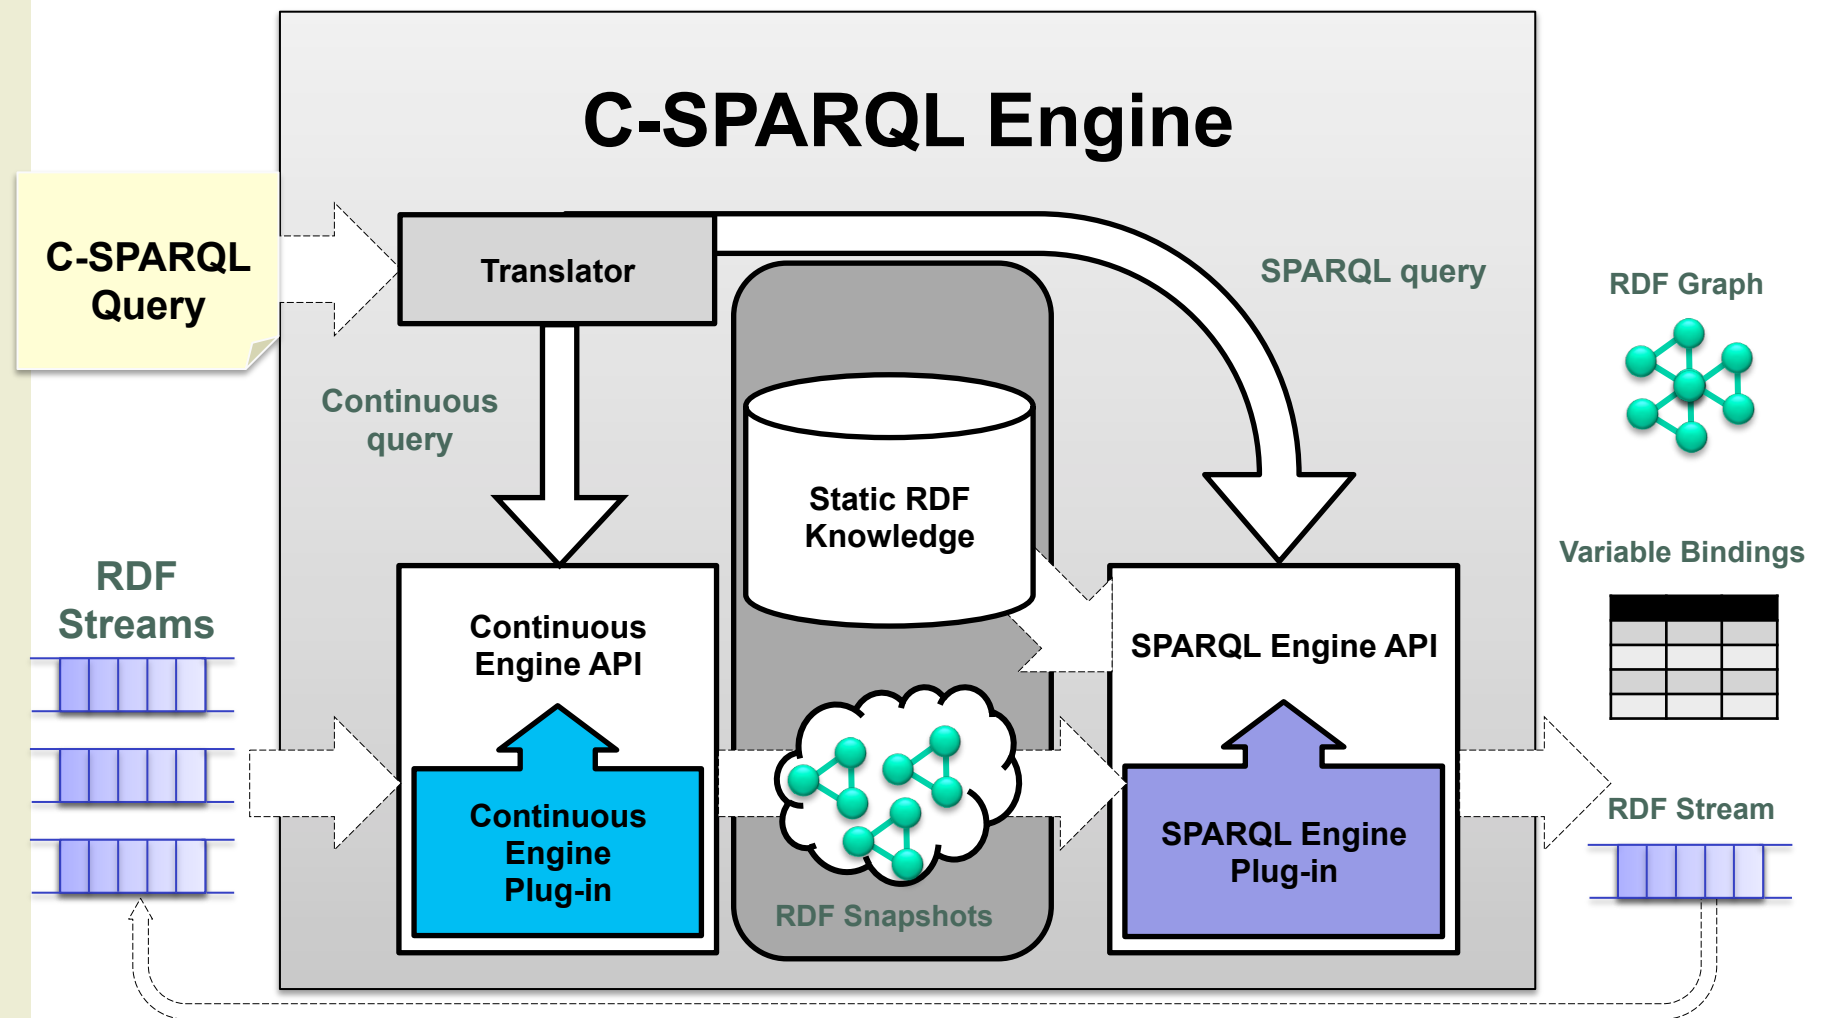
```
FROM STREAM <http://streams.org/gates.trdf> [RANGE 30m STEP 5m]
```

```
WHERE { ?tollgate t:registers ?car .  
        ?district c:contains ?street .  
        ?tollgate c:placedIn ?street . }
```

```
GROUP BY ?district
```

Subject	Predicate	Object	Timestamp
district-1	has-entering-cars	21	T ₁
district-2	has-entering-cars	34	T ₁
district-1	has-entering-cars	15	T ₂
district-2	has-entering-cars	7	T ₂

C-SPARQL Engine Architecture [4]

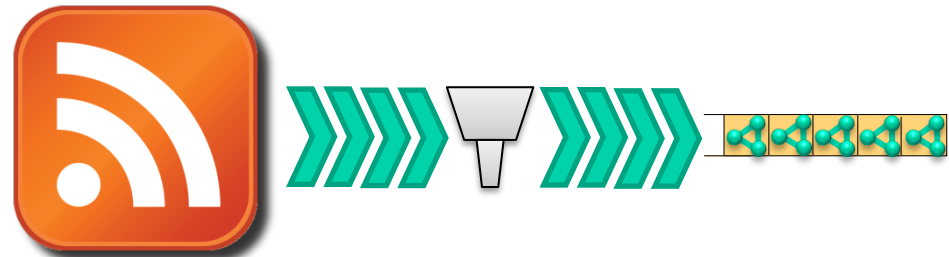


[4] D. Barbieri, D. Braga, S. Ceri and M. Grossniklaus: An Execution Environment for C-SPARQL Queries International Conference on Extending Database Technology (EDBT) 2010

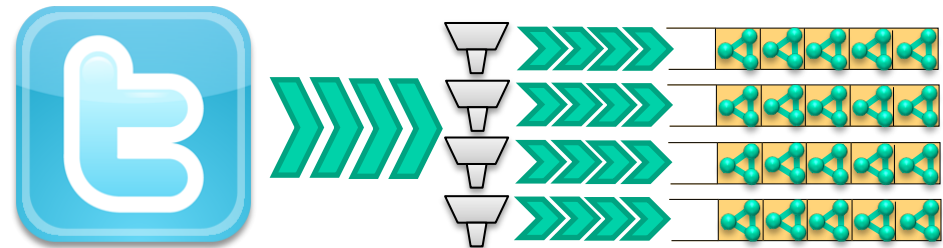
Use cases and APIs

- Some RDF-Stream generators have been developed:

RSS2RDFStream



Twitter2RDFStream [7] [8]



A ReST API has been designed in order to allow the invocation of C-SPARQL Engine as a service [8]

- [7] D. Barbieri, D. Braga, S. Ceri, E. Della Valle and M. Grossniklaus: ***Continuous Queries and Real-time Analysis of Social Semantic Data with C-SPARQL***, (SDoW2009)
- [8] D. Barbieri, and E. Della Valle: **A Proposal for Publishing Data Streams as Linked Data - A Position Paper - Linked Data on the Web (LDOW2010)**
- [9] D. Barbieri et al.: **Deductive and Inductive Stream Reasoning for Semantic Social Media Analytics - IEEE Intelligent Systems 2010**

Future Work

- Extending reasoning:
 - techniques that can exploit the strong temporal connotation of the triples in a RDF stream
- Optimizations:
 - Multi-query optimization
 - Distribution and parallelism
 - Low-level improvements (e.g. anticipation of aggregates)
- Extending the implementation:
 - Dynamic adaptation of query computation to load/source availability