



Vector Processing => Multimedia

2002

Vector Processors

- Initially developed for super-computing applications, today important for multimedia.
- Vector processors have high-level operations that work on linear arrays of numbers: "vectors"



Properties of Vector Processors

- Single vector instruction implies lots of work (loop)
 - Fewer instruction fetches
- Each result independent of previous result
 - Multiple operations can be executed in parallel
 - Simpler design, high clock rate
 - Compiler (programmer) ensures no dependencies
- Reduces branches and branch problems in pipelines
- Vector instructions access memory with known pattern
 - Effective prefetching
 - Amortize memory latency of over large number of elements
 - Can exploit a high bandwidth memory system
 - No (data) caches required!

Styles of Vector Architectures

- Memory-memory vector processors
 - All vector operations are memory to memory
- Vector-register processors
 - All vector operations between vector registers (except vector load and store)
 - Vector equivalent of load-store architectures
 - Includes all vector machines since late 1980s
 - We assume vector-register for rest of the lecture

Historical Perspective

- Mid-60s fear perf. stagnates
- SIMD processor arrays actively developed during late 60's - mid 70's
 - bit-parallel machines for image processing
 - pepe, staran, mpp
 - word-parallel for scientific
 - Illiac IV
- Cray develops fast scalar
 - CDC 6600, 7600
- CDC bets of vectors with Star-100
- Amdahl argues against vector

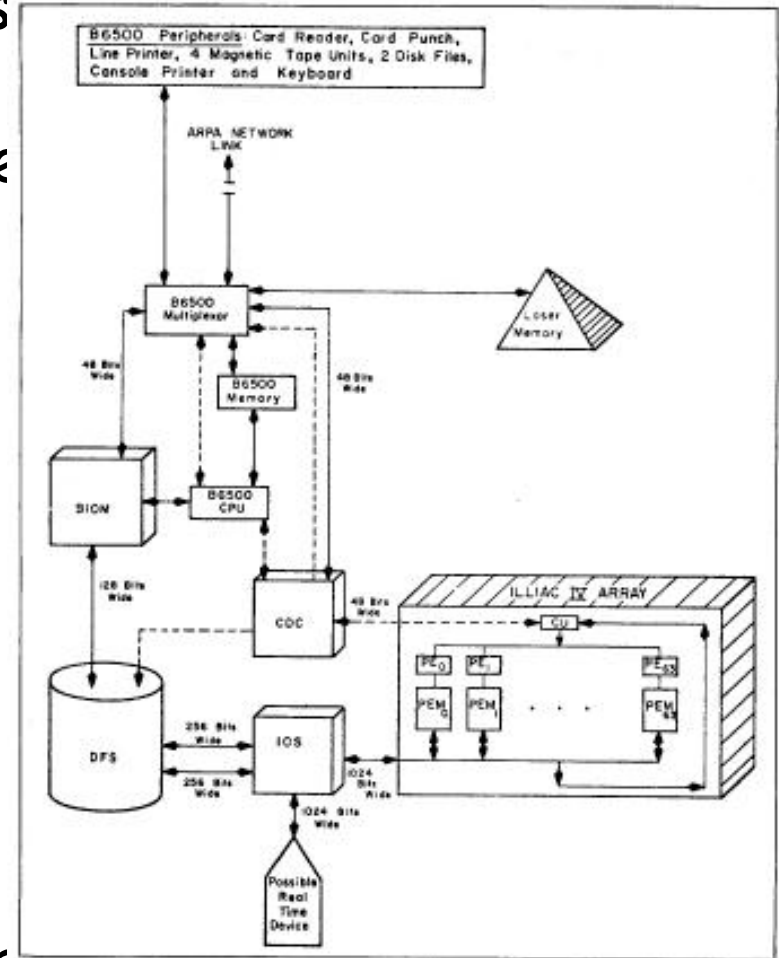



Fig. 11. Illiac IV system.

Cray-1 Breakthrough

- Fast, simple scalar processor
 - 80 MHz!
 - single-phase, latches
- Exquisite electrical and mechanical design
- Semiconductor memory
- Vector register concept
 - vast simplification of instruction set
 - reduced necc. memory bandwidth
- Tight integration of vector and scalar
- Piggy-back off 7600 stacklib
- Later vectorizing compilers developed
- Owned high-performance computing for a decade
 - what happened then?
 - VLIW competition



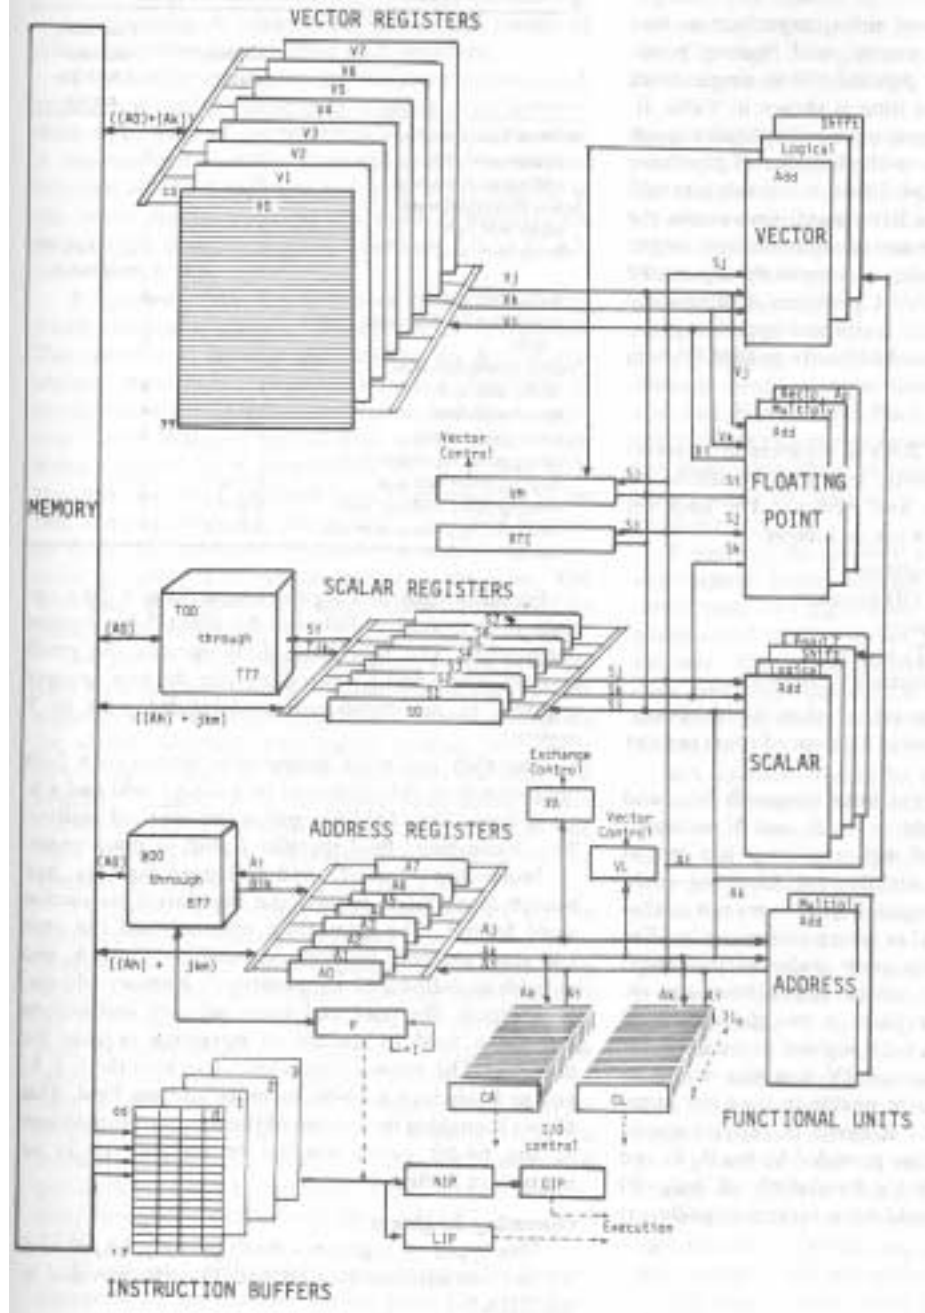
Components of a Vector Processor

- **Scalar CPU:** registers, datapaths, instruction fetch logic
- **Vector register**
 - Fixed length memory bank holding a single vector
 - Typically 8-32 vector registers, each holding 1 to 8 Kbits
 - Has at least 2 read and 1 write ports
 - MM: Can be viewed as array of 64b, 32b, 16b, or 8b elements
- **Vector functional units (FUs)**
 - Fully pipelined, start new operation every clock
 - Typically 2 to 8 FUs: integer and FP
 - Multiple datapaths (pipelines) used for each unit to process multiple elements per cycle
- **Vector load-store units (LSUs)**
 - Fully pipelined unit to load or store a vector
 - Multiple elements fetched/stored per cycle
 - May have multiple LSUs
- **Cross-bar** to connect FUs , LSUs, registers

Cray-1 Block Diagram

- Simple 16-bit RR instr
- 32-bit with immed
- Natural combinations of scalar and vector
- Scalar bit-vectors match vector length
- Gather/scatter M-R
- Cond. merge

Fig. 5. Block diagram of registers.



Basic Vector Instructions

<u>Instr.</u>	<u>Operands</u>	<u>Operation</u>	<u>Comment</u>
VADD.VV	V1, V2, V3	$V1 = V2 + V3$	vector + vector
VADD.SV	V1, R0, V2	$V1 = R0 + V2$	scalar + vector
VMUL.VV	V1, V2, V3	$V1 = V2 \times V3$	vector x vector
VMUL.SV	V1, R0, V2	$V1 = R0 \times V2$	scalar x vector
VLD	V1, R1	$V1 = M[R1..R1+63]$	load, stride=1
VLD _S	V1, R1, R2	$V1 = M[R1..R1+63 \times R2]$	load, stride=R2
VLD _X	V1, R1, V2	$V1 = M[R1+V2i, i=0..63]$	indexed("gather")
VST	V1, R1	$M[R1..R1+63] = V1$	store, stride=1
VST _S	V1, R1, R2	$V1 = M[R1..R1+63 \times R2]$	store, stride=R2
VST _X	V1, R1, V2	$V1 = M[R1+V2i, i=0..63]$	indexed("scatter")

+ all the regular scalar instructions (RISC style)...

Vector Memory Operations

- Load/store operations move groups of data between registers and memory
- Three types of addressing
 - Unit stride
 - Fastest
 - Non-unit (constant) stride
 - Indexed (gather-scatter)
 - Vector equivalent of register indirect
 - Good for sparse arrays of data
 - Increases number of programs that vectorize
 - compress/expand variant also
- Support for various combinations of data widths in memory
 - $\{.L., .W., .H., .B\} \times \{64b, 32b, 16b, 8b\}$

Vector Code Example

$$Y[0:63] = Y[0:653] + a * X[0:63]$$

64 element SAXPY: scalar

```
LD      R0,a
ADDI    R4,Rx,#512
loop: LD      R2, 0(Rx)
      MULTD   R2,R0,R2
      LD      R4, 0(Ry)
      ADDD    R4,R2,R4
      SD      R4, 0(Ry)
      ADDI    Rx,Rx,#8
      ADDI    Ry,Ry,#8
      SUB     R20,R4,Rx
      BNZ     R20,loop
```

64 element SAXPY: vector

```
LD      R0,a      #load scalar a
VLD     V1,Rx      #load vector X
VMUL.SV V2,R0,V1   #vector mult
VLD     V3,Ry      #load vector Y
VADD.VV V4,V2,V3   #vector add
VST     Ry,V4      #store vector Y
```

Vector Length

- A vector register can hold some maximum number of elements for each data width (maximum vector length or MVL)
- What to do when the application vector length is not exactly MVL?
- **Vector-length (VL)** register controls the length of any vector operation, including a vector load or store
 - E.g. vadd.vv with VL=10 is

```
for (I=0; I<10; I++) V1[I]=V2[I]+V3[I]
```
- VL can be anything from 0 to MVL
- How do you code an application where the vector length is not known until run-time?

Strip Mining

- Suppose application vector length $> \text{MVL}$
- Strip mining
 - Generation of a loop that handles MVL elements per iteration
 - A set operations on MVL elements is translated to a single vector instruction
- Example: vector saxpy of N elements
 - First loop handles $(N \bmod \text{MVL})$ elements, the rest handle MVL

```
VL = (N mod MVL);           // set VL = N mod MVL
for (I=0; I<VL; I++)         // 1st loop is a single set of
    Y[I]=A*X[I]+Y[I];       // vector instructions
low = (N mod MVL);
VL = MVL;                   // set VL to MVL
for (I=low; I<N; I++)        // 2nd loop requires N/MVL
    Y[I]=A*X[I]+Y[I];       // sets of vector instructions
```

Optimization 1: Chaining

- Suppose:

`vmul.vv` `v1`, `v2`, `v3`


`vadd.vv` `v4`, `v1`, `v5` # RAW hazard

- Chaining

- Vector register (V1) is not as a single entity but as a group of individual registers


- Pipeline forwarding can work on individual vector elements

- Flexible chaining: allow vector to chain to any other active vector operation => more read/write ports




**Cray X-mp
introduces
memory chaining**

Optimization 2: Multi-lane Implementation



- Elements for vector registers interleaved across the lanes
- Each lane receives identical control
- Multiple element operations executed per cycle
- Modular, scalable design
- No need for inter-lane communication for most vector instructions

Chaining & Multi-lane Example



- VL=16, 4 lanes, 2 FUs, 1 LSU, chaining -> 12 ops/cycle
- Just one new instruction issued per cycle !!!!

Optimization 3: Conditional Execution

- Suppose you want to vectorize this:

```
for (I=0; I<N; I++)  
    if (A[I]!= B[I]) A[I] -= B[I];
```

- Solution: vector conditional execution
 - Add vector flag registers with single-bit elements
 - Use a vector compare to set the a flag register
 - Use flag register as mask control for the vector sub
 - Addition executed only for vector elements with corresponding flag element set
- Vector code

```
vld          V1, Ra  
vld          V2, Rb  
vcmp.neq.vv  F0, V1, V2      # vector compare  
vsub.vv      V3, V2, V1, F0  # conditional vadd  
vst          V3, Ra
```

–Cray uses vector mask & merge

Two Ways to View Vectorization


- Inner loop vectorization (Classic approach)
 - Think of machine as, say, 32 vector registers each with 16 elements
 - 1 instruction updates 32 elements of 1 vector register
 - Good for vectorizing single-dimension arrays or regular kernels (e.g. saxpy)
- Outer loop vectorization (post-CM2)
 - Think of machine as 16 "virtual processors" (VPs) each with 32 scalar registers! (- multithreaded processor)
 - 1 instruction updates 1 scalar register in 16 VPs
 - Good for irregular kernels or kernels with loop-carried dependences in the inner loop
- These are just two compiler perspectives
 - The hardware is the same for both

Vectorizing Matrix Mult

```
// Matrix-matrix multiply:
//   sum a[i][t] * b[t][j] to get c[i][j]
for (i=1; i<n; i++)
{
    for (j=1; j<n; j++)
    {
        sum = 0;
        for (t=1; t<n; t++)
        {
            sum += a[i][t] * b[t][j]; // loop-carried
        }                             // dependence
        c[i][j] = sum;
    }
}
```

Parallelize Inner Product

Sum of Partial Products



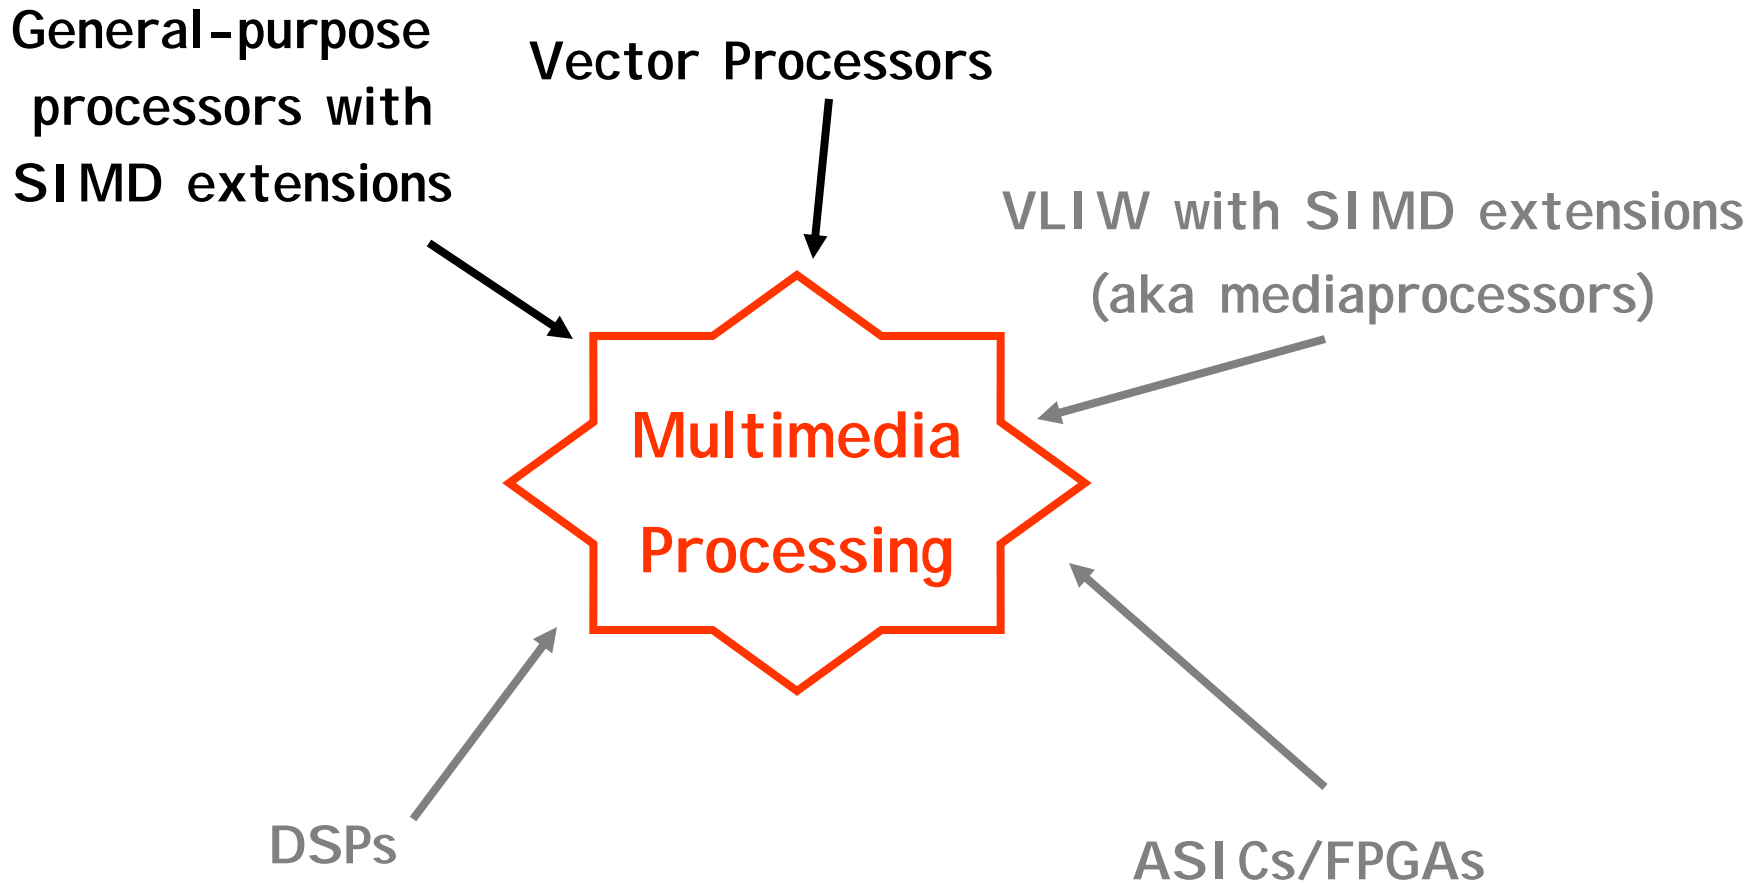
Outer-loop Approach

```
// Outer-loop Matrix-matrix multiply:
//      sum a[i][t] * b[t][j] to get c[i][j]
//      32 elements of the result calculated in parallel
//      with each iteration of the j-loop (c[i][j:j+31])
for (i=1; i<n; i++) {

    for (j=1; j<n; j+=32) {      // loop being vectorized
        sum[0:31] = 0;
        for (t=1; t<n; t++) {
            ascalar = a[i][t]; // scalar load
            bvector[0:31] = b[t][j:j+31]; // vector load
            prod[0:31] = b_vector[0:31]*ascalar; // vector mul
            sum[0:31] += prod[0:31];      // vector add
        }

        c[i][j:j+31] = sum[0:31]; // vector store
    }
}
```

Approaches to Mediaprocessing



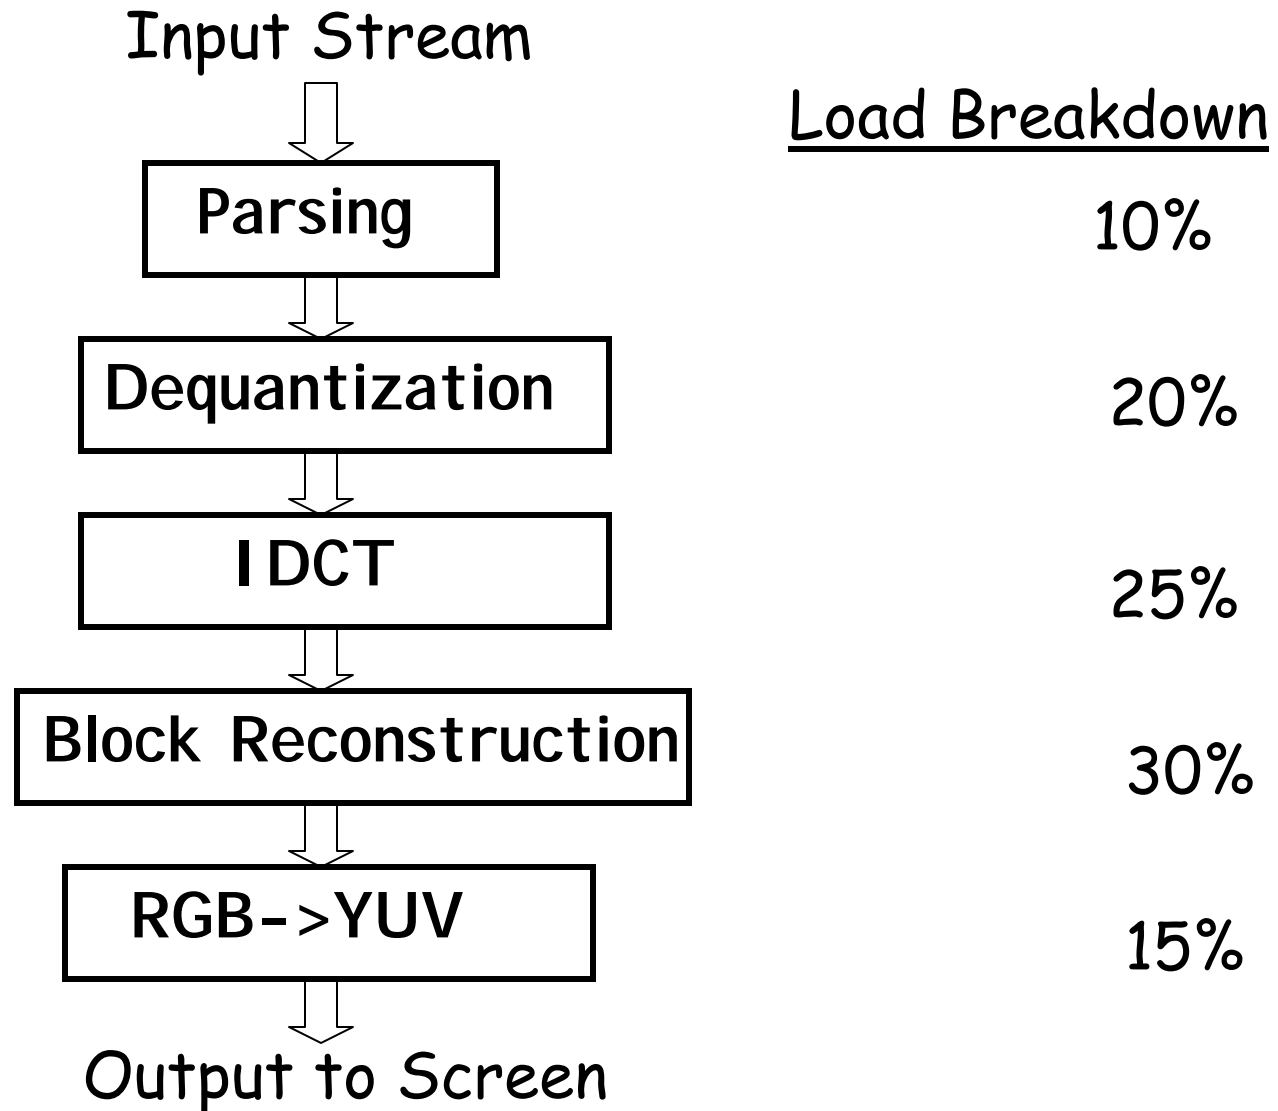
What is Multimedia Processing?

- Desktop:
 - 3D graphics (games)
 - Speech recognition (voice input)
 - Video/audio decoding (mpeg-mp3 playback)
- Servers:
 - Video/audio encoding (video servers, IP telephony)
 - Digital libraries and media mining (video servers)
 - Computer animation, 3D modeling & rendering (movies)
- Embedded:
 - 3D graphics (game consoles)
 - Video/audio decoding & encoding (set top boxes)
 - Image processing (digital cameras)
 - Signal processing (cellular phones)

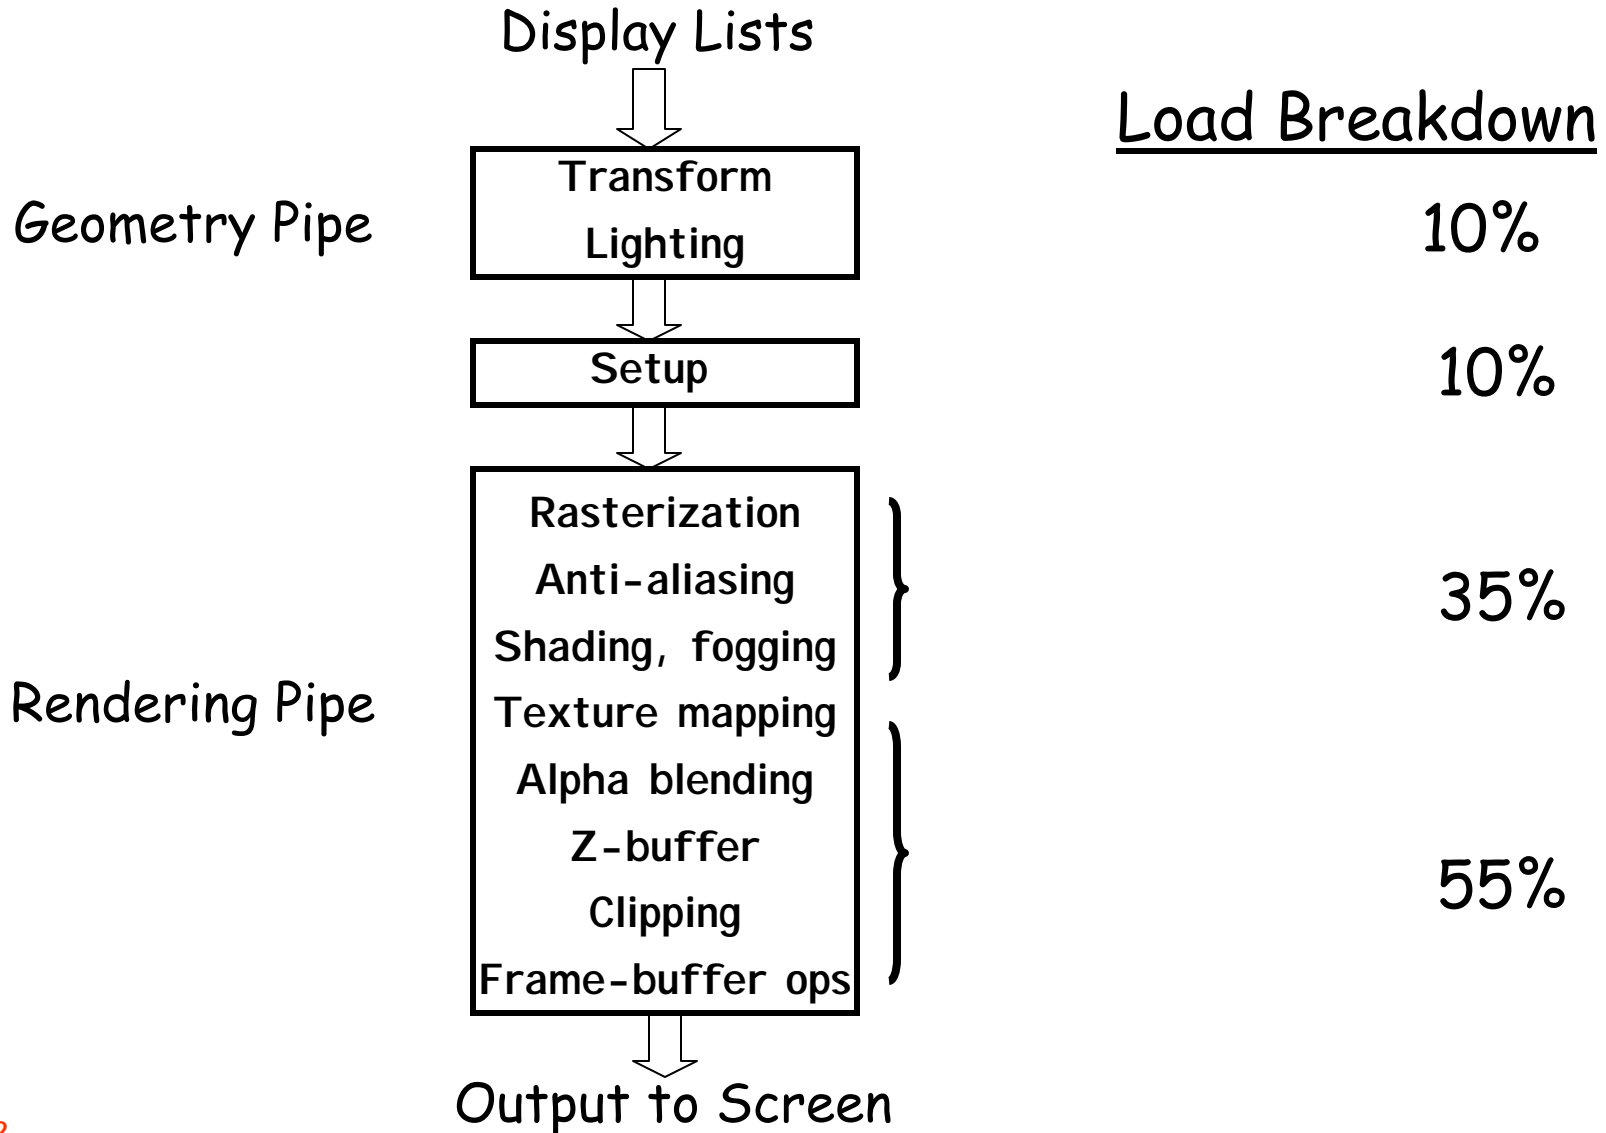
The Need for Multimedia ISAs

- Why aren't general-purpose processors and ISAs sufficient for multimedia (despite Moore's law)?
- Performance
 - A 1.2GHz Athlon can do MPEG-4 encoding at 6.4fps
 - One 384Kbps W-CDMA channel requires 6.9 GOPS
- Power consumption
 - A 1.2GHz Athlon consumes ~60W
 - Power consumption increases with clock frequency and complexity
- Cost
 - A 1.2GHz Athlon costs ~\$62 to manufacture and has a list price of ~\$600 (module)
 - Cost increases with complexity, area, transistor count, power, etc

Example: MPEG Decoding



Example: 3D Graphics



Characteristics of Multimedia Apps (1)

- Requirement for real-time response
 - “Incorrect” result often preferred to slow result
 - Unpredictability can be bad (e.g. dynamic execution)
- Narrow data-types
 - Typical width of data in memory: 8 to 16 bits
 - Typical width of data during computation: 16 to 32 bits
 - 64-bit data types rarely needed
 - Fixed-point arithmetic often replaces floating-point
- Fine-grain (data) parallelism
 - Identical operation applied on streams of input data
 - Branches have high predictability
 - High instruction locality in small loops or kernels

Characteristics of Multimedia Apps (2)

- Coarse-grain parallelism
 - Most apps organized as a pipeline of functions
 - Multiple threads of execution can be used
- Memory requirements
 - High bandwidth requirements but can tolerate high latency
 - High spatial locality (predictable pattern) but low temporal locality
 - Cache bypassing and prefetching can be crucial

Examples of Media Functions

- Matrix transpose/multiply (3D graphics)
- DCT/FFT (Video, audio, communications)
- Motion estimation (Video)
- Gamma correction (3D graphics)
- Haar transform (Media mining)
- Median filter (Image processing)
- Separable convolution (Image processing)
- Viterbi decode (Communications, speech)
- Bit packing (Communications, cryptography)
- Galois-fields arithmetic (Communications, cryptography)
- ...

SIMD Extensions for GPP

- Motivation
 - Low media-processing performance of GPPs
 - Cost and lack of flexibility of specialized ASICs for graphics/video
 - Underutilized datapaths and registers
- Basic idea: sub-word parallelism
 - Treat a 64-bit register as a vector of 2 32-bit or 4 16-bit or 8 8-bit values (short vectors)
 - Partition 64-bit datapaths to handle multiple narrow operations in parallel
- Initial constraints
 - No additional architecture state (registers)
 - No additional exceptions
 - Minimum area overhead

Overview of SIMD Extensions

Vendor	Extension	Year	# Instr	Registers
HP	MAX-1 and 2	94,95	9,8 (int)	Int 32x64b
Sun	VIS	95	121 (int)	FP 32x64b
Intel	MMX	97	57 (int)	FP 8x64b
AMD	3DNow!	98	21 (fp)	FP 8x64b
Motorola	AltiVec	98	162 (int,fp)	32x128b (new)
Intel	SSE	98	70 (fp)	8x128b (new)
MIPS	MIPS-3D	?	23 (fp)	FP 32x64b
AMD	E 3DNow!	99	24 (fp)	8x128 (new)
Intel	SSE-2	01	144 (int,fp)	8x128 (new)

Summary of SIMD Operations (1)

- Integer arithmetic
 - Addition and subtraction with saturation
 - Fixed-point rounding modes for multiply and shift
 - Sum of absolute differences
 - Multiply-add, multiplication with reduction
 - Min, max
- Floating-point arithmetic
 - Packed floating-point operations
 - Square root, reciprocal
 - Exception masks
- Data communication
 - Merge, insert, extract
 - Pack, unpack (width conversion)
 - Permute, shuffle


Summary of SIMD Operations (2)

- Comparisons
 - Integer and FP packed comparison
 - Compare absolute values
 - Element masks and bit vectors
- Memory
 - No new load-store instructions for short vector
 - No support for strides or indexing
 - Short vectors handled with 64b load and store instructions
 - Pack, unpack, shift, rotate, shuffle to handle alignment of narrow data-types within a wider one
 - Prefetch instructions for utilizing temporal locality

Programming with SIMD Extensions

- Optimized shared libraries
 - Written in assembly, distributed by vendor
 - Need well defined API for data format and use
- Language macros for variables and operations
 - C/C++ wrappers for short vector variables and function calls
 - Allows instruction scheduling and register allocation optimizations for specific processors
 - Lack of portability, non standard
- Compilers for SIMD extensions
 - No commercially available compiler so far
 - Problems
 - Language support for expressing fixed-point arithmetic and SIMD parallelism
 - Complicated model for loading/storing vectors
 - Frequent updates
- Assembly coding


SIMD Performance



Limitations

- Memory bandwidth
- Overhead of handling alignment and data width adjustments

A Closer Look at MMX/SSE




- Higher speedup for kernels with narrow data where 128b SSE instructions can be used
- Lower speedup for those with irregular or strided accesses

Choosing the Data Type Width

- Alternatives for selecting the width of elements in a vector register (64b, 32b, 16b, 8b)
- Separate instructions for each width
 - E.g. vadd64, vadd32, vadd16, vadd8
 - Popular with SIMD extensions for GPPs
 - Uses too many opcodes
- Specify it in a control register
 - Virtual-processor width (VPW)
 - Updated only on width changes
- NOTE
 - MVL increases when width (VPW) gets narrower
 - E.g. with 2Kbits for register, MVL is 32,64,128,256 for 64-,32-,16-,8-bit data respectively
 - Always pick the narrowest VPW needed by the application

Other Features for Multimedia

- Support for fixed-point arithmetic
 - Saturation, rounding-modes etc
- Permutation instructions of vector registers
 - For reductions and FFTs
 - Not general permutations (too expensive)
- Example: permutation for reductions
 - Move 2nd half of a vector register into another one
 - Repeatedly use with vadd to execute reduction
 - Vector length halved after each step



Designing a Vector Processor

- Changes to scalar core
- How to pick the maximum vector length?
- How to pick the number of vector registers?
- Context switch overhead?
- Exception handling?
- Masking and flag instructions?

Changes to Scalar Processor

- Decode vector instructions
- Send scalar registers to vector unit (vector-scalar ops)
- Synchronization for results back from vector register, including exceptions
- Things that don't run in vector don't have high ILP, so can make scalar CPU simple

How to Pick Max. Vector Length?

- Vector length => Keep all VFUs busy:
- Vector length $\geq \frac{(\# \text{ lanes}) \times (\# \text{ VFUs})}{\# \text{ Vector instr. issued/cycle}}$
- Notes:
 - Single instruction issue is always the simplest
 - Don't forget you have to issue some scalar instructions as well
 - Cray get mileage from VL \leq word length

How to Pick # of Vector Registers?

- More vector registers:
 - Reduces vector register “spills” (save/restore)
 - Aggressive scheduling of vector instructions: better compiling to take advantage of ILP
- Fewer
 - Fewer bits in instruction format (usually 3 fields)
- 32 vector registers are usually enough

Context Switch Overhead?

- The vector register file holds a huge amount of architectural state
 - Too expensive to save and restore all on each context switch
 - Cray: exchange packet
- Extra dirty bit per processor
 - If vector registers not written, don't need to save on context switch
- Extra valid bit per vector register, cleared on process start
 - Don't need to restore on context switch until needed
- Extra tip:
 - Save/restore vector state only if the new context needs to issue vector instructions

Exception Handling: Arithmetic

- Arithmetic traps are hard
- Precise interrupts => large performance loss
 - Multimedia applications don't care much about arithmetic traps anyway
- Alternative model
 - Store exception information in vector flag registers
 - A set flag bit indicates that the corresponding element operation caused an exception
 - Software inserts trap barrier instructions from SW to check the flag bits as needed
 - IEEE floating point requires 5 flag registers (5 types of traps)

Exception Handling: Page Faults

- Page faults must be precise
 - Instruction page faults not a problem
 - Data page faults harder
- Option 1: Save/restore internal vector unit state
 - Freeze pipeline, (dump all vector state), fix fault, (restore state and) continue vector pipeline
- Option 2: expand memory pipeline to check all addresses before send to memory
 - Requires address and instruction buffers to avoid stalls during address checks
 - On a page-fault on only needs to save state in those buffers
 - Instructions that have cleared the buffer can be allowed to complete

Exception Handling: Interrupts

- Interrupts due to external sources
 - I/O, timers etc
- Handled by the scalar core
- Should the vector unit be interrupted?
 - Not immediately (no context switch)
 - Only if it causes an exception or the interrupt handler needs to execute a vector instruction

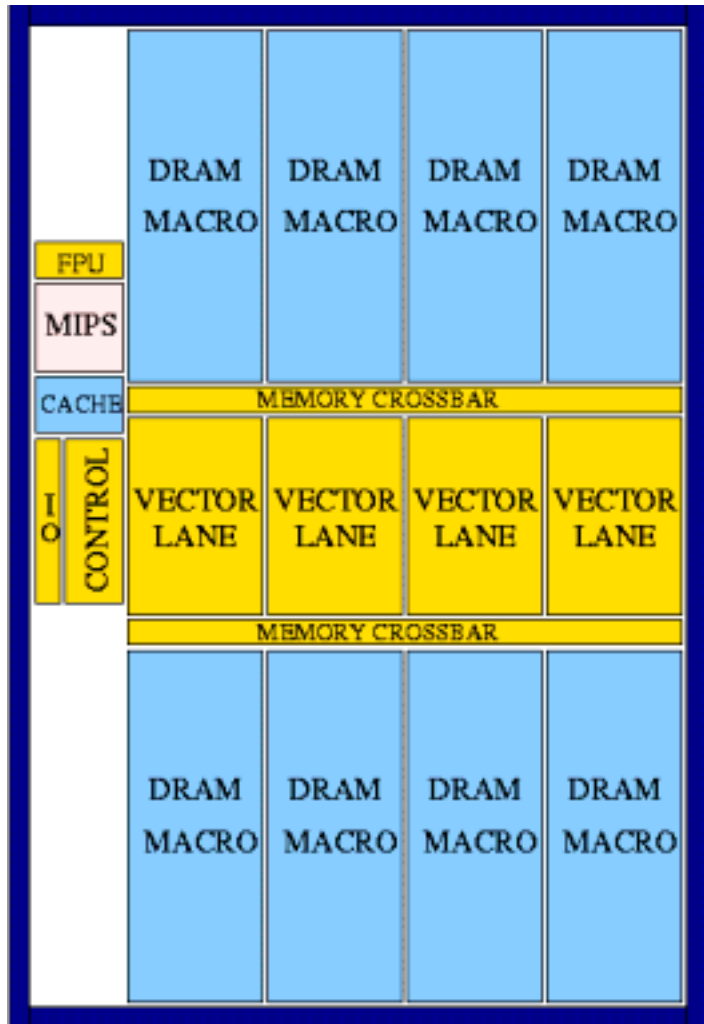
Vector Power Consumption

- Can trade-off parallelism for power
 - $\text{Power} = C * V_{dd}^2 * f$
 - If we double the lanes, peak performance doubles
 - Halving f restores peak performance but also allows halving of the V_{dd}
 - $\text{Power}_{\text{new}} = (2C) * (V_{dd}/2)^2 * (f/2) = \text{Power}/4$
- Simpler logic
 - Replicated control for all lanes
 - No multiple issue or dynamic execution logic
- Simpler to gate clocks
 - Each vector instruction explicitly describes all the resources it needs for a number of cycles
 - Conditional execution leads to further savings

Why Vectors for Multimedia?

- Natural match to parallelism in multimedia
 - Vector operations with VL the image or frame width
 - Easy to efficiently support vectors of narrow data types
- High performance at low cost
 - Multiple ops/cycle while issuing 1 instr/cycle
 - Multiple ops/cycle at low power consumption
 - Structured access pattern for registers and memory
- Scalable
 - Get higher performance by adding lanes without architecture modifications
- Compact code size
 - Describe N operations with 1 short instruction (v. VLIW)
- Predictable performance
 - No need for caches, no dynamic execution
- Mature, developed compiler technology

A Vector Media-Processor: VIRAM




- Technology: IBM SA-27E
 - 0.18mm CMOS, 6 copper layers
- 280 mm² die area
 - 158 mm² DRAM, 50 mm² logic
- Transistor count: ~115M
 - 14 Mbytes DRAM
- Power supply & consumption
 - 1.2V for logic, 1.8V for DRAM
 - 2W at 1.2V
- Peak performance
 - 1.6/3.2 /6.4 Gops (64/32/16b ops)
 - 3.2/6.4/12.8 Gops (with madd)
 - 1.6 Gflops (single-precision)
- Designed by 5 graduate students

Performance Comparison


	VIRAM	MMX
iDCT	0.75	3.75 (5.0x)
Color Conversion	0.78	8.00 (10.2x)
Image Convolution	1.23	5.49 (4.5x)
QCIF (176x144)	7.1M	33M (4.6x)
CIF (352x288)	28M	140M (5.0x)

- QCIF and CIF numbers are in clock cycles per frame
- All other numbers are in clock cycles per pixel
- MMX results assume no first level cache misses

FFT (1)



FFT (2)



SIMD Summary

- Narrow vector extensions for GPPs
 - 64b or 128b registers as vectors of 32b, 16b, and 8b elements
- Based on sub-word parallelism and partitioned datapaths
- Instructions
 - Packed fixed- and floating-point, multiply-add, reductions
 - Pack, unpack, permutations
 - Limited memory support
- 2x to 4x performance improvement over base architecture
 - Limited by memory bandwidth
- Difficult to use (no compilers)

Vector Summary

- Alternative model for explicitly expressing data parallelism
- If code is vectorizable, then simpler hardware, more power efficient, and better real-time model than out-of-order machines with SIMD support
- Design issues include number of lanes, number of functional units, number of vector registers, length of vector registers, exception handling, conditional operations
- Will multimedia popularity revive vector architectures?