



Using the proc filesystem

Lecturer:

Ing. Luca Pizzamiglio

Politecnico di Milano, DEI

luca.pizzamiglio@gmail.com

Overview



- Introduction
- Managing procfs entries
- Communicating with userland
- Tips and tricks

The proc filesystem



- **proc** is a virtual filesystem
 - ▶ It is not associated with a block device but exists only in **memory**
- Files in procfs allow **userland** programs to **access** certain information from the kernel
 - ▶ **Debug** purposes
- **/proc/sys** are sysctl files
 - ▶ They don't belong to procfs and are governed by a completely different API

Managing procfs entries (1/3)



- If you want to use any of the procfs functions, be sure to include the correct header file

```
#include <linux/proc_fs.h>
```

- Creating a regular file

```
struct proc_dir_entry* create_proc_entry(const char*  
    name, mode_t mode, struct proc_dir_entry* parent);
```

- ▶ NULL as *parent* parameter to create the file in the root of the proc

- ▶ It returns a pointer to the freshly created struct

- NULL if some error has happened

- ▶ It is possible to pass a path that spans multiple directories

```
create_proc_entry("foo/bar/info")
```

- It creates the bar directory, if necessary, with 755 permissions

- ▶ `create_proc_read_entry` to be able only to read the file

Managing procfs entries (2/3)



```
struct proc_dir_entry* proc_symlink(const char* name,  
    struct proc_dir_entry* parent, const char* dest)
```

- It creates a symlink in the procfs directory *parent* that points from *name* to *dest*

```
ln -s name dest
```

```
struct proc_dir_entry* proc_mknod(const char* name,  
    mode_t mode, struct proc_dir_entry* parent, kdev_t  
    rdev)
```

- It creates a device file *name* with mode *mode* in the procfs directory *parent*
 - ▶ *Mode* parameter must contain S_IFBLK or S_IFCHR

```
mknod --mode=mode
```

Managing procfs entries (3/3)



```
struct proc_dir_entry* proc_mkdir(const char* name,  
    struct proc_dir_entry* parent)
```

- It creates a directory *name* in *parent*

```
void remove_proc_entry(const char*name, struct  
    proc_dir_entry* parent)
```

- It removes the entry *name* in the directory *parent*
 - ▶ Entries are removed by name
 - ▶ This function doesn't recursively remove entries
 - Be sure to free the *data* entry from the struct `proc_dir_entry`

Communicating with userland



- Procfs works with callback functions for files

```
struct proc_dir_entry* entry;
```

```
entry -> read_proc = read_proc_foo;
```

```
entry -> write_proc = write_proc_foo;
```

- `create_proc_read_entry` creates and initialize the procfs entry in one single call

Reading and writing



```
int read_func(char* page, char** start, off_t off, int
count, int* eof, void *data)
```

- ▶ It is used to read data from the kernel
- ▶ Writes its information to *page*
 - It should start writing at *off* and write at most *count* bytes
 - *eof* used to signal that the end of the file has been reached (1 in the memory location it points to)
 - It returns the number of bytes written into the page

```
int write_func(struct file* file, const char* buffer,
long count, void *data)
```

- ▶ It should read *count* bytes at maximum from the *buffer*
- ▶ Buffer doesn't live in kernel memory space
 - `copy_from_user` to copy it

A single callback



- Useful when a large number of almost identical files is used
 - ▶ It distinguishes between files using the *data* field in struct `proc_dir_entry`
 - It is a `void*`, so it can be initialized with anything
 - Be sure to free data when removing the `procfs` entry

```
struct proc_dir_entry* entry;  
struct my_file_data* file_data;  
file_data = kmalloc(sizeof(struct my_file_data),  
                    GFP_KERNEL);  
entry-> data = file_data
```

Inside the callback



```
int foo_read_func(char* page, char** start, off_t off,
    int count, int* eof, void* data)
{
    int len;
    if (data == file_data) {
        /* special case */
    } else {
        /* normal processing */
    }
    return len;
}
```

Tips and tricks



- `create_proc_read_entry` to create and initialize an entry with a single call
- If `procfs` is being used from within a module, be sure to set the owner field in the struct `proc_dir_entry` to `THIS_MODULE`

```
struct proc_dir_entry* entry;  
entry->owner = THIS_MODULE;
```
- Change the mode and/or ownership of the entry

```
struct proc_dir_entry* entry;  
entry->uid = 0;  
entry->mode = S_IWUSR | S_IRUSR | S_IRGRP  
            | S_IROTH;
```