



Parallel Processing

Prof. William Fornaciari

fornacia@elet.polimi.it
www.elet.polimi.it/~fornacia

Multiple Processor Organization



- Single instruction, single data stream - SISD
- Single instruction, multiple data stream - SIMD
- Multiple instruction, single data stream - MISD
- Multiple instruction, multiple data stream- MIMD

Single Instruction, Single Data Stream - SISD



- Single processor
- Single instruction stream
- Data stored in single memory
- Uni-processor

Single Instruction, Multiple Data Stream - SIMD



- Single machine instruction
- Controls simultaneous execution
- Number of processing elements
- Lockstep basis
- Each processing element has associated data memory
- Each instruction executed on different set of data by different processors
- Vector and array processors

Multiple Instruction, Single Data Stream - MISD



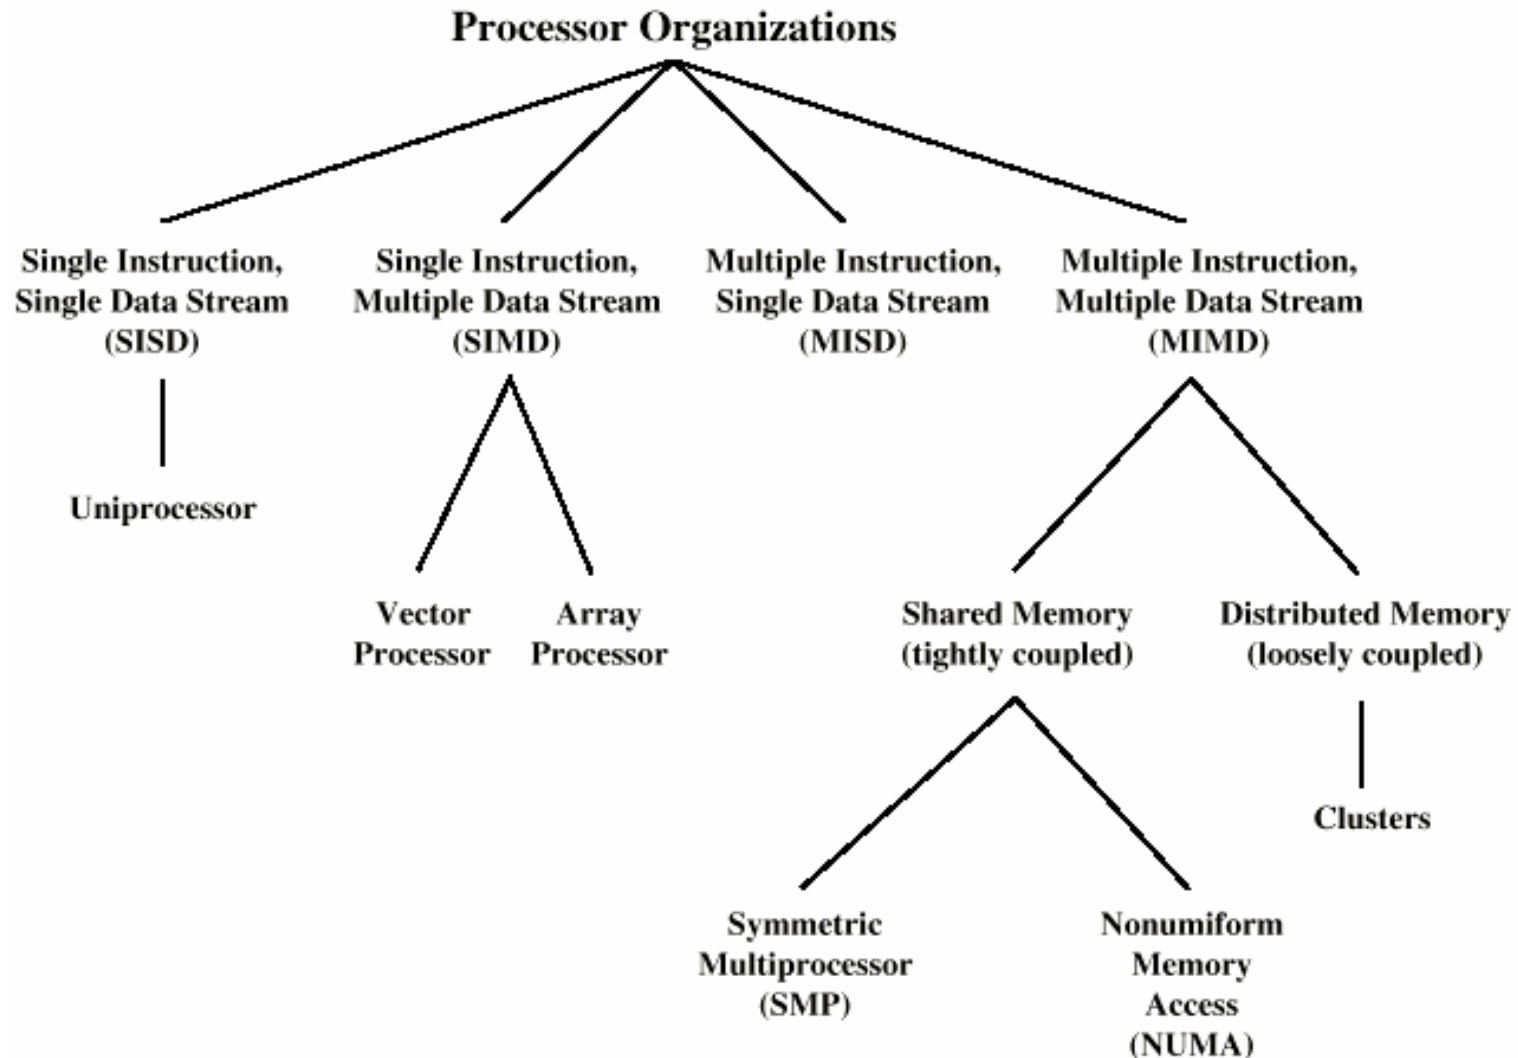
- Sequence of data
- Transmitted to set of processors
- Each processor executes different instruction sequence
- Never been implemented

Multiple Instruction, Multiple Data Stream- MIMD



- Set of processors
- Simultaneously execute different instruction sequences
- Different sets of data
- SMPs, clusters and NUMA systems

Taxonomy of Parallel Processor Architectures



MIMD - Overview



- General purpose processors
- Each can process all instructions necessary
- Further classified by method of processor communication

Tightly Coupled - SMP



- Processors share memory
- Communicate via that shared memory
- Symmetric Multiprocessor (SMP)
 - ▶ Share single memory or pool
 - ▶ Shared bus to access memory
 - ▶ Memory access time to given area of memory is approximately the same for each processor

Tightly Coupled - NUMA




- NonUniform Memory Access
- Access times to different regions of memory may differ

Loosely Coupled - Clusters




- Collection of independent uniprocessors or SMPs
- Interconnected to form a cluster
- Communication via fixed path or network connections

Parallel Organizations - SISD




Parallel Organizations - SIMD




Parallel Organizations - MIMD

Shared Memory



Parallel Organizations - MIMD

Distributed Memory



Symmetric Multiprocessors




- A stand alone computer with the following characteristics
 - ▶ Two or more similar processors of comparable capacity
 - ▶ Processors share same memory and I/O
 - ▶ Processors are connected by a bus or other internal connection
 - ▶ Memory access time is approximately the same for each processor
 - ▶ All processors share access to I/O
 - Either through same channels or different channels giving paths to same devices
 - ▶ All processors can perform the same functions (hence symmetric)
 - ▶ System controlled by integrated operating system
 - providing interaction between processors
 - Interaction at job, task, file and data element levels

SMP Advantages



- Performance
 - ▶ If some work can be done in parallel
- Availability
 - ▶ Since all processors can perform the same functions, failure of a single processor does not halt the system
- Incremental growth
 - ▶ User can enhance performance by adding additional processors
- Scaling
 - ▶ Vendors can offer range of products based on number of processors

Block Diagram of Tightly Coupled Multiprocessor



Organization Classification



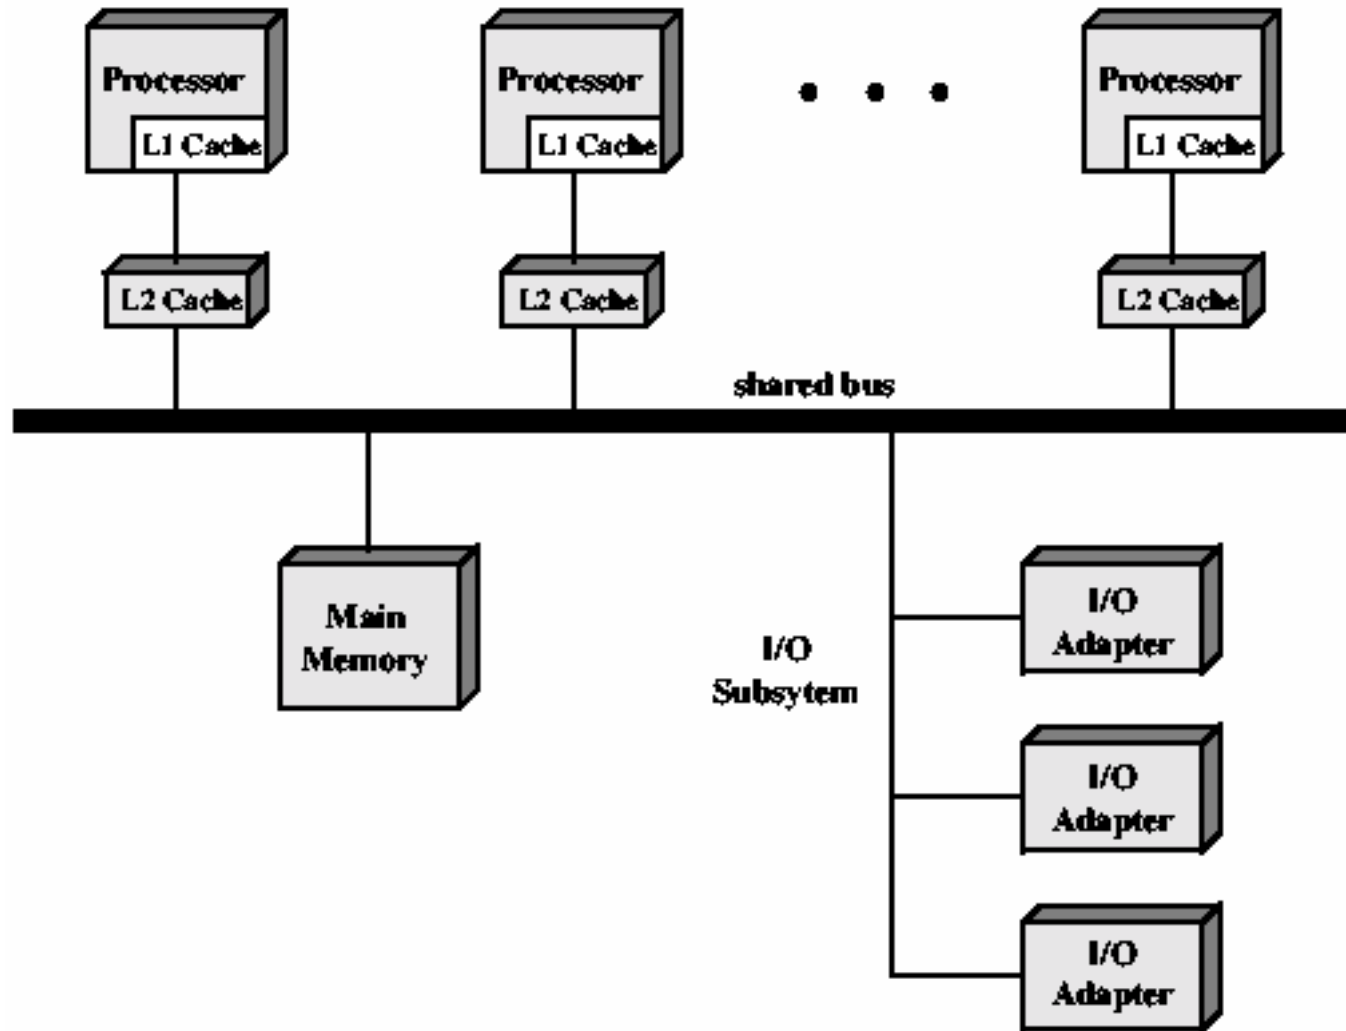
- Time shared or common bus
- Multiport memory
- Central control unit

Time Shared Bus



- Simplest form
- Structure and interface similar to single processor system
- Following features provided
 - ▶ Addressing - distinguish modules on bus
 - ▶ Arbitration - any module can be temporary master
 - ▶ Time sharing - if one module has the bus, others must wait and may have to suspend
- Now have multiple processors as well as multiple I/O modules

Shared Bus



Time Share Bus - Advantages



- Simplicity
- Flexibility
- Reliability

Time Share Bus - Disadvantage



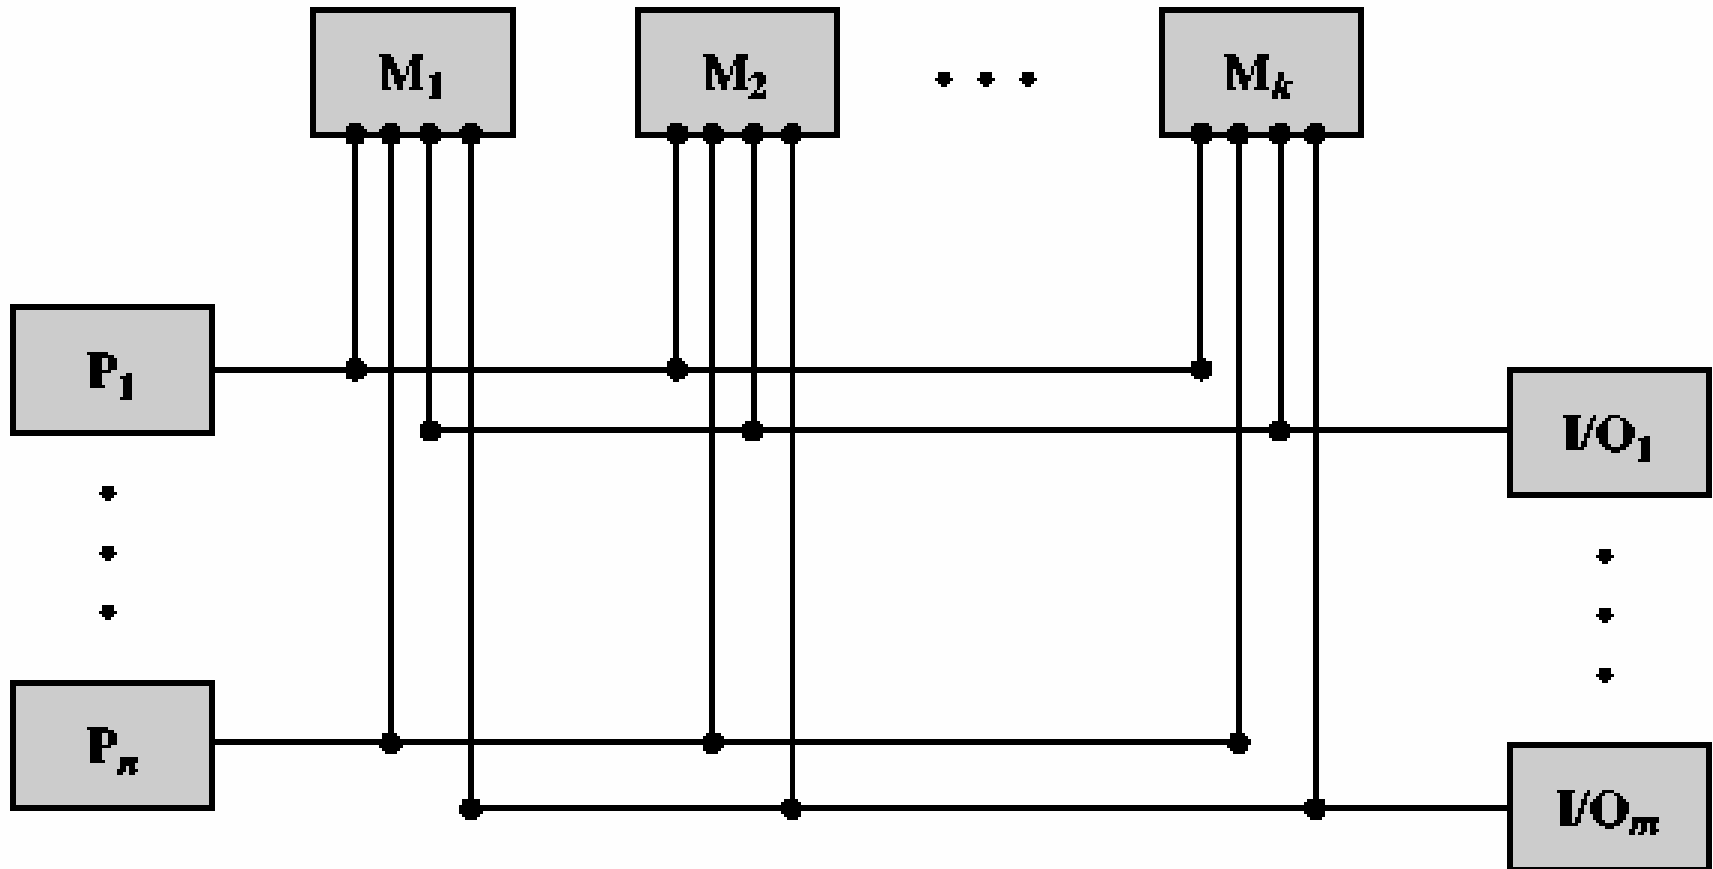
- Performance limited by bus cycle time
- Each processor should have local cache
 - ▶ Reduce number of bus accesses
- Leads to problems with cache coherence
 - ▶ Solved in hardware - see later

Multiport Memory



- Direct independent access of memory modules by each processor
- Logic required to resolve conflicts
- Little or no modification to processors or modules required

Multiport Memory Diagram



Multiport Memory - Advantages and Disadvantages



- More complex
 - ▶ Extra login in memory system
- Better performance
 - ▶ Each processor has dedicated path to each module
- Can configure portions of memory as private to one or more processors
 - ▶ Increased security
- Write through cache policy

Central Control Unit



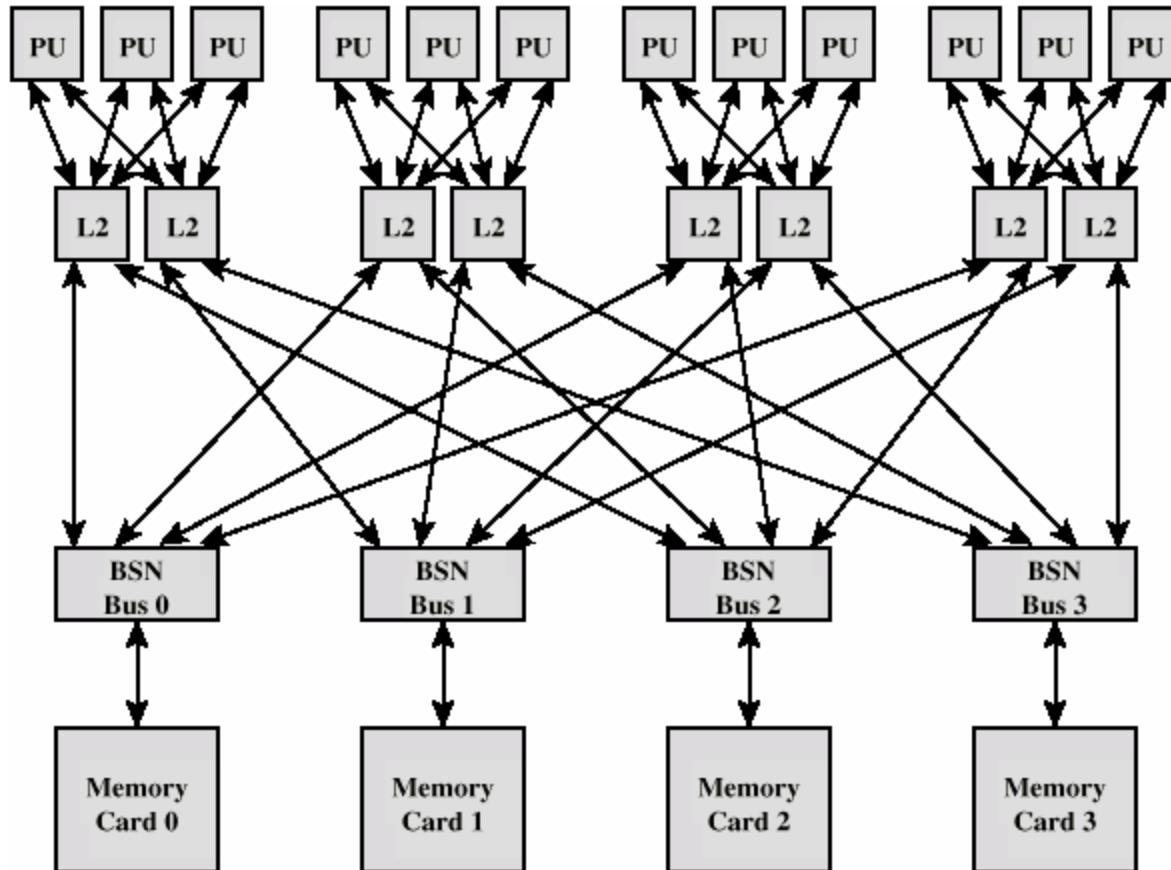
- Funnels separate data streams between independent modules
- Can buffer requests
- Performs arbitration and timing
- Pass status and control
- Perform cache update alerting
- Interfaces to modules remain the same
- e.g. IBM S/370

Operating System Issues



- Simultaneous concurrent processes
- Scheduling
- Synchronization
- Memory management
- Reliability and fault tolerance

IBM S/390 Mainframe SMP



S/390 - Key components



- Processor unit (PU)
 - ▶ CISC microprocessor
 - ▶ Frequently used instructions hard wired
 - ▶ 64k L1 unified cache with 1 cycle access time
- L2 cache
 - ▶ 384k
- Bus switching network adapter (BSN)
 - ▶ Includes 2M of L3 cache
- Memory card
 - ▶ 8G per card

Cache Coherence and MESI Protocol



- Problem - multiple copies of same data in different caches
- Can result in an inconsistent view of memory
- Write back policy can lead to inconsistency
- Write through can also give problems unless caches monitor memory traffic

Software Solutions



- Compiler and operating system deal with problem
- Overhead transferred to compile time
- Design complexity transferred from hardware to software
- However, software tends to make conservative decisions
 - ▶ Inefficient cache utilization
- Analyze code to determine safe periods for caching shared variables

Hardware Solution



- Cache coherence protocols
- Dynamic recognition of potential problems
- Run time
- More efficient use of cache
- Transparent to programmer
- Directory protocols
- Snoopy protocols

Directory Protocols



- Collect and maintain information about copies of data in cache
- Directory stored in main memory
- Requests are checked against directory
- Appropriate transfers are performed
- Creates central bottleneck
- Effective in large scale systems with complex interconnection schemes

Snoopy Protocols



- Distribute cache coherence responsibility among cache controllers
- Cache recognizes that a line is shared
- Updates announced to other caches
- Suited to bus based multiprocessor
- Increases bus traffic

Write Invalidate




- Multiple readers, one writer
- When a write is required, all other caches of the line are invalidated
- Writing processor then has exclusive (cheap) access until line required by another processor
- Used in Pentium II and PowerPC systems
- State of every line is marked as modified, exclusive, shared or invalid
- MESI

Write Update




- Multiple readers and writers
- Updated word is distributed to all other processors
- Some systems use an adaptive mixture of both solutions

MESI State Transition Diagram



(a) Line in cache at initiating processor



(b) Line in snooping cache

RH	Read hit		Dirty line copyback
RMS	Read miss, shared		Invalidate transaction
RME	Read miss, exclusive		Read-with-intent-to-modify
WH	Write hit		Cache line fill
WM	Write miss		
SHR	Snoop hit on read		
SHW	Snoop hit on write or read-with-intent-to-modify		

Clusters




- Alternative to SMP
 - High performance
 - High availability
 - Server applications
-
- A group of interconnected whole computers
 - Working together as unified resource
 - Illusion of being one machine
 - Each computer called a node

Cluster Benefits




- Absolute scalability
- Incremental scalability
- High availability
- Superior price/performance

Cluster Configurations - Standby Server, No Shared Disk



Cluster Configurations - Shared Disk



Operating Systems Design Issues



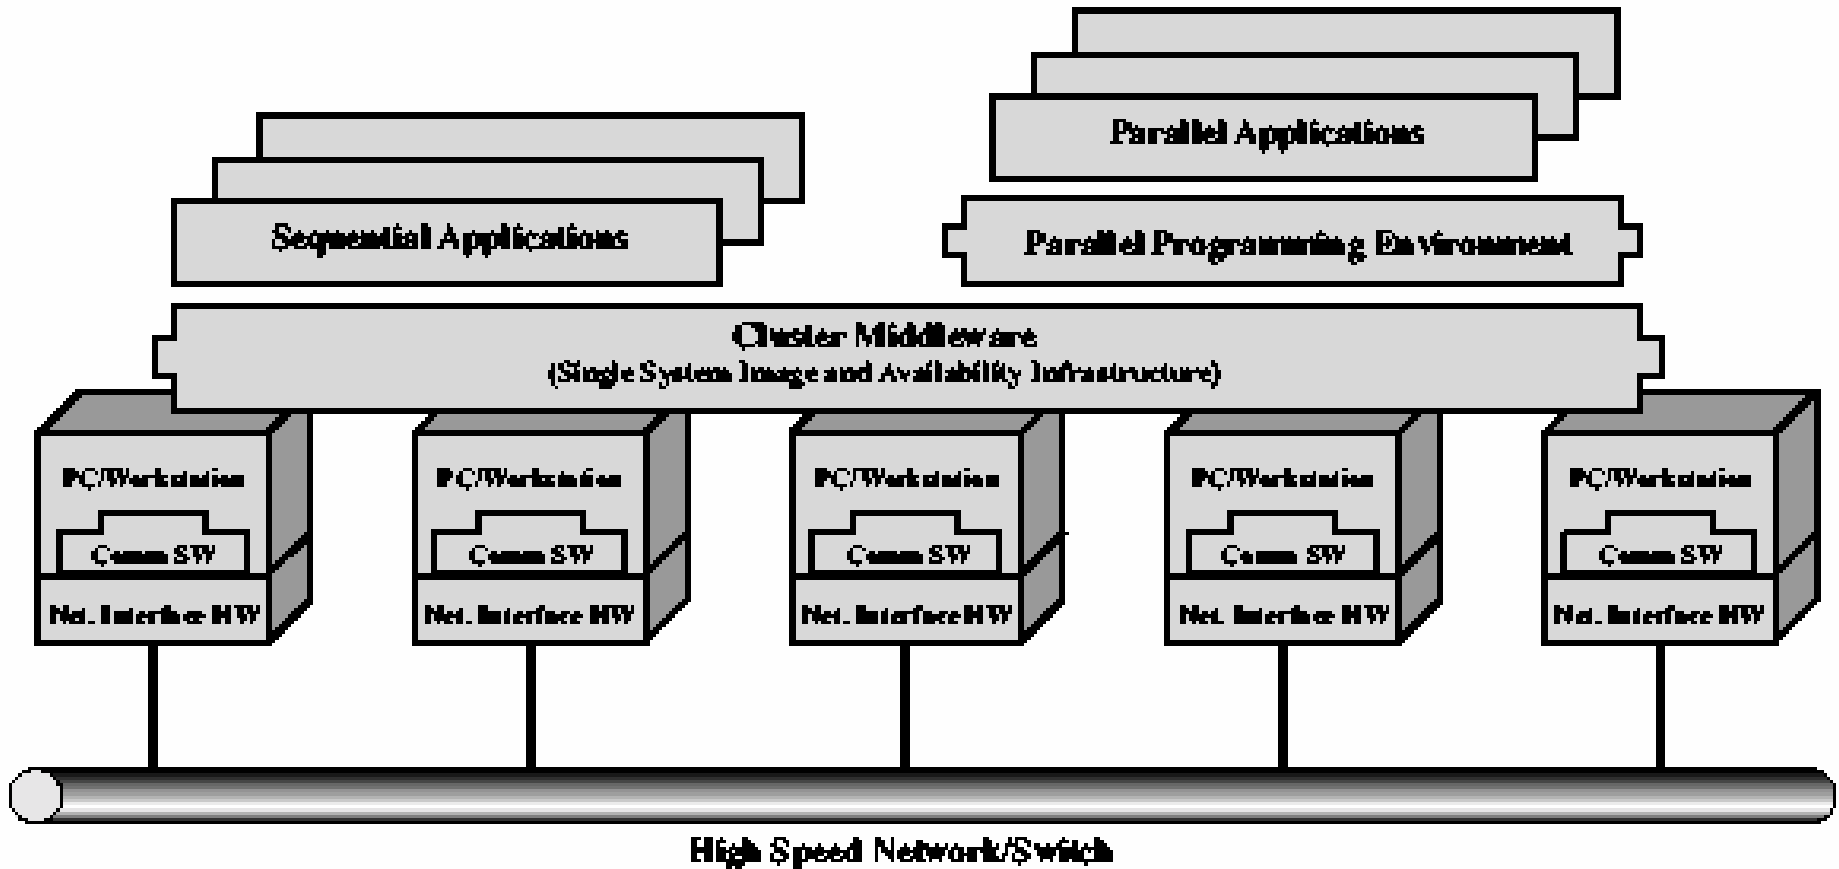
- Failure Management
 - ▶ High availability
 - ▶ Fault tolerant
 - ▶ Failover
 - Switching applications & data from failed system to alternative within cluster
 - ▶ Failback
 - Restoration of applications and data to original system
 - After problem is fixed
- Load balancing
 - ▶ Incremental scalability
 - ▶ Automatically include new computers in scheduling
 - ▶ Middleware needs to recognise that processes may switch between machines

Parallelizing



- Single application executing in parallel on a number of machines in cluster
 - ▶ Compiler
 - Determines at compile time which parts can be executed in parallel
 - Split off for different computers
 - ▶ Application
 - Application written from scratch to be parallel
 - Message passing to move data between nodes
 - Hard to program
 - Best end result
 - ▶ Parametric computing
 - If a problem is repeated execution of algorithm on different sets of data
 - e.g. simulation using different scenarios
 - Needs effective tools to organize and run

Cluster Computer Architecture



Cluster Middleware



- Unified image to user
 - ▶ Single system image
- Single point of entry
- Single file hierarchy
- Single control point
- Single virtual networking
- Single memory space
- Single job management system
- Single user interface
- Single I/O space
- Single process space
- Checkpointing
- Process migration

Cluster v. SMP



- Both provide multiprocessor support to high demand applications.
- Both available commercially
 - ▶ SMP for longer
- SMP:
 - ▶ Easier to manage and control
 - ▶ Closer to single processor systems
 - Scheduling is main difference
 - Less physical space
 - Lower power consumption
- Clustering:
 - ▶ Superior incremental & absolute scalability
 - ▶ Superior availability
 - Redundancy



Nonuniform Memory Access (NUMA)

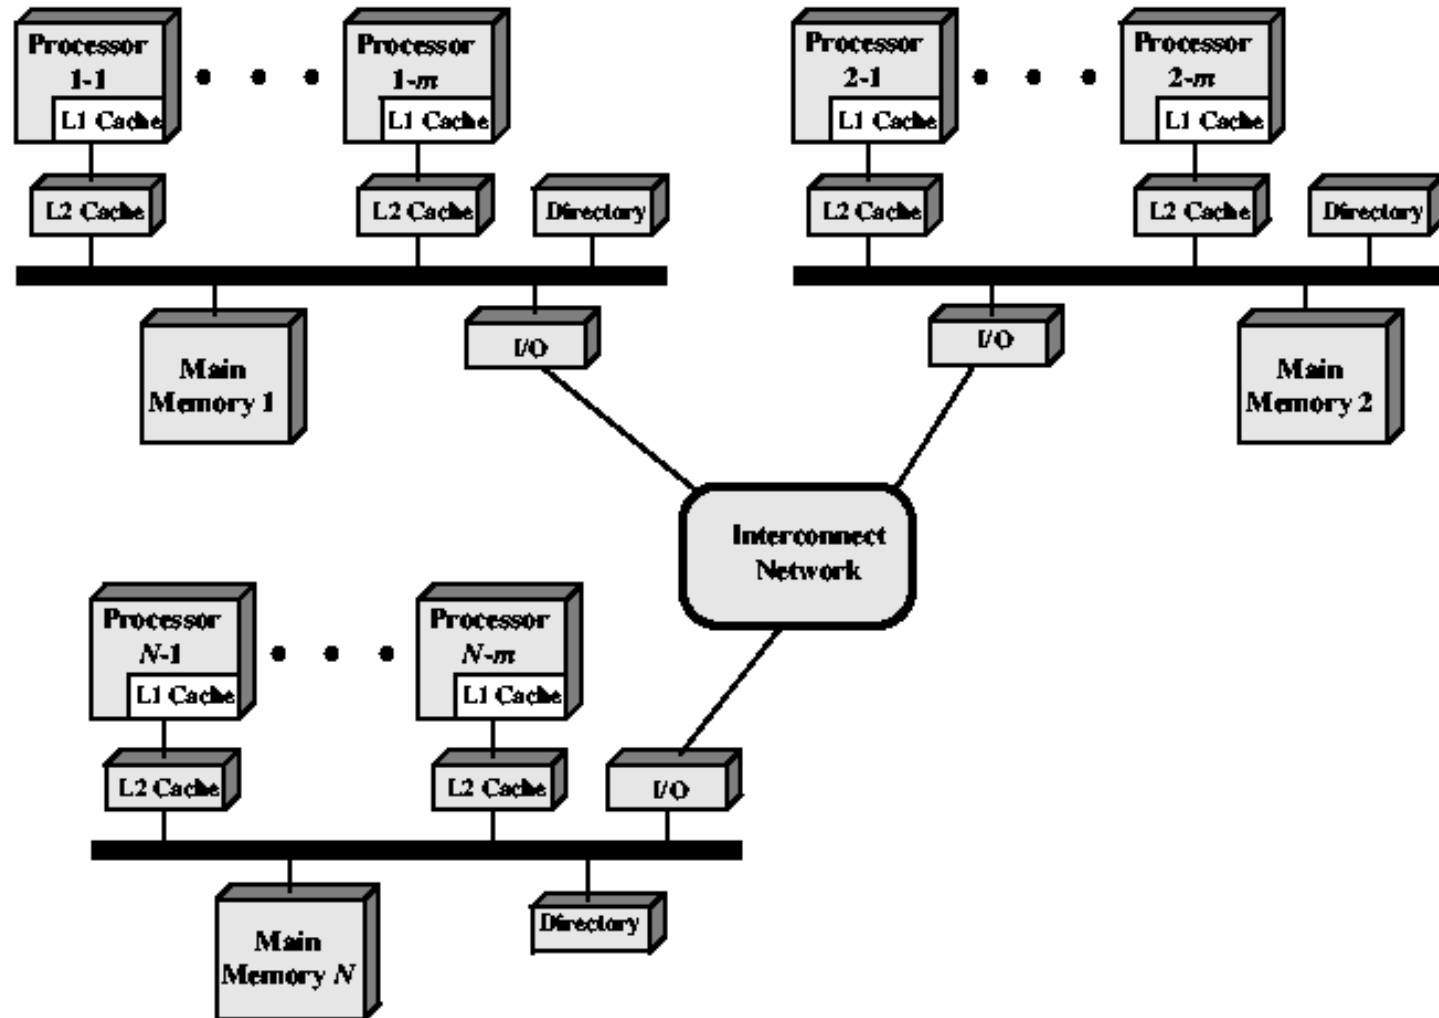
- Alternative to SMP & clustering
- Uniform memory access
 - ▶ All processors have access to all parts of memory
 - Using load & store
 - ▶ Access time to all regions of memory is the same
 - ▶ Access time to memory for different processors same
 - ▶ As used by SMP
- Nonuniform memory access
 - ▶ All processors have access to all parts of memory
 - Using load & store
 - ▶ Access time of processor differs depending on region of memory
 - ▶ Different processors access different regions of memory at different speeds
- Cache coherent NUMA
 - ▶ Cache coherence is maintained among the caches of the various processors
 - ▶ Significantly different from SMP and clusters

Motivation



- SMP has practical limit to number of processors
 - ▶ Bus traffic limits to between 16 and 64 processors
- In clusters each node has own memory
 - ▶ Apps do not see large global memory
 - ▶ Coherence maintained by software not hardware
- NUMA retains SMP flavour while giving large scale multiprocessing
 - ▶ e.g. Silicon Graphics Origin NUMA 1024 MIPS R10000 processors
- Objective is to maintain transparent system wide memory while permitting multiprocessor nodes, each with own bus or internal interconnection system

CC-NUMA Organization



CC-NUMA Operation



- Each processor has own L1 and L2 cache
- Each node has own main memory
- Nodes connected by some networking facility
- Each processor sees single addressable memory space
- Memory request order:
 - ▶ L1 cache (local to processor)
 - ▶ L2 cache (local to processor)
 - ▶ Main memory (local to node)
 - ▶ Remote memory
 - Delivered to requesting (local to processor) cache
- Automatic and transparent

Memory Access Sequence



- Each node maintains directory of location of portions of memory and cache status
- e.g. node 2 processor 3 (P2-3) requests location 798 which is in memory of node 1
 - ▶ P2-3 issues read request on snoopy bus of node 2
 - ▶ Directory on node 2 recognises location is on node 1
 - ▶ Node 2 directory requests node 1's directory
 - ▶ Node 1 directory requests contents of 798
 - ▶ Node 1 memory puts data on (node 1 local) bus
 - ▶ Node 1 directory gets data from (node 1 local) bus
 - ▶ Data transferred to node 2's directory
 - ▶ Node 2 directory puts data on (node 2 local) bus
 - ▶ Data picked up, put in P2-3's cache and delivered to processor

Cache Coherence



- Node 1 directory keeps note that node 2 has copy of data
- If data modified in cache, this is broadcast to other nodes
- Local directories monitor and purge local cache if necessary
- Local directory monitors changes to local data in remote caches and marks memory invalid until writeback
- Local directory forces writeback if memory location requested by another processor

NUMA Pros & Cons



- Effective performance at higher levels of parallelism than SMP
- No major software changes
- Performance can breakdown if too much access to remote memory
 - ▶ Can be avoided by:
 - L1 & L2 cache design reducing all memory access
 - Need good temporal locality of software
 - Good spatial locality of software
 - Virtual memory management moving pages to nodes that are using them most
- Not transparent
 - ▶ Page allocation, process allocation and load balancing changes needed
- Availability?

Vector Computation



- Maths problems involving physical processes present different difficulties for computation
 - ▶ Aerodynamics, seismology, meteorology
 - ▶ Continuous field simulation
- High precision
- Repeated floating point calculations on large arrays of numbers
- Supercomputers handle these types of problem
 - ▶ Hundreds of millions of flops
 - ▶ \$10-15 million
 - ▶ Optimised for calculation rather than multitasking and I/O
 - ▶ Limited market
 - Research, government agencies, meteorology
- Array processor
 - ▶ Alternative to supercomputer
 - ▶ Configured as peripherals to mainframe & mini
 - ▶ Just run vector portion of problems

Vector Addition Example



$$\begin{bmatrix} 1.5 \\ 7.1 \\ 6.9 \\ 100.5 \\ 0 \\ 59.7 \end{bmatrix} + \begin{bmatrix} 2.0 \\ 39.7 \\ 1000.003 \\ 11 \\ 21.1 \\ 19.7 \end{bmatrix} = \begin{bmatrix} 3.5 \\ 46.8 \\ 1006.903 \\ 111.5 \\ 21.1 \\ 79.4 \end{bmatrix}$$

$$A + B = C$$

Approaches



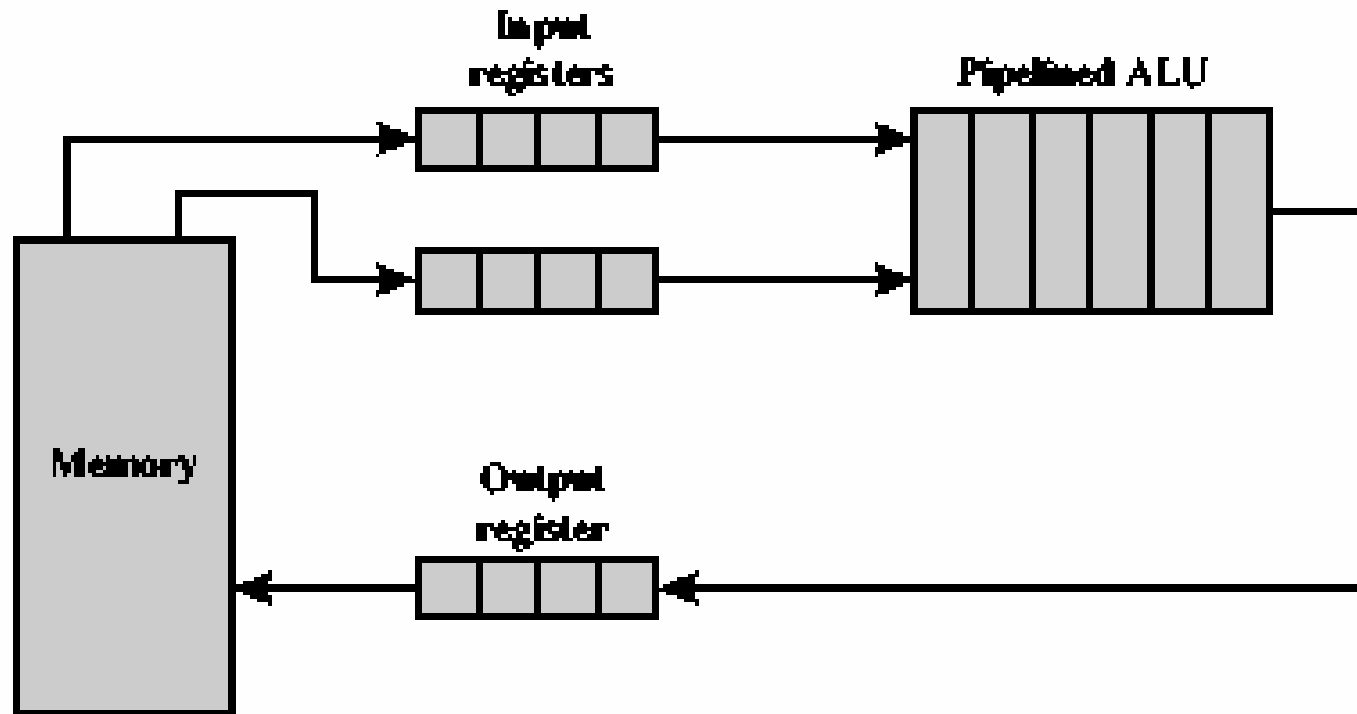
- General purpose computers rely on iteration to do vector calculations
- In example this needs six calculations
- Vector processing
 - ▶ Assume possible to operate on one-dimensional vector of data
 - ▶ All elements in a particular row can be calculated in parallel
- Parallel processing
 - ▶ Independent processors functioning in parallel
 - ▶ Use FORK N to start individual process at location N
 - ▶ JOIN N causes N independent processes to join and merge following JOIN
 - O/S Co-ordinates JOINS
 - Execution is blocked until all N processes have reached JOIN

Processor Designs



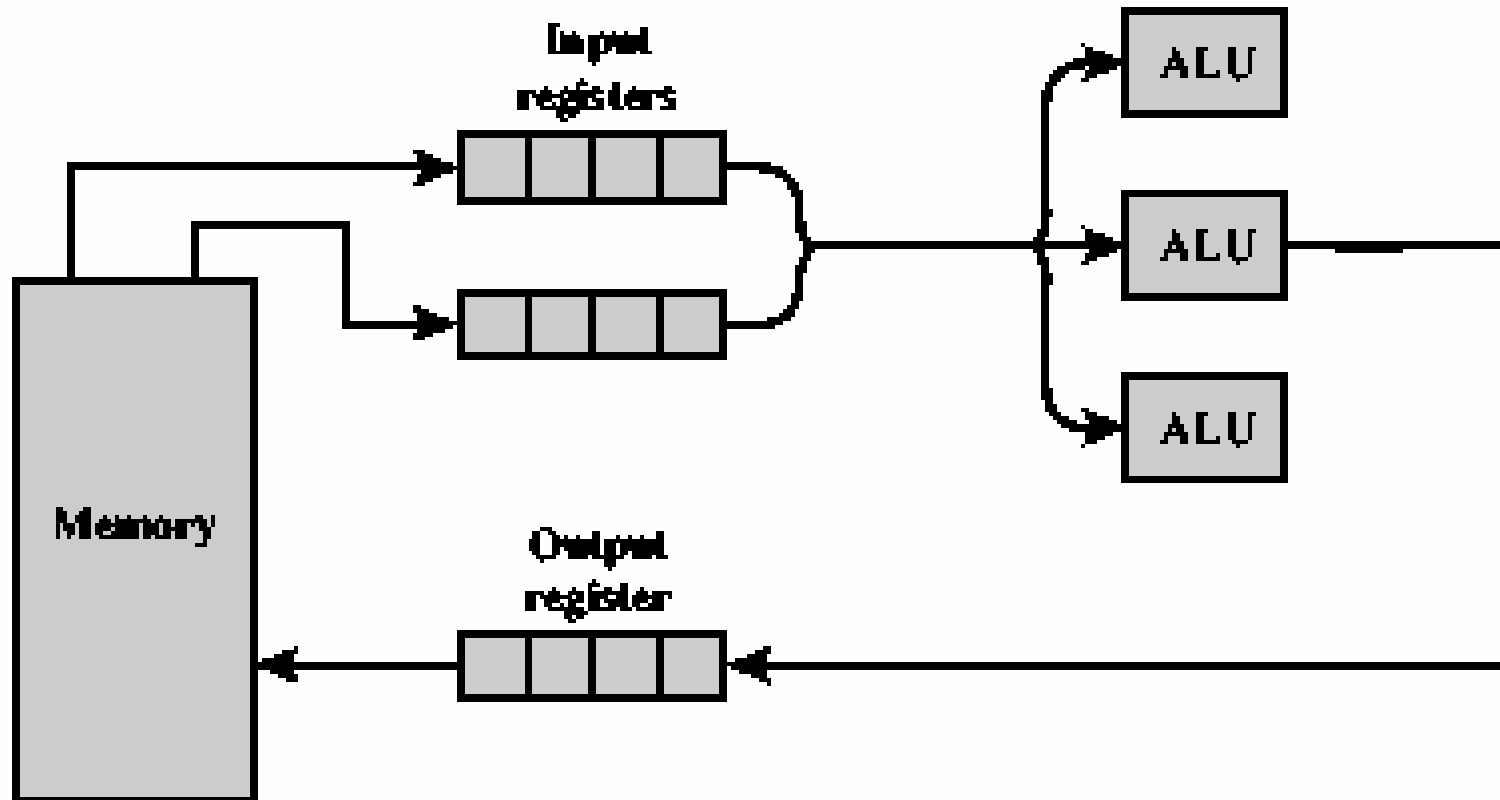
- Pipelined ALU
 - ▶ Within operations
 - ▶ Across operations
- Parallel ALUs
- Parallel processors

Approaches to Vector Computation



(a) Pipelined ALU

Approaches to Vector Computation



(b) Parallel ALUs

Chaining




- Cray Supercomputers
- Vector operation may start as soon as first element of operand vector available and functional unit is free
- Result from one functional unit is fed immediately into another
- If vector registers used, intermediate results do not have to be stored in memory

Computer Organizations



Single Control Unit



Multiple Control Units

