



# Real Time Operating Systems

RT features in commercial products

Lecturer:

Prof. William Fornaciari

Politecnico di Milano

[fornacia@elet.polimi.it](mailto:fornacia@elet.polimi.it)

[www.elet.polimi.it/~fornacia](http://www.elet.polimi.it/~fornacia)

# Outline

---



- *Introduction to Real Time Systems*
- *The role of a RTOS*
- *Time constraints*
- *The problem of scheduling RT activities*
- *Classification of scheduling strategies*
- *Review of RT scheduling approaches*
- Linux, Unix SVR4, Win2000
- VxWork real-time features
- Windows CE real-time features



- In addition to traditional UNIX-like scheduling, two scheduling classes for soft-RT are added
  - ▶ SCHED\_FIFO: First-in-first-out real-time threads
  - ▶ SCHED\_RR: Round-robin real-time threads
  - ▶ SCHED\_OTHER: Other, non-real-time threads
- Within each class multiple priorities may be used, with priority for RT classes higher than for SCHED\_OTHER class



# FIFO threads: rules

---

- The system can interrupt a FIFO thread only if:
  - ▶ Another FIFO thread of higher priority becomes ready
  - ▶ The running FIFO thread blocks waiting for an event
  - ▶ The running FIFO thread voluntarily gives up the processor (sys call *sched\_yield*)
- An interrupted FIFO thread is suspended in the queue with its priority
- When a FIFO thread becomes ready, it can preempt current thread if its priority is higher
- SCHED\_RR is similar to FIFO, except for the addition of a time quote associated with each thread

# RR and OTHER threads: scheduling rules

---



- SCHED\_RR is similar to FIFO, except for the addition of a time quote associated with each thread
- When SCHED\_RR thread executed for its quote, it is suspended and a RT thread of  $\geq$  priority is selected
- A SCHED\_OTHER thread can only execute if there are no RT threads ready to execute, using traditional UNIX scheduling



# FIFO vs RR: Example

- Three processes with relative priority, where all waiting threads are ready to execute and no higher priority thread is awakened while a thread is running

A	minimum
B	middle
C	middle
D	maximum

(a) Relative thread priorities

D → B → C → A →

(b) Flow with FIFO scheduling

D → B → C → B → C → A →

(c) Flow with RR scheduling

# UNIX SVR4 Scheduling

---



- Complete overhaul of the UNIX scheduler
- Highest preference to real-time processes
- Next-highest to kernel-mode processes
- Lowest preference to other user-mode processes (time shared processes)

# UNIX SVR4: major modifications

---




- Addition of a preemptable static priority scheduler, with 160 priorities partitioned in three classes
  - ▶ Real Time (159-100)
    - Processes are selected to run before any Kernel or TS process
  - ▶ Kernel (99-60)
    - Processes are selected to run before TS processes, but must defer to RT processes
  - ▶ Time Shared (59-0)
    - Intended for user other than no RT applications
- Insertion of preemption points
  - ▶ Since basic kernel is not preemptive, safe points to interrupt kernel and schedule new processes have been identified
  - ▶ Safe place: region of code where all kernel data structures are either updated and consistent or locked via a semaphore





# SVR4 Scheduling: Dispatch Queues

- A dispatch queue is associated with each priority level and processes at a given level execute RR
- *dqactmap*: bit-map vector, one bit per priority level
  - ▶ 1 for nonempty queues
  - ▶ The scheduler checks *dqactmap* and dispatch a ready process from the highest priority nonempty queue
- In preemption points, the kernel checks a flag (*kprunrun*), indicating the existence of RT processes in the Ready state. If it exists, a preemption will



# SVR4 Scheduling: Dispatch Queues

---



- The priority of TS class is variable
  - ▶ *Reduce* if it uses time quantum
  - ▶ *Rise* if Process is blocked on event
- Variable time quantum for TS class
  - ▶ 100 ms for priority 0
  - ▶ 10 ms for priority 59
- Each RT process has fixed priority and fixed time quantum

# Windows 2000 Scheduling


---



- Flexible priority-driven preemptive scheduler
  - ▶ RR scheduling within each level
  - ▶ For some levels variable priority based on current thread activity
- Priorities organized into two bands or classes, each consisting of 16 priority levels
  - ▶ Real-time
  - ▶ Variable
- RT threads have precedence over other threads
- Lower priority thread is preempted by higher priority



# W2k: Thread Dispatching Priorities



# Windows 2000 Priorities

---




- RT priority class
  - ▶ all threads have fixed priority, each level is RR
- Variable priority class
  - ▶ Priority can move up/down during thread's lifetime, but not move to RT class
  - ▶ Initial priority of a thread
    - Process base priority (0..15) attribute of process object
    - Thread base priority. Relative to that of the process
  - ▶ The initial priority of a thread is in a range around that of process (base process + base thread)
  - ▶ The dynamic priority may never fall below lower range of thread base priority and over 15
  - ▶ W2K executive tends to raise I/O bound (interactive) threads



# W2K Priorities Relationship: Example

- Process base priority is 4
- Each threads objects have base priority 2-6
- Dynamic priority may fluctuate from 2 to 15



# Outline

---



- VxWork Real-Time Features
  - ▶ Introduction
  - ▶ Kernel
  - ▶ Memory Architecture
  - ▶ Real-Time Measurement Tools

# WxWorks Definition

---



- Unix and Windows are not appropriate for real time systems
- VxWorks has to be used with another operating system: it handles the critical real-time chores, while the host machine is used for program development and non real-time applications





- All applications consist of a tasks
- A task is the basic unit to which the operating system allocates processor time
- Tasks are essential to organize applications into independent, though cooperating, programs

# Processes

---



- WxWorks creates a context for each task
- A task's context includes
  - ▶ The CPU register
  - ▶ A stack for dynamic variables and function calls
  - ▶ I/O assignments for standard input, output and error
  - ▶ A delay timer
  - ▶ A timeslice timer
  - ▶ Signal handlers

# Scheduling

---



- VxWorks uses, by default, a priority-base preemptive scheduling algorithm
- 256 priority levels:
  - 0: highest priority level for time critical activities
  - 1-254
  - 255: lower priority level
- Higher priority levels are assigned to real-time applications
- Preemptive scheduling can be augmented with round-robin scheduling

# Scheduling

---



- Threads with higher priority are scheduled to run first
- Threads with the same priority run in a round-robin fashion
- Lower priority threads do not run until higher priority threads have finished
- Lower priority threads can be suspended by higher priority threads



- Time Slice:
  - ▶ Threads run for a time slice
  - ▶ A run-time counter is kept for each task, and incremented on every clock cycle
  - ▶ When the time slice is completed, the task is placed at the tail of the queue of tasks at its priority
  - ▶ If a task is preempted by a higher priority task during its interval, its run-time count is saved

# Interrupts

---



- VxWorks supplies routines for handling hardware interrupts and software traps without having to resort to assembly language coding
- Routines are provided to connect C routines to hardware interrupt vectors, and to manipulate the processor interrupt level

# Interrupts

---



- For the fastest possible response to interrupts, interrupt service routines (ISRs) in VxWorks run in a special context outside of any task's context
- Interrupt handling involves no task context switch
- All ISRs use the same interrupt stack

# Interrupts

---



- When an interrupt occurs, the connected C function is called at interrupt level with the specified argument
- When the interrupt handling is finished, the connected function returns
- A routine connected to an interrupt in this way is called an interrupt service routine (ISR)



# Intertask Communications

---



- Intertask communications can be realized through the following mechanism:
  - ▶ Shared Data Structures (global variables, linear buffers, ring buffers, linked lists, and pointers)
  - ▶ Mutual exclusion
  - ▶ Semaphores
  - ▶ Message Queues
  - ▶ Pipes
  - ▶ Signals

# Memory Architecture

---



- VxWorks supplies management facility for dynamically allocating, freeing, and reallocating blocks of memory from a memory pool
- Blocks of arbitrary size can be allocated
- The size of the memory pool can be specified
- VxWorks manages several separate memory pools

# Virtual Memory

---



- VxWorks provides two levels of virtual memory support
- The basic level is bundled with VxWorks and provides caching on a per-page basis
- The full level is unbundled, and requires an optional component, VxVMI
- VxVMI provides:
  - ▶ Write protection of text segments
  - ▶ VxWorks exception vector table
  - ▶ Architecture-independent interface to the CPU's memory management unit (MMU)

# Real-Time Measurement Tools

---



- VxWorks provides various timing facilities to understand and optimize the performance of a real-time system
- VxWorks execution timer can time any subroutine or group of subroutines
- The timer can also repeatedly execute a group of routines until the time of a single iteration is known to a reasonable accuracy



## Spy utility

---

- There is also a spy utility which provides CPU utilization information for each task:
  - ▶ CPU time consumed
  - ▶ Time spent at interrupt level
  - ▶ Amount of idle time
- Time is displayed in ticks and in percentages

# Outline

---



- Windows CE Real-Time Features
  - ▶ Introduction
  - ▶ Kernel
  - ▶ Memory Architecture
  - ▶ Real-Time Measurement Tools



# Windows CE Definition

---

- Windows CE is the modular real-time embedded operating system for small footprint and mobile 32-bit intelligent and connected device. (Microsoft)



- All applications consist of a process and one or more threads
- Each process starts with a single thread: primary thread
- A thread is the basic unit to which the operating system allocates processor time



# Processes

---



- Windows CE can run up to 32 processes at one time
- When a Windows CE-based device starts:
  - ▶ The OS creates a single 4-GB virtual address space
  - ▶ It is divided into 33 slots of 32 MB each
  - ▶ When a process start the OS assigns it one slot
  - ▶ Slot zero is for the currently running process



- Threads are implemented in kernel space
- Problems to avoid in multi-thread applications:
  - ▶ *Deadlocks*: when all threads are in the blocked state, Windows CE enters in the idle mode
  - ▶ *Race conditions*: thread sharing common resources must coordinate their work by using a method of synchronization

# Scheduling

---



- Windows CE uses a priority-base time-slice algorithm
- 256 priority levels:
  - 0: highest priority level for time critical activities
  - 1-254
  - 255: lower priority level
- Higher priority levels are assigned to real-time applications
- Process priority classes are not supported

# Scheduling

---



- Threads with higher priority are scheduled to run first
- Threads with the same priority run in a round-robin fashion
- Lower priority threads do not run until higher priority threads have finished
- Lower priority threads can be suspended by higher priority threads



- Time Quantum:
  - ▶ Threads run for a time slice
  - ▶ By setting the time quantum of a thread to zero, a round-robin scheduling algorithm can change to a run-to-completion algorithm
  - ▶ Threads set to run-to-completion can be preempted only by higher priority threads
  - ▶ The default quantum is 100 milliseconds

# Scheduling

---



- Priority Inversion:
  - ▶ Thread priorities are fixed and do not change
  - ▶ Priority inversion occurs when a high priority thread is blocked by a lower priority thread
  - ▶ Windows CE guarantees the handling of priority inversion only to a depth of one level

# Interrupts

---



- The use of interrupts requires balancing performance against ease of use
- Interrupt processing is split into two steps:
  - ▶ ISR: interrupt service routine
    - It directs the kernel to launch the appropriate IST
  - ▶ IST: interrupt service thread
    - It contains the interrupt handling



- Nested Interrupts:
  - ▶ Windows CE handling of nested ISR calls:
    - The kernel disables all IRQ lines for the same and lower priorities
    - If a higher priority IRQ arrives, the current ISR state is stored
    - The kernel calls the higher-priority ISR
    - When the request is completed, the kernel loads the original ISR



# Interrupts

---




- Interrupt latencies:
  - ▶ They are bounded for threads locked in memory if paging does not occur
  - ▶ It is possible to calculate the worst-case latencies

# Interrupts



- Interrupt latencies:
  - ▶ ISR latency
    - Time from when an IRQ is set at the CPU to when the ISR begins to run

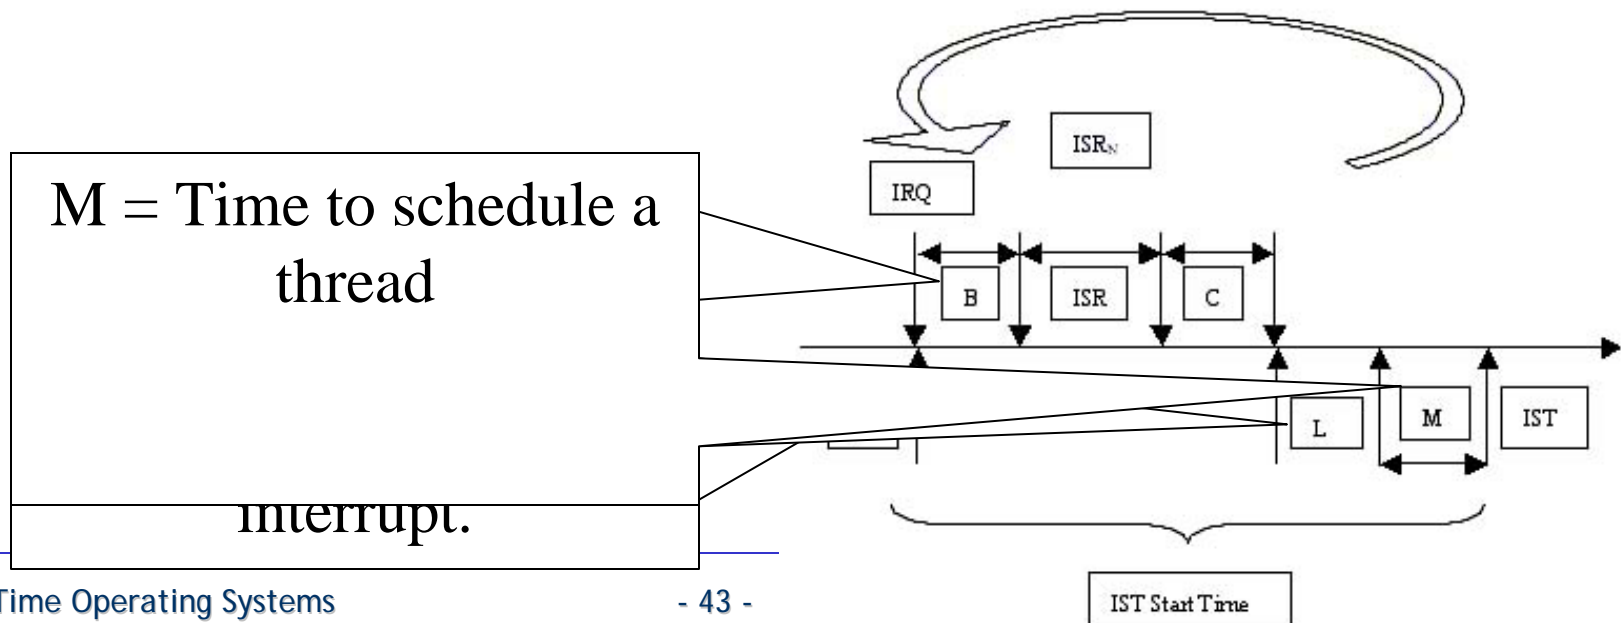
C = Time between when an ISR returns to the kernel and the kernel stops processing



# Interrupts



- Interrupt latencies:
  - ▶ IST latency:
    - Time from when an ISR finishes execution to when IST begins execution



# Interrupts

---



- Interrupt latencies:
  - ▶ ISR latency:
    - Windows CE is in control of the variables A, B, C, all of which are bounded
  - ▶ IST latency:
    - Windows CE is in control of the variables B, C, L, and M all of which are bounded.

# Synchronization

---



- Synchronization solutions:
  - ▶ Wait functions
  - ▶ Interlocked functions
  - ▶ Synchronization objects:
    - Event objects
    - Mutex
    - Critical sections
    - Semaphores
  - ▶ Messages

# Memory Architecture

---



- The RAM is divided into two areas:
  - ▶ The **object store** resembles a permanent, virtual RAM disk
  - ▶ The **program memory** works like a RAM, storing heaps and stacks
- The maximum size for the RAM file system is 256 MB
- The boundary between the two areas is movable

# Virtual Memory

---



- Windows CE implements a paged virtual memory management system
- Each process has 32 MB of virtual address space.
- This space is used for: processes, DLLs, heaps, stacks and virtual memory allocation
- Allocating virtual memory may affect real-time performance

# Heap Memory

---



- Windows CE only supports allocating fixed blocks in the heap
- It can lead to the heap becoming fragmented
- Fragmentation can affect performance
- To reduce the impact on allocating heap memory, a process should allocate it before proceeding with normal processing



# Stack Memory

---



- When a thread is created, Windows CE reserves enough memory space to accommodate the maximum stack size
- If no memory space is available, the thread is suspended until the space is granted
- To prevent the initial commitment from affecting performance, ensure that the thread is scheduled at least once before real-time processing

# Real-Time Measurement Tools

---



- Two kernel-level tools to test real-time performance:
  - ▶ Interrupt Timing Analysis (IntrTime)
    - IntrTime Command Prompt
    - Extnal Interrupt Response
  - ▶ Scheduler Timing Analysis (CEBench)



# Interrupt Timing Analysis

---

- IntrTime Command Prompt:
  - ▶ The aim is the measurements of ISR and IST latencies
  - ▶ The measurements are done using the system clock timer
  - ▶ The following variations can be introduced:
    - Setting the IST to run on various priorities
    - Flushing or not flushing the cache after each interrupt
    - Changing the ISR rate and number of interrupts captured
    - Printing or outputting to file the collected results



# Interrupt Timing Analysis

---

- External Interrupt Response
  - ▶ The previous method relies on the timer on the device itself, which may affect the measurements
  - ▶ In this analysis two machine can be set up:
    - A workstation that generates an external interrupt and measures the time it takes to receive acknowledgements from ISR and IST routines
    - Device-under-test that receives the external interrupt and toggles output lines when ISR and IST routines are reached



# Scheduler Timing Analysis

---

- It measures the time required to perform basic kernel operations
- Timing samples are collected for the following performance metrics:
  - ▶ Acquire/release critical section
  - ▶ Wait/signal an event
  - ▶ Semaphores
  - ▶ Mutexes
  - ▶ Voluntary yield using `Sleep(0)`

# Scheduler Timing Analysis

---



- The QueryPerformanceCounter is used to obtain timing information
- Using the function call to get time stamps is not free
- In cases where the operation takes a very short time to complete, the overhead of the function call becomes significant