



# Real Time Operating Systems

## Scheduling

Lecturer:

Prof. William Fornaciari

Politecnico di Milano

[fornacia@elet.polimi.it](mailto:fornacia@elet.polimi.it)

[www.elet.polimi.it/~fornacia](http://www.elet.polimi.it/~fornacia)

# Outline

---



- *Introduction to Real Time Systems*
- *The role of a RTOS*
- *Time constraints*
- The problem of scheduling RT activities
- Classification of scheduling strategies
- Review of RT scheduling approaches
- *VxWork real-time features*
- *Windows CE real-time features*
- *Unix SVR4, Win2000, Linux*

# Short-Term Scheduler enables RT

---

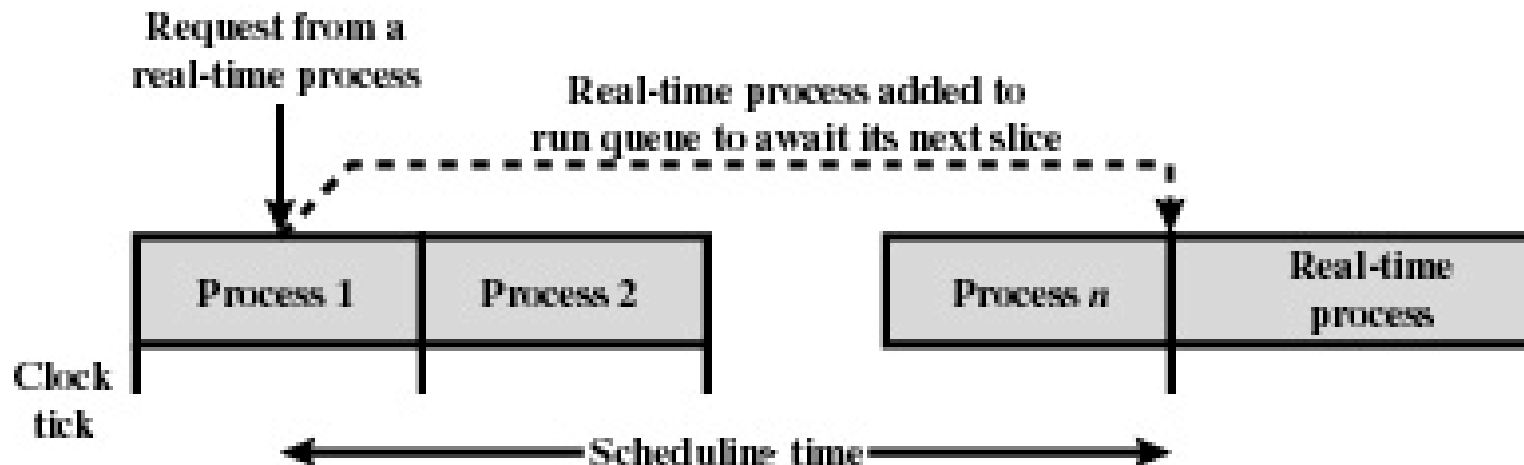


- The heart of a RTOS is the short-term scheduler
  - ▶ Fairness and min avg response time are not paramount
  - ▶ crucial: all hard-RT tasks must complete (or start) by their deadline and as many as possible soft-RT tasks should also complete (or start) meeting their deadlines
- Most current RTOSs are unable to deal with deadlines
  - ▶ they are designed to be as responsive as possible to RT tasks, so that, when deadline approaches, they can be quickly scheduled
  - ▶ this approach requires deterministic response time sometimes below milliseconds



# Towards a RT scheduler (1)

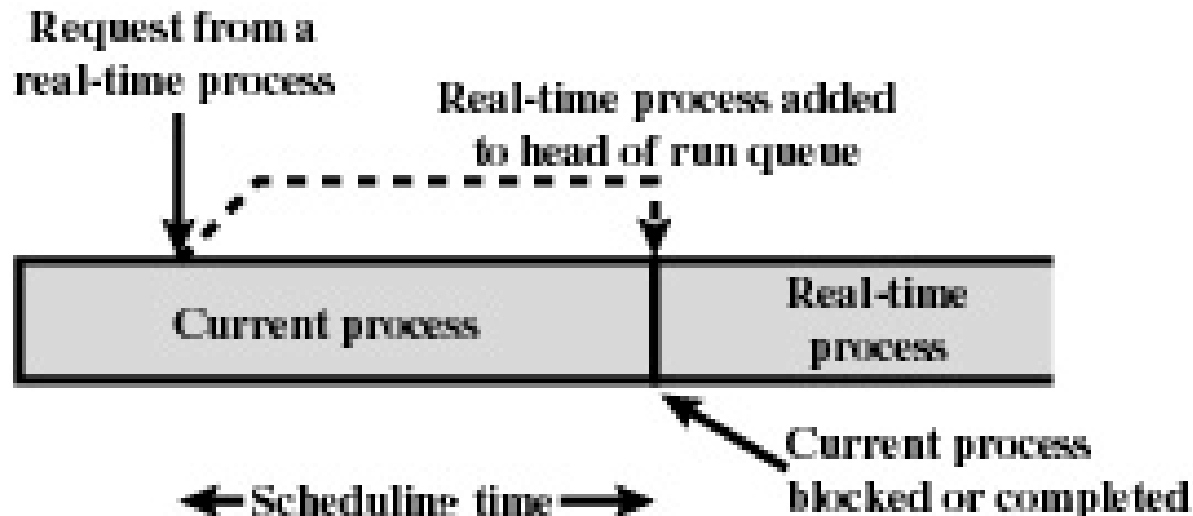
- Round Robin preemptive scheduler
  - ▶ The RT task is appended to the ready queue to await its next timeslice
  - ▶ The delay can be unacceptable for RT applications





## Towards a RT scheduler (2)

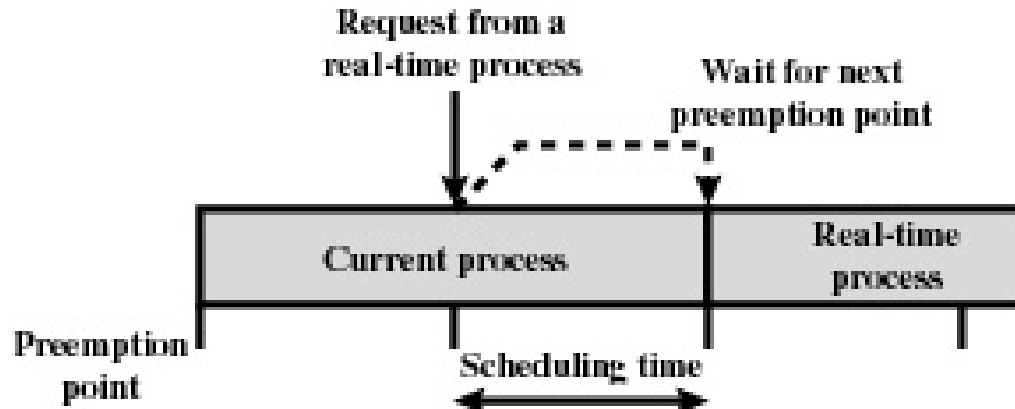
- Priority driven nonpreemptive scheduler
  - ▶ RT tasks have higher priority
  - ▶ A RT task is scheduled when the current P is blocked or runs to completion (even if with low priority)
  - ▶ Possible delay of seconds, unacceptable for RT



# Towards a RT scheduler (3)



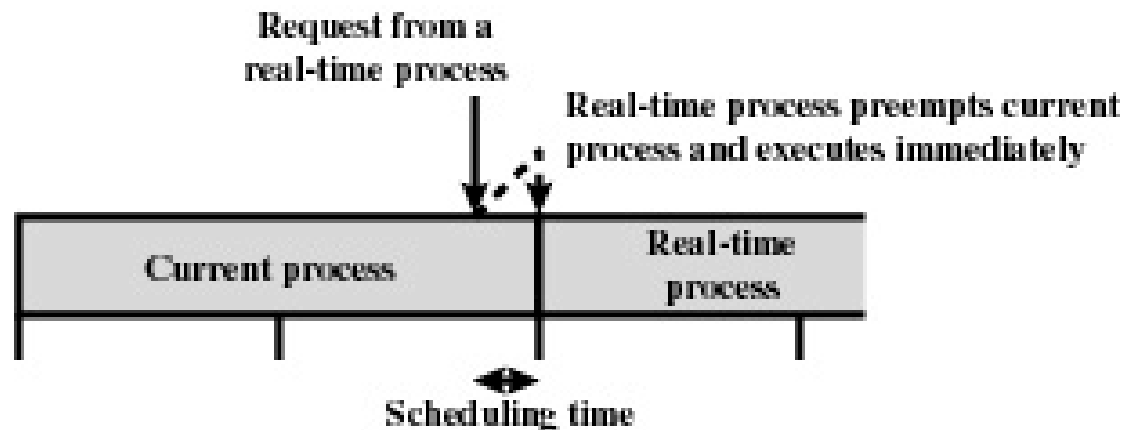
- Priority driven preemptive scheduler on preemption points
  - ▶ preemption takes place at the end of some regular intervals
  - ▶ in those points, the highest priority task is scheduled (including kernel tasks)
  - ▶ delays in the order of some ms, adequate for some applications not for more demanding ones





# Towards a RT scheduler (4)

- Immediate preemptive scheduling
  - ▶ Apart from the case OS is executing a critical region, the service to the interrupt is almost immediate
  - ▶ Scheduling delays fall down to 100  $\mu$ s, or less
  - ▶ Good for critical systems



(d) Immediate Preemptive Scheduler

# Factors influencing RT scheduling

---



- The system performs (or not) schedulability analysis
- Static vs dynamic schedulability analysis
- The result of the analysis can be
  - ▶ a clear scheduling
  - ▶ a strategy (plan) to be followed at run-time for task dispatching
- Classes of algorithms
  - ▶ static table-driven
  - ▶ static priority-driven preemptive
  - ▶ dynamic planning-based
  - ▶ dynamic best-effort



# Static table driven

---



- Through a static feasibility analysis of schedule, determines, at run-time, when a task must begin execution
- Applicable to periodic tasks
- Analysis inputs:
  - ▶ periodic arrival time
  - ▶ execution time
  - ▶ periodic ending deadline
  - ▶ relative priority of each task
- Predictable but inflexible approach: any change in the requirements of tasks imply the computation of a new schedule
- Example: Earliest-deadline-first



# Static priority-driven preemptive

---

- Uses the traditional priority-driven preemptive scheduler
- Static analysis is performed but no schedule is drawn-up, it is used to assign priorities to task
  - ▶ in no RT time-sharing systems, typ priority can change depending on I/O vs CPU bound process nature
  - ▶ in RT systems depends on time constraints associated with tasks
- Example: Rate Monotonic assigns static priorities to tasks based on the lengths of their periods

# Dynamic planning-based

---



- Feasibility is determined at run-time rather than offline
- A task is accepted for execution iff it is feasible to meet its time constraints
  - ▶ before its execution an attempt is made to create a schedule including previous tasks and the new one
  - ▶ the deadline of the extended task set must be met

# Dynamic best-effort

---



- Used in many commercial RT systems, easy to implement
- No feasibility analysis is performed
- Typ the tasks are aperiodic so that no static scheduling analysis is possible
- When a task arrives, the system assign a priority based on its characteristics (e.g. based on earliest deadlines)
- The system tries to meet deadlines and aborts any started process whose deadline is missed
- Until the task is completed (or deadline arrives), it is unknow if time constraints will be met



# RT Scheduling

Deadline based  
Rate monotonic

# Deadline scheduling (1)

---



- Real-time applications are not concerned with speed but with completing tasks
- Priorities provide a crude tool and do not capture the requirement of completion (or initiation) at the most valuable time
- Other deadline related information should be taken into account

# Deadline scheduling - information used

---



- Ready time
  - ▶ time at which task is ready for execution. For periodic task it is a sequence of times known in advance
- Starting deadline
- Completion deadline
  - ▶ typical RT application will have either starting or completion deadlines, but not both
- Processing time
  - ▶ in some cases it is supplied, in others OS measures an exponential average
- Resource requirements
  - ▶ in addition to microprocessor

# Deadline scheduling - information used

---



- Priority
  - ▶ measures relative importance of tasks. Hard-RT tasks have “absolute” priority
- Subtask structure
  - ▶ task possibly decomposed in mandatory (the only with hard RT deadlines) and optional subtasks



# Design issues considering deadlines

---



- Which task to schedule next?
  - ▶ For a given preemption strategy, using either starting or completion deadlines, scheduling tasks with the earliest deadline minimized the fraction of tasks that miss their deadlines (true for single and multiprocessor configurations)
- What sort of preemption is allowed?
  - ▶ When starting deadlines are specified, nonpreemptive scheduler makes sense
  - ▶ RT task has the responsibility to block itself after executing the critical portion, allowing starting other RT deadlines

# Executing profile of two periodic tasks

---

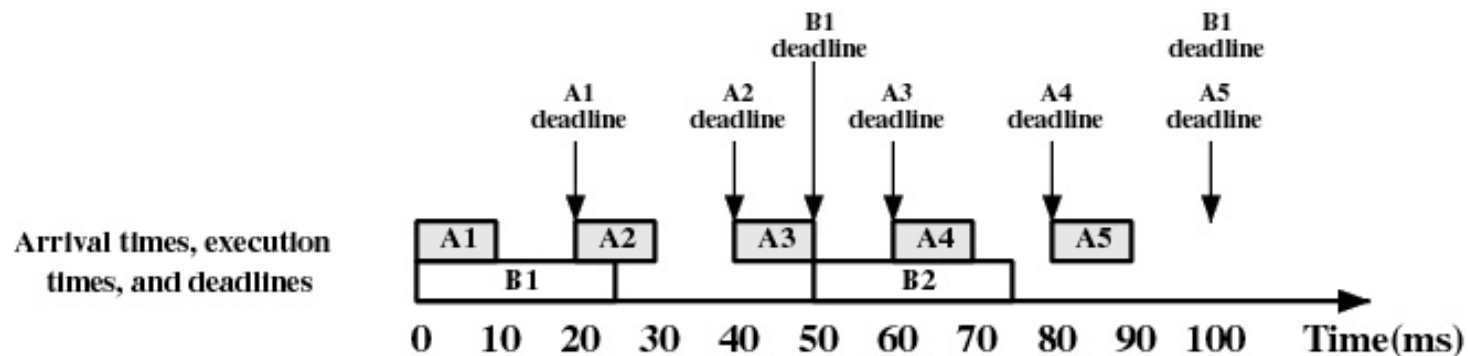


- System collecting and processing data from two sensors, A and B
  - ▶ Deadlines: A must be read every 20 ms, B every 50 ms
  - ▶ Processing time (including OS overhead): 10 ms for A samples, 25 ms for data from B
  - ▶ The computer makes scheduling decision every 10 ms

# Executing profile of two periodic tasks



Process	Arrival Time	Execution Time	Ending Deadline
A(1)	0	10	20
A(2)	20	10	40
A(3)	40	10	60
A(4)	60	10	80
A(5)	80	10	100
•	•	•	•
•	•	•	•
•	•	•	•
B(1)	0	25	50
B(2)	50	25	100
•	•	•	•
•	•	•	•
•	•	•	•



# Scheduling of periodic RT tasks with Completion deadlines (1)

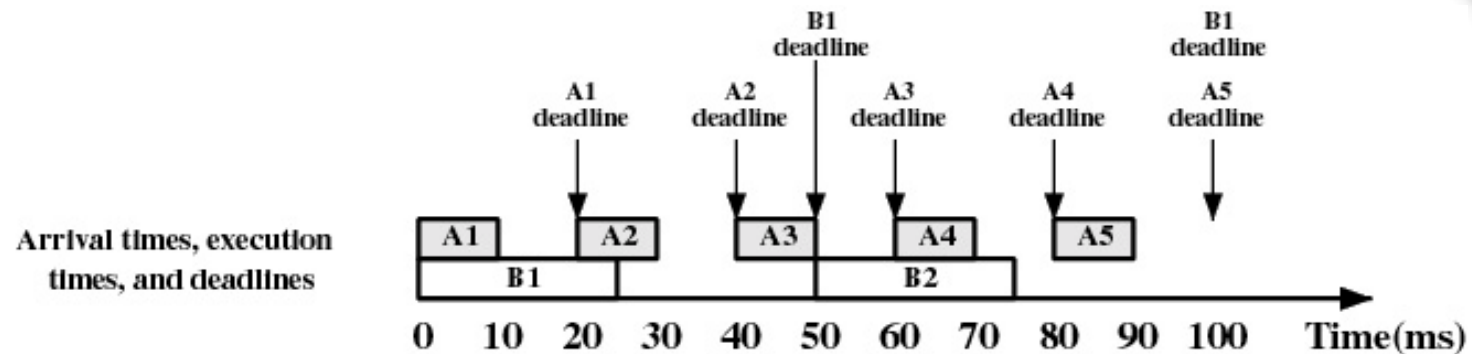
---




- Priority scheduling
  - ▶ A has higher priority -> B fails deadline
  - ▶ B has higher priority -> A fails deadline
- Earliest deadline schema
  - ▶ the scheduler gives priority at any preemption point to the task with the nearest deadline
  - ▶ All the deadlined can be met




# Scheduling of PERIODIC RT tasks with COMPLETION deadlines (2)



Fixed-priority scheduling;  
A has priority



Fixed-priority scheduling;  
B has priority



Earliest deadline scheduling  
using completion deadlines



# Scheduling of APERIODIC RT tasks with STARTING deadlines (1)

---



- Five tasks each having execution time of 20 ms
- Earliest deadline
  - ▶ schedule the ready task with the earliest deadline and let the task run to completion
  - ▶ The immediate service required by B is denied. Typical case of aperiodic task with starting deadline
- FCFS
  - ▶ Tasks B and E do not meet their deadlines


# Scheduling of APERIODIC RT tasks with STARTING deadlines (2)

---



- Earliest deadline with unforced idle times
  - ▶ Refinement possible if deadlines can be known in advance of the time when the task is ready
  - ▶ Always schedule the *elegible* task with the earliest deadline and let the task run to completion
  - ▶ Eligible tasks may not be ready -> processor can remain idle though there are ready tasks
  - ▶ All the requirements are met, even with non optimal processor exploitation

# Scheduling of APERIODIC RT tasks with STARTING deadlines (3)

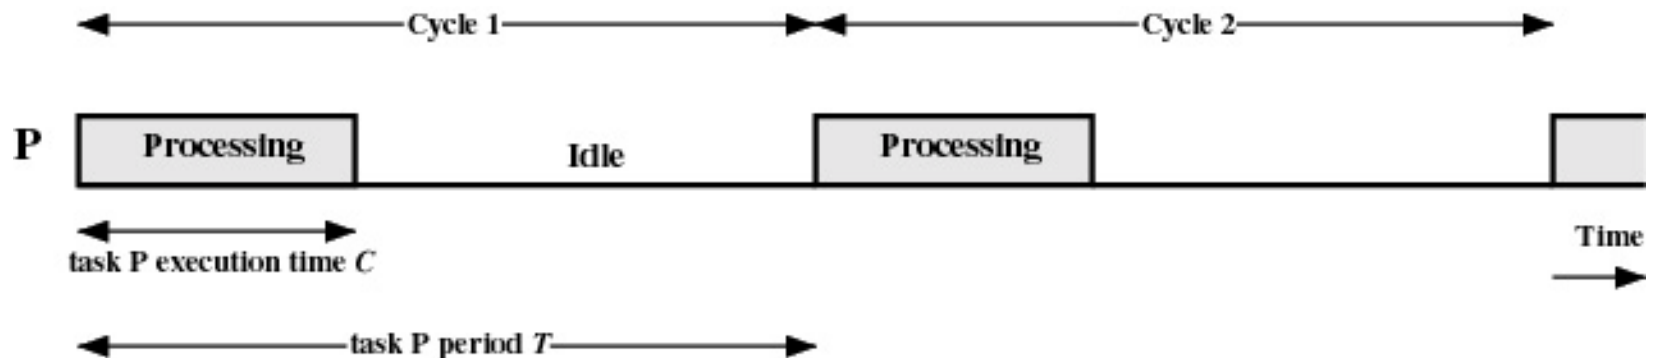






# Rate Monotonic Scheduling (RMS)

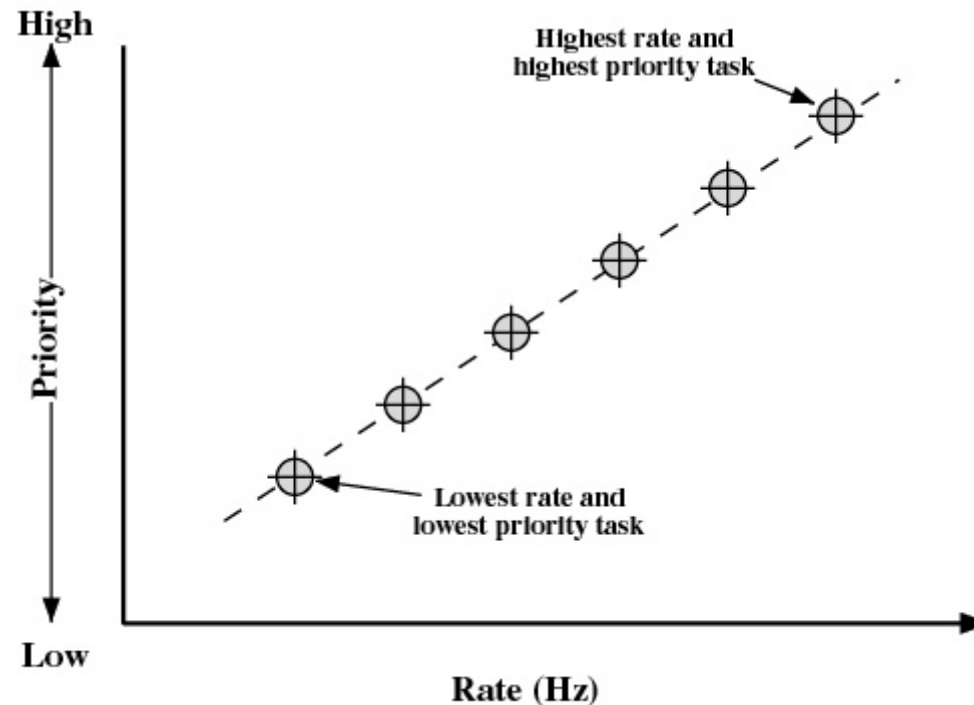
- Each Task has a period  $T$  and an Execution time  $C$  for each occurrence of the task
- In uniprocessor systems  $C < T$
- Processor Utilization  $U = C/T (< 100\%)$





# Rate Monotonic Scheduling (RMS)

- Assigns priorities to tasks on the basis of their periods
- Highest-priority task is the one with the shortest period



# Scheduling of periodic tasks: evaluation

---



- Effectiveness measure of a periodic scheduling algorithm: capability to meet deadlines
- n tasks with a fixed period and execution time, to meet all deadlines, the processor utilizations of individual tasks must not exceeds the available computational power

$$\frac{C_1}{T_1} + \frac{C_2}{T_2} + \dots + \frac{C_n}{T_n} \leq 1$$

- For RMS, it can be shown

$$\frac{C_1}{T_1} + \frac{C_2}{T_2} + \dots + \frac{C_n}{T_n} \leq n(2^{1/n} - 1)$$

# RMS - schedulability analysis

---



- The total utilization of the tasks must be less than the upper bound
- Examples of upper bounds
  - ▶  $n=2$ : 0.82
  - ▶  $n=4$ : 0.75
  - ▶  $n=6$ : 0.73
  - ▶ ...
  - ▶  $n \rightarrow \infty$ : 0.693 ( $\ln 2$ )
- The same constraint also holds for Earliest Deadline scheduling
- It is possible to achieve greater processor utilization with the EDF scheduling, nevertheless RMS is widely used for industrial applications...why?

# RMS wins over EDF?

---



- Performance difference is small in practice
  - ▶ upper bound is conservative, usage of 90% frequent
- Most Hard RT systems also have soft RT components
  - ▶ e.g. built-in self test, displays, ... can execute with low priority to absorb processor time left by RMS
- Stability is easier to achieve with RMS
  - ▶ in case system problems(e.g. due to transient errors), deadlines of essential tasks must be guaranteed
  - ▶ in static priority assignment one only need to ensure that essential tasks have relatively high priorities
  - ▶ in RMS is sufficient to organize essential tasks to have short periods or to modify RMS priority
  - ▶ in EDF a periodic task's priority can change from one period to another, difficult to meet deadlines



# RT Scheduling - Addendum

Cyclic Scheduling  
Deterministic Scheduling  
Capacity-Based Scheduling  
Dynamic Priority Scheduling  
Scheduling Tasks with Imprecise Results

# Scheduling

---



- Cyclic Scheduling
- Deterministic Scheduling
- Capacity-Based Scheduling
- Dynamic Priority Scheduling
- Scheduling Tasks with Imprecise Results

# Cyclic Executive

---



- A cyclic executive is a supervisory control program
- It schedules tasks according to schedule constructed during the system design phase
- The schedule consists of a sequence of action to be taken
- Long-term external conditions are used to choose a schedule for execution



# Cyclic Scheduling(1)

---



- The cyclic executive provides a practical means for executing a *cyclic schedule*
- The cyclic schedule is a timed sequence of computations which is to be repeated indefinitely, in a cyclic manner

# Cyclic Scheduling(2)

---



- Event-based processing can be handled with the cyclic schedule in different ways
  - ▶ Higher priority to the processes that are reacting to events
  - ▶ Allocate slots in scheduling where aperiodic, event-based processes can be serviced
  - ▶ Service aperiodic events in the background

# Cyclic Scheduling (3)

---



- Deterministic scheduling theory may be used to help find schedules
- Optimal deterministic scheduling algorithms require *a priori* knowledge
- Optimal non-preemptive scheduling of computations with timing constraints is NP-Hard

# Cyclic Scheduling(4)

---



- Advantages
  - ▶ Simple to implement, efficient, predictable
- Critical issues
  - ▶ Design
  - ▶ Runtime: the system cannot adapt to a dynamically changing environment
  - ▶ Maintenance: the code may reflect job splitting and sequencing details of the schedule

# Deterministic Scheduling(1)

---



- It provides methods for constructing schedules in which the assignment of tasks to processors is known exactly for each point in time
- Information needed a priori:
  - Vector processing time
  - Arrival time
  - Deadline
  - Priority
  - Task splitting (preemption vs non-preemption)

# Deterministic Scheduling(2)

---



- A schedule is an assignment of processors to tasks

- Each task in a schedule has:
  - Completion time
  - Flow time
  - Lateness
  - Tardiness
  - Unit penalty
- Schedules are evaluated using:
  - Schedule length
  - Mean flow time
  - Mean weighted flow time
  - Maximum lateness
  - Mean tardiness
  - Mean weighted tardiness
  - Number of tardy tasks

# Deterministic Scheduling(3)

---



- Properties of scheduling:
  - ▶ Not only measures for evaluating but also criteria for optimization
- Many of optimization problems are NP-hard
- Approximate solutions are found by relaxing some assumption

# Deterministic Scheduling(4)

---



- Limits of deterministic schedule:
  - ▶ The computation times are not known in advance
  - ▶ Time to produce a schedule is larger than the time it takes to service the tasks
  - ▶ Tasks arrive dynamically, updated schedule is needed



# Capacity-Based Scheduling(1)

---



- Requirements:
  - ▶ Information about the amount of computation
  - ▶ The amount of computation available
- For a restricted class of real-time activities it's possible to determine if a task is schedulable

# Capacity-Based Scheduling(2)

---



- Assumptions:
  - ▶ Each task in the task set must be periodic
  - ▶ The deadline is the end of the period
  - ▶ The computation time must be constant
  - ▶ Neither communication nor synchronization between tasks
  - ▶ No critical regions in any of the computation

# Capacity-Based Scheduling(3)

---



- Schedule by using fixed priority preemptive scheduling
  - ▶ Order the tasks according to their frequency
  - ▶ Assign integer priorities to the tasks
  - ▶ The higher priority assigned to highest frequency task
- This is called: Rate Monotonic Priority Assignment

# Dynamic Priority Scheduling(1)

---



- Task priorities may change
- The approach:
  - ▶ Define a selection discipline
  - ▶ Make on-line scheduling decisions
  - ▶ Consider the instantaneous state of the system

# Dynamic Priority Scheduling(2)

---



- Earliest-deadline-first:
  - ▶ All tasks are ready at time  $t=0$
  - ▶ The computations are non-preemptive
  - ▶ If the algorithm can schedule without missing any deadlines it minimizes the maximum lateness in the task set.

# Dynamic Priority Scheduling(3)

---



- Least-slack-time:
  - ▶ The slack-times the amount of time that the task can be delayed without missing its deadline
  - ▶ Order tasks according to nondecreasing slack-time
  - ▶ It maximizes minimum task lateness and minimum task tardiness

# Scheduling Tasks with Imprecise Results(1)

---



- Considerations:
  - ▶ A certain amount of processor time will produce a reasonably accurate result.
  - ▶ Any additional processing time will increase the accuracy
- Each task is considered to be divided into:
  - ▶ A mandatory subtask
  - ▶ An optional subtask

# Scheduling Tasks with Imprecise Results(2)

---



- The schedule guarantees:
  - ▶ All mandatory subtasks will be completed by their deadlines.
  - ▶ The optional subtasks are scheduled in the remaining processor time.
- Scheduling algorithms:
  - ▶ Mandatory subtasks: traditional deterministic scheduling theory
  - ▶ Optional subtasks: various scheduling criteria