



Power estimation of digital systems

System level power estimation

Outline



- Target
 - ▶ mixed Hw/Sw systems
- Power estimation of the Hw part
 - ▶ Based on constant additive model
 - ▶ Based on power FSMs
- Power estimation of the software part of the system
 - ▶ Gate and RT-level techniques
 - ▶ Bus activity based techniques
 - ▶ Instruction level techniques
- The POET project
 - ▶ Towards source-level power aware design and optimization

Constant additive models



- PowerPlay
 - ▶ Spreadsheet framework with WWW interface
 - ▶ Power estimation is simply based on manipulation of power, timing and area of the models available for the functional blocks of the systems
 - ▶ Models are parametrized and scalable w.r.t. supply voltage and technology
 - ▶ The estimation engine gathers data from the models and computes the overall sum (constant additive model)
- Basic model template, power dissipated by a component

$$P = C_{sw} V_{dd}^2 f + I V_{dd}$$

- ▶ C_{sw} average switched capacitance (different model for various components); I static current drawn from the supply
- ▶ No dynamic (profiling) parameters in the model -> lower accuracy than behavioral estimation

Constant additive models



- Computational block model
 - ▶ Relates switched capacitance to complexity (e.g., bit width)
 - ▶ The more complex the block, the larger the number of coefficients that must be used
 - ▶ Example: model with N capacitive coefficients
 - ▶ More accurate, analytical (closed) models can be derived if layout of the block is available
- Storage unit model
 - ▶ Registers can be treated as functional units
 - ▶ The model relates switched capacitance C_{sw} to complexity
 - ▶ Example: N words of size W

$$C_{sw} = C_0 + C_1N + C_2W + C_3NW$$

Constant additive models



- Controller model

- ▶ Different models for the switched capacitance C_{sw} are used depending on the implementation style (ROM-based vs random logic)
- ▶ Example: Random logic (N_1 inputs, N_0 outputs, N_m minterms)

$$C_{sw} = C_0 \alpha_0 N_1 N_0 + C_1 \alpha_1 N_1 N_0$$

- Interconnect model

- ▶ Derived from estimate of interconnect area
 - Rent's rule relates # of connection to a module to the number of gates (param. Feature size, Tr size, #layers, unit gate area, avg gate fanout)
- ▶ Line capacitance is then parametrized by capacitance per unit of area

Constant additive models



- Other available models
 - ▶ Programmable processor
 - ▶ Analog circuit
 - ▶ DC-DC converter
- Constant additive models of PowerPlay is limited by the fact that workload is left to the designer
 - ▶ Sizable estimation errors can be foreseen
- The estimation approach based on Power FSM tries to overcome such limitation
 - ▶ It is close to an implementation model, ready to compare alternatives




- Event-driven system model
 - ▶ The system is composed by a power manager, a set of resources and a battery
 - ▶ Users (environment) interacts with resources by feeding requests to the power manager
- Key feature
 - ▶ No overhead for long inactivity (no events)
- Resource model: PFSM
 - ▶ States have peak power annotation
 - ▶ Transition have performance annotation
 - ▶ Transitions are driven by the power manager

Power FSMs (PFSMs)



- Issues
 - ▶ State transition takes time
 - ▶ Transition times depend on head and tail states
 - ▶ Power depends on state
- Example: LCD display unit



Power FSMs (PFSMs)



- Improvement resource model: activity-level PFSM
 - ▶ Power annotation on a state indicates the peak power
 - ▶ Actually power depends on the workload (activity) of the resource
 - ▶ Use workload dependent power annotation
$$P_s = WL \text{ PeakPower}$$
 - ▶ Where WL is a Workload factor
- Power manager
 - ▶ Receives service requests
 - ▶ Dispatches requests to resources
 - ▶ Controls the power state of resources according to the selected power management policy
 - ▶ Can be implemented either Hw or Sw

Power FSMs (PFSMs)



- Additional models
 - ▶ User/environment (non deterministic FSM)
 - ▶ Power models for battery and DC-DC converters
- Power estimation engine
 - ▶ Based on simulation
 - ▶ Allows capturing non-uniform workloads
 - Activity bursts, inactivity periods
 - ▶ Enables to account for device latency
- Results
 - ▶ Benchmarked on real-life portable systems
 - ▶ Accuracy 10%-15% against measured data
 - ▶ High-simulation speed
 - ▶ Provides instantaneous power and performance estimates

Software Power Estimation



- Software-related power is becoming a significant component of the overall power budget
- Fast design space exploration allows
 - ▶ Simulation of the specifications, possibly in conjunction with hardware-bound parts of the system
 - ▶ Analysis of the power consumption at different levels of abstraction, to tradeoff accuracy and simulation speed
 - ▶ Comparison of alternative software codings
 - ▶ Comparison of different microprocessors
- To estimate software dissipation it is necessary to identify
 - ▶ Sources of software power consumption
 - ▶ Power models for the software and the
- Note that the terms *software* encompasses both system and application code

Sources of Sw power



- Memory system (caches and main memory)
 - ▶ Memory accesses are power consuming due to capacitive address line switching, control/decoding logic, internal word and data lines
 - ▶ Memory accesses are affected by software (e.g., locality of references)
- Bus (data, instruction, address)
 - ▶ Bus accesses are expensive because of line capacitances
 - ▶ Instruction and address buses are affected by SW; for data based accesses mostly decided at runtime
- Data-path (ALUS, FPU's)
 - ▶ Datapath accesses are data-dependent, but operation sequence and data dependencies are imposed by software

Approaches to Sw power estimation



- Different abstraction/accuracy pairs
 - ▶ Gate-level
 - ▶ RT-level
 - ▶ Bus activity based
 - ▶ Instruction level
- Gate level
 - ▶ Straight forward simulation on gate-level description of processor
 - ▶ Useful for benchmarking of other approaches
 - ▶ Characteristics
 - Slow but accurate
 - Unavailability of gate level description of non proprietary designs

Sw power estimation: RT-Level



- Requires a model of the processor in terms of high-level blocks (ALUs and memory) and their relative power dissipation estimates
- Each instruction is characterized in terms of the block that it activates
- The average power is obtained as

$$\text{Power} = (1/N_{\text{inst}})(\sum_I \sum_M \text{power}_{IM})$$

- Where M is the module activated by instruction I
- Faster than Gate-Level

Sw power estimation: Instruction-Level



- Possible approaches for instruction power characterization
 - ▶ Direct measurement of the currents drawn from the power supply while executing the instructions
 - ▶ HDL simulation
 - The instruction are simulated on a processor model in some HDL
 - The processor is plugged into a tester machine and simulation traces are applied. The current drawn is measured by the tester
 - ▶ Use simulation of a gate-level description of the processor

Sw power estimation: Instruction-Level



- A power cost is assigned to each instruction
- Two component of the cost
 - ▶ Base cost: static component, without the notion of processor state
 - ▶ Dynamic component: account for the circuit state effects, namely the previous processor state
- Dynamic cost modeling allows to capture events depending on sequence of events, like cache misses and pipeline stalls

Sw power estimation: Instruction-Level



- Base cost is computed
 - ▶ By repeating and infinite loop containing a total of N copies of the target instruction to be executed
 - ▶ The average current is then measured and the power cost is obtained from the current, supply voltage and cycle/instruction
 - ▶ N should be large enough to amortize loop control overhead
- Dynamic effects are currently modeled limiting the analysis windows to pairs of instructions
 - ▶ Marginally account for cache misses and pipelines stalls

Sw power estimation: Instruction-Level



- Circuit state effects are computed as follows
 - ▶ Execution of an infinite loop containing N copies for an alternating sequence of target pair of instruction $I_1 \rightarrow I_2$
 - ▶ The power cost is obtained as for the base cost
 - ▶ The averaging process shields the costs for $I_1 I_2$ and $I_2 I_1$
- The cost of a program can be so summarized
$$\text{Cost (program)} = \sum_i (B_i N_i) + \sum_{i,j} (O_{ij} N_{ij}) + \sum_k E_k$$
 - ▶ B_i Base cost of instruction i
 - ▶ N_i Number of occurrences of instruction i
 - ▶ O_{ij} Dynamic cost of sequence $i \rightarrow j$
 - ▶ N_{ij} Number of occurrences of sequence $i \rightarrow j$
 - ▶ E_k Other effects obtained from program profiling

A new functional approach working at instruction level



- Problem definition
- Model for functional decomposition
- Single processor model
- Multi processor model
- Experimental results
 - ▶ Identification of functionalities
 - ▶ Instruction characterization
 - ▶ Generalization properties
- Concluding remarks

Goals



- Definition of a formal instruction-level model, based on a *functional* decomposition of the μ P activities
- Intra-processor generalization
 - ▶ Capability to cope with incomplete instruction Set (IS) power characterization
- Inter-processor generalization
 - ▶ Capability to predict power characteristics of a μ P based on the data of other μ Ps and its IS
- Simplification of the measurements
- Statistical validation of the measurement data

The model: definitions



- *Functionality* F_i : set of activities involving -partially or totally- one or more μP units
- *Space-disjoint*: if F_i and F_j involve different units
- *Time-disjoint*: if F_i and F_j are accomplished at different times
- *Compatible model*: A set of functionalities forms a compatible model iff the current i_s absorbed by each instruction can be expressed as a linear combination of the currents if_j associated with F_j

$$i_s = \sum_{j=1..k} (if_j \cdot a_{s,j})$$

- ▶ The coefficients $a_{s,j}$ characterize each instruction in terms of the different functionalities
- ▶ Adherence to physical reality imposes $if_j > 0$

The model: single processor



- The energy e_s absorbed by instruction s is:

$$e_s = \sum_{j=1..k} e_{s,j} = V_{dd} \cdot i_s \cdot n_{ck,s} \cdot \tau$$

- Let $IN = \{i_s \cdot n_{ck,s}\}$

$$IN = A \times IF + R$$

- Solving this equation in the least square sense, gives estimate of the functionality currents

$$IF_{est} = IF + A^* \times R$$

$$\text{where } A^* = (A' \times A)^{-1} \times A'$$

The model: single processor



- A statistical analysis proves the formal correctness
- Applicability to the available measures can be justified through a $Z_{0.95}$ or $Z_{0.99}$ test
- The linearity of the model and its statistical properties allows to
 - ▶ Derive the model parameters IF_{est} from a limited set of instruction measures (*learning-set*)
 - ▶ Extrapolate the energy of other uncharacterized instructions (*generalization-set*)

The model: multiple processors



- For the same type of instruction
 - ▶ the *absolute* current strongly depends on the considered μP ,
 - ▶ the *relative* currents ($i_{rel,s}$) are *nearly independent* of the μP
- A general model based on a set P of p μP s for learning has been defined using the $i_{rel,s}$ defined as:

$$i_{rel,s} = i_s / i_{ref} = \sum_{j=1..k} if_{rel,j} \cdot a_{s,j}$$

k = # of functionalities considered

- For the generic q-th μP of P, characterized by $IN_{rel,q} = \{i_{s,rel} \cdot n_{ck,s}\}$ and A_q

$$IN_{rel,q} = A_q \times IF_{rel,q} + R_{rel,q}$$

The model: multiple processors



- Solving the above equation in the least square sense and considering that the model should depend on a general set of parameters IF_{rel} instead of microprocessor-specific:

$$\begin{aligned} IF_{est\ rel,q} &= A_q^* \times IN_{rel,q} = \\ &= IF_{rel} + A_q^* \times R_{rel,q} \end{aligned}$$

- Adding up the relations for all the p μP of the set

$$\begin{aligned} (1/p) \cdot \sum_{q=1..p} IF_{est\ rel,q} &= \\ = IF_{rel} + (1/p) \sum_{q=1..p} A_q^* \times R_{rel,q} \end{aligned}$$

The model: multiple processors



Result

- ▶ average of the estimated parameters of each μP composing P is an adequate estimator for the parameters of the general model
- ▶ Statistical analysis proves the formal correctness of this choice
- ▶ Increasing the number of considered processors the variance of the parameters decreases as:

$$\text{VAR}[IF_{\text{est}}] = (1/p^2) \sum_{q=1..p} \text{VAR}[IF_{\text{est,rel},q}]$$

Functionality identification



- Identification of the functionalities: compromise between accuracy and specific knowledge on μP architecture
- Selection of five functionalities, based on measured power consumption figures
 - ▶ Principal component analysis shows that no functionality can be neglected without affecting model accuracy

Functionality	Activity	Normalized contribution (typ)
F&D	Fetch and Decode	0.08
Br	Branch	0.08
WrReg	Write Register	0.13
A&L	Arithmetic and Logic	0.20
Ld&St	Load and Store	0.51

Functionality identification



- The values of $a_{s,j}$ have to be determined according to the identified functionalities
- Example: decomposition in F&D and Exec

$$a_{s,F\&D} \cdot if_{F\&D} + a_{s,Exec} \cdot if_{Exec} = i_s \cdot n_{ck,s}$$

- ▶ where, a reasonable choice for $a_{s,F\&D}$ and $a_{s,Exec}$ is:

$$a_{s,F\&D} = n_{ck,s,F\&D}, \quad a_{s,Exec} = n_{ck,s,Exec}$$

- General case: F_j is involved in the execution of the instruction s if the *activation coefficients* $b_{s,j} = 1$
- The relation between $b_{s,j}$ and $a_{s,j}$ is

$$a_{s,j} = \begin{cases} b_{s,j} n_{ck,s,F\&D} & j = 1 \\ b_{s,j} w_s & j = 2 \dots k \end{cases} \quad w_s = \begin{cases} 0 & \sum_{i=2}^k b_{s,i} = 0 \\ n_{ck,s,Exec} / \sum_{i=2}^k b_{s,i} & \text{otherwise} \end{cases}$$

Estimation on a single processor



- The stimulated functionalities depend on the **operation class** and the **addressing mode**


Class	Functionalities
Arithmetic & Logic	F&D, A&L
Data Transfer	F&D
Control	F&D, Br
System	F&D, Br
Floating Point	F&D, A&L
Decimal	F&D, Br
String	F&D, A&L, WrReg

Addressing mode	Sample	Functionalities
Register	R2	WrReg
Relative	10(R2)	Ld&St
Auto-increment	(R2)+	Ld&St, A&L, WrReg

Estimation on a single processor



- 6 μ Ps considered, with similar results
 - ▶ ARM7, StrongARM, i960JF, i960HD, SPARC, i80486DX
 - **Example I486DX**. Gaussian noise Hypothesis holds:
 $\mu_R = 9.94$ falls in the range $z_{0.95} = \pm 40$.
 - Comparison with 18 energy characterized instructions




Functionality	F&D	Br	WrReg	A&L	Ld&St
Current [mA]	421	355	228	228	505
STD	48	26	46	38	39

Generalization on single processor



- Procedure


- ◎ Select learning-set from characterized instructions
 1. Estimate the currents of the remaining instructions (generalization set)
 2. Compute the learning and generalization errors
 - Errors tend to compensate (mean $\sim 10^{-10}$): absolute error values have been used to assess the methodology accuracy, that is always $< 9\%$



Generalization over different μ Ps



- Combinations of μ Ps are used to generate different μ P learning-sets, whose estimated model parameters allow prediction of the parameters of other μ Ps




- According to the theoretical model, general mean error decreases as the μ P learning-set size increases, and similarly the variance

Generalization over different μ Ps



- Example: estimation of model parameters for the i486 based on those of SPARClite and ARM7TDMI





- Identification of a widely applicable general model to estimate power consumption of 32-bits μ Ps
 - ▶ Single μ P: copes with partially characterized Instruction Sets
 - ▶ Multiple μ Ps: considering relative currents, allows prediction of the characteristics of a μ P capitalizing the knowledge of other characterized μ Ps
 - ▶ Multiple μ Ps: if needed, the current per clock cycle of a **single** specific instruction might be used as reference to compute absolute values

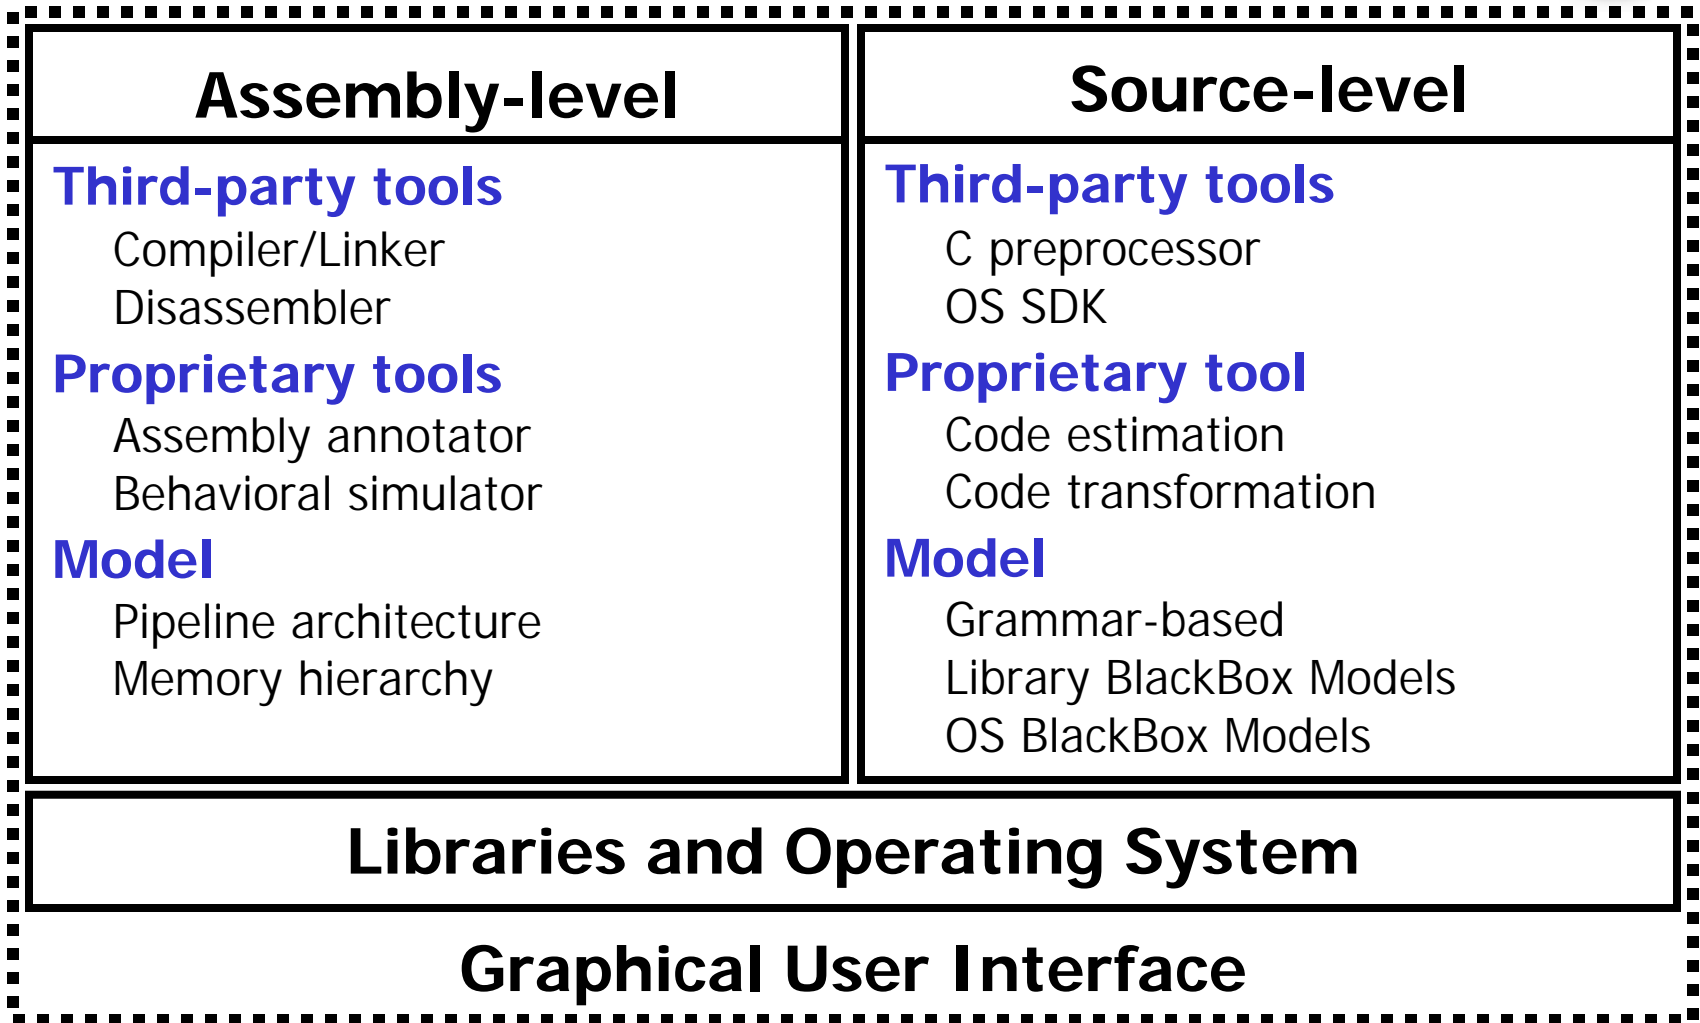
Current enhancements



- The presence of dynamic effects (assembly level) has been added according to other models considering higher-level features of code execution (see POET project)
- Execution time on a pipelined and/or superscalar architecture depends on
 - ▶ The **nominal execution time** $n(s)$ of instruction s
 - Is the nominal latency of an instruction (static measure)
 - Can be found in the documentation of the processor
 - ▶ The **execution time overhead** $oh(s)$ due to stalls
 - Depends on the status of the pipeline(s)
 - Is a dynamic measure
 - ▶ The **actual parallelism** $p(s)$ that instruction s exploits
 - Depends on the status of the pipeline(s)
 - Is a dynamic measure
- This is summarized by the relation

$$CPI(s) = p(s) \times [n(s) + oh(s)]$$

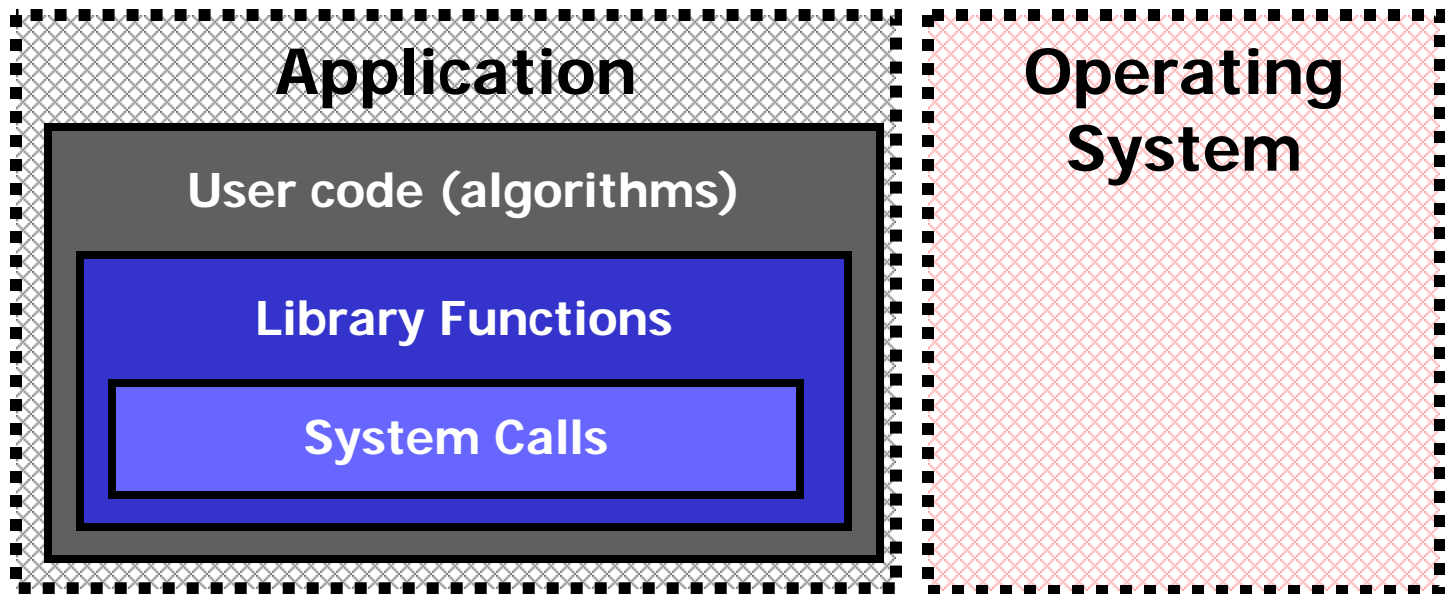
POET tools and methodology



Source level - Models



- Timing model
 - ▶ **Deterministic** execution time contribution
 - ▶ **Statistical** execution time correction
 - ▶ **Statistical black-box** library function and OS models




Source level - Models



- Source code
 - ▶ A **tree** of atoms
- Atom
 - ▶ Statement
 - Entire
 - Part
 - ▶ Defined by a **grammar**
 - ▶ Atom name
 - Left-hand side symbol
 - ▶ Atom structure
 - Right-hand side terminal symbols

```
...  
    if( a>0 || b<t[2] ) {  
        b++;  
    } else {  
        while( ... ) {  
        }  
    }  
...
```



Source level - Models



- Each atom γ_k is characterized by
 - ▶ A **determinisc** time contribution t
 - Local to the specific atom
 - Based on a virtual architecture/compiler pair
 - ▶ A **statistical** time correction δ
 - Global to more atoms
 - Based on the results of actual compilation
 - ▶ An **average current absorption** per clock cycle
- Thus:

$$t_{\text{est}}(\gamma_k) = t(\gamma_k) + \delta(\gamma_k)$$

$$e_{\text{est}}(\gamma_k) = t_{\text{est}}(\gamma_k) \times i_{\text{ave}}(\gamma_k) \times V_{\text{dd}} \times T_{\text{ck}}$$

Library analysis - Models



- Concentrates on execution timing
 - ▶ Highly data-dependent
 - ▶ Huge amount of data to be processed
- Energy is derived based on average power dissipation data of the target architecture
- Based on statistical analysis
 - ▶ Library functions used in different contexts
 - Computation-dominated applications
 - Control-dominated application
 - Text processing
 - ▶ Construction of a function timing catalog



- To derive a black-box model
 - ▶ A function must be **executed** with **different data**
- This requires
 - ▶ Describing the **arguments semantics** in a formal way
 - The C typesystem is not very tight
 - The same C type is used for data with different semantics
 - The semantic of a formal argument gives hints on the way the datum is manipulated by the function
 - ▶ Describing the statistical **distribution of the data** that the function is supposed to process in the specific application

Library analysis - Models



```
int strlen( char* s )
```

`char*`

Formal type

`string`

Semantic type: `string`
D-parameter: `length`

`length`

Distribution: `gaussian`
Parameters: `avg`, `std`

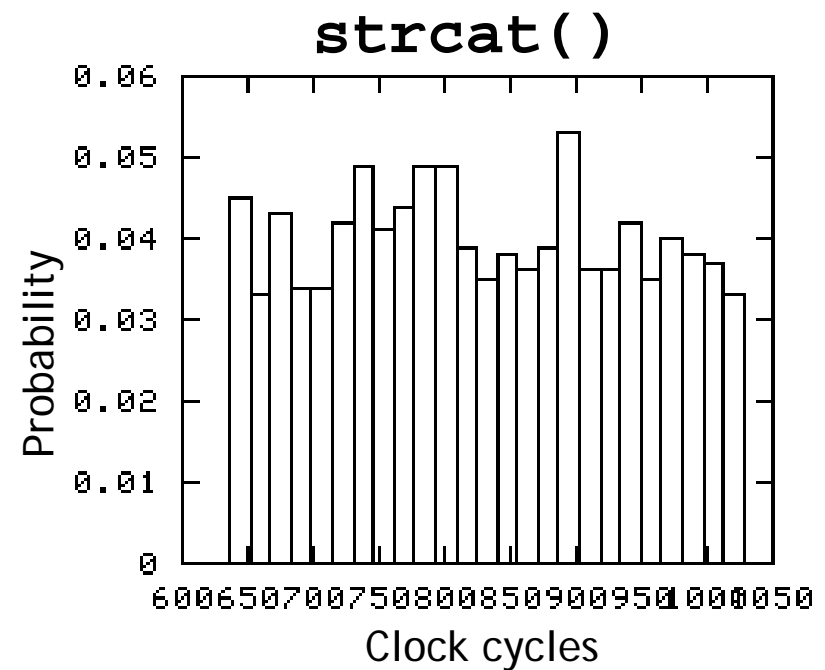
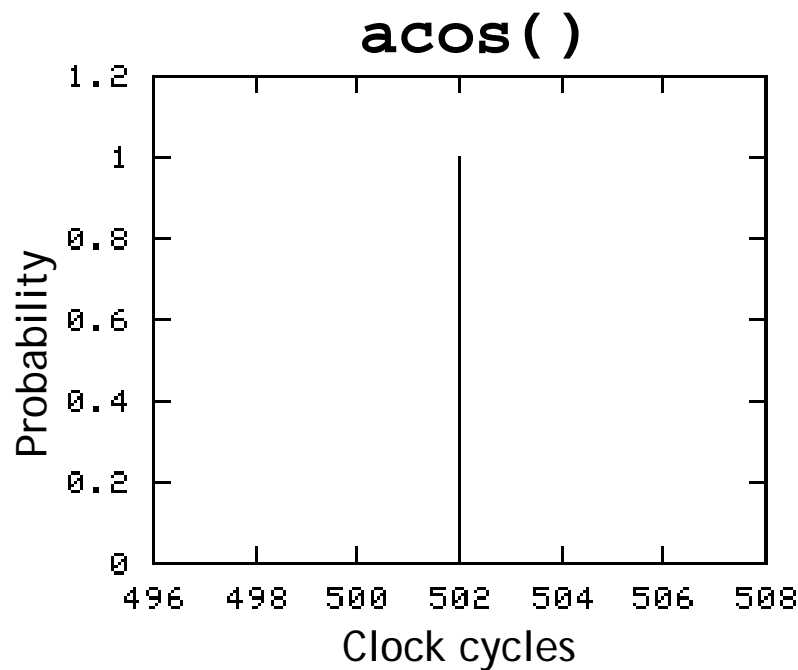
`avg=60`
`std=12`

Distribution
parameters

Library analysis - Results



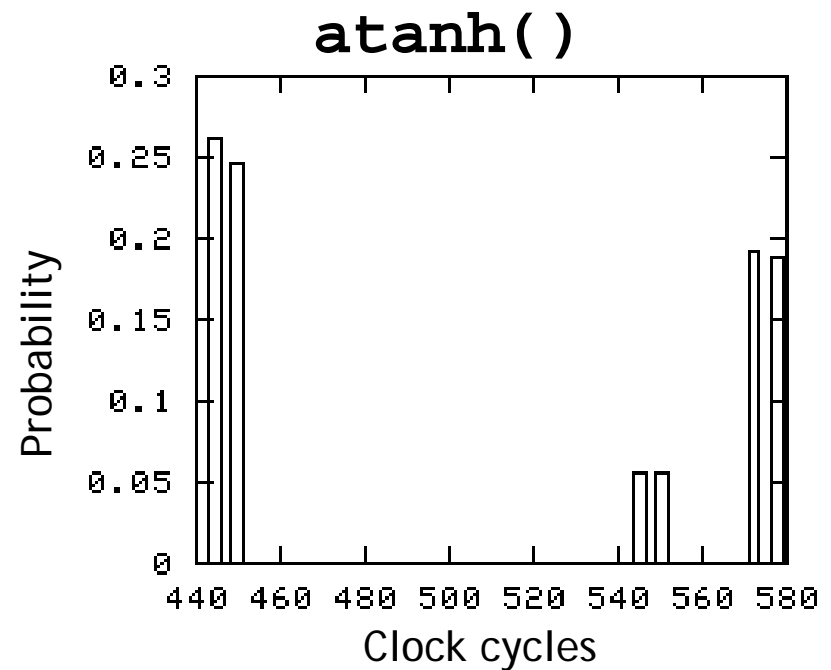
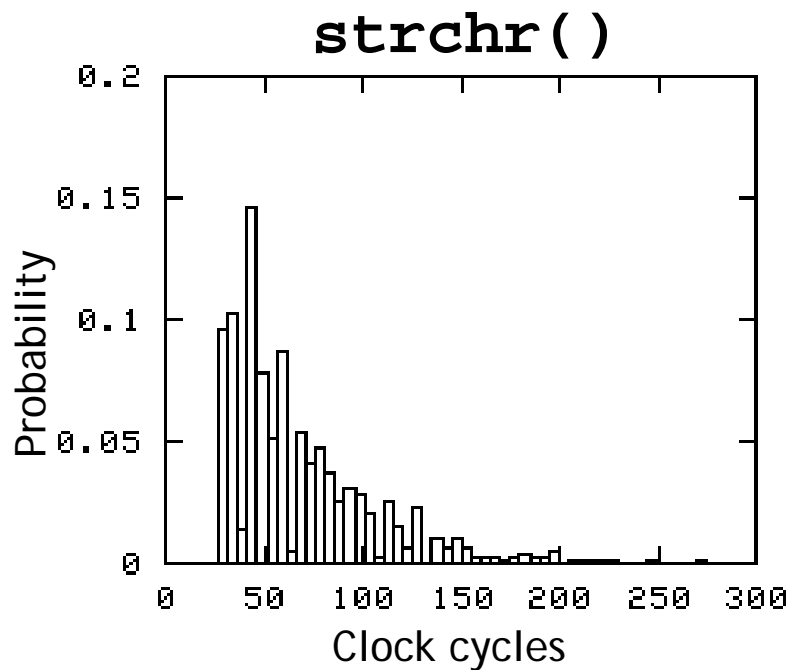
- Routines with simple behavior
 - ▶ Post-processing **can extract** analytical models



Library analysis - Results



- Routines with complex behavior
 - ▶ Post-processing **can hardly extract** analytical models



Source level - Results



- Based on the deterministic contribution only

Program	Time (Clock cycles)			Energy (μ J)		
	Real	Estimated	Error (%)	Real	Estimated	Error (%)
loop	25256	25394	-0.54	1867	1888	-1.11
factorial	56940	61030	-6.70	4169	5137	-18.84
mcd3	32193	32094	+0.31	2379	2373	+0.27
arith	48294	49949	-3.31	3553	3678	-3.41
crc16	226093	246012	-8.10	17510	18539	-5.55
base64enc	162445	149211	+8.88	11501	11098	+3.63
real2frac	46735	48688	-4.01	3407	3533	-3.57
pfactor	84058	84158	-0.12	6264	6500	-3.63
adpcm	41384	37629	-9.07	3213	3275	+1.93
Average	4.55			4.65		

Source level - Results



- Effect of the statistical contribution
 - ▶ A linear error can be corrected assuming $\delta = \text{const}$

